
The art of service discovery at scale

Nitesh Kant,
Software Engineer,
Netflix Edge Engineering.



@NiteshKant







Nitesh Kant

 @NiteshKant

Who Am I?

- ❖ Engineer, Edge Engineering, Netflix.
- ❖ Core contributor, RxNetty*
- ❖ Contributor, Zuul**
- ❖ Ran Netflix's Service Discovery for a year.
- ❖ Conceptualized and designed Eureka v2***.

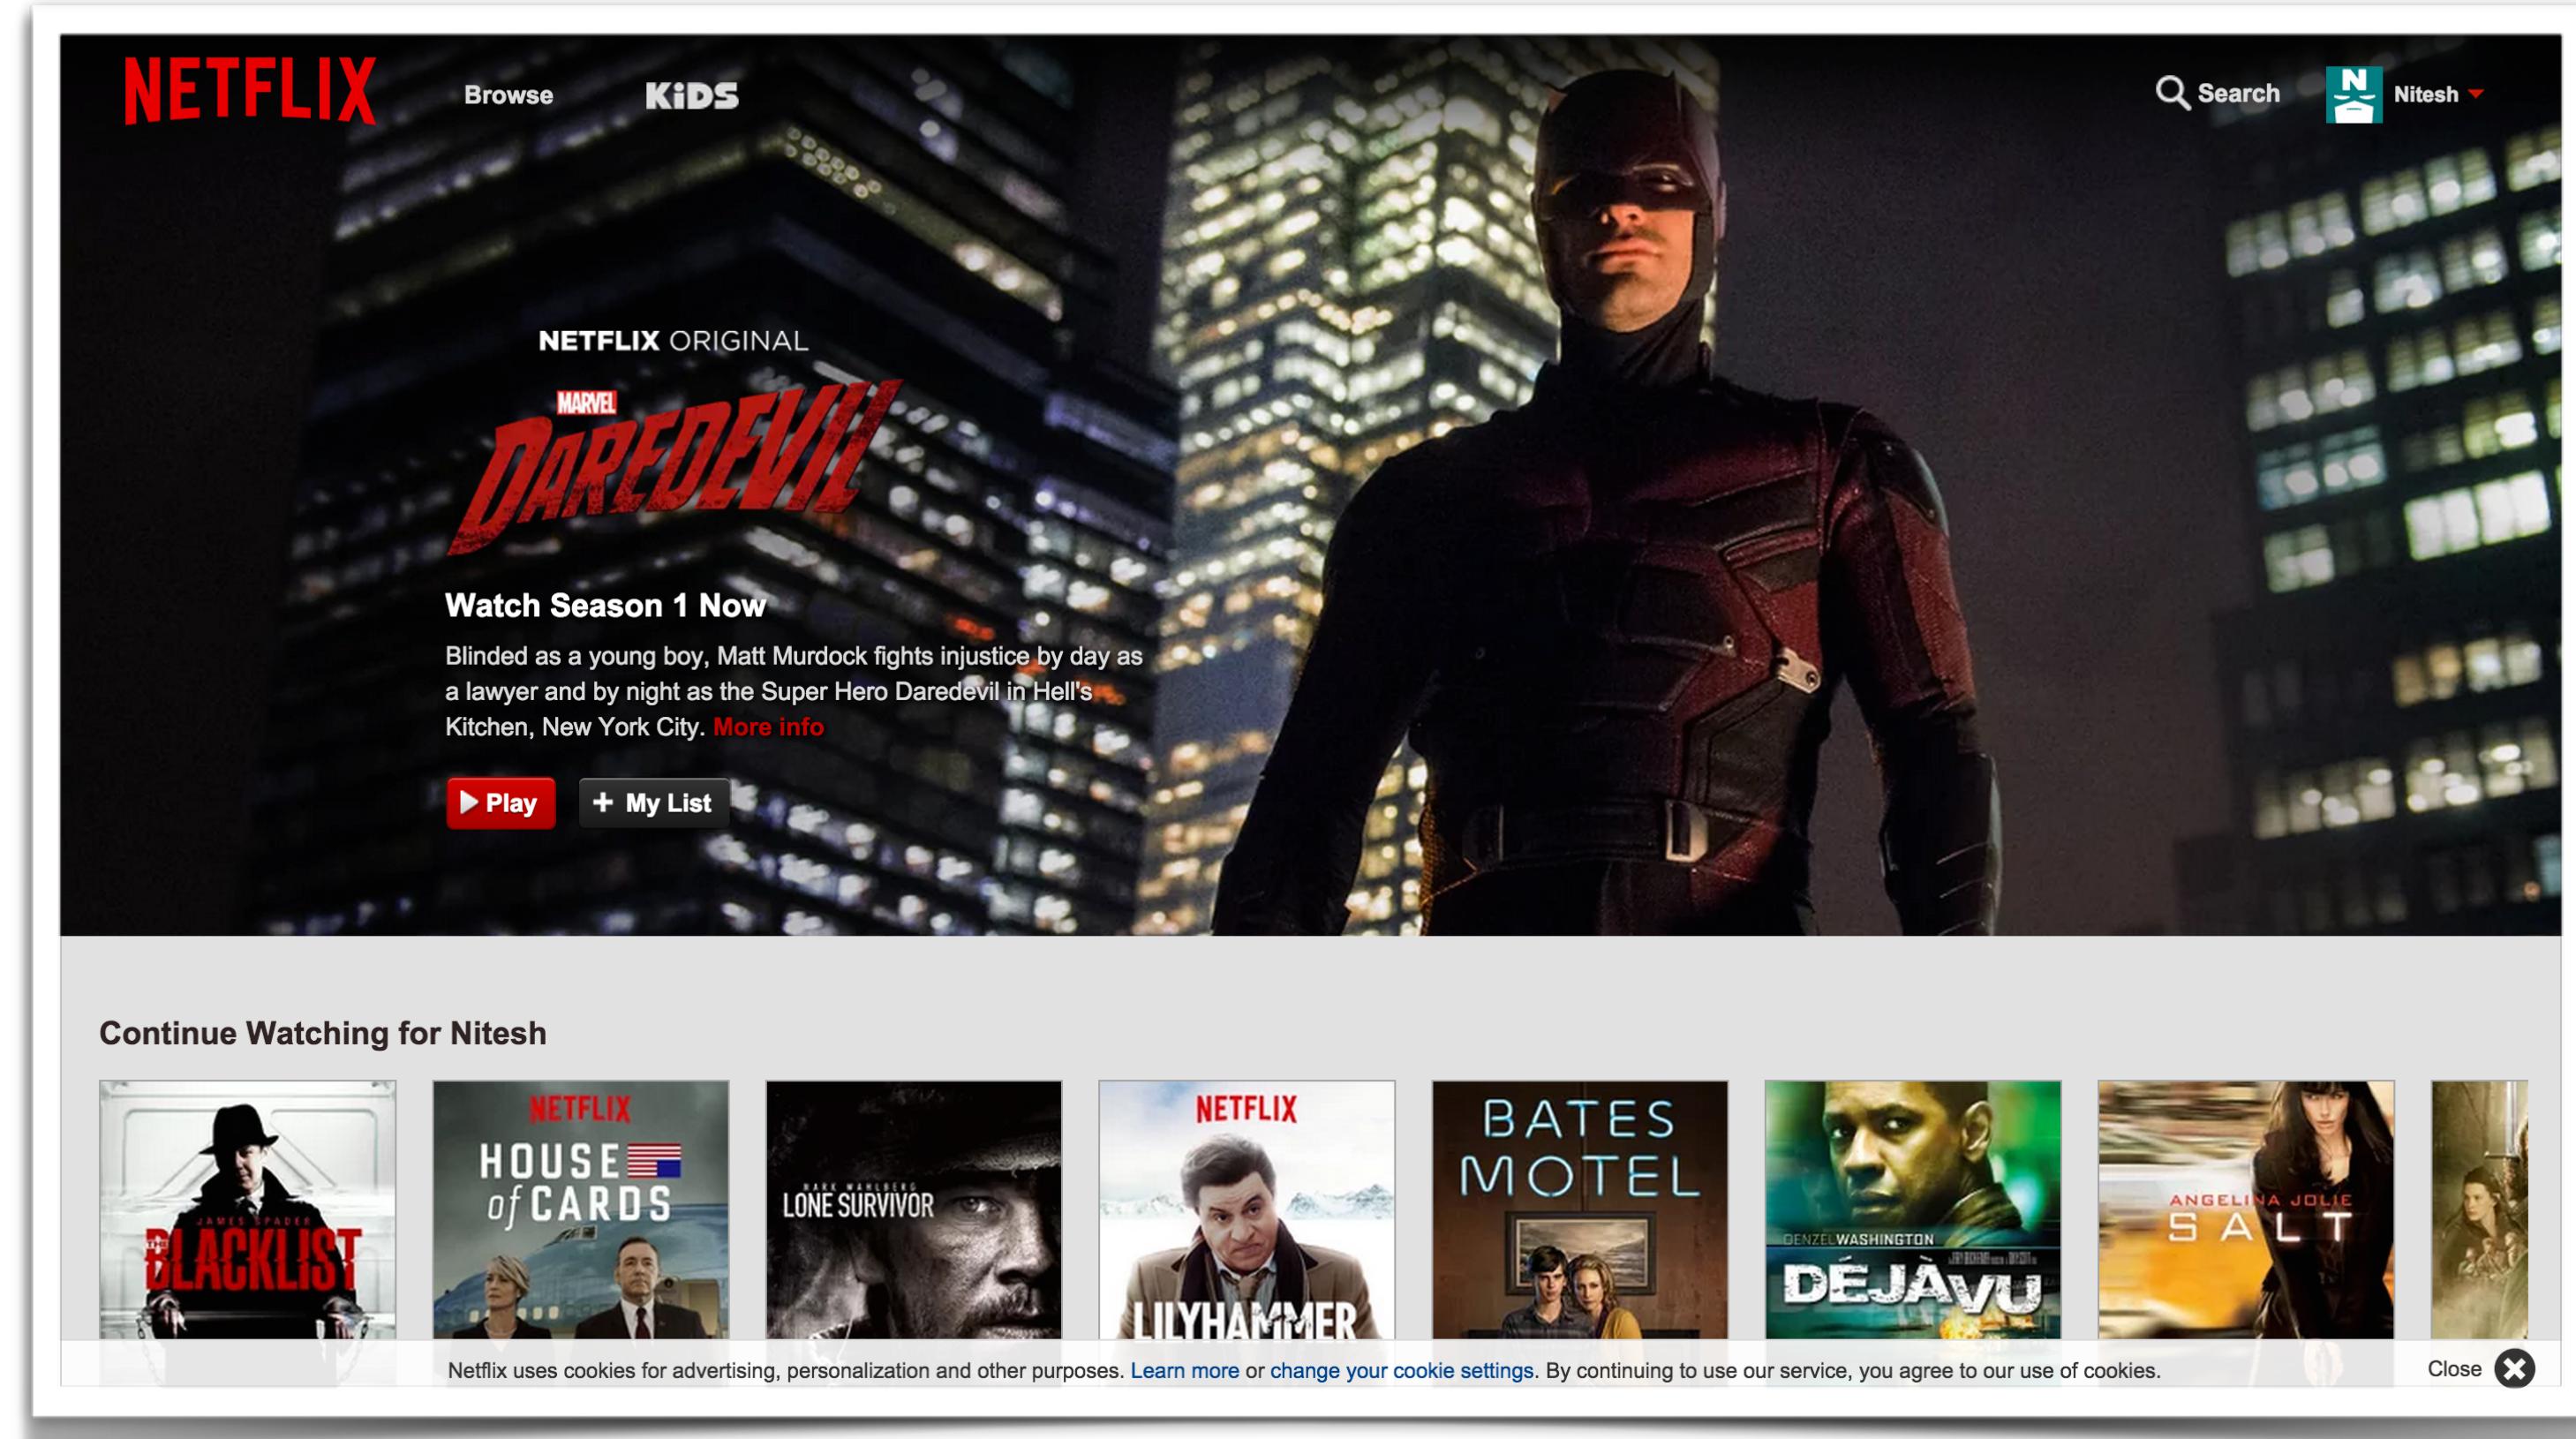


* <https://github.com/ReactiveX/RxNetty>

** <https://github.com/Netflix/zuul>

*** <https://github.com/Netflix/eureka>

What is Service Discovery?



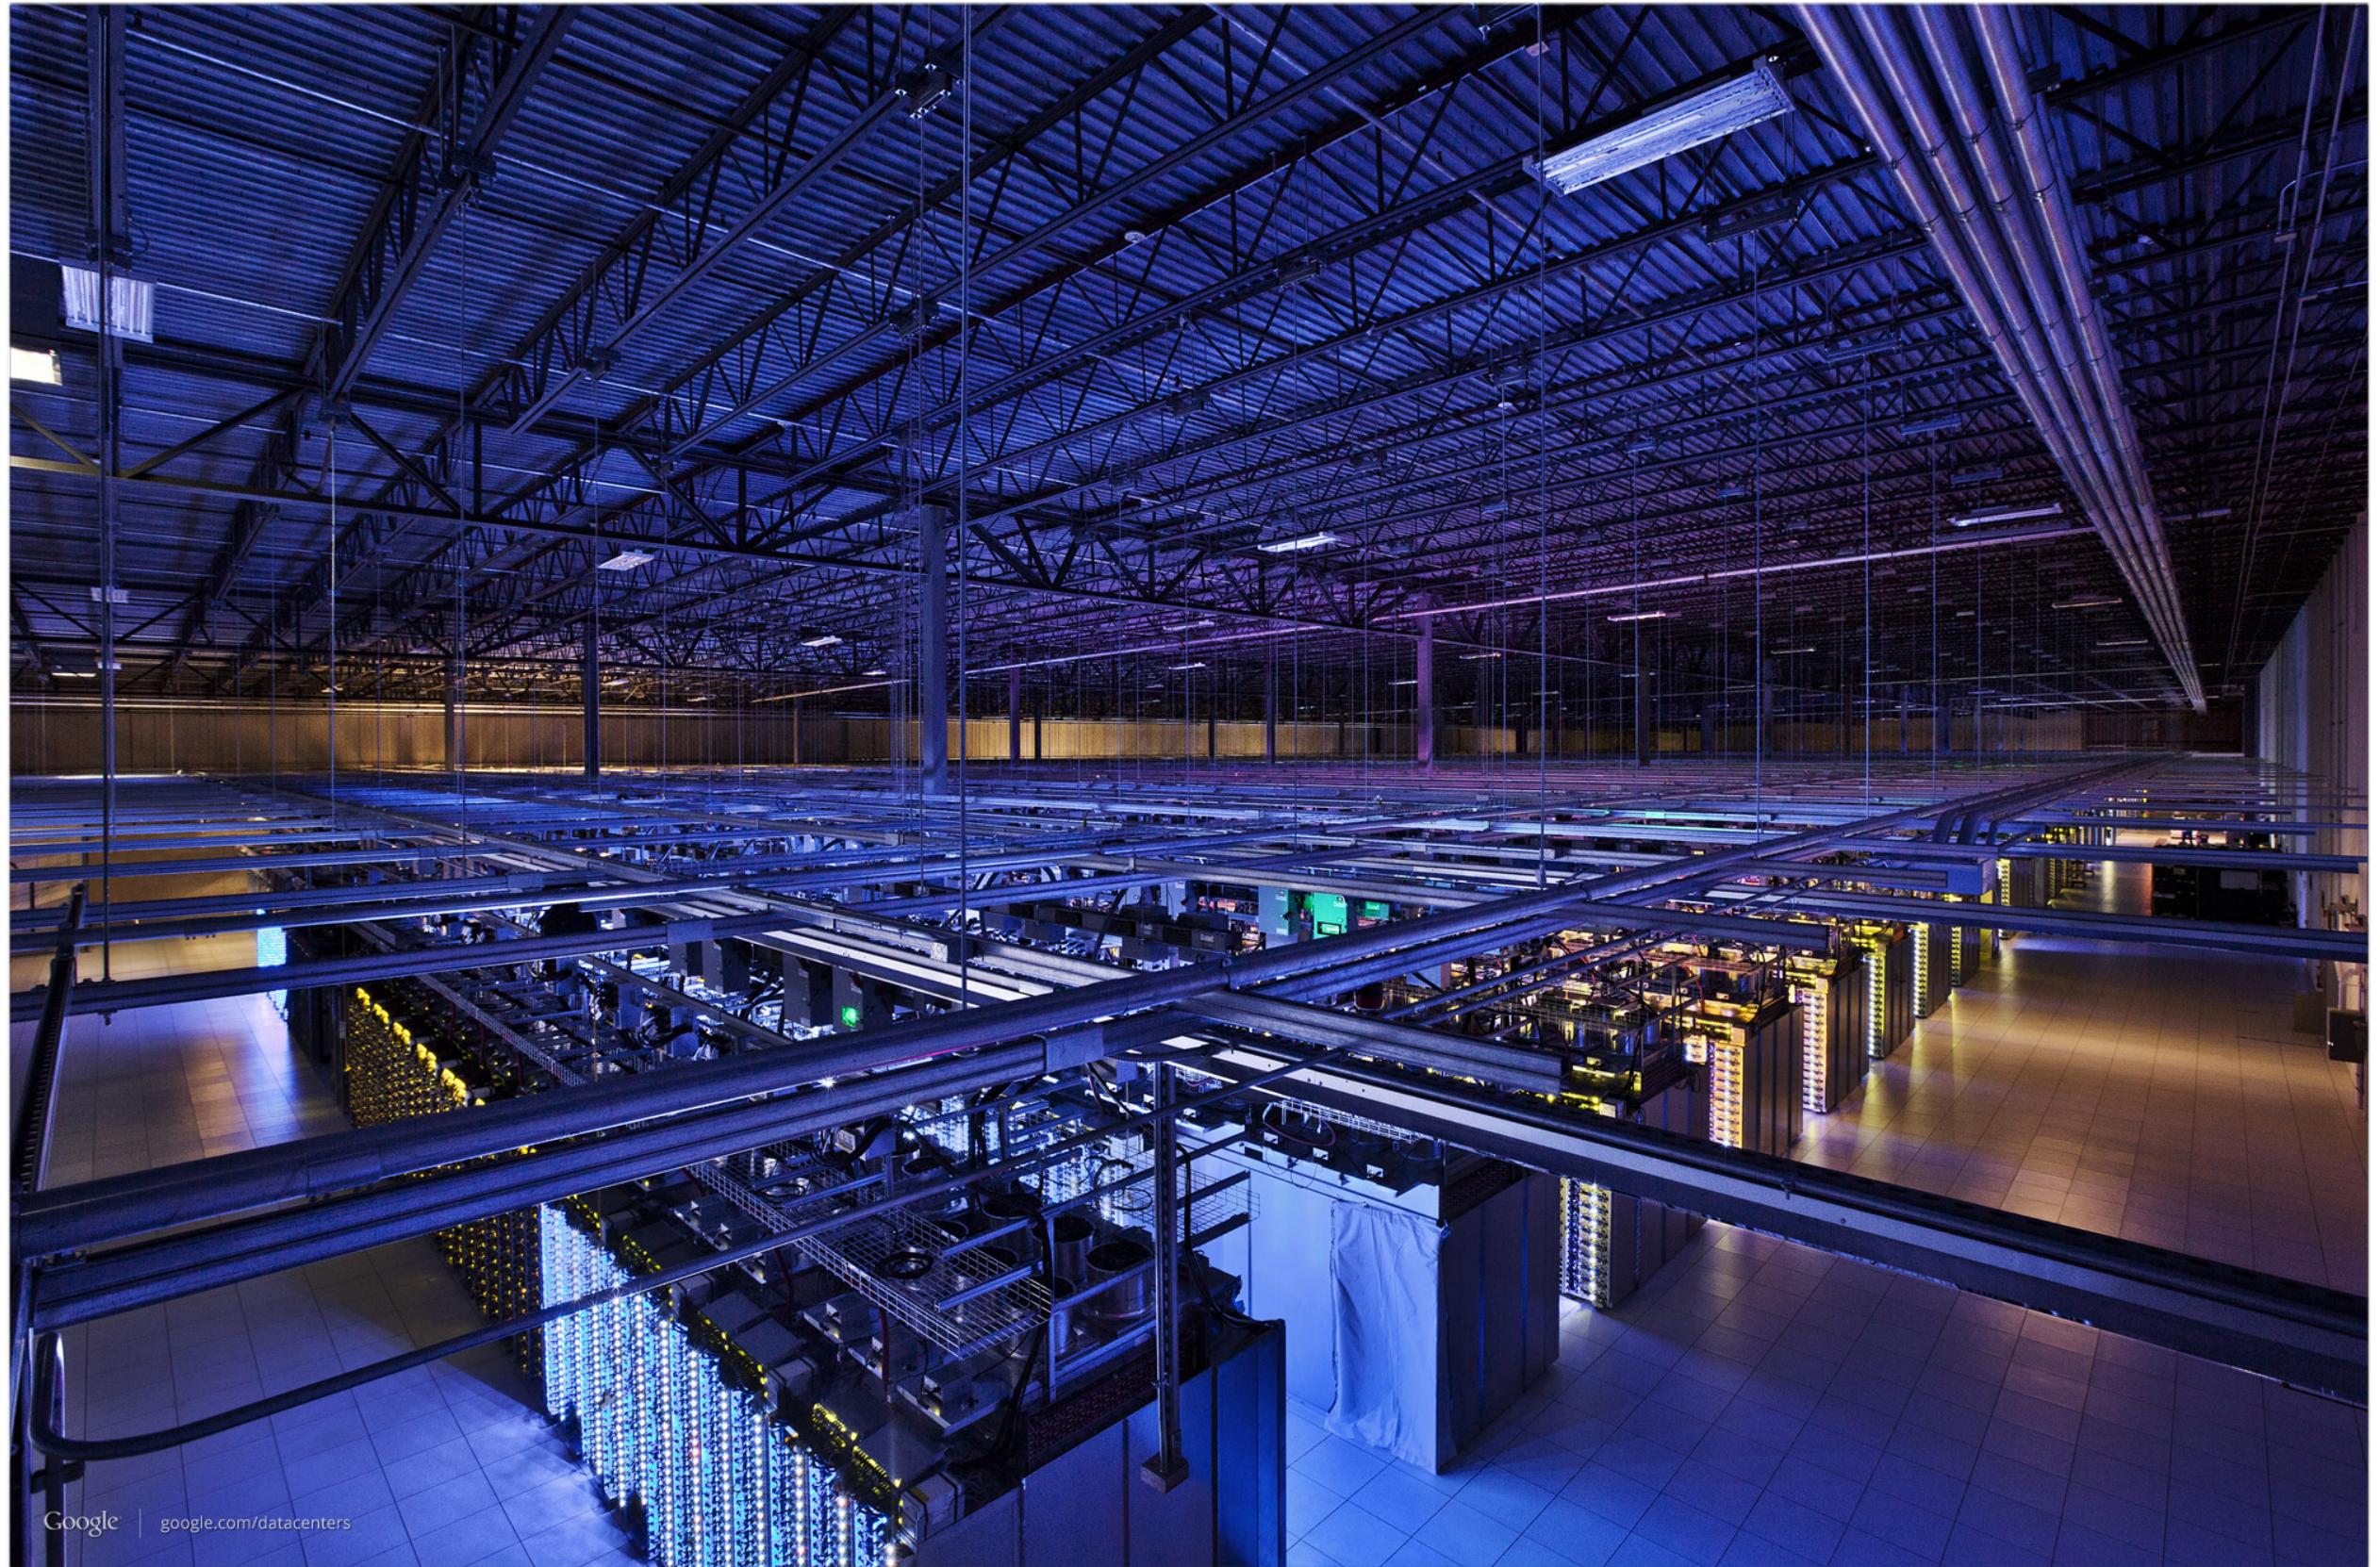
Delivering Netflix

What goes behind the scenes?

Pipe Dream

Always available, infinitely
scalable machine.



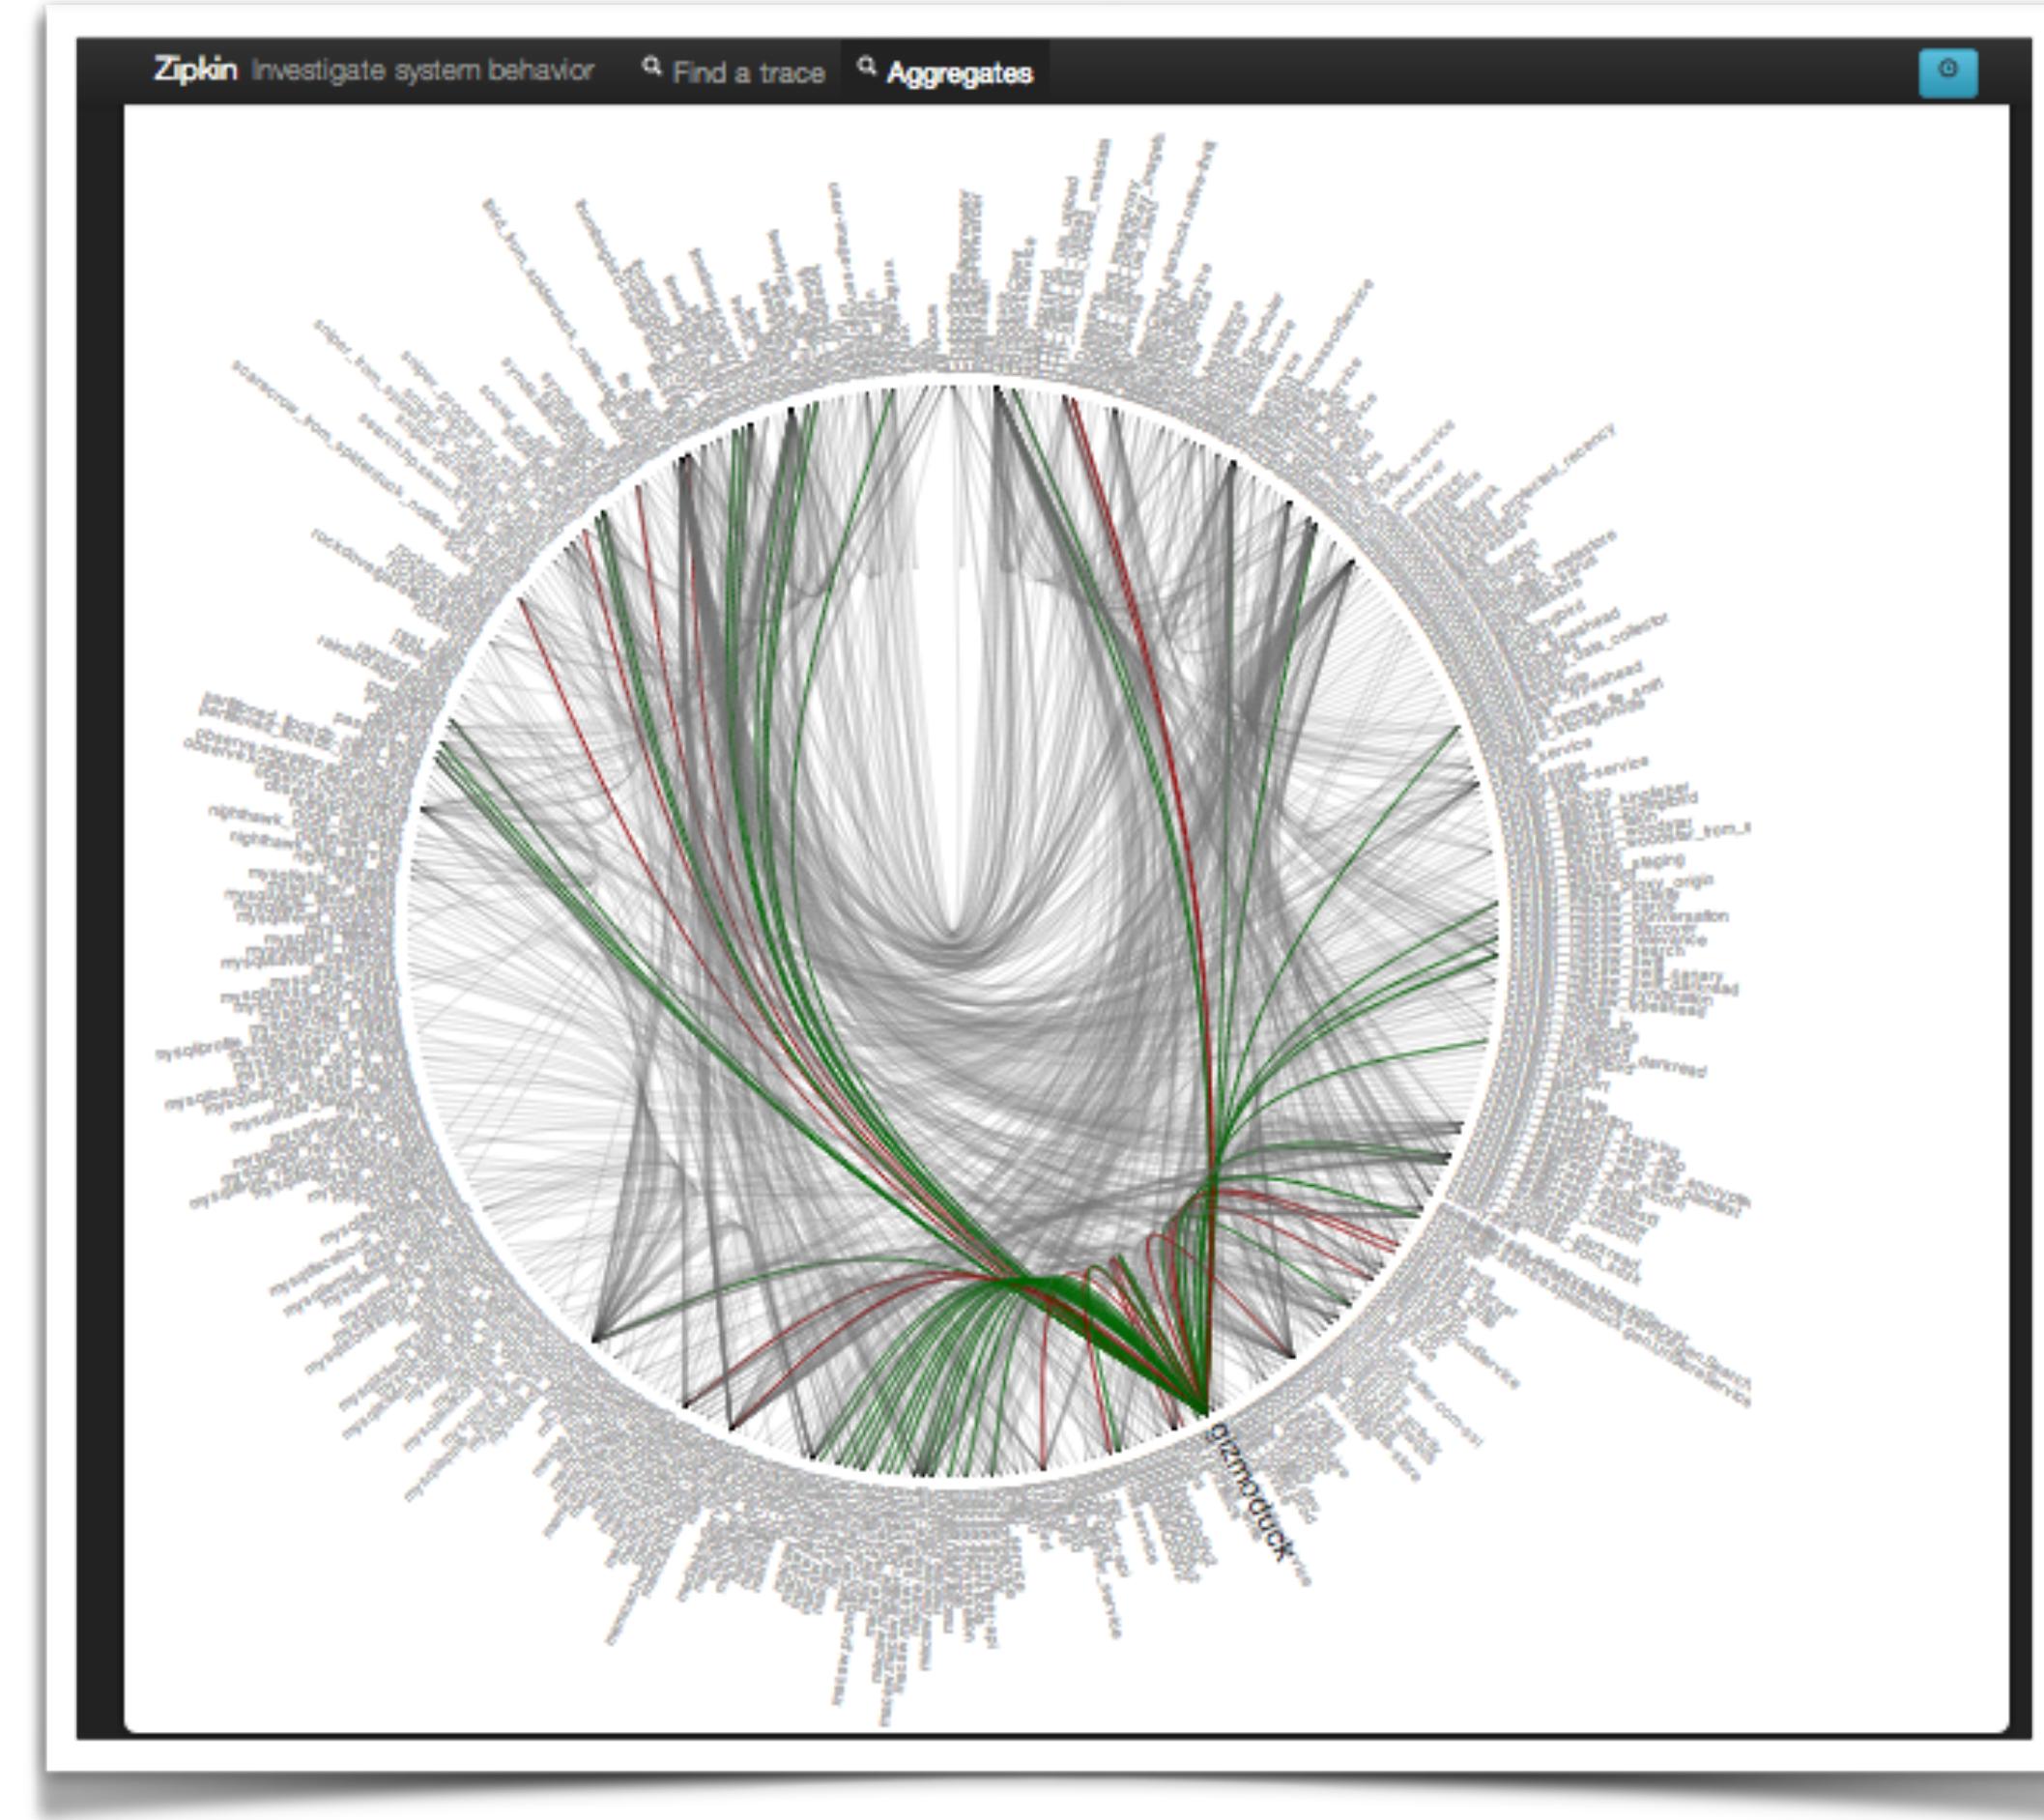


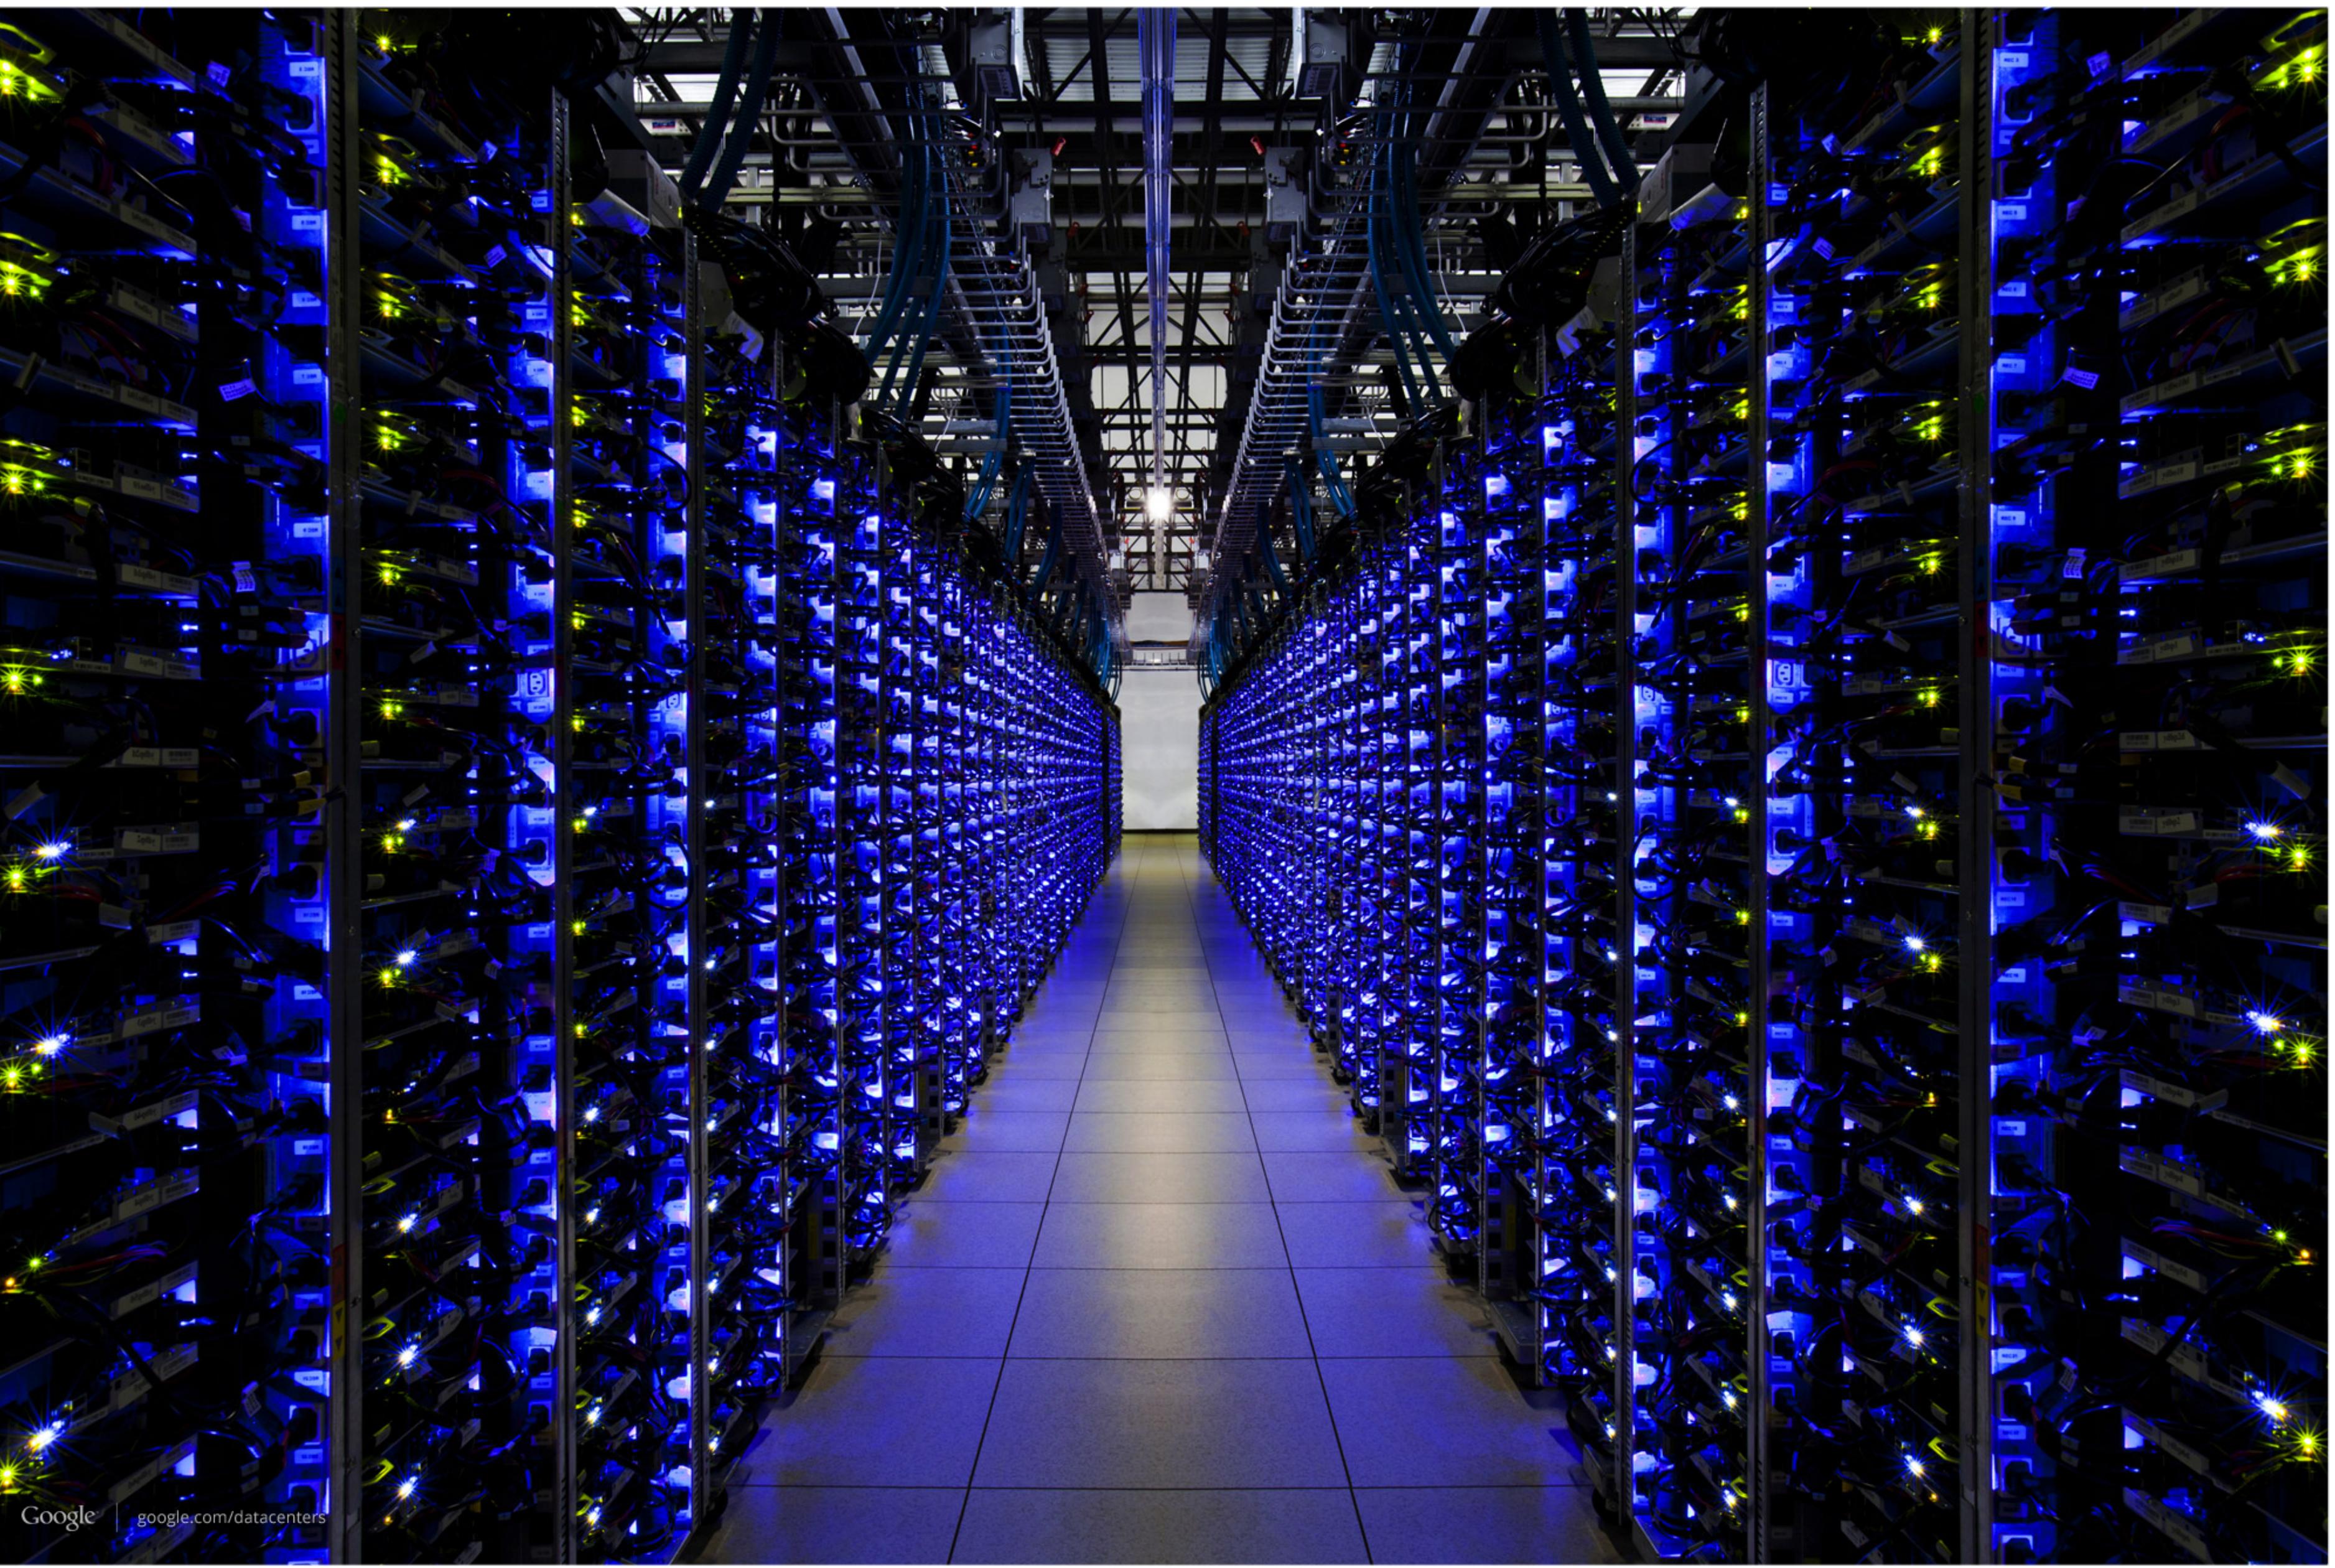
Reality

Thousands of machines with
hundreds of microservices

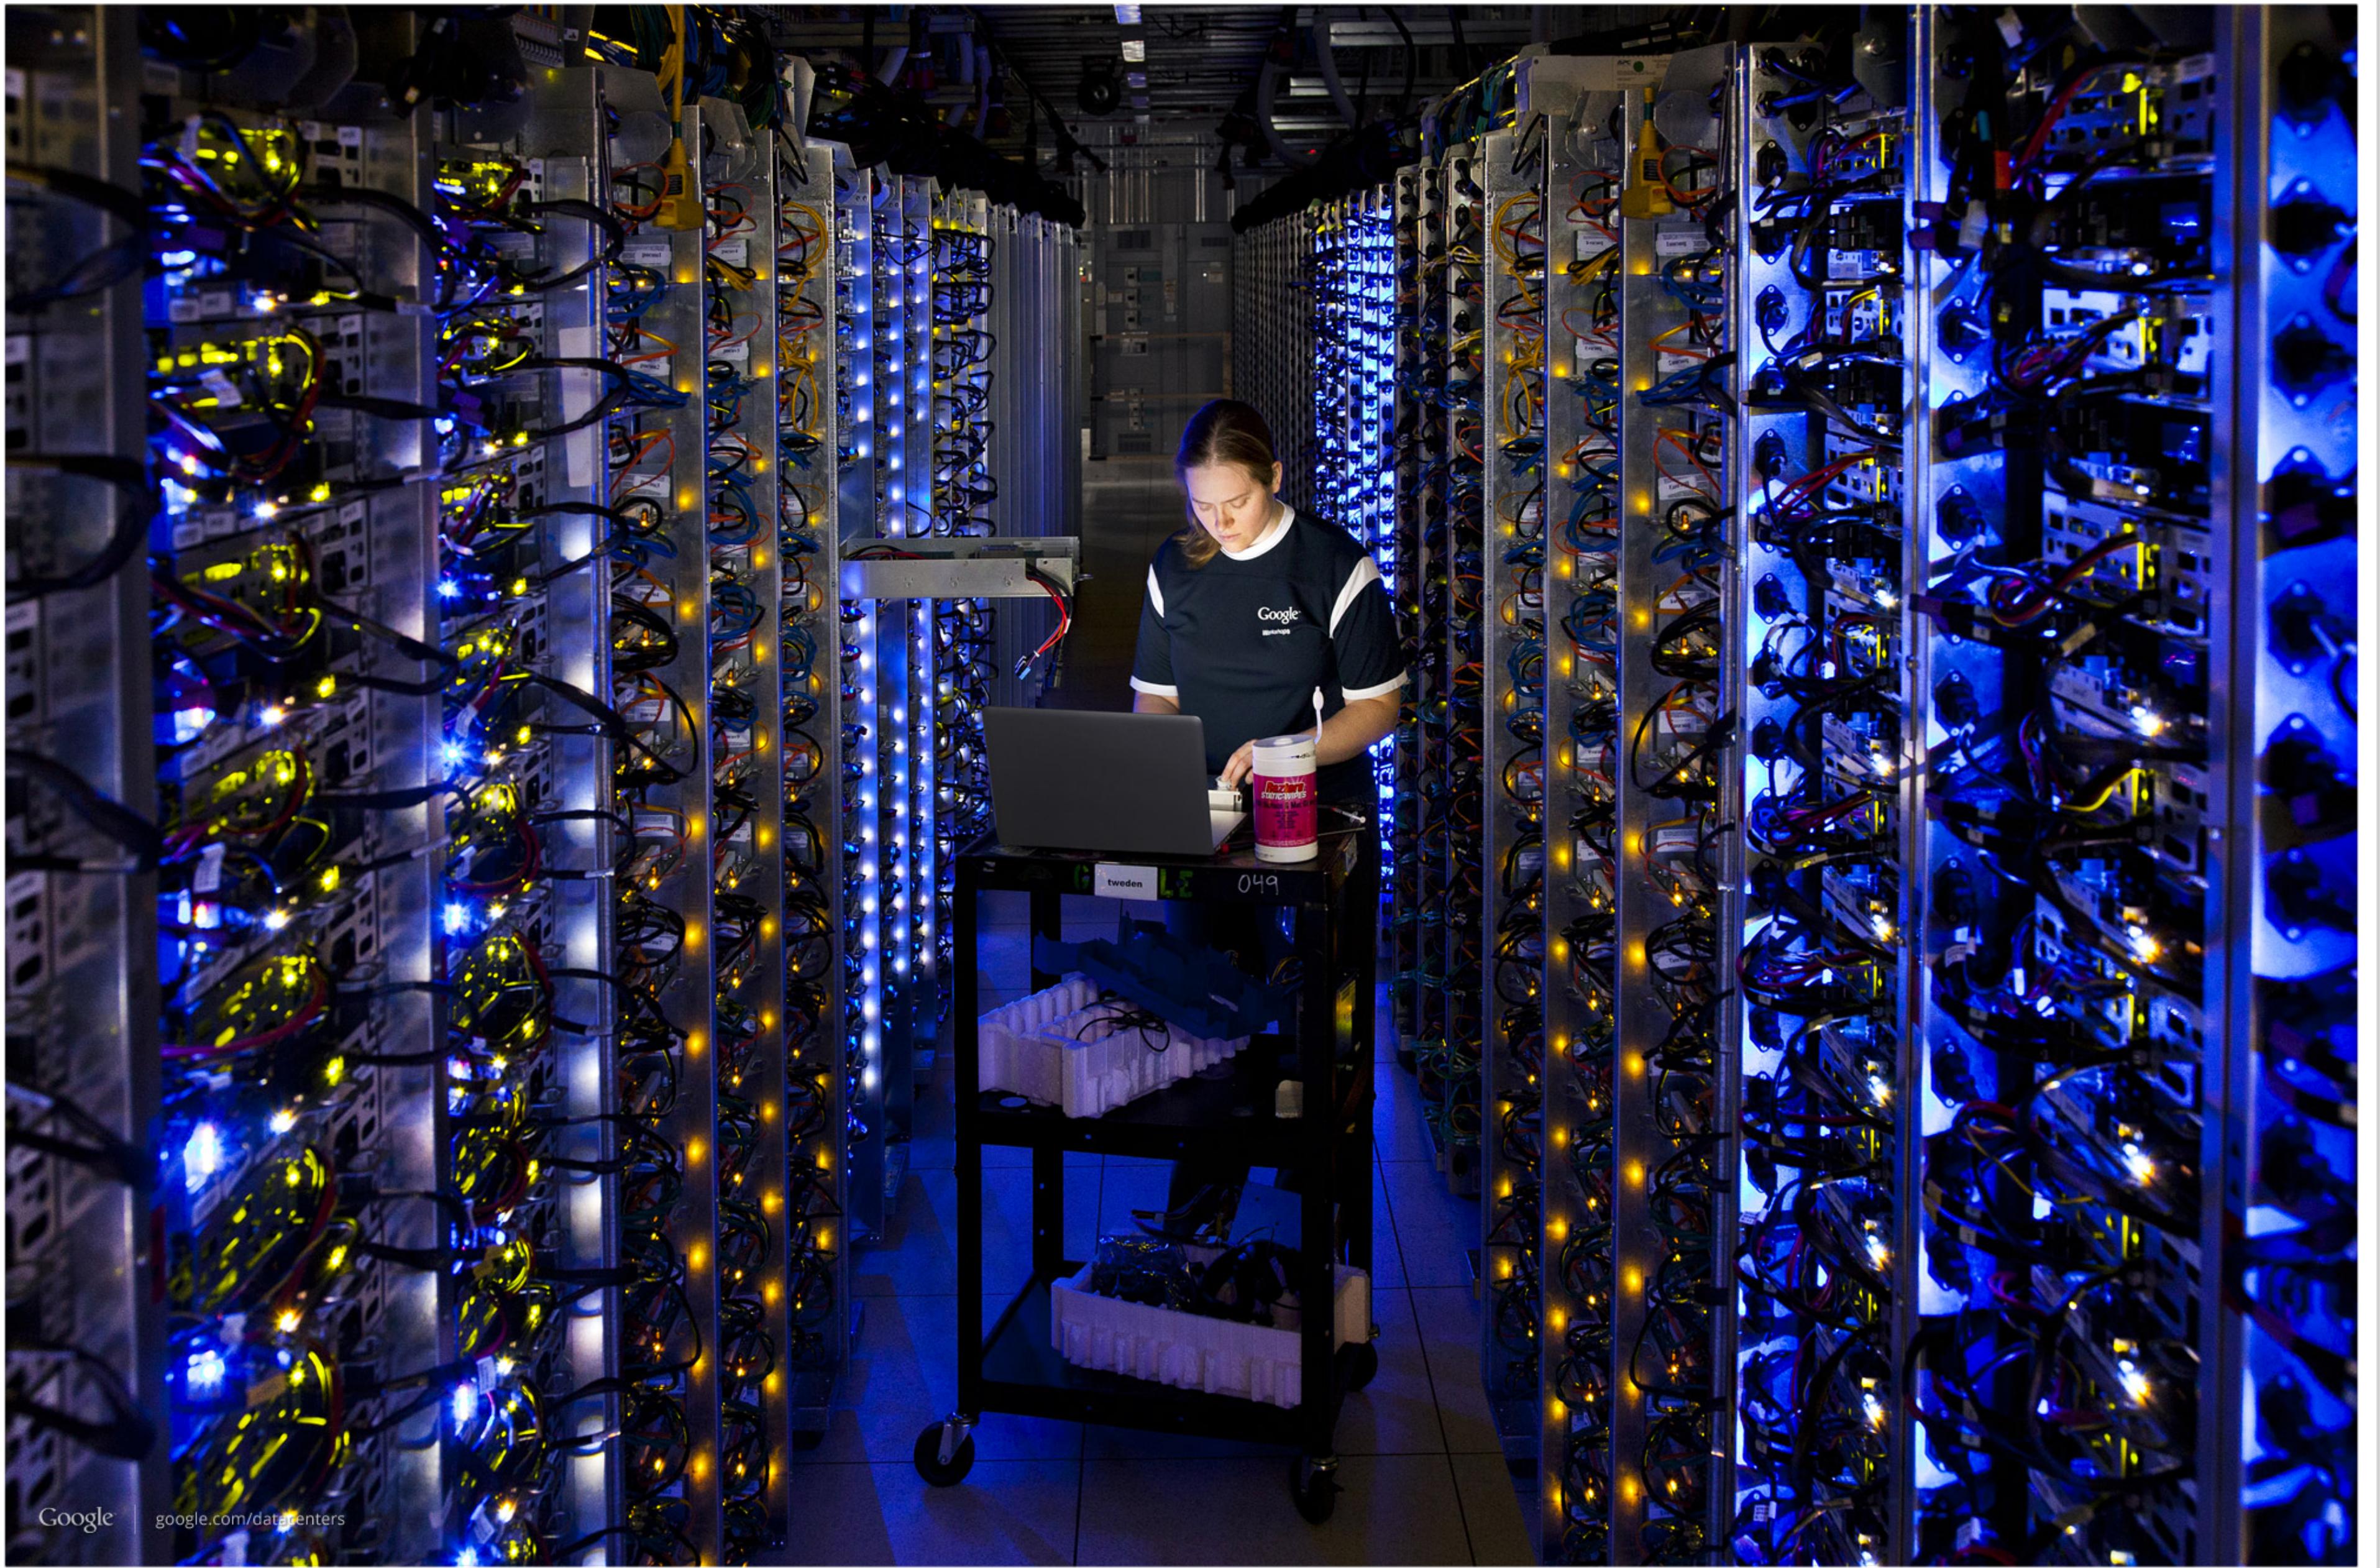
Microservices

A typical relationship between different microservices.

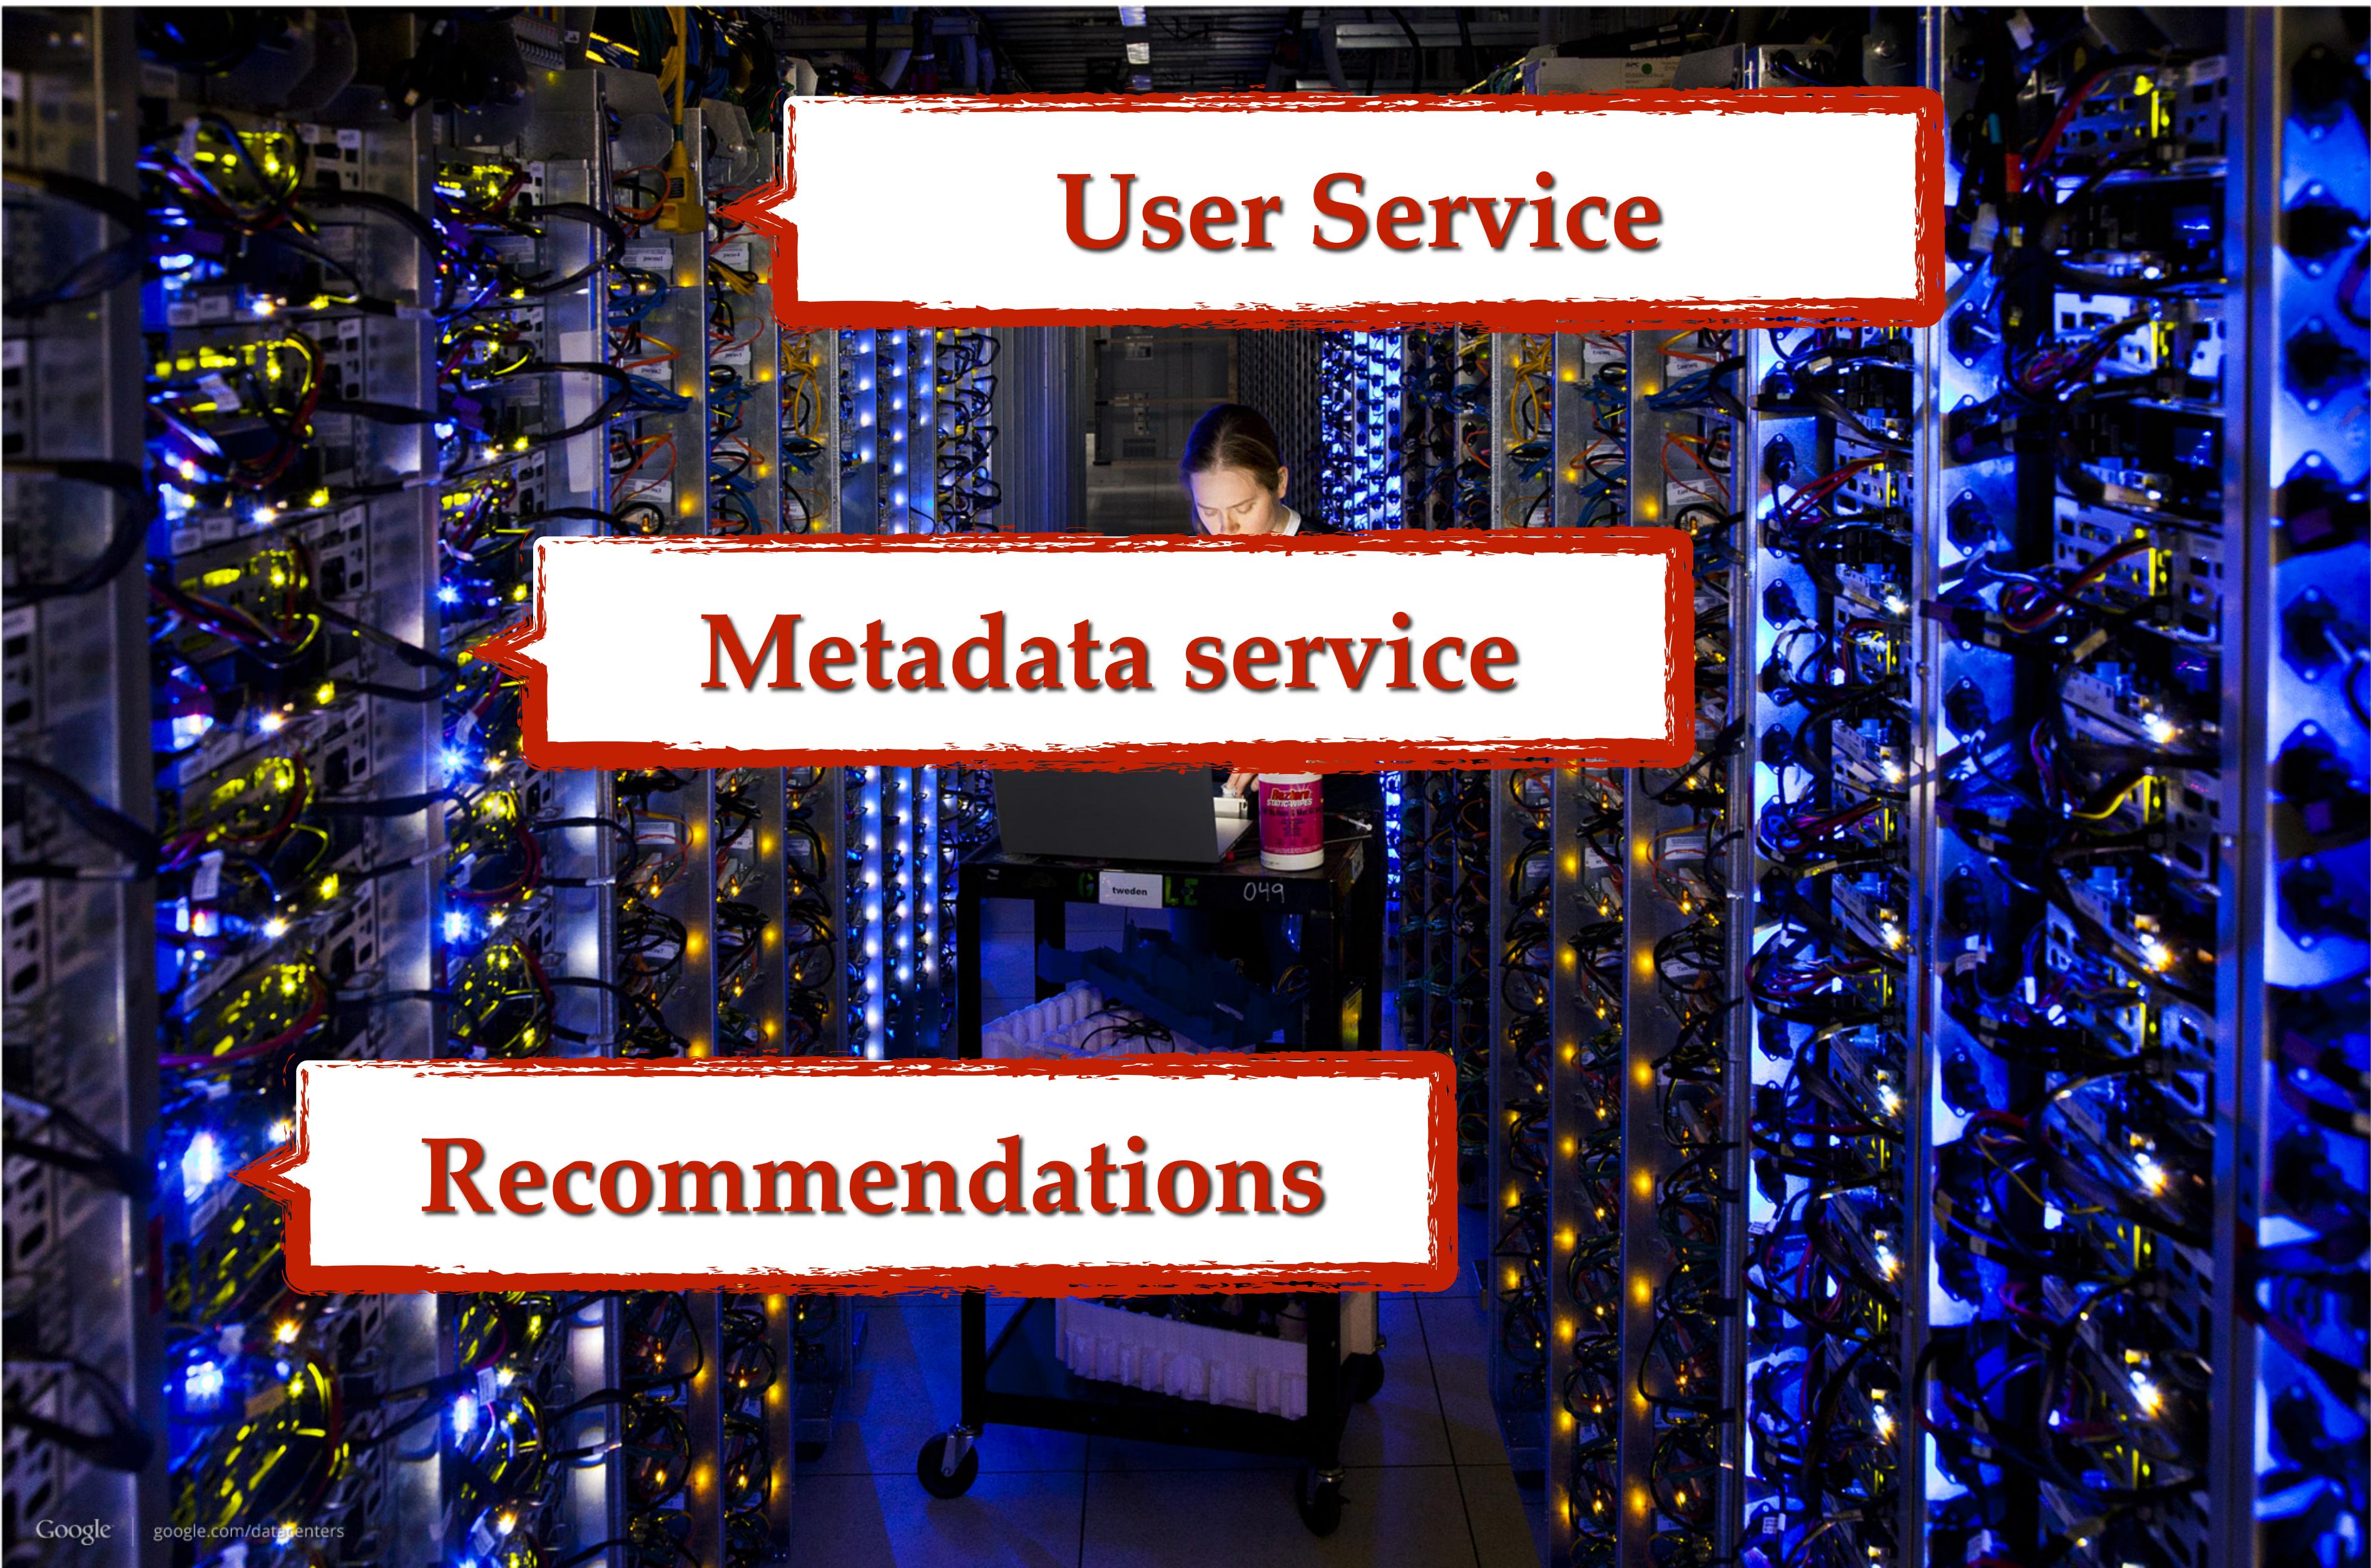




Google | google.com/datacenters



Google | google.com/datacenters



User Service

Metadata service

Recommendations

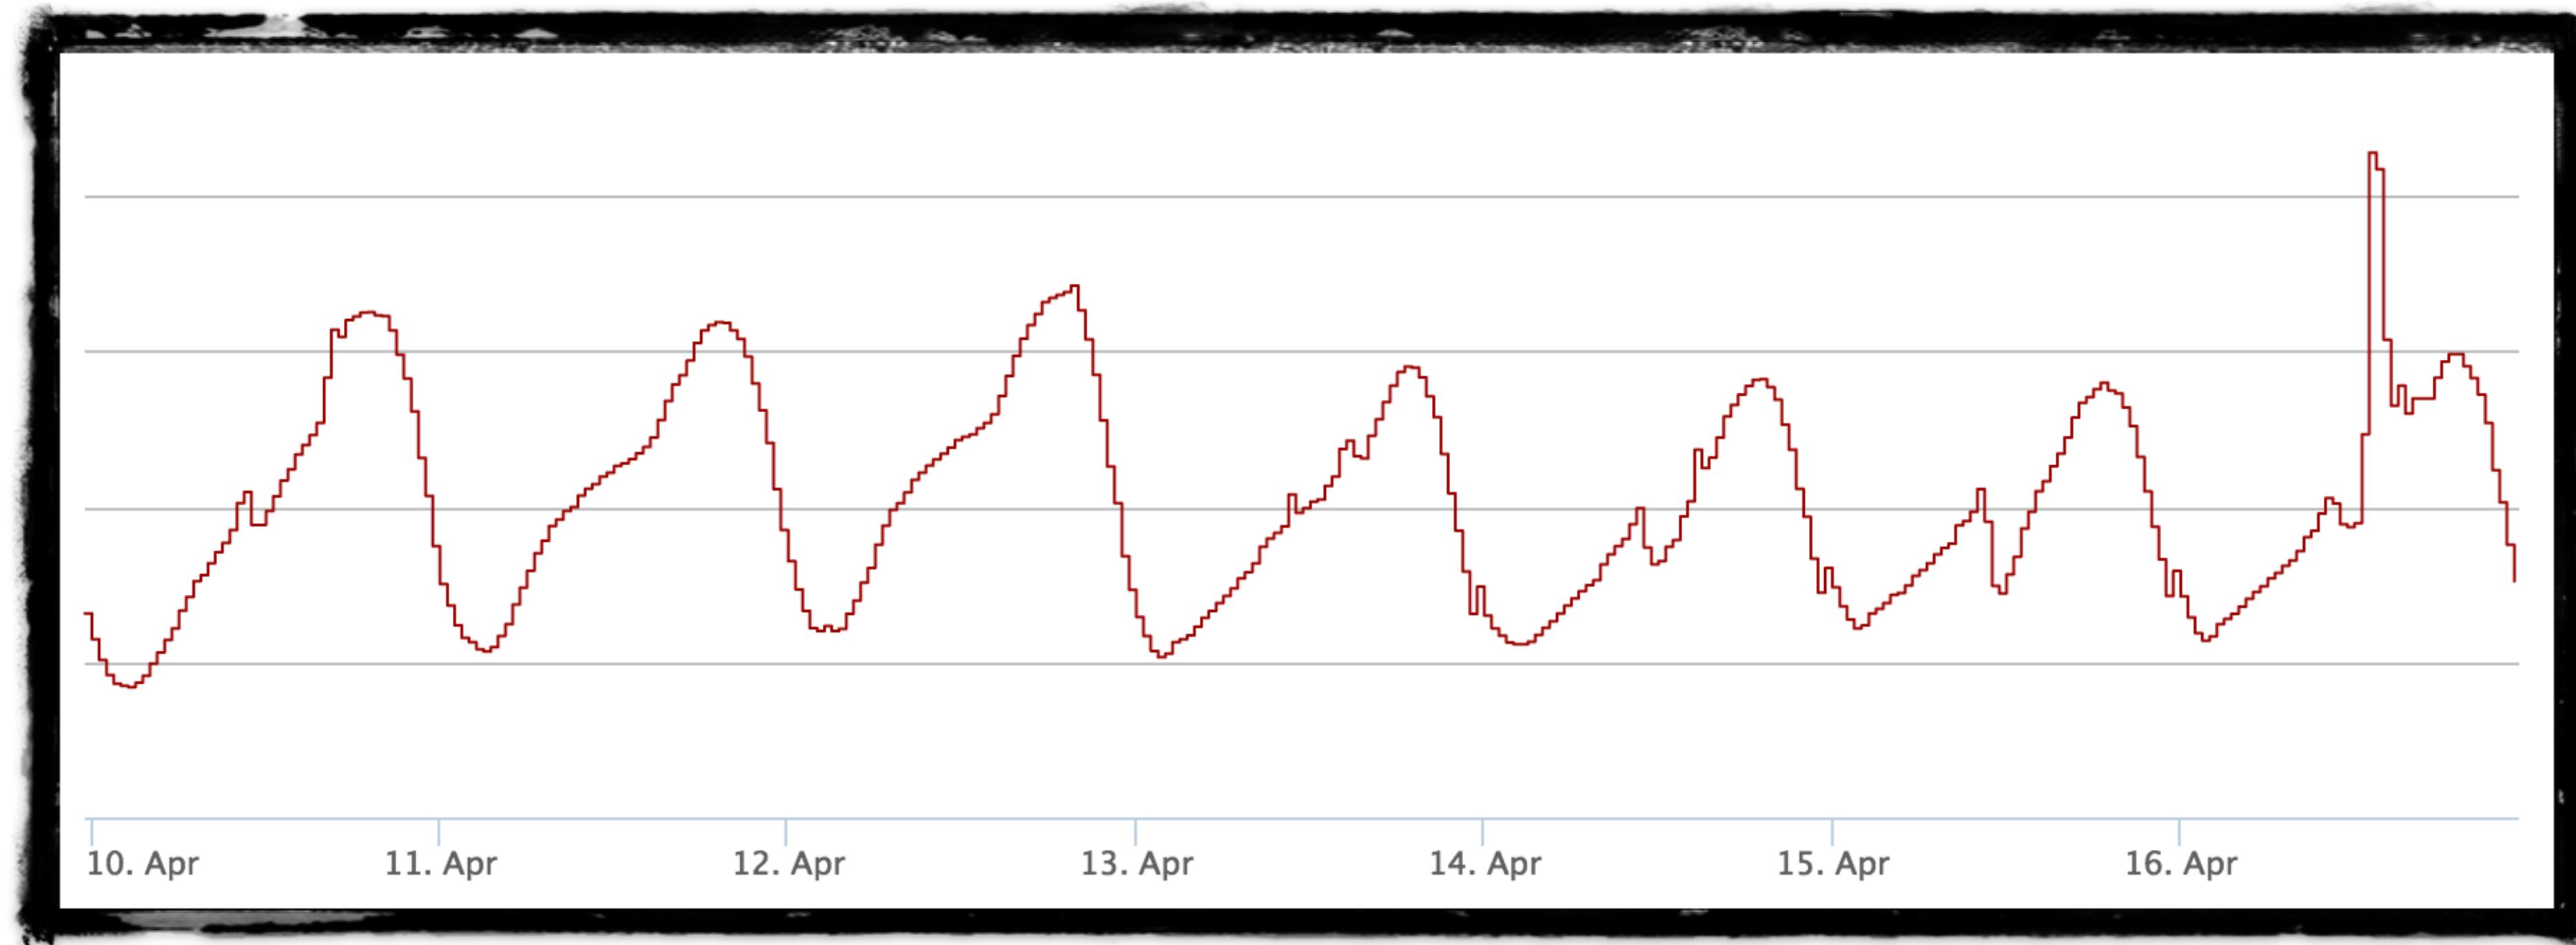
Service Discovery!



HEH

WHAT ABOUT THE CLOUD?

memegenerator.net

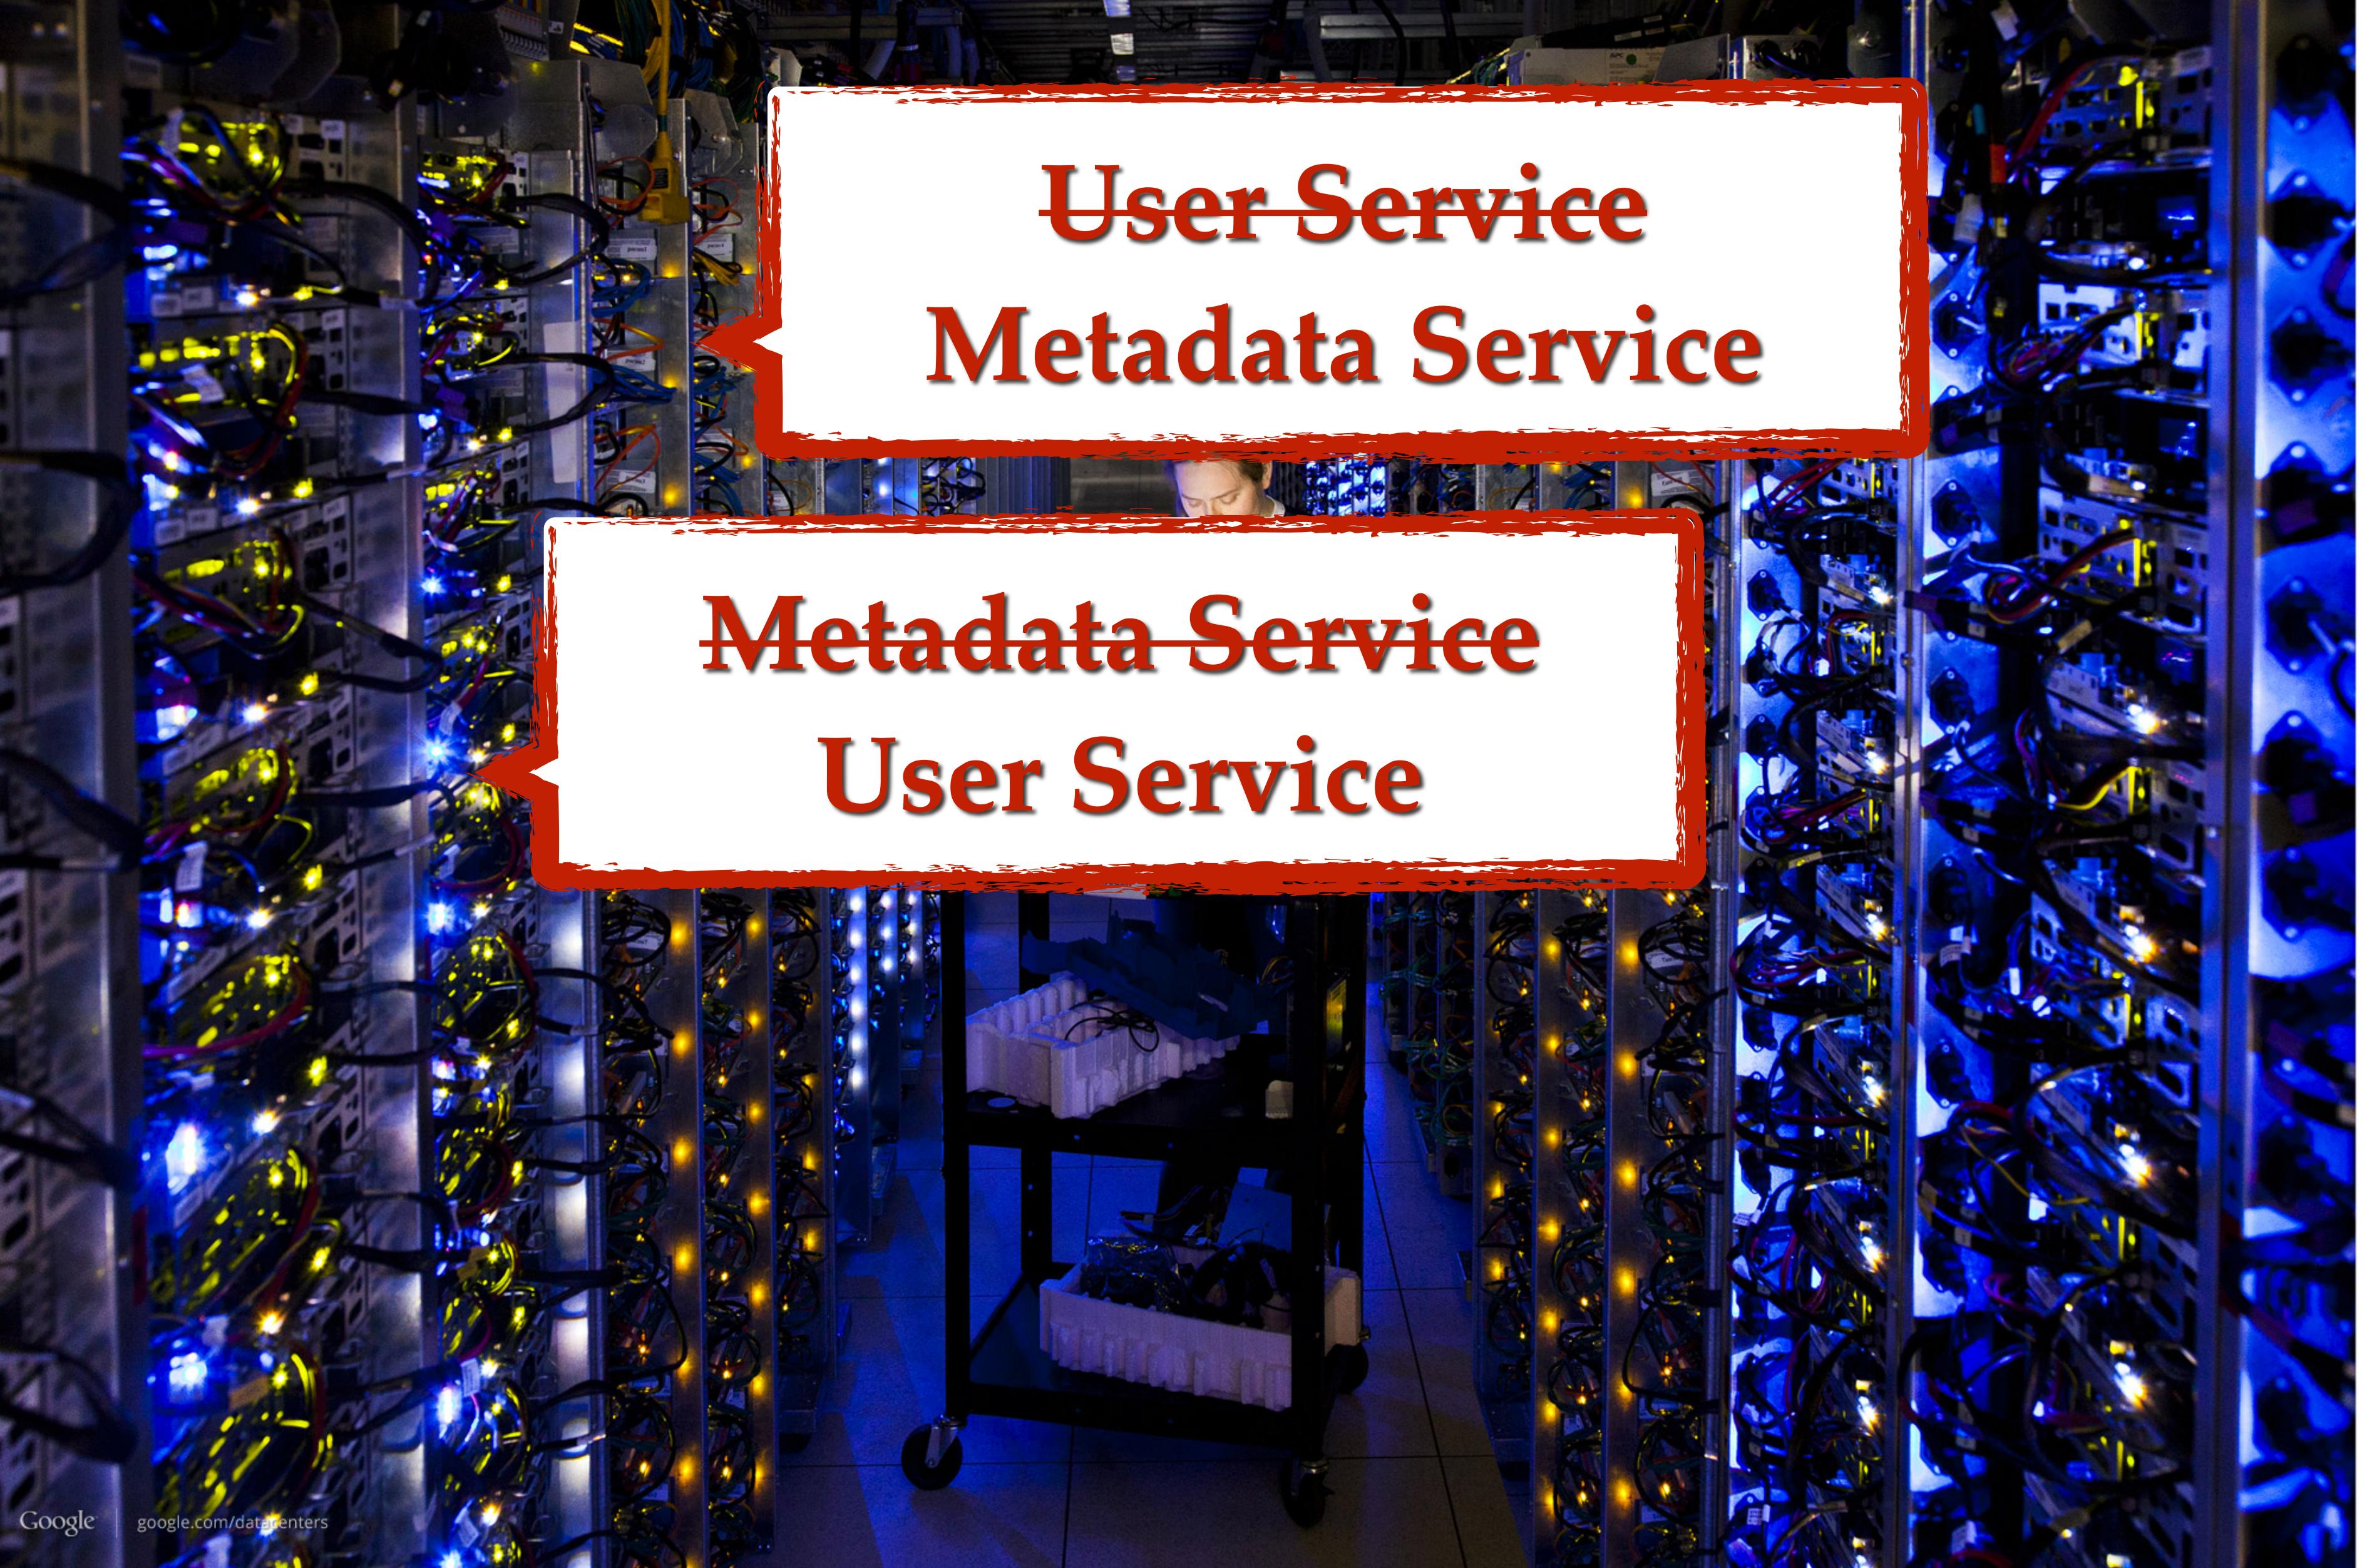


Instance variations

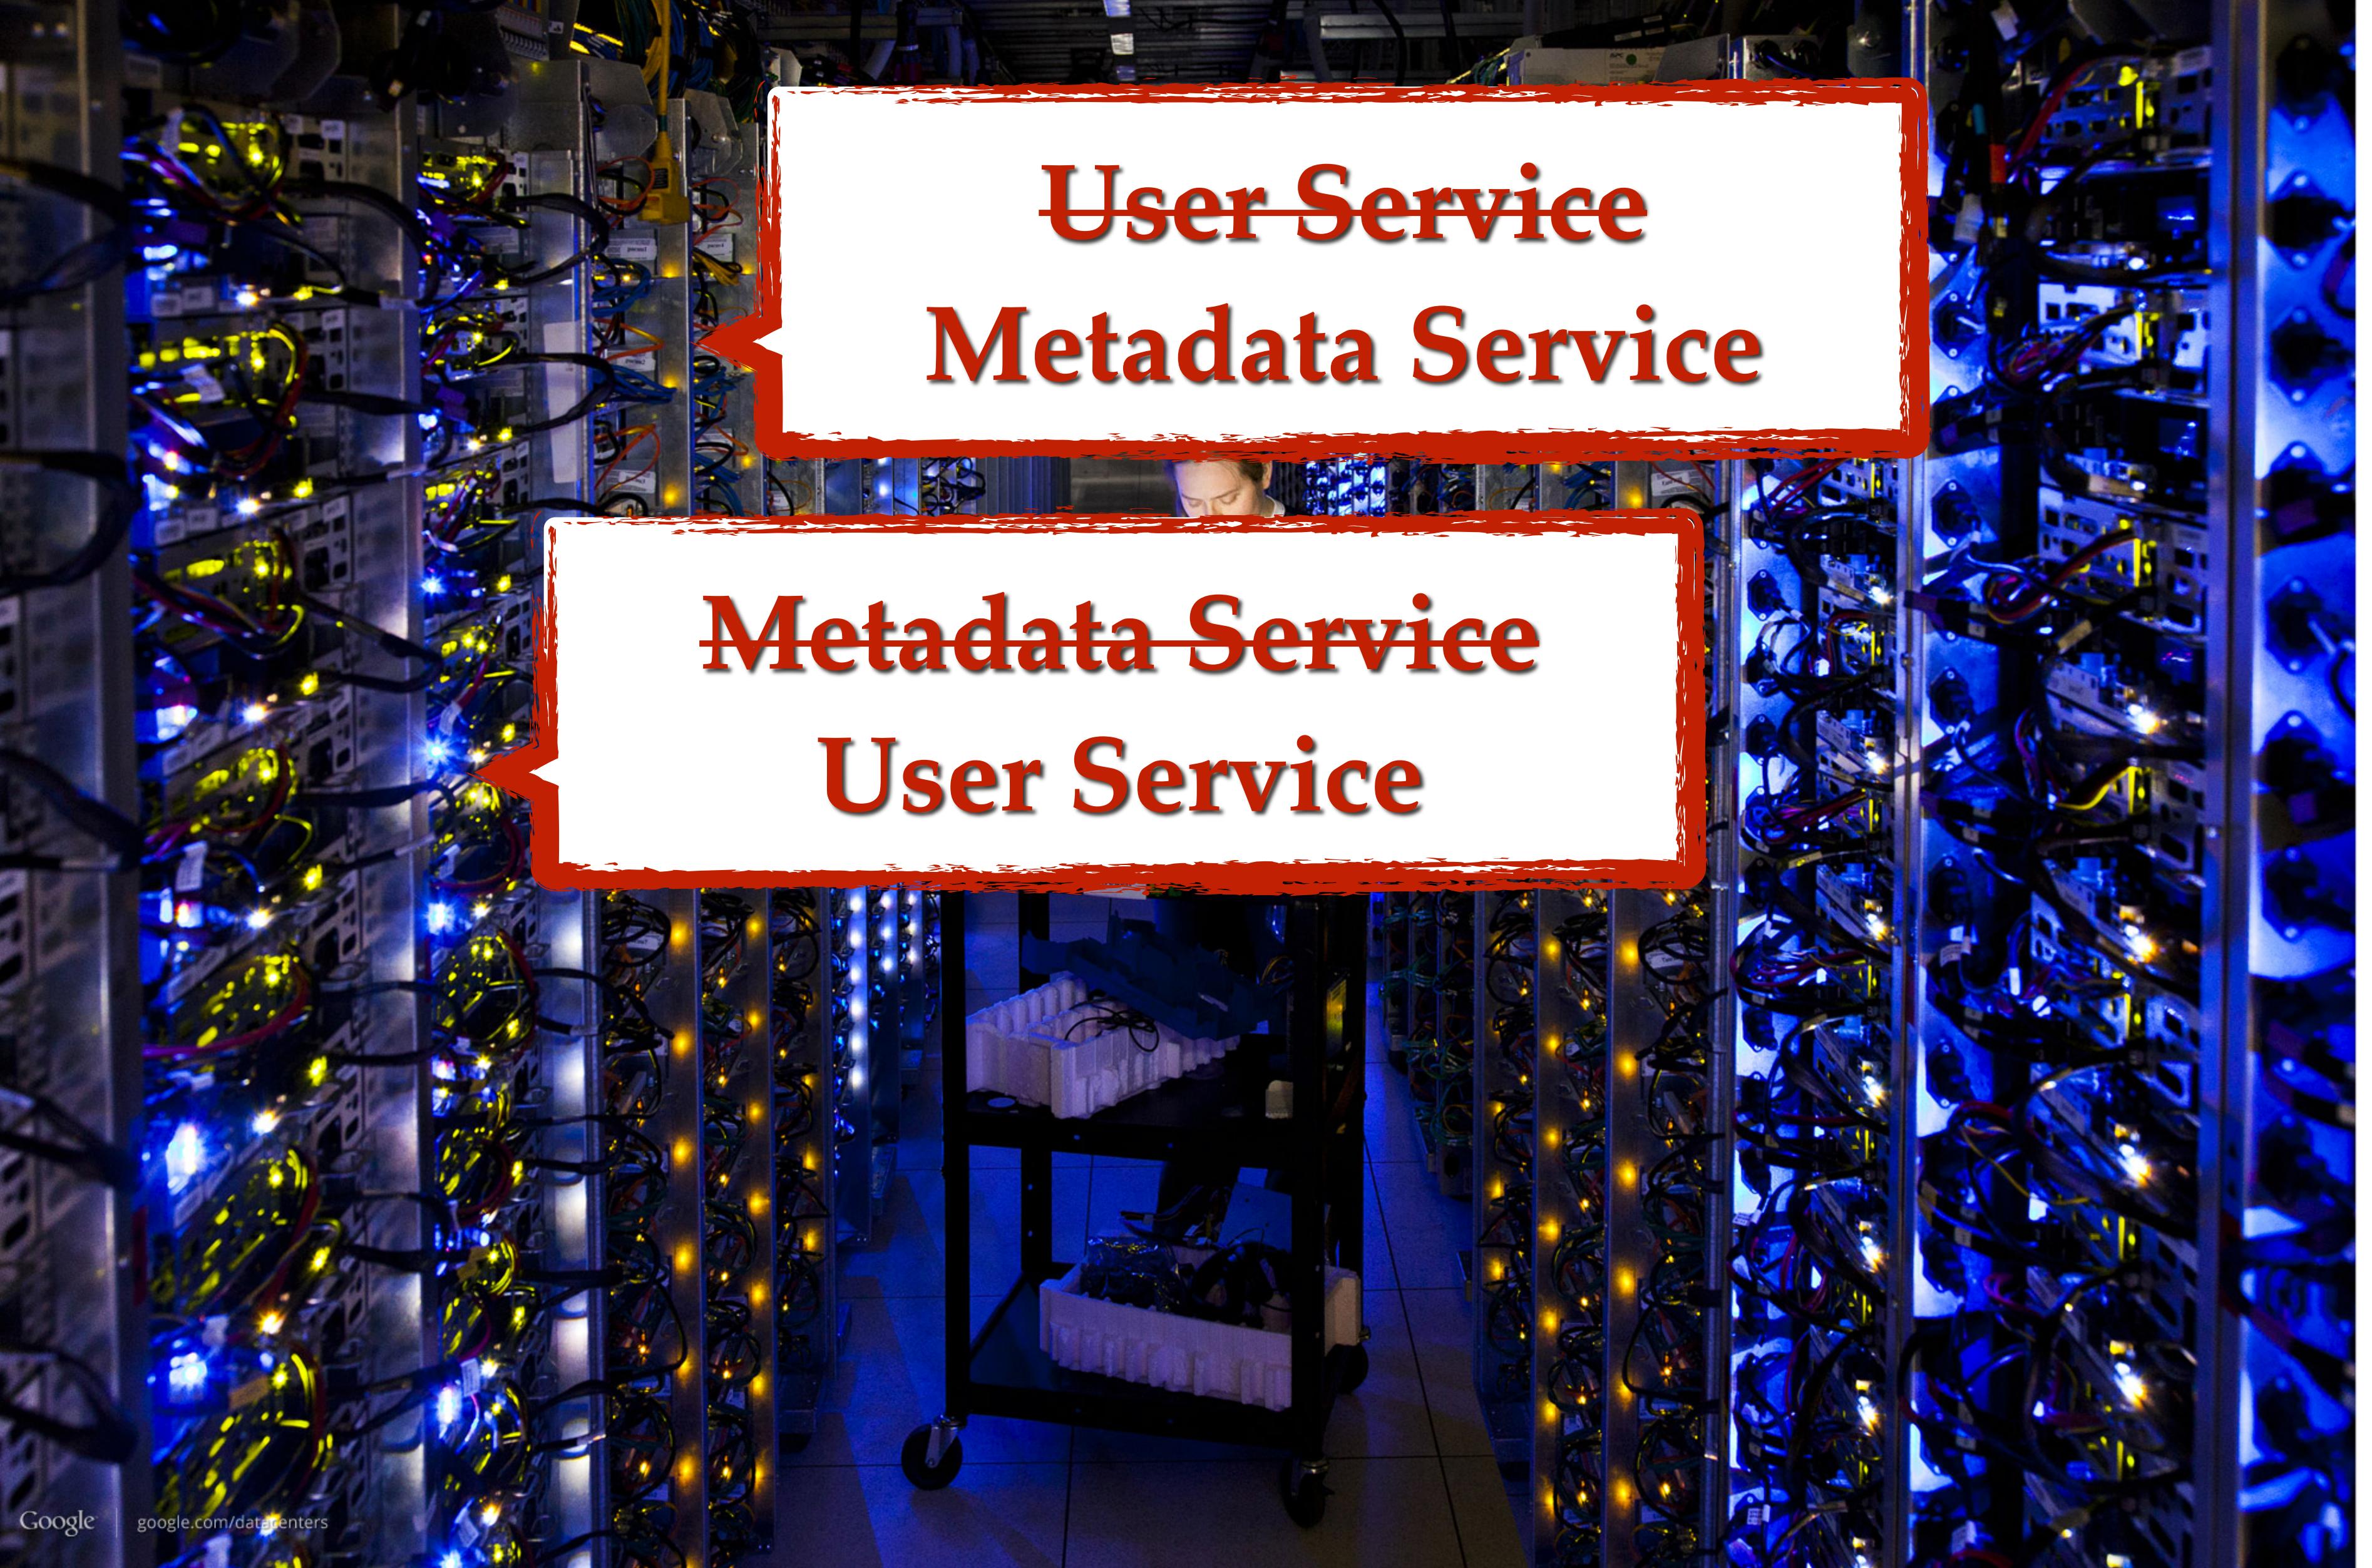
We auto-scale up / down all the time.



VMs crash



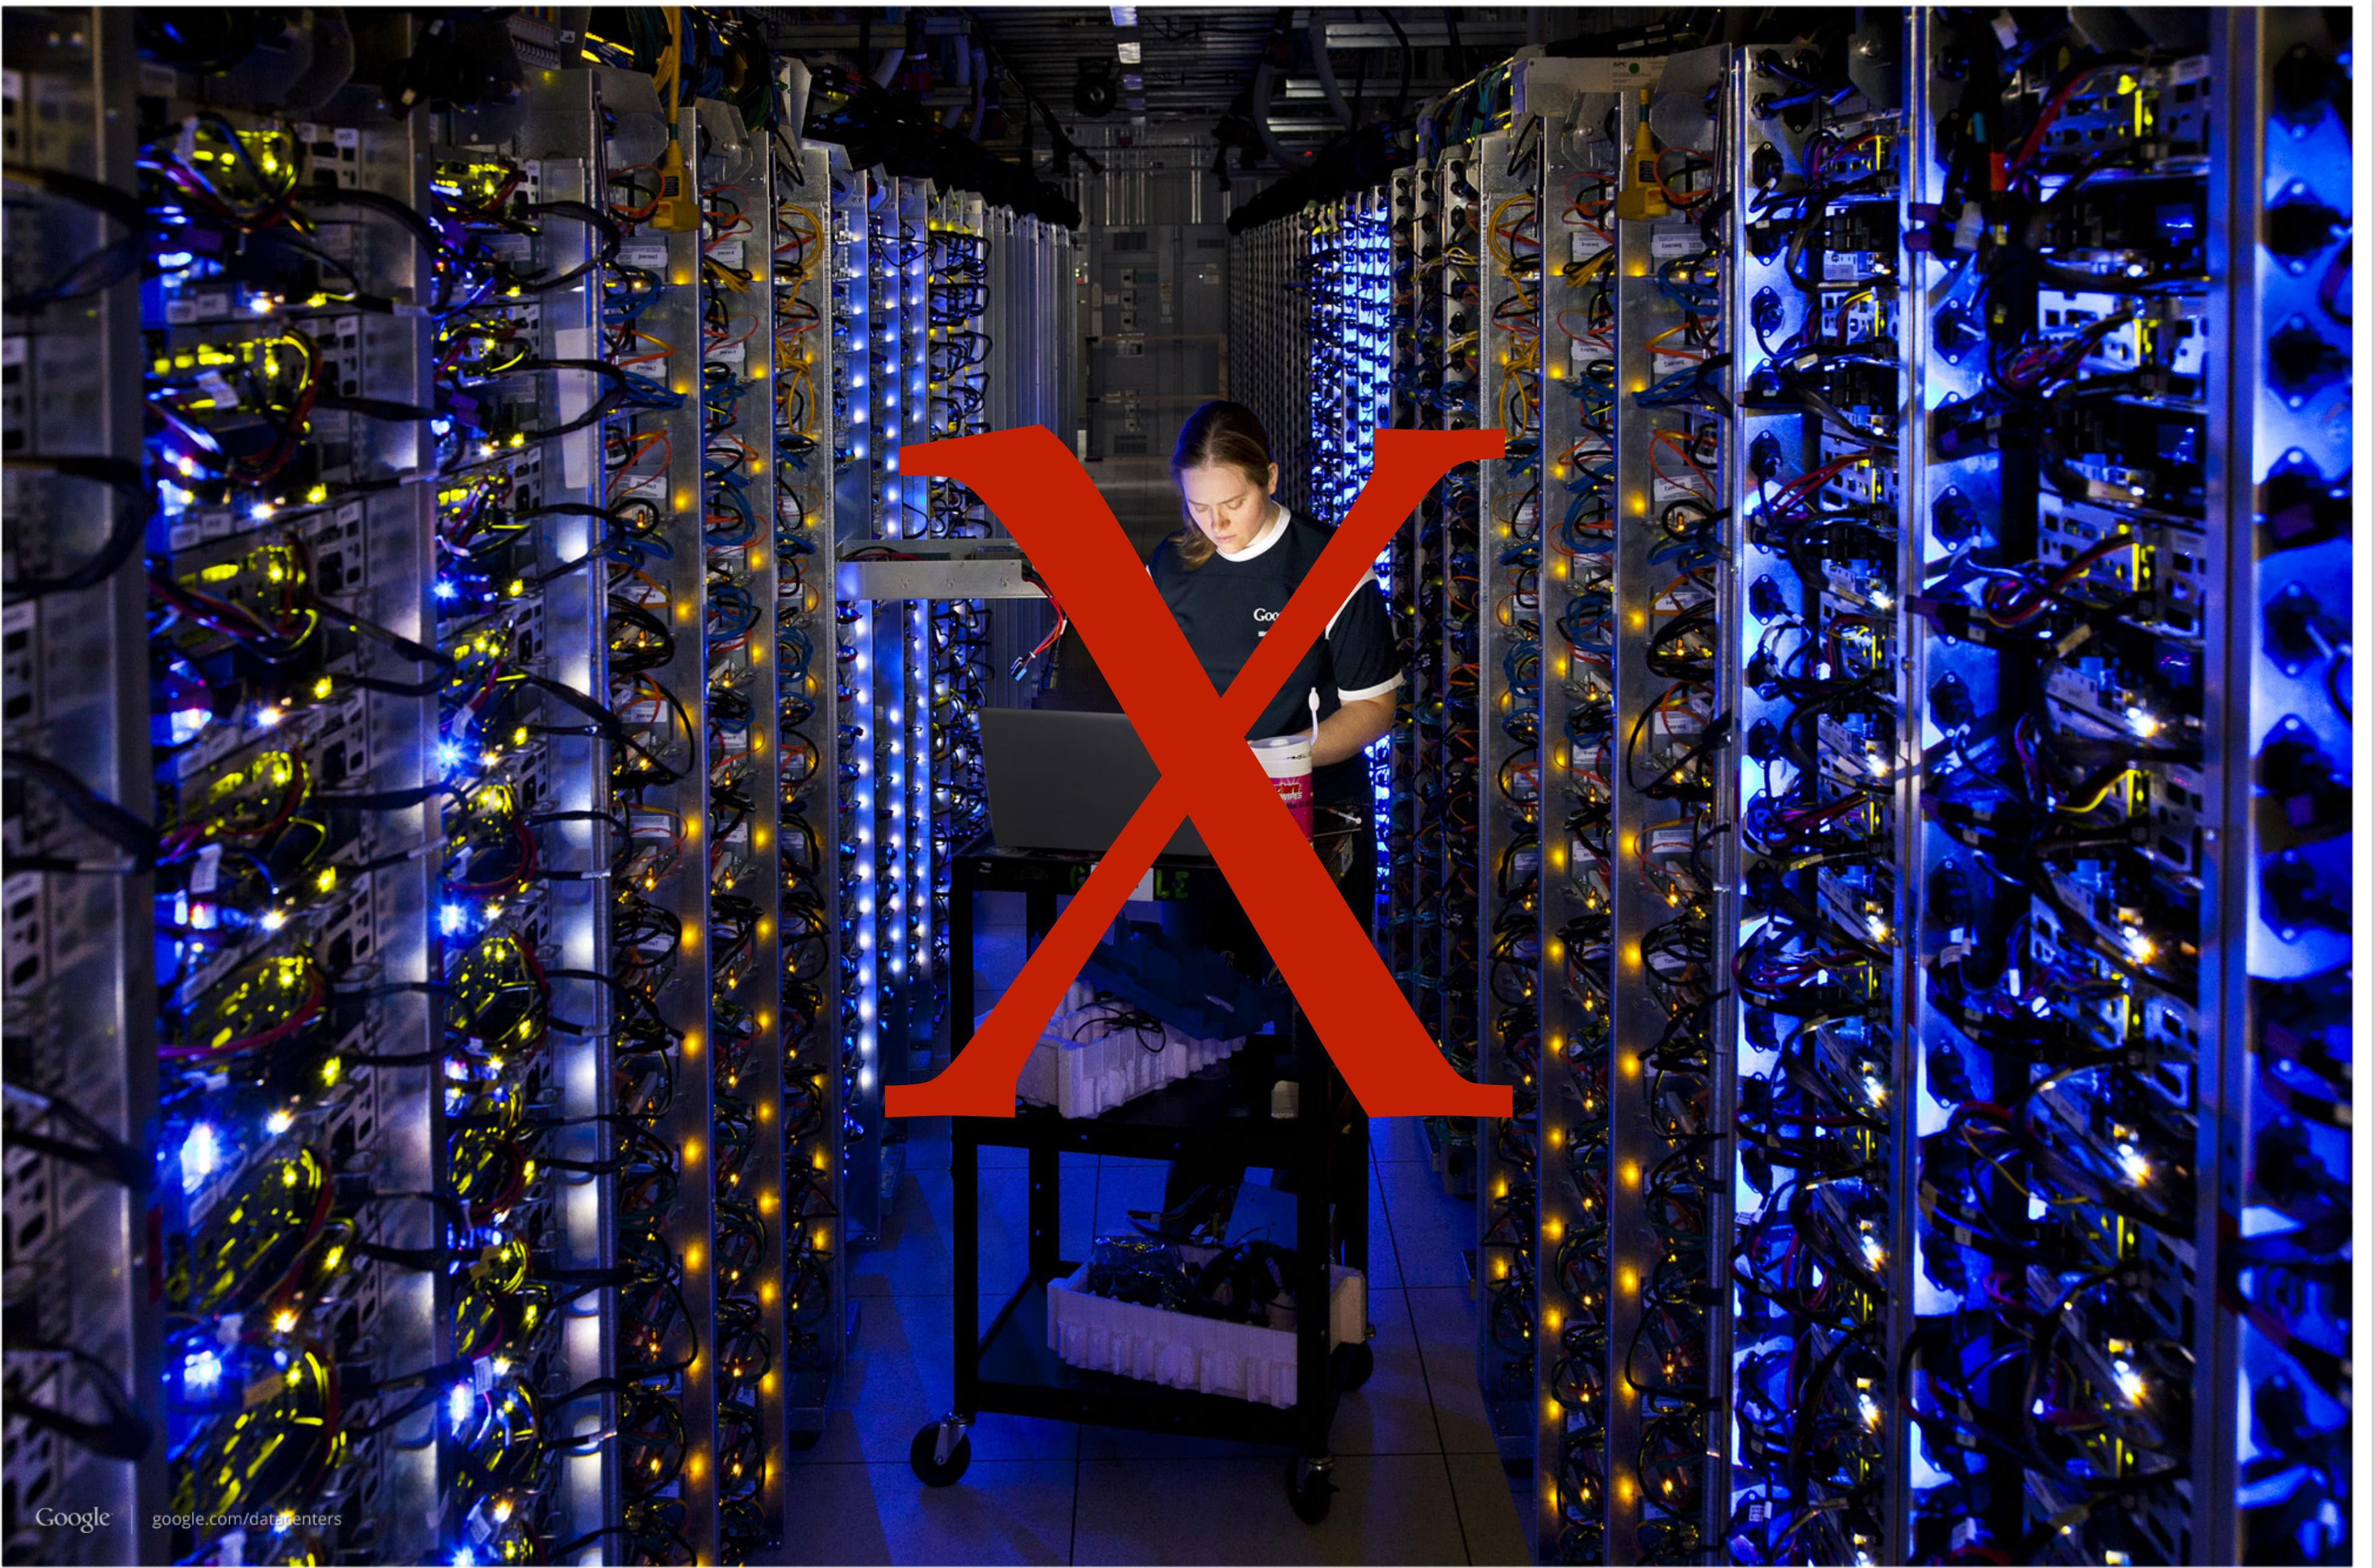
User Service
Metadata Service



Metadata Service
User Service

It's an ever changing eco-system,
static mappings, don't work.

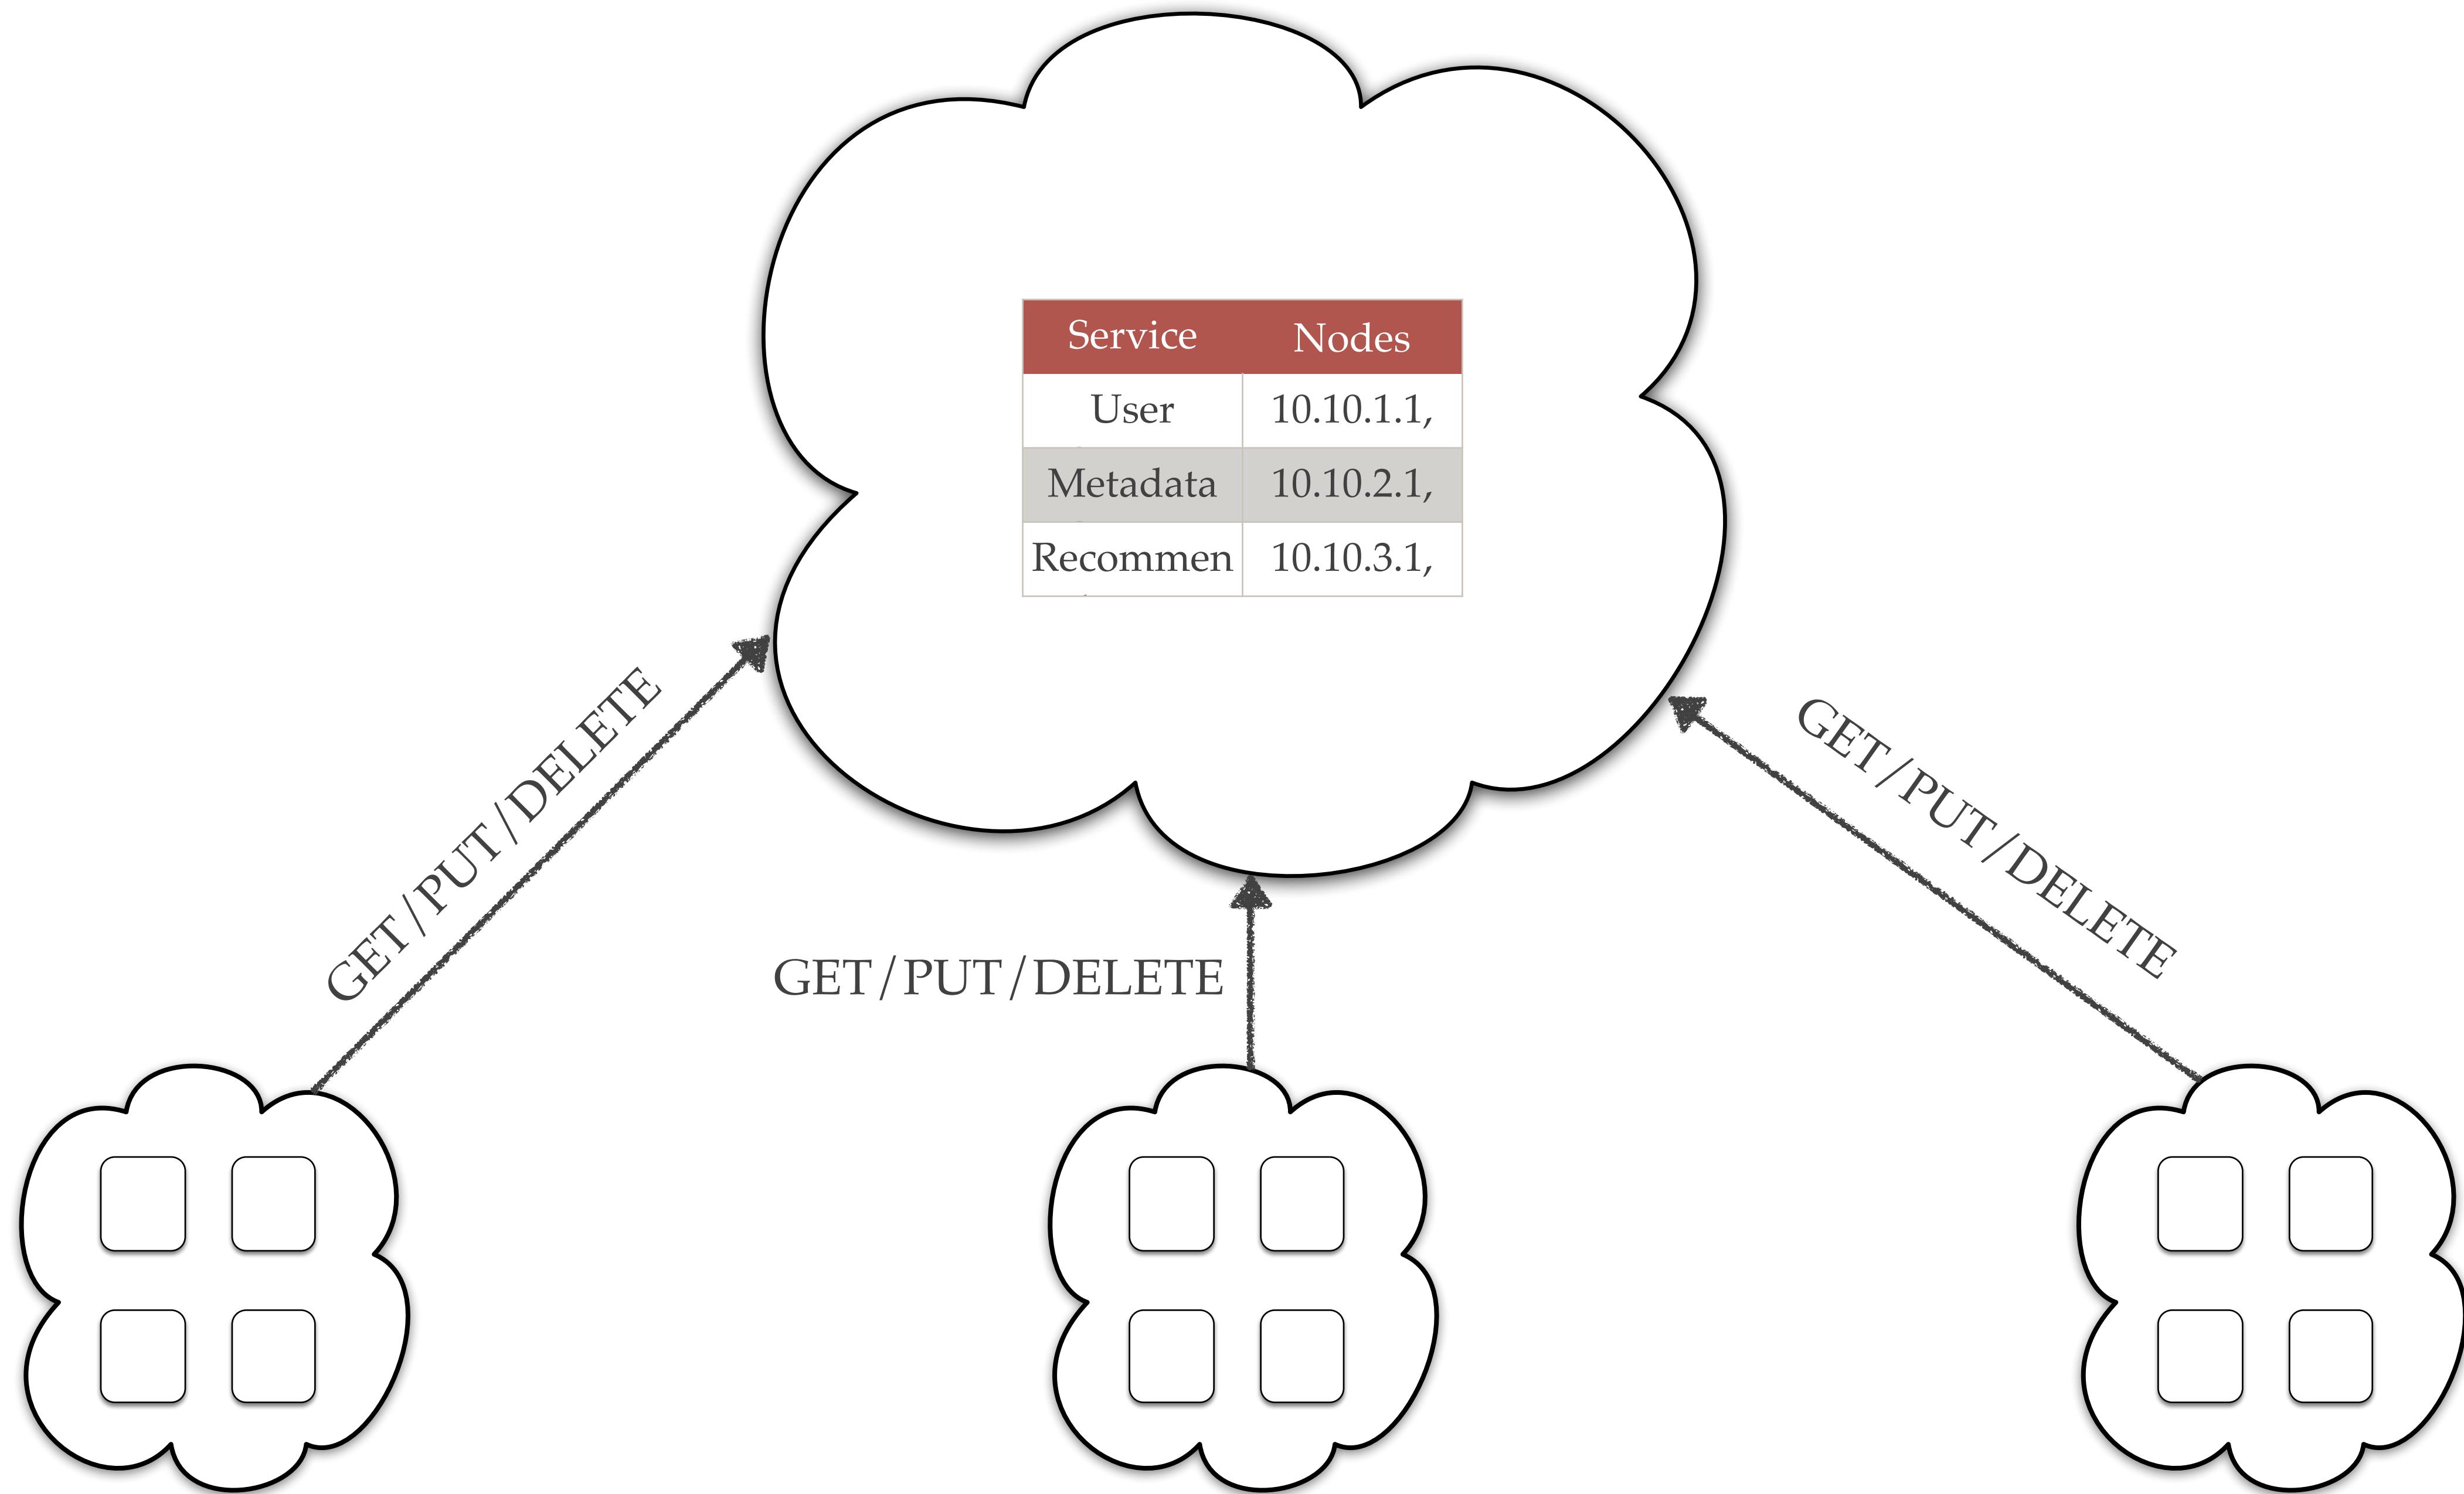
Service discovery is a problem of
cloud environments.



Google | google.com/datacenters

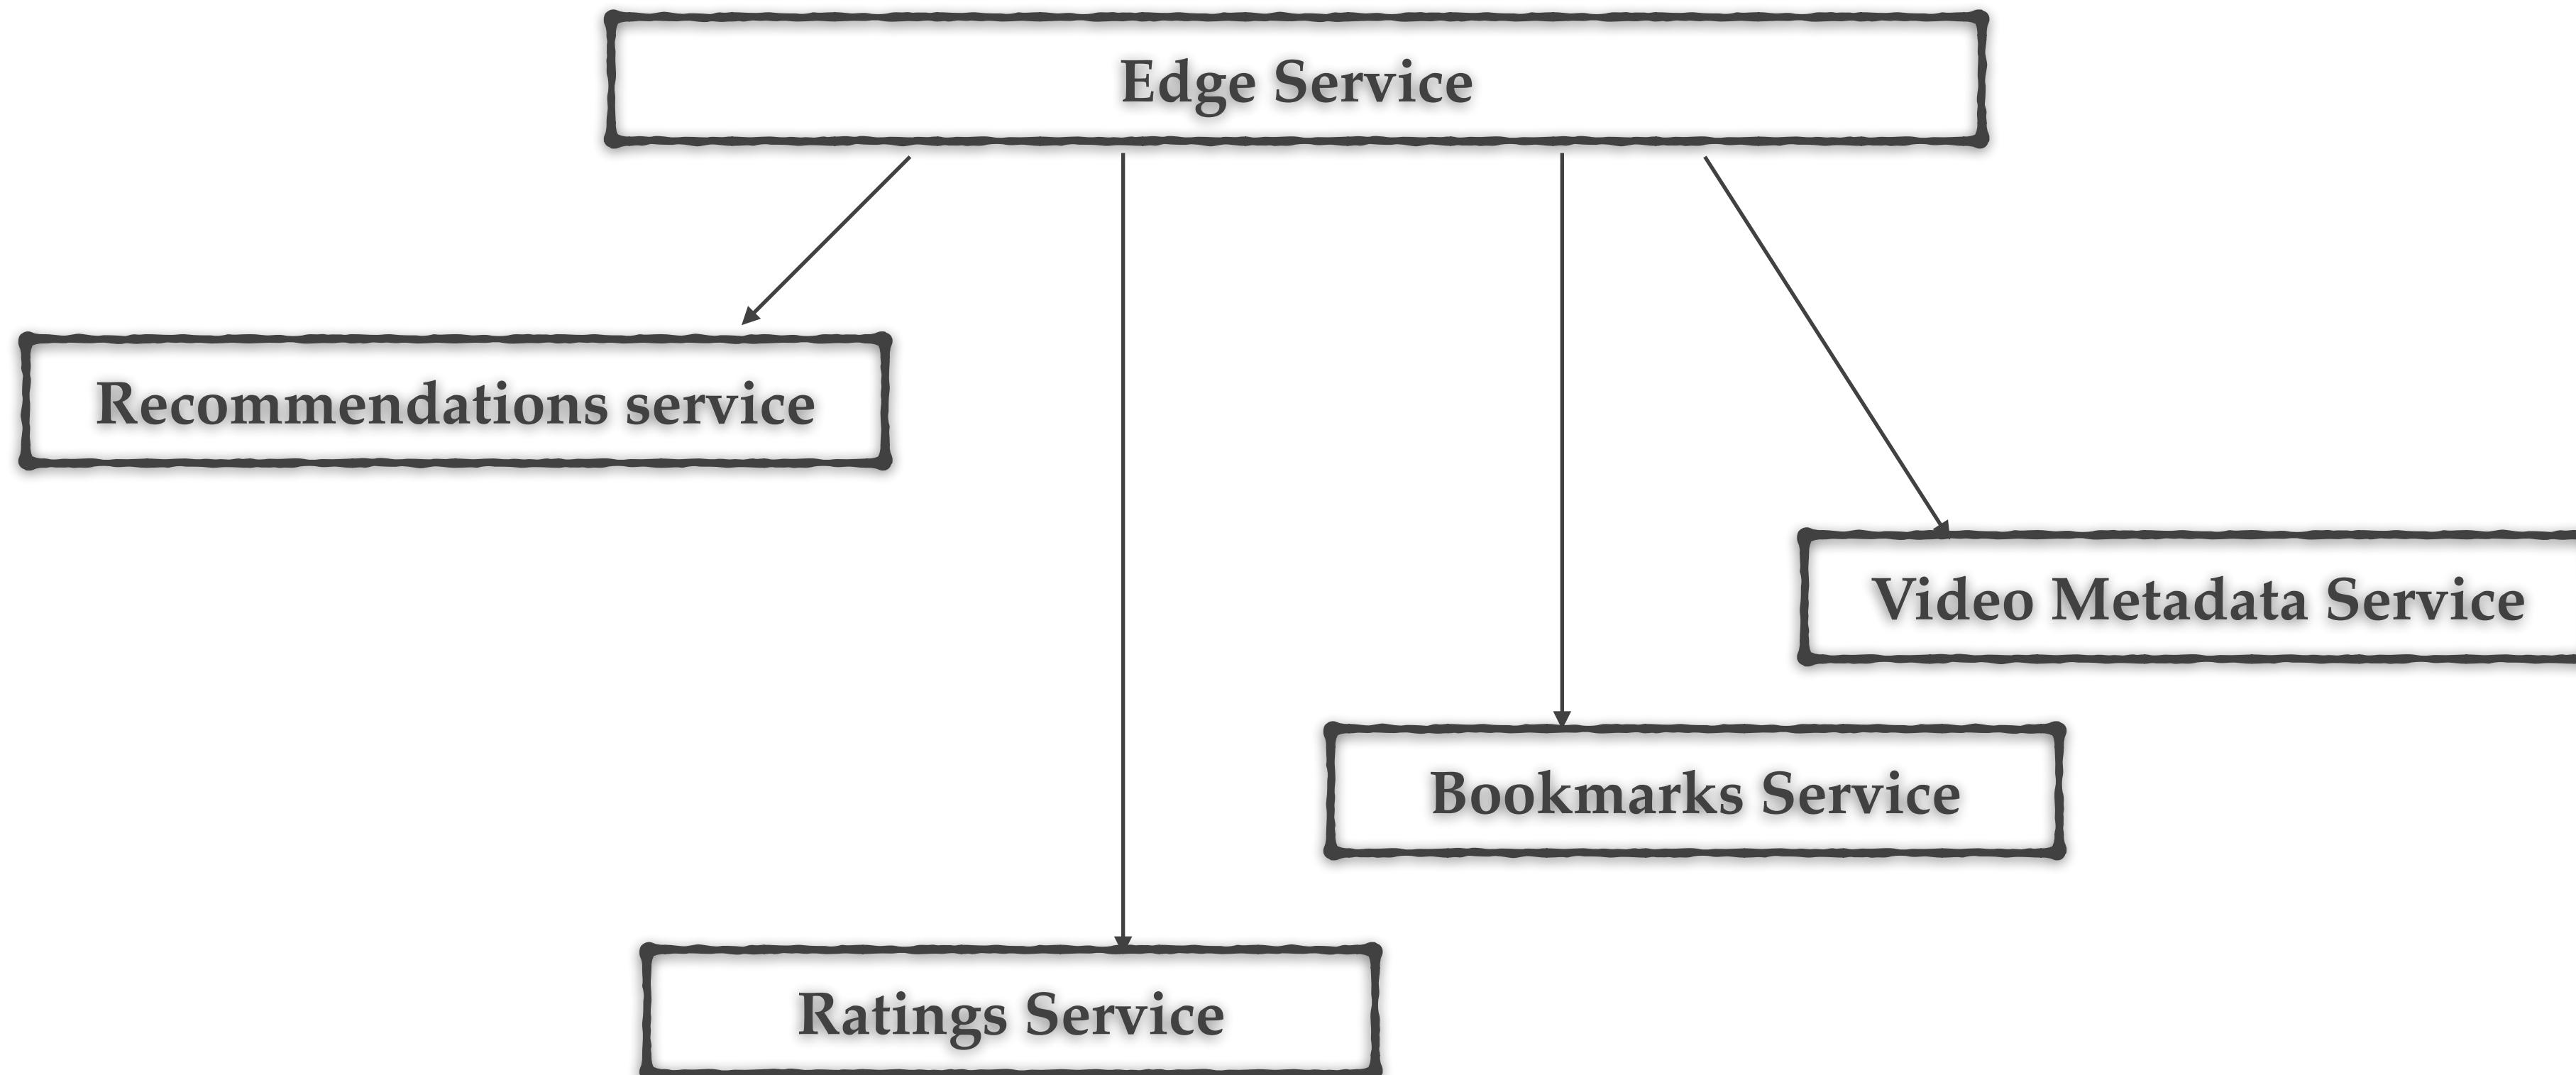
Service Name	Nodes
User Service	10.10.1.1, 10.10.1.2, 10.10.1.3, 10.10.1.4
Metadata Service	10.10.2.1
Recommendations Service	10.10.3.1, 10.10.3.2, 10.10.3.3

Service Name	Nodes
User Service	10.10.1.1, 10.10.1.2
Metadata Service	10.10.2.1, 10.10.2.2
Recommendations Service	10.10.3.1, 10.10.3.2



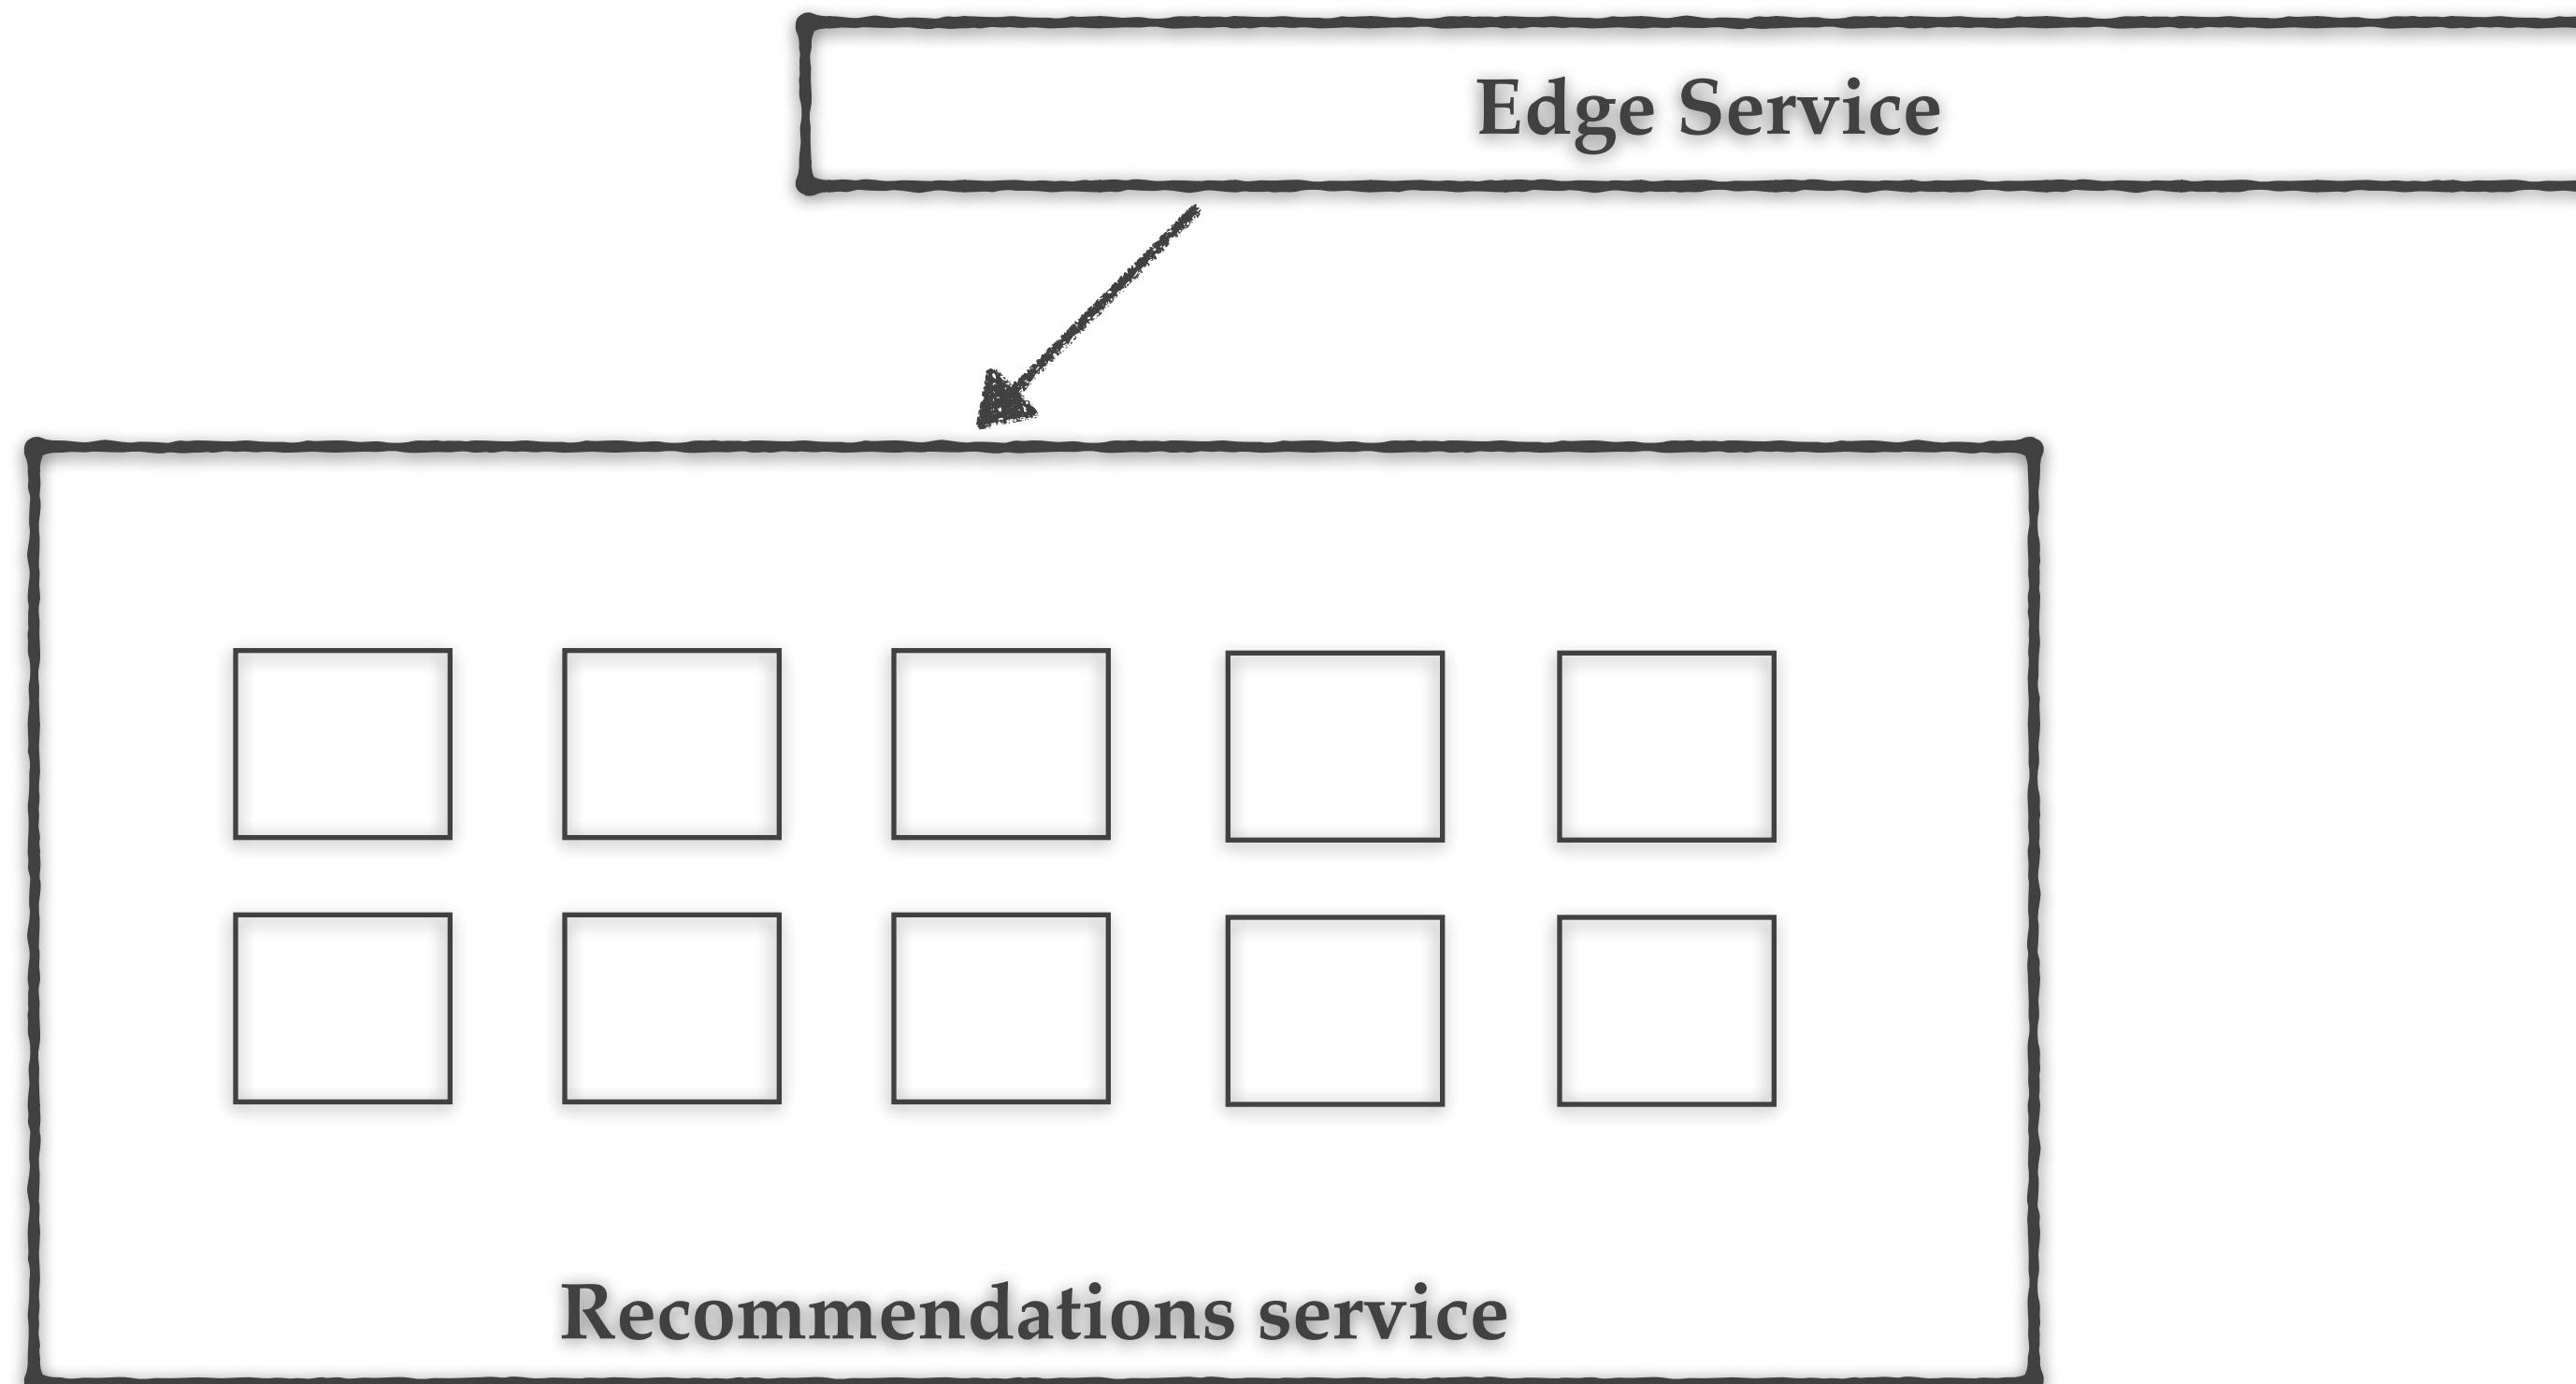
How is service discovery used?

Delivering Netflix



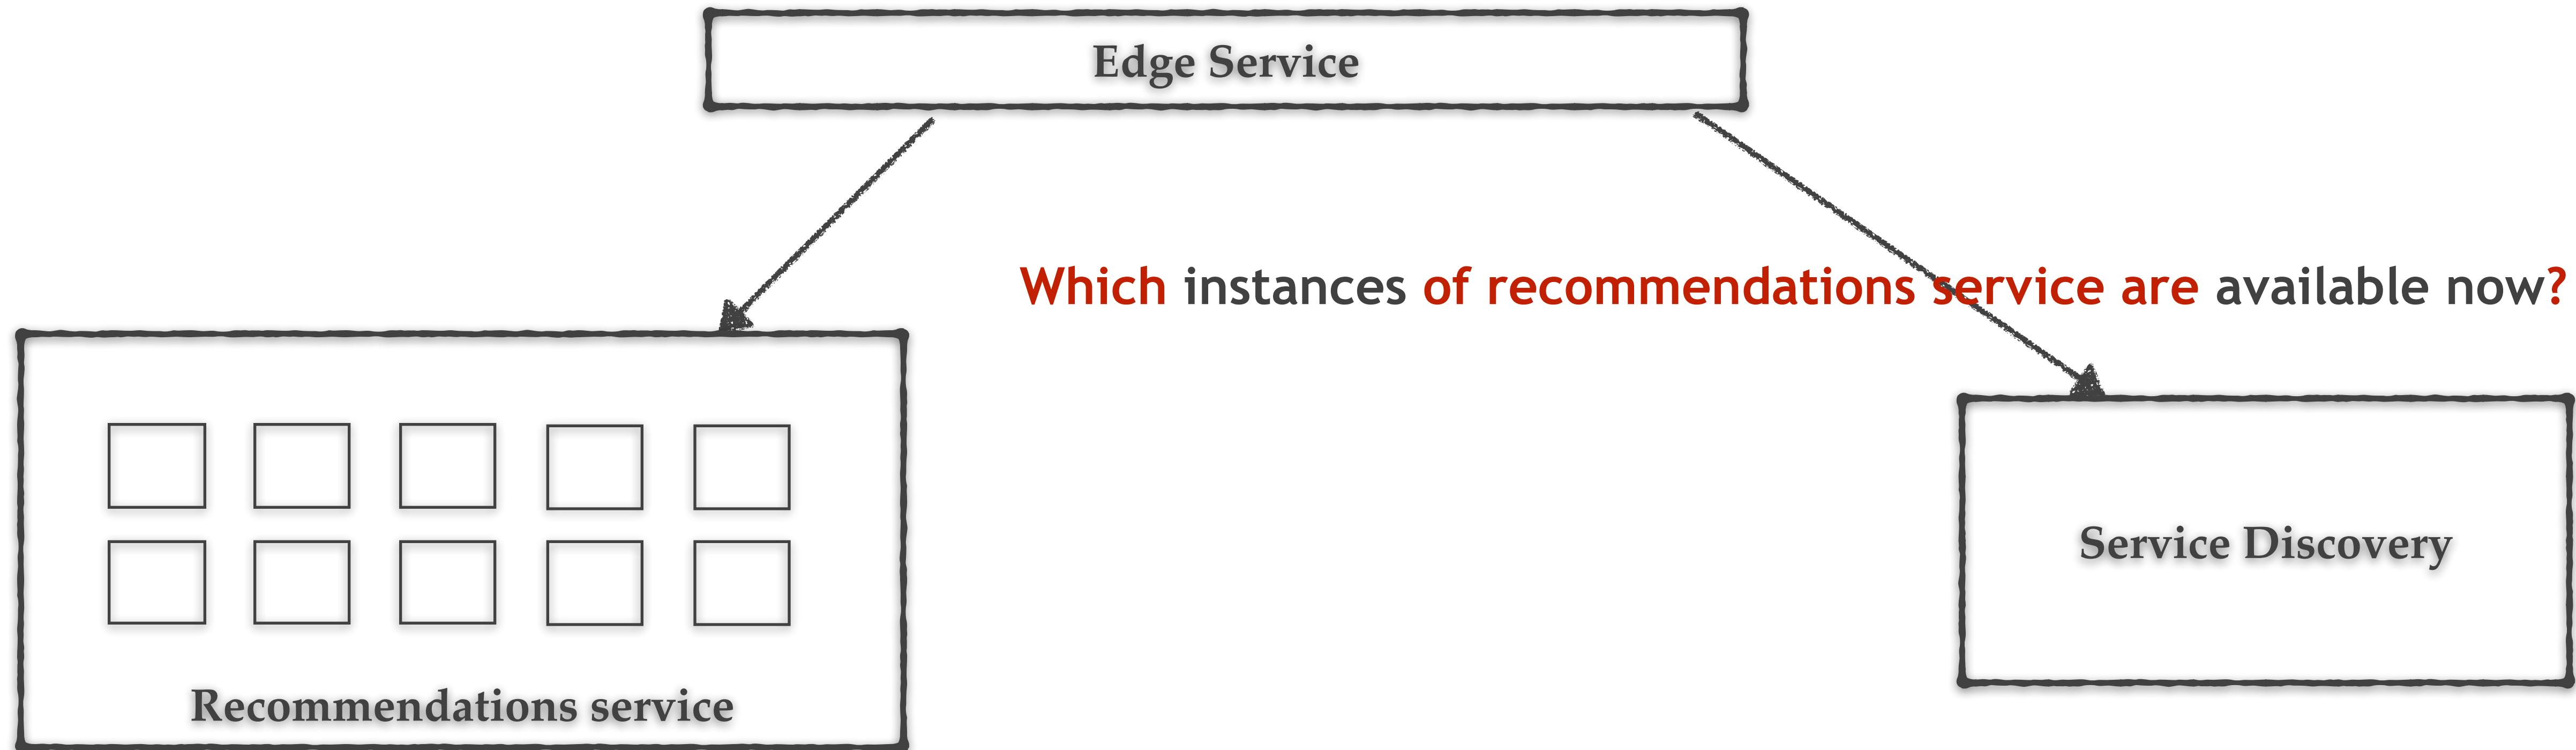
Disclaimer: This is an example and not an exact representation of the processing

Delivering Netflix



Disclaimer: This is an example and not an exact representation of the processing

Delivering Netflix

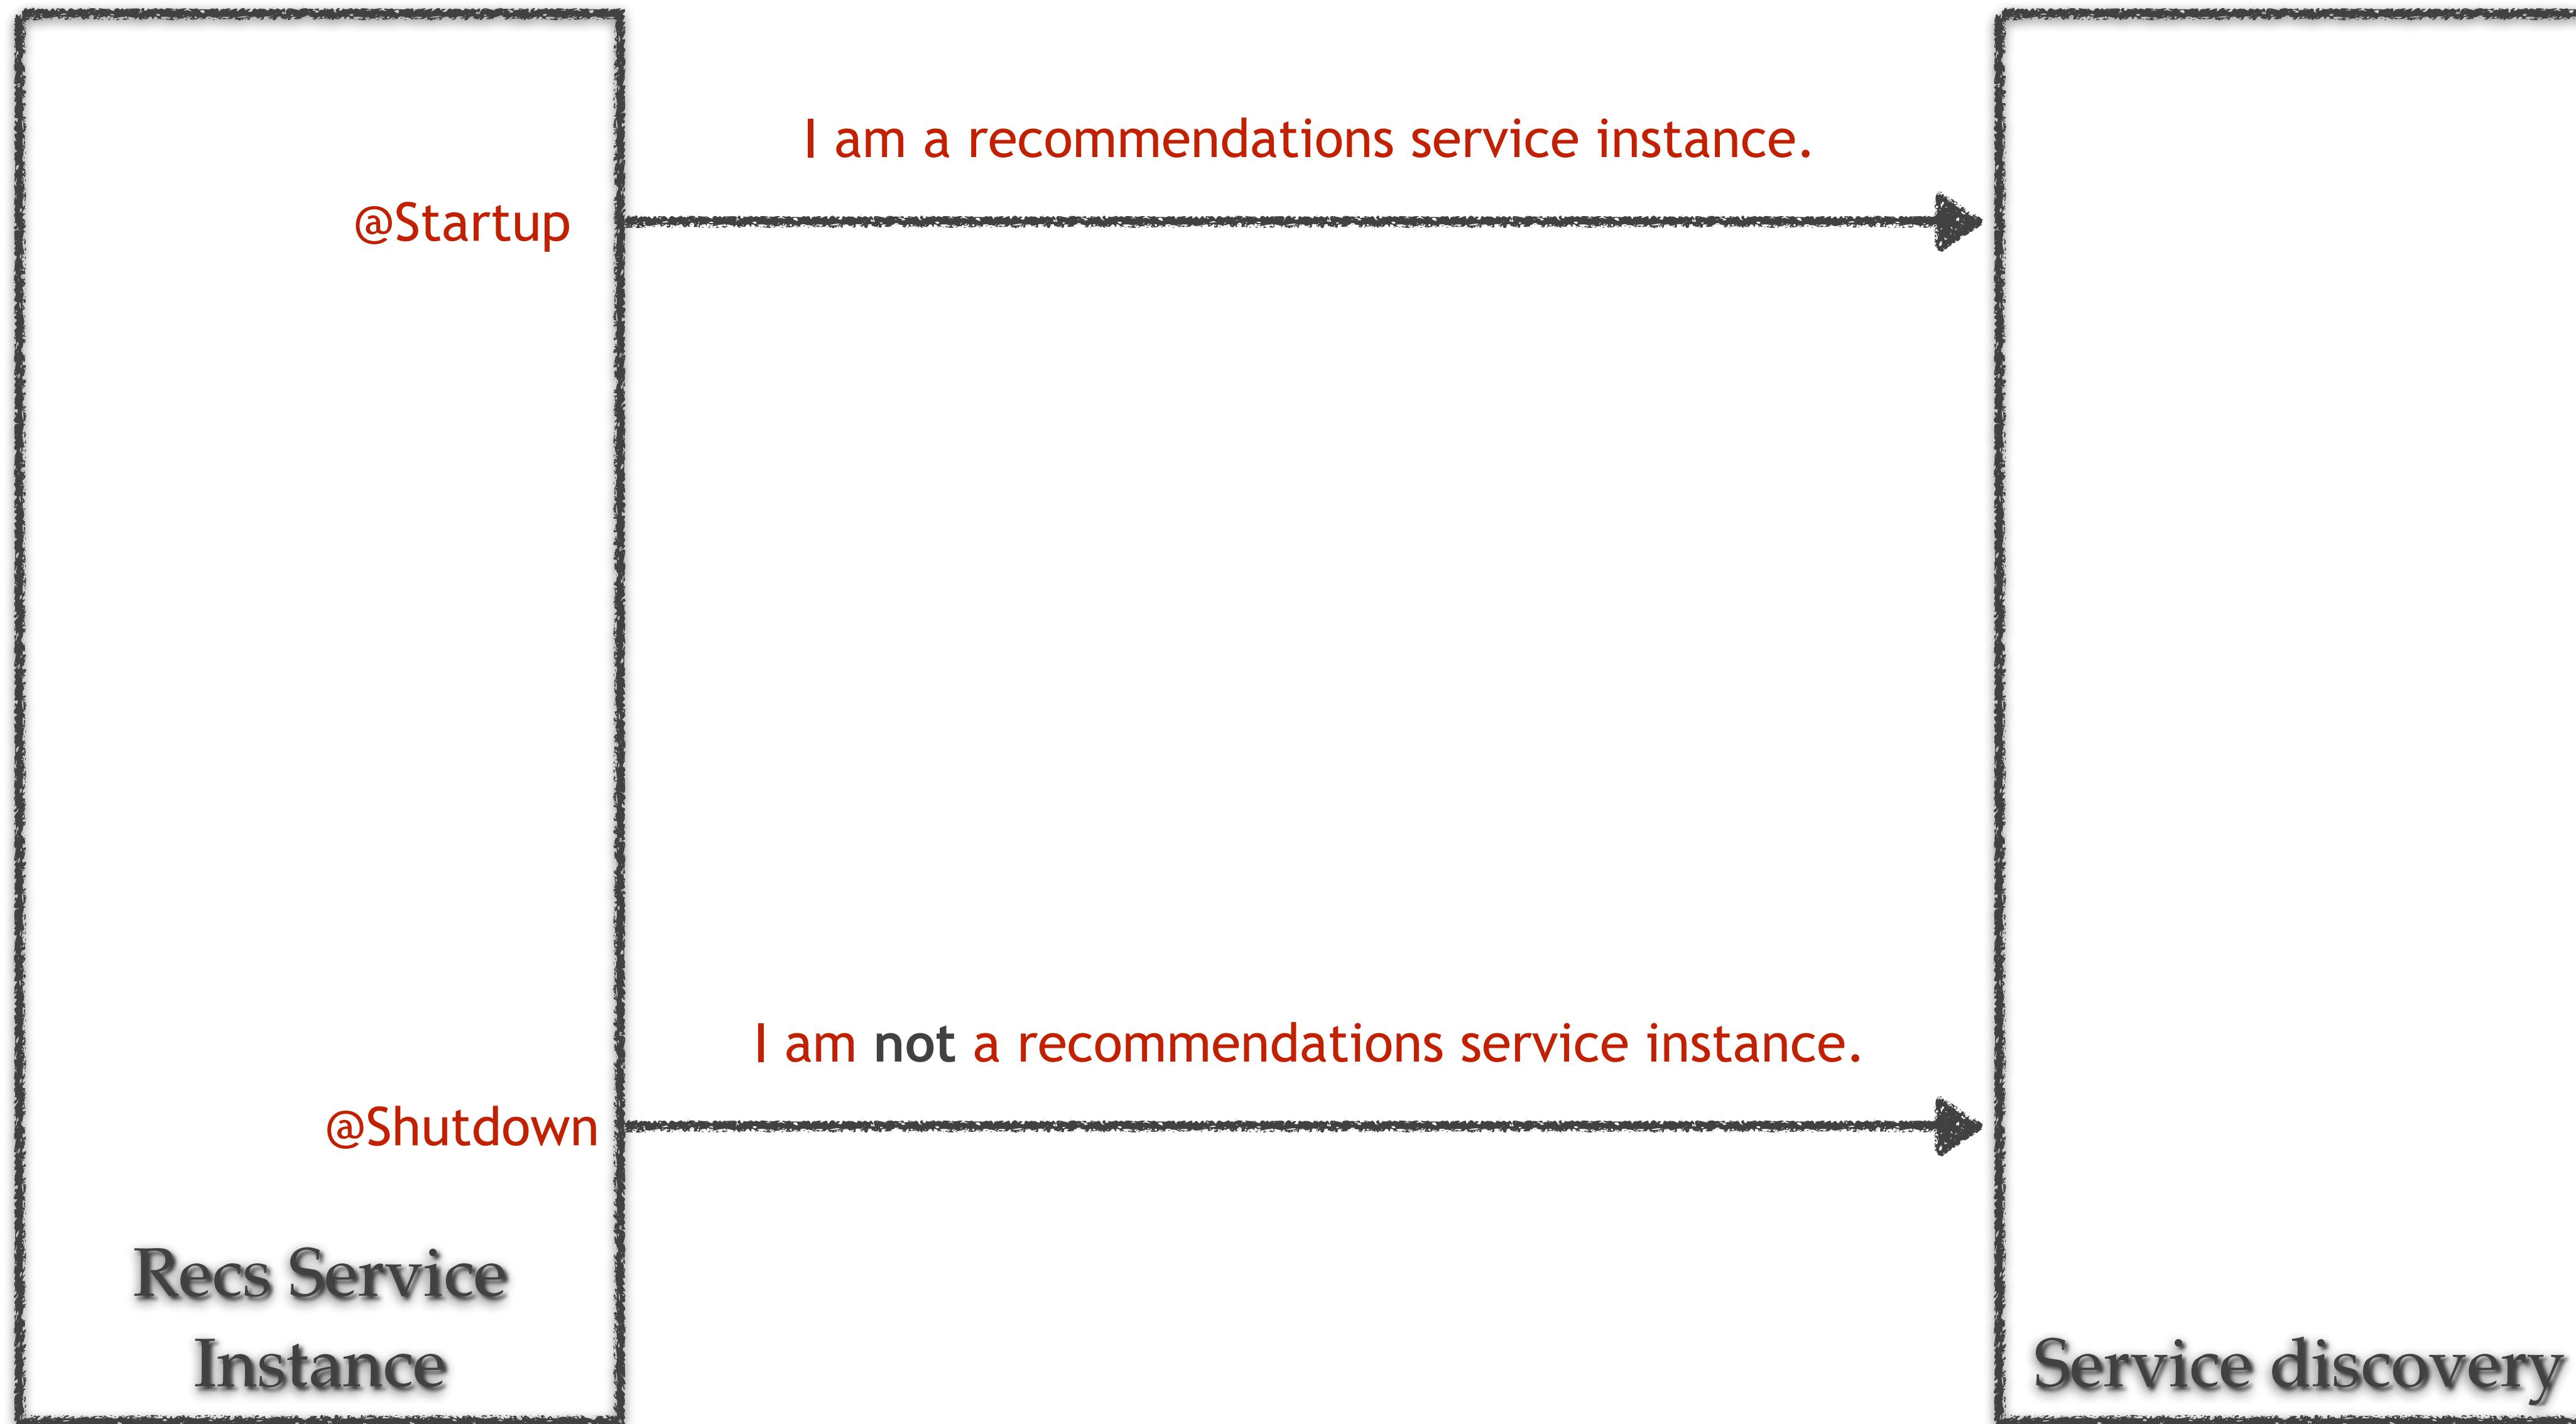


Disclaimer: This is an example and not an exact representation of the processing

Which instances of recommendations service are available now?

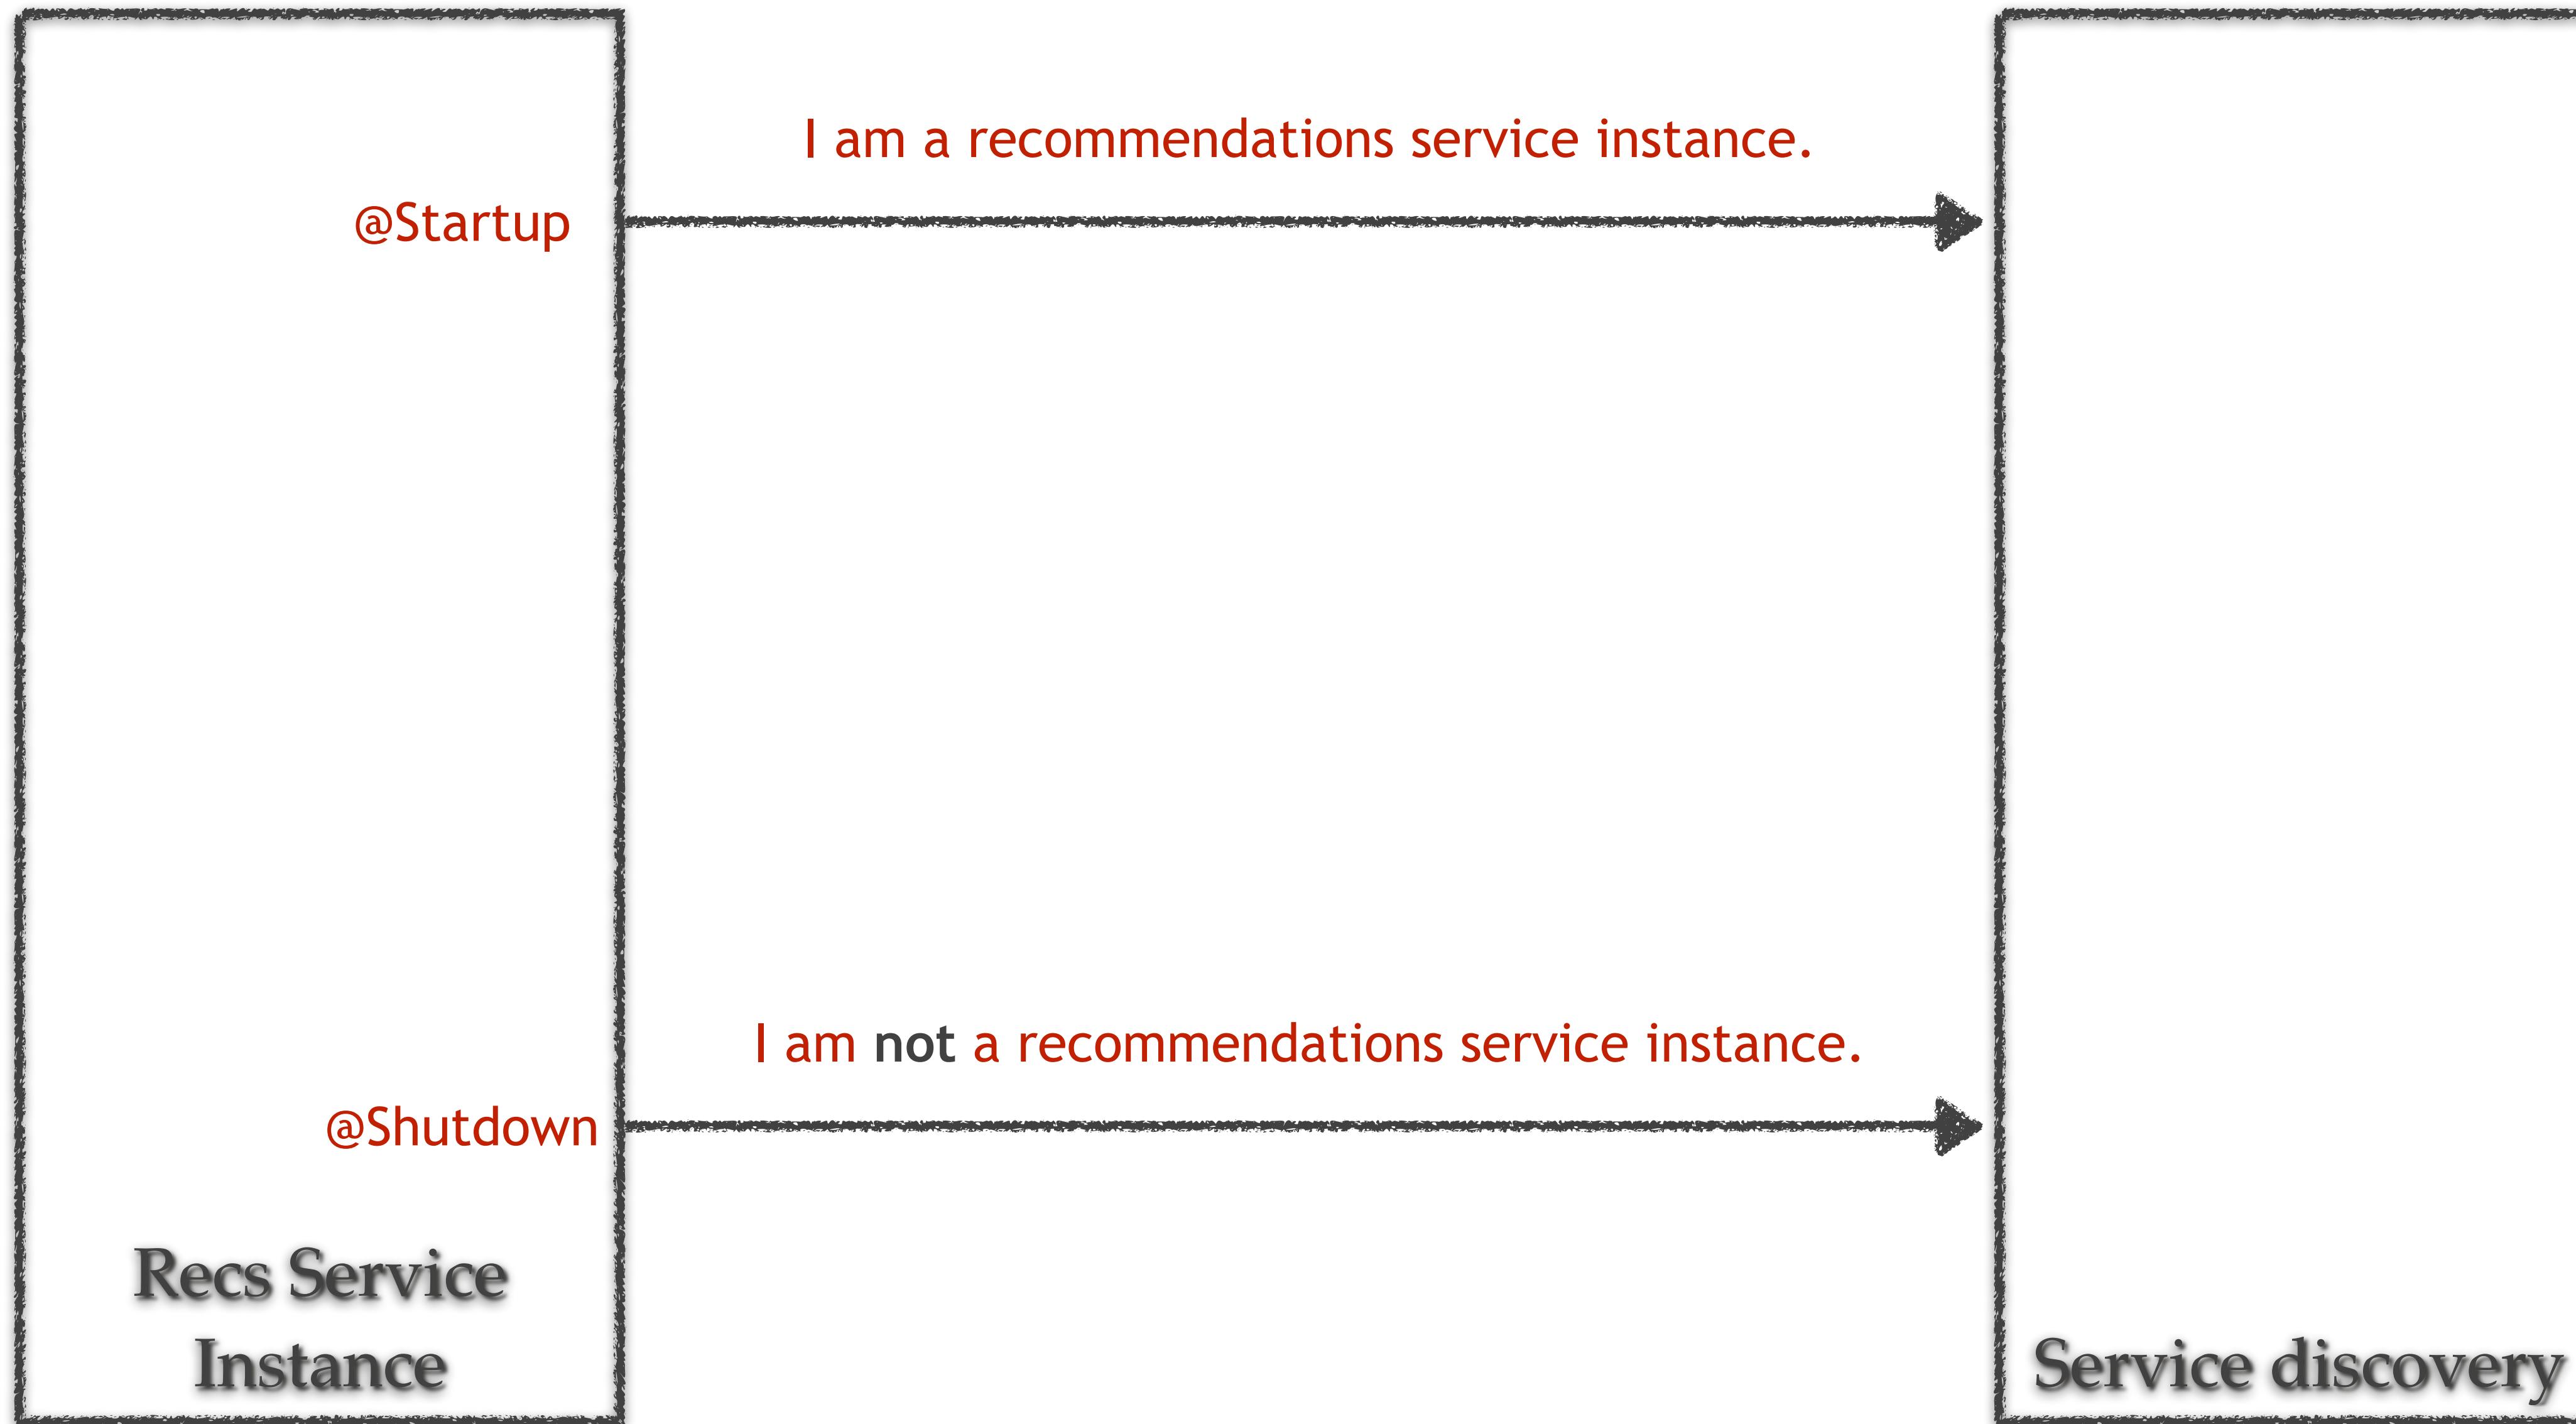
Which of the machines in the datacenter have the recommendations service software deployed?

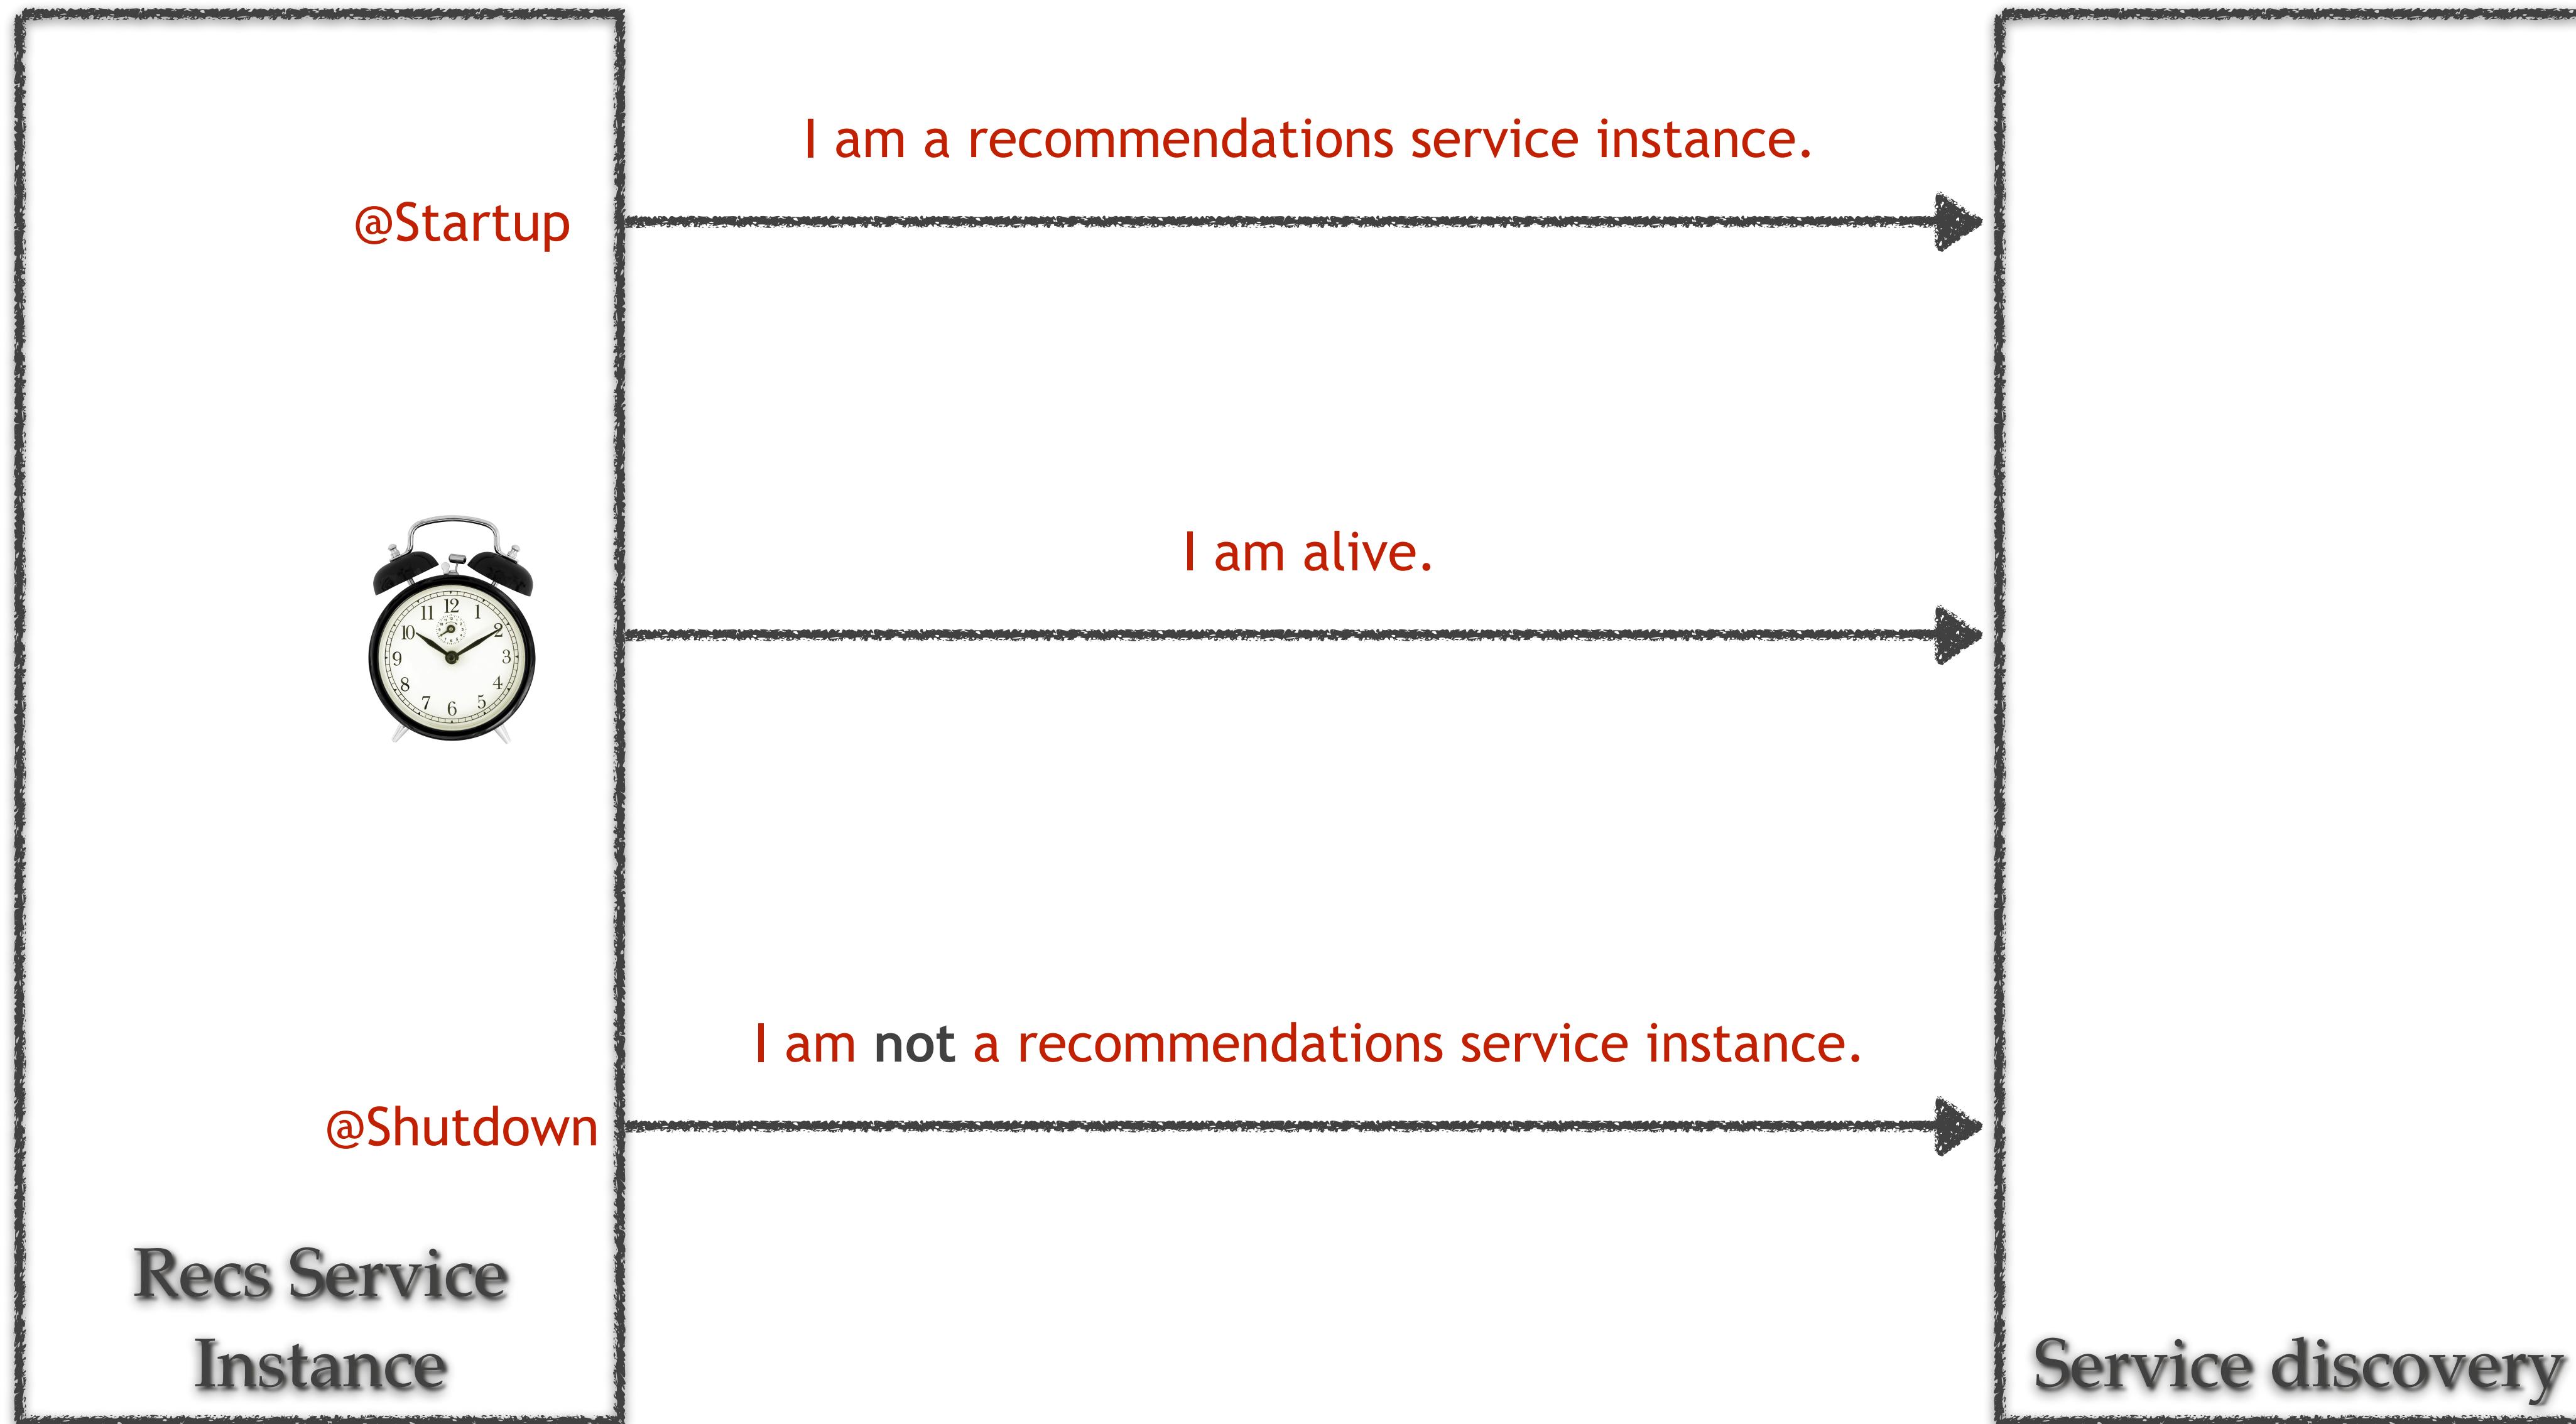
Service Name	Nodes
User Service	10.10.1.1, 10.10.1.2
Metadata Service	10.10.2.1, 10.10.2.2
Recommendations Service	10.10.3.1, 10.10.3.2



Which instances of recommendations service are available now?

available now?





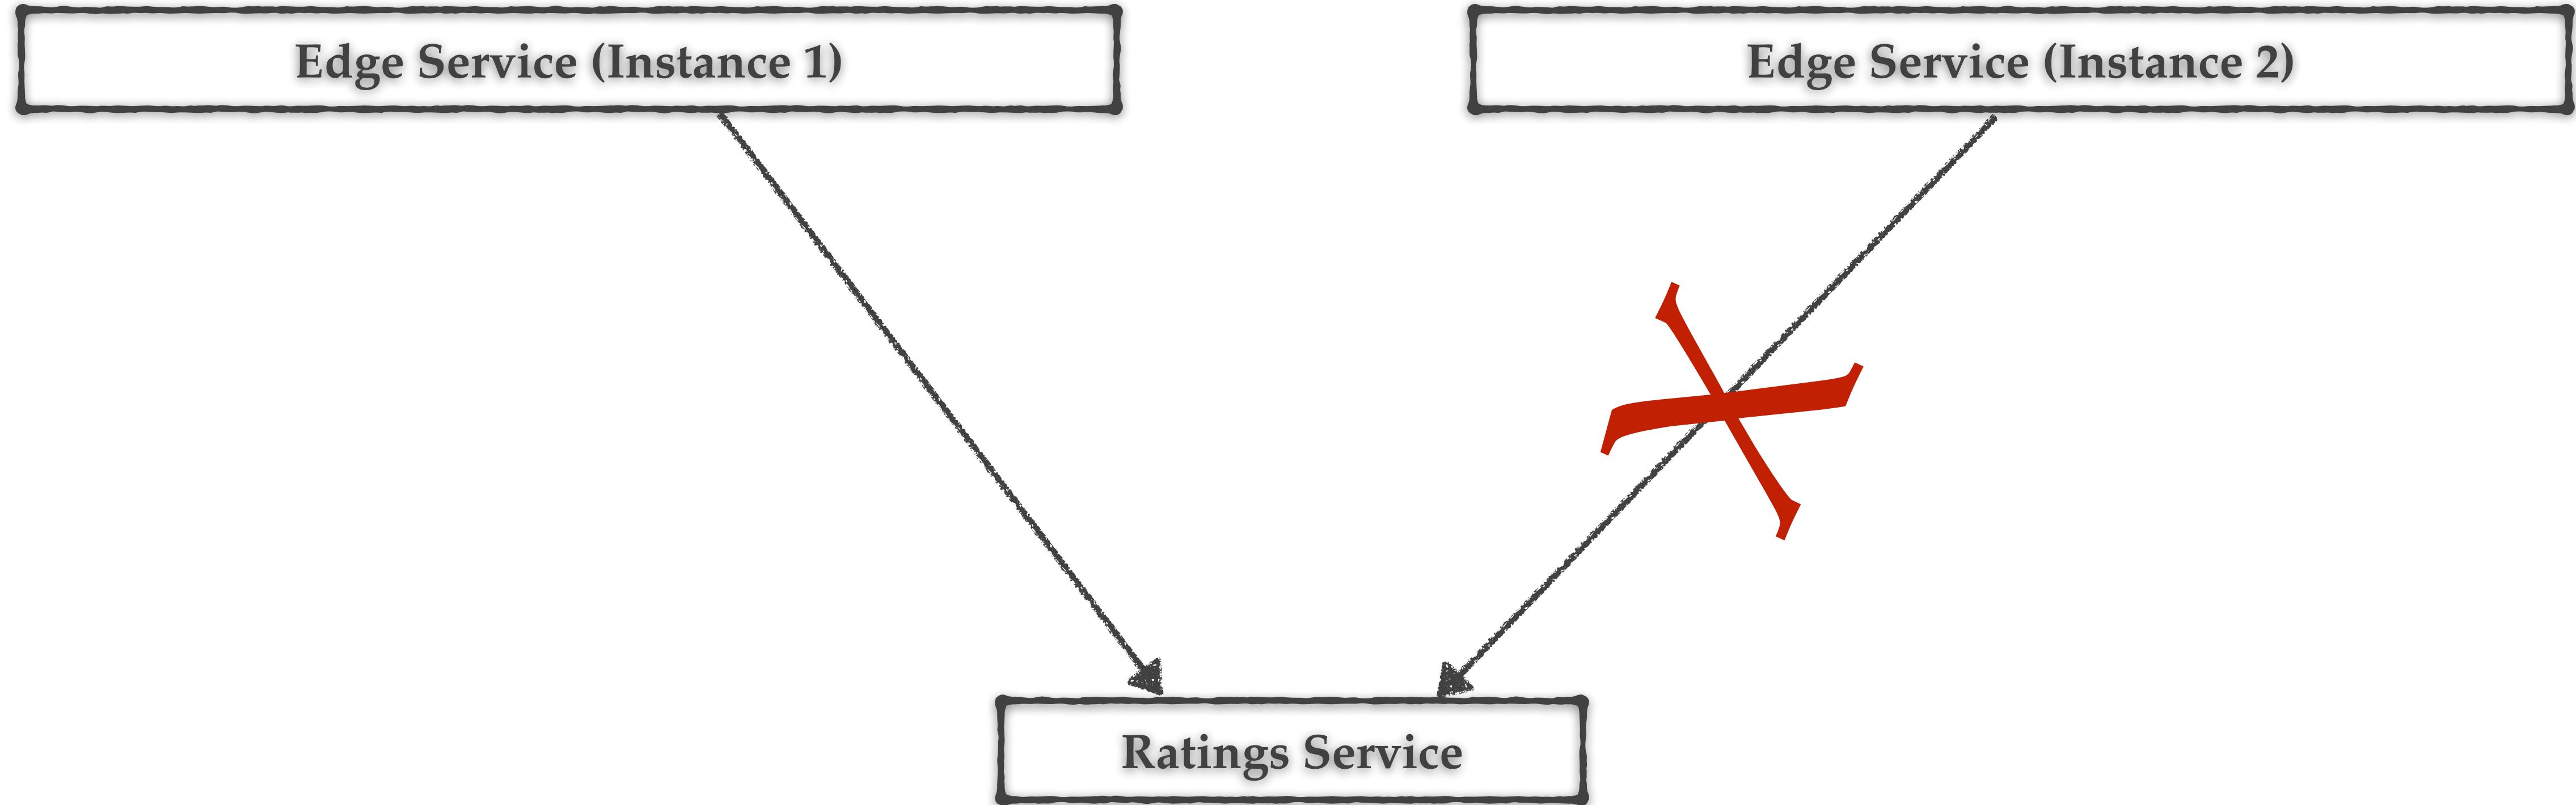
available now?

**YEAH, IF YOU CAN ANSWER THIS
QUESTION**

**THAT WOULD BE
GREAT**

MEMES & FUNNY STUFF funnyjunk.com

Network Partitions



Network Partitions

Being available in a distributed system, is very subjective.

Network Partitions

Being available in a distributed system, is very subjective.

What is available to one can be unavailable to other.

Network Partitions

Being available in a distributed system, is very subjective.

What is available to one can be unavailable to other.

A node's availability decision is best when it is local.

Service discovery does not
guarantee node availability.

Why should Service discovery store
node status?

Why should Service discovery store node status?

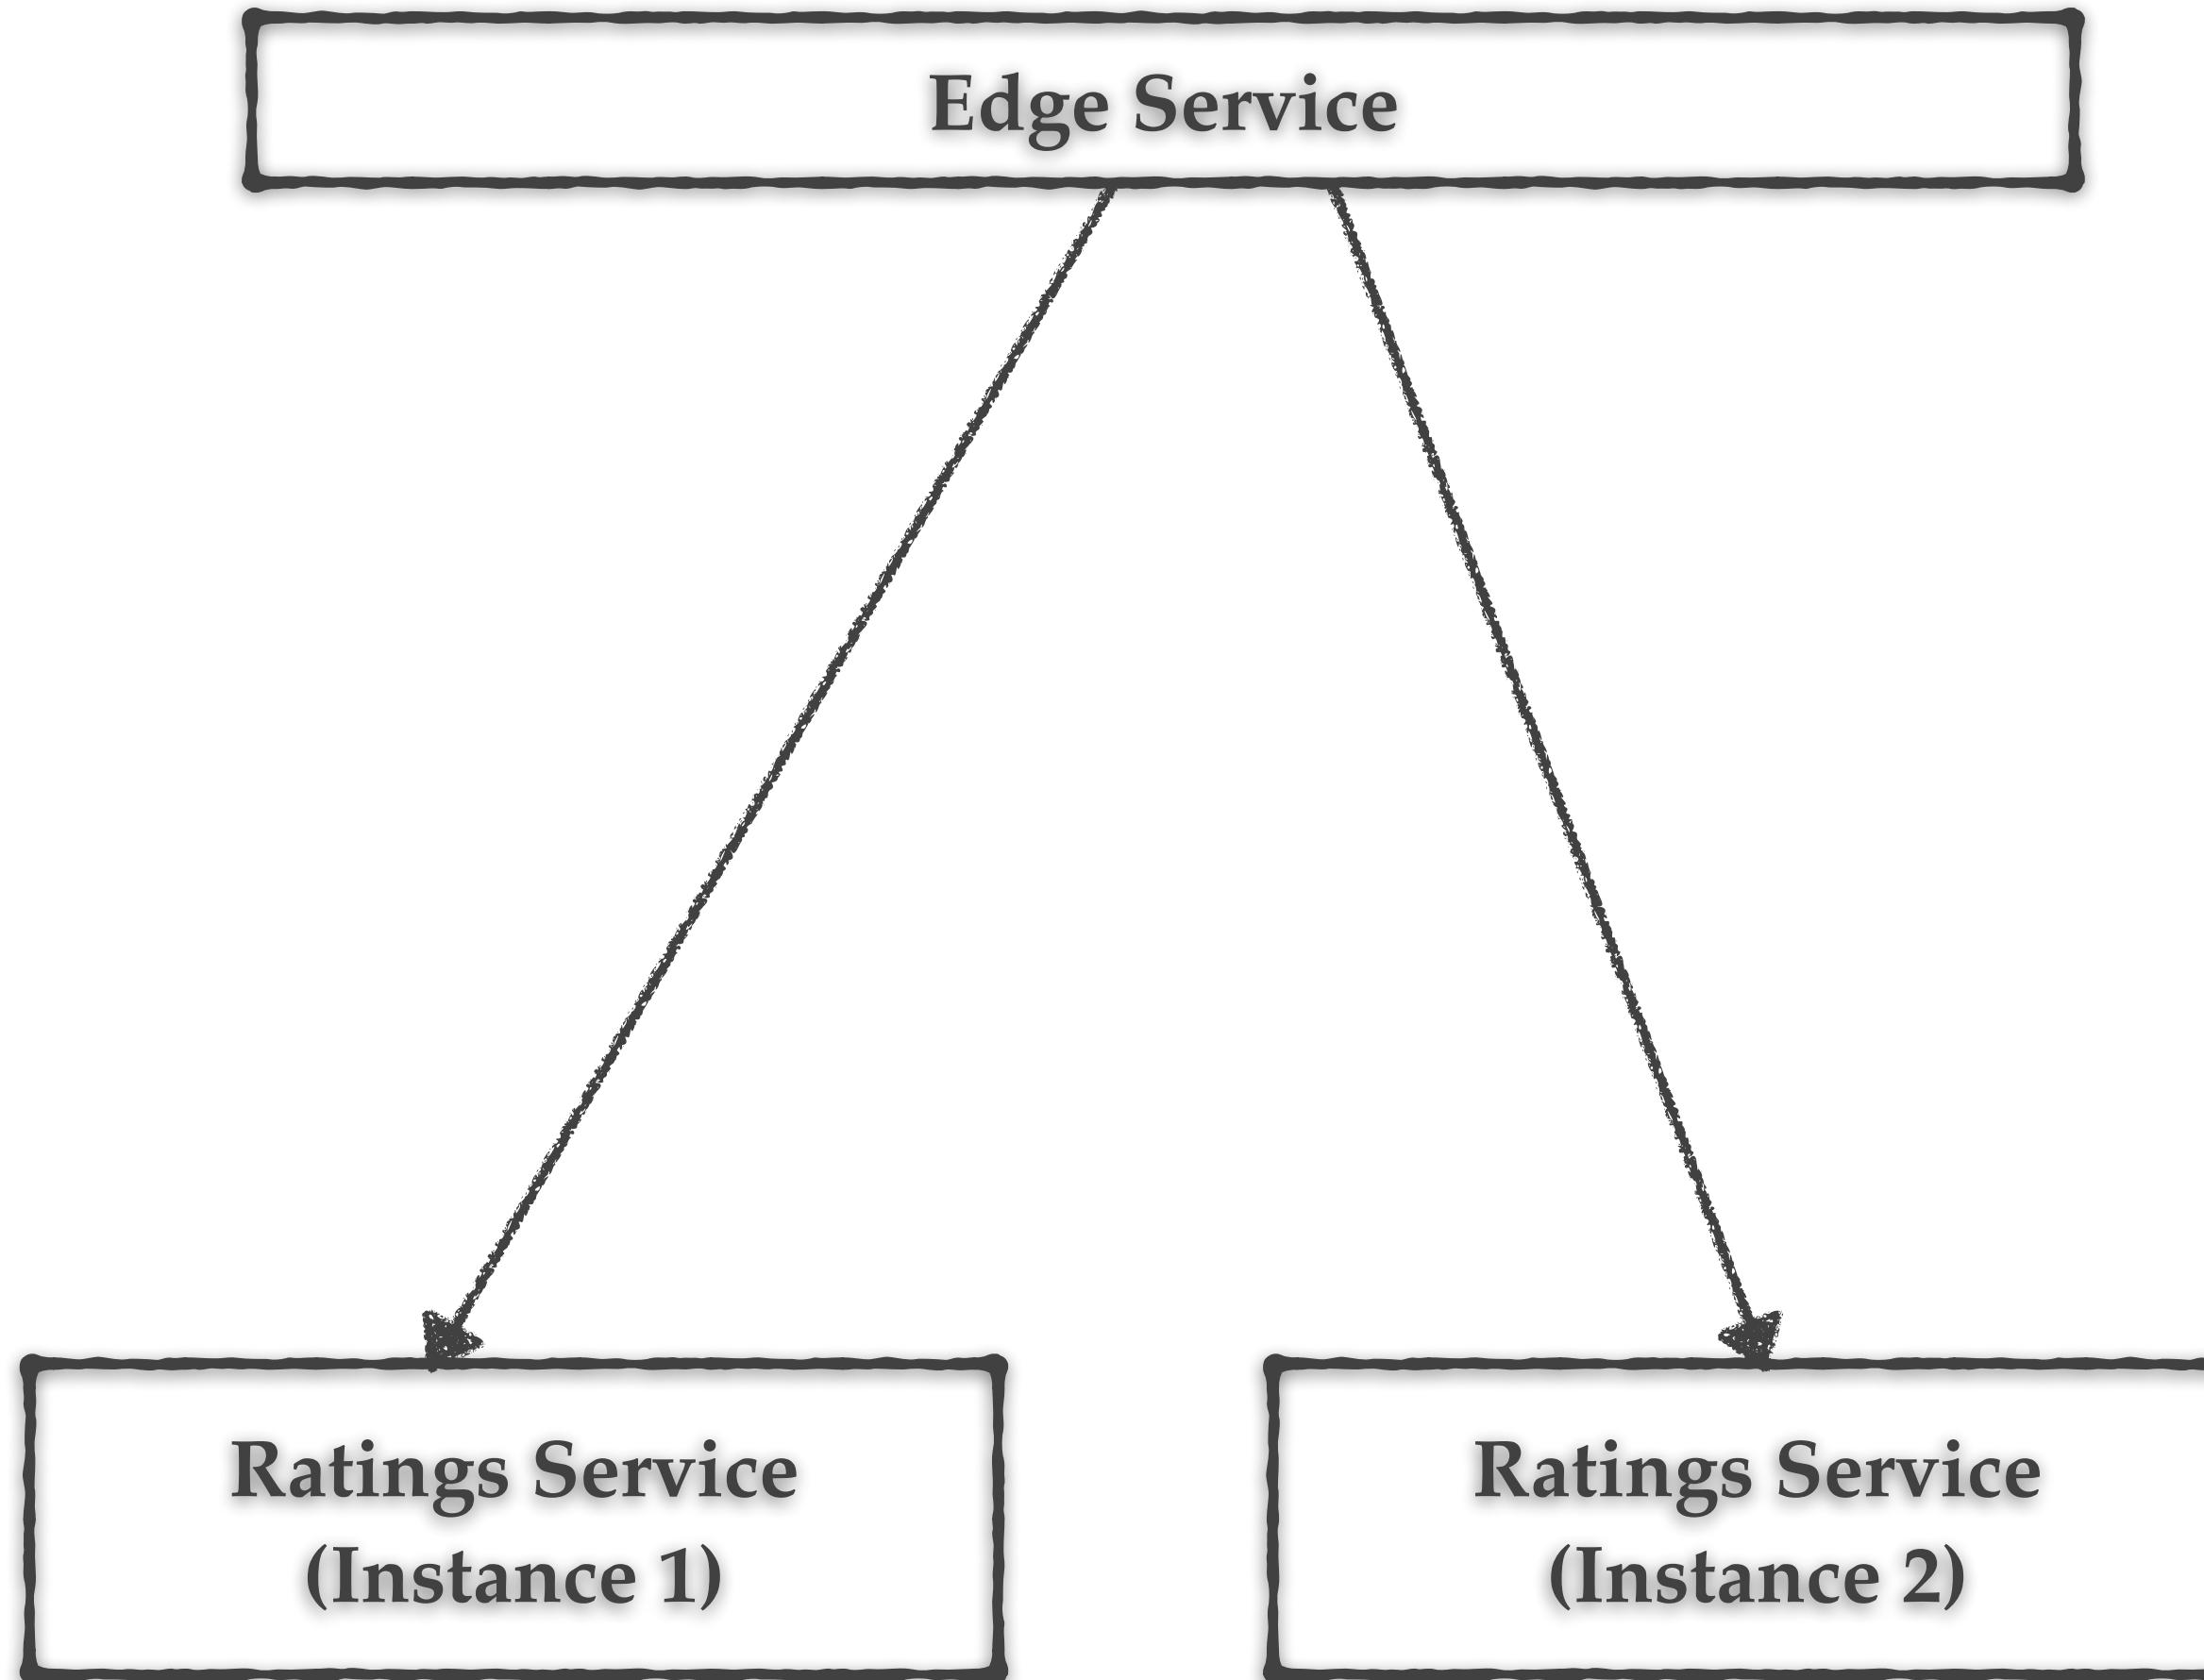
Availability is doubtful but Unavailability can be trusted.

Unavailability can be trusted

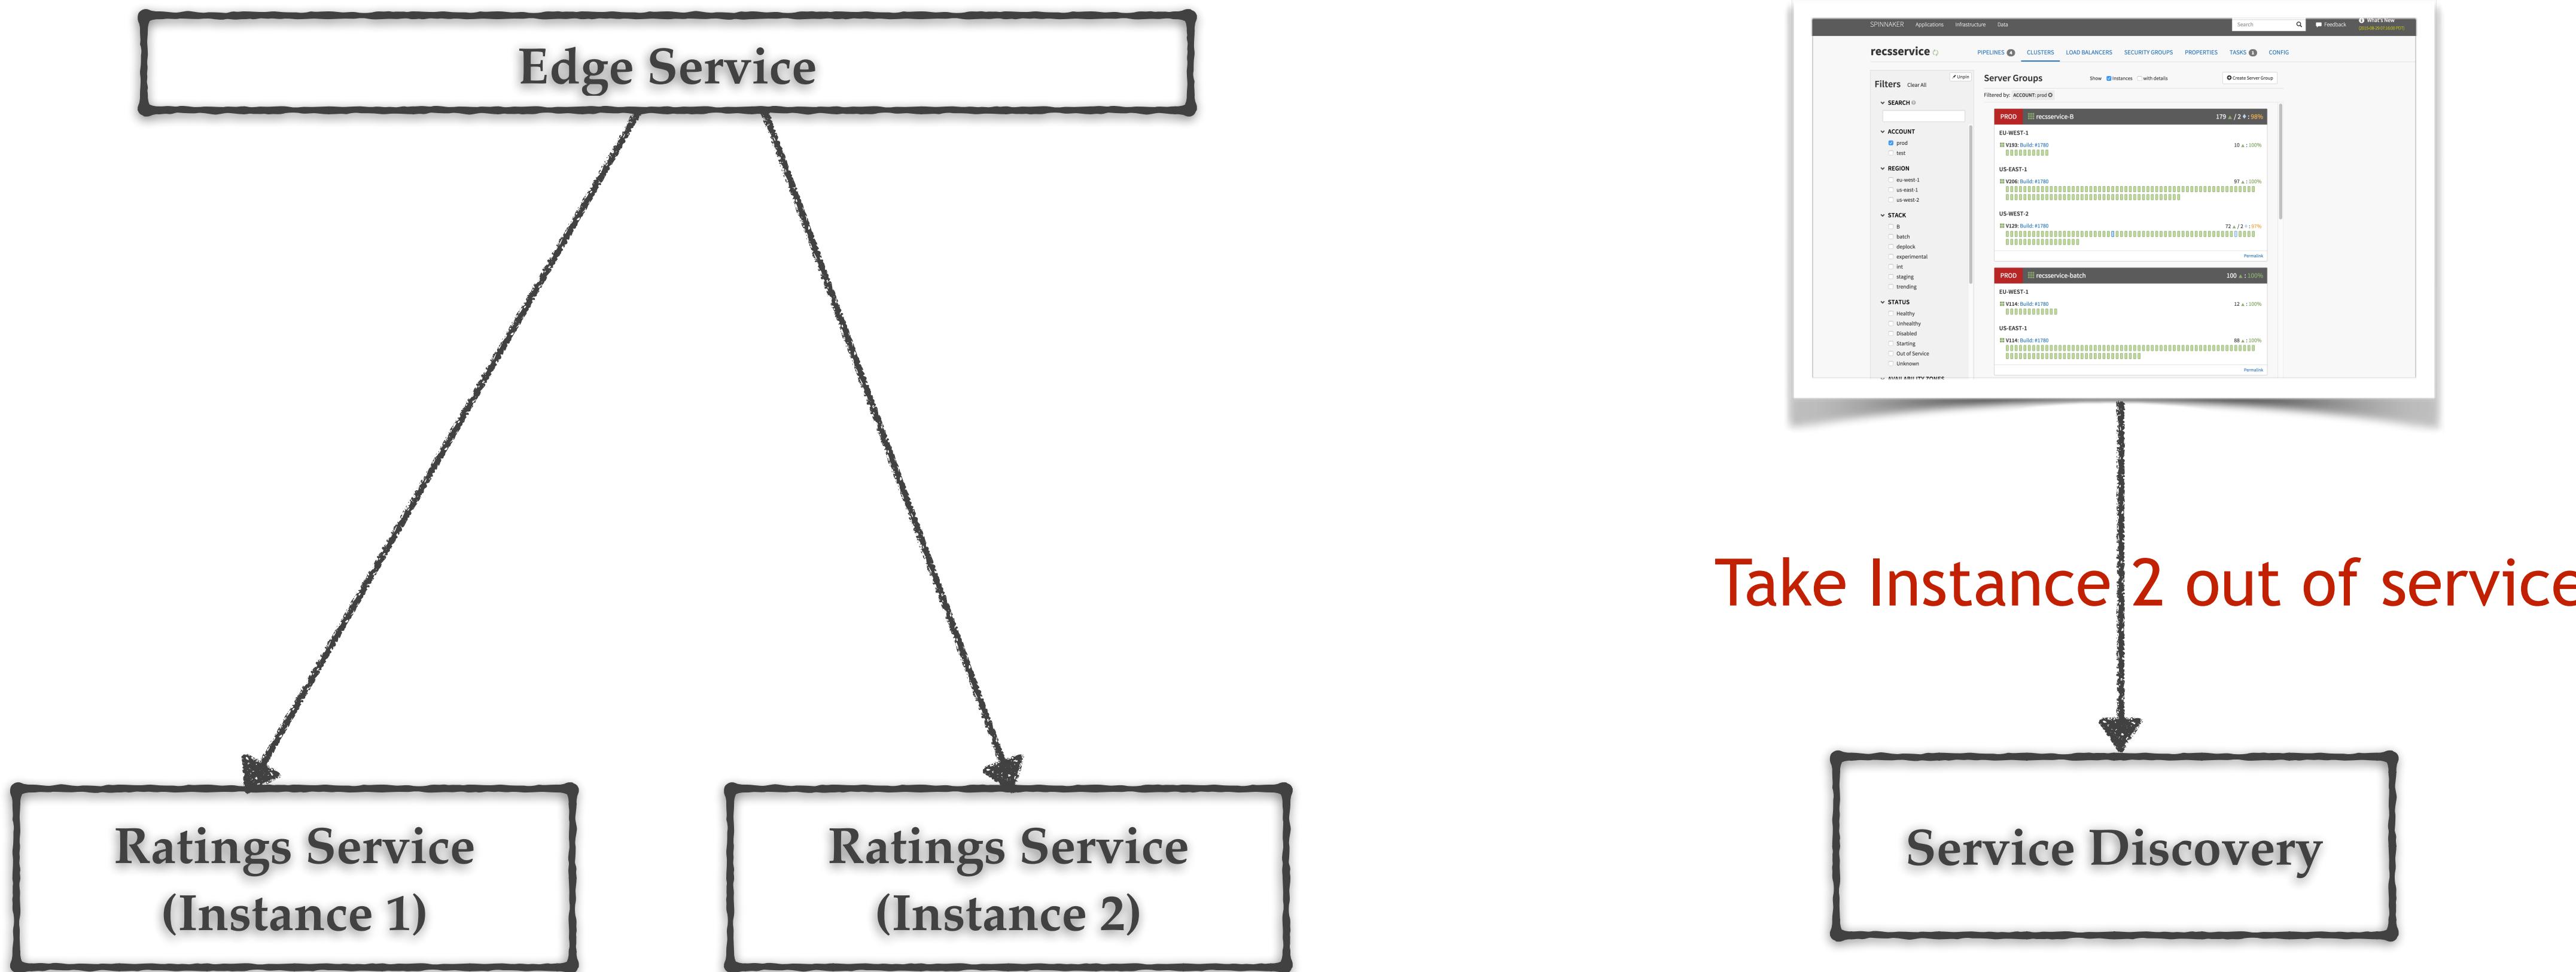
Unavailability can be trusted

Node status override

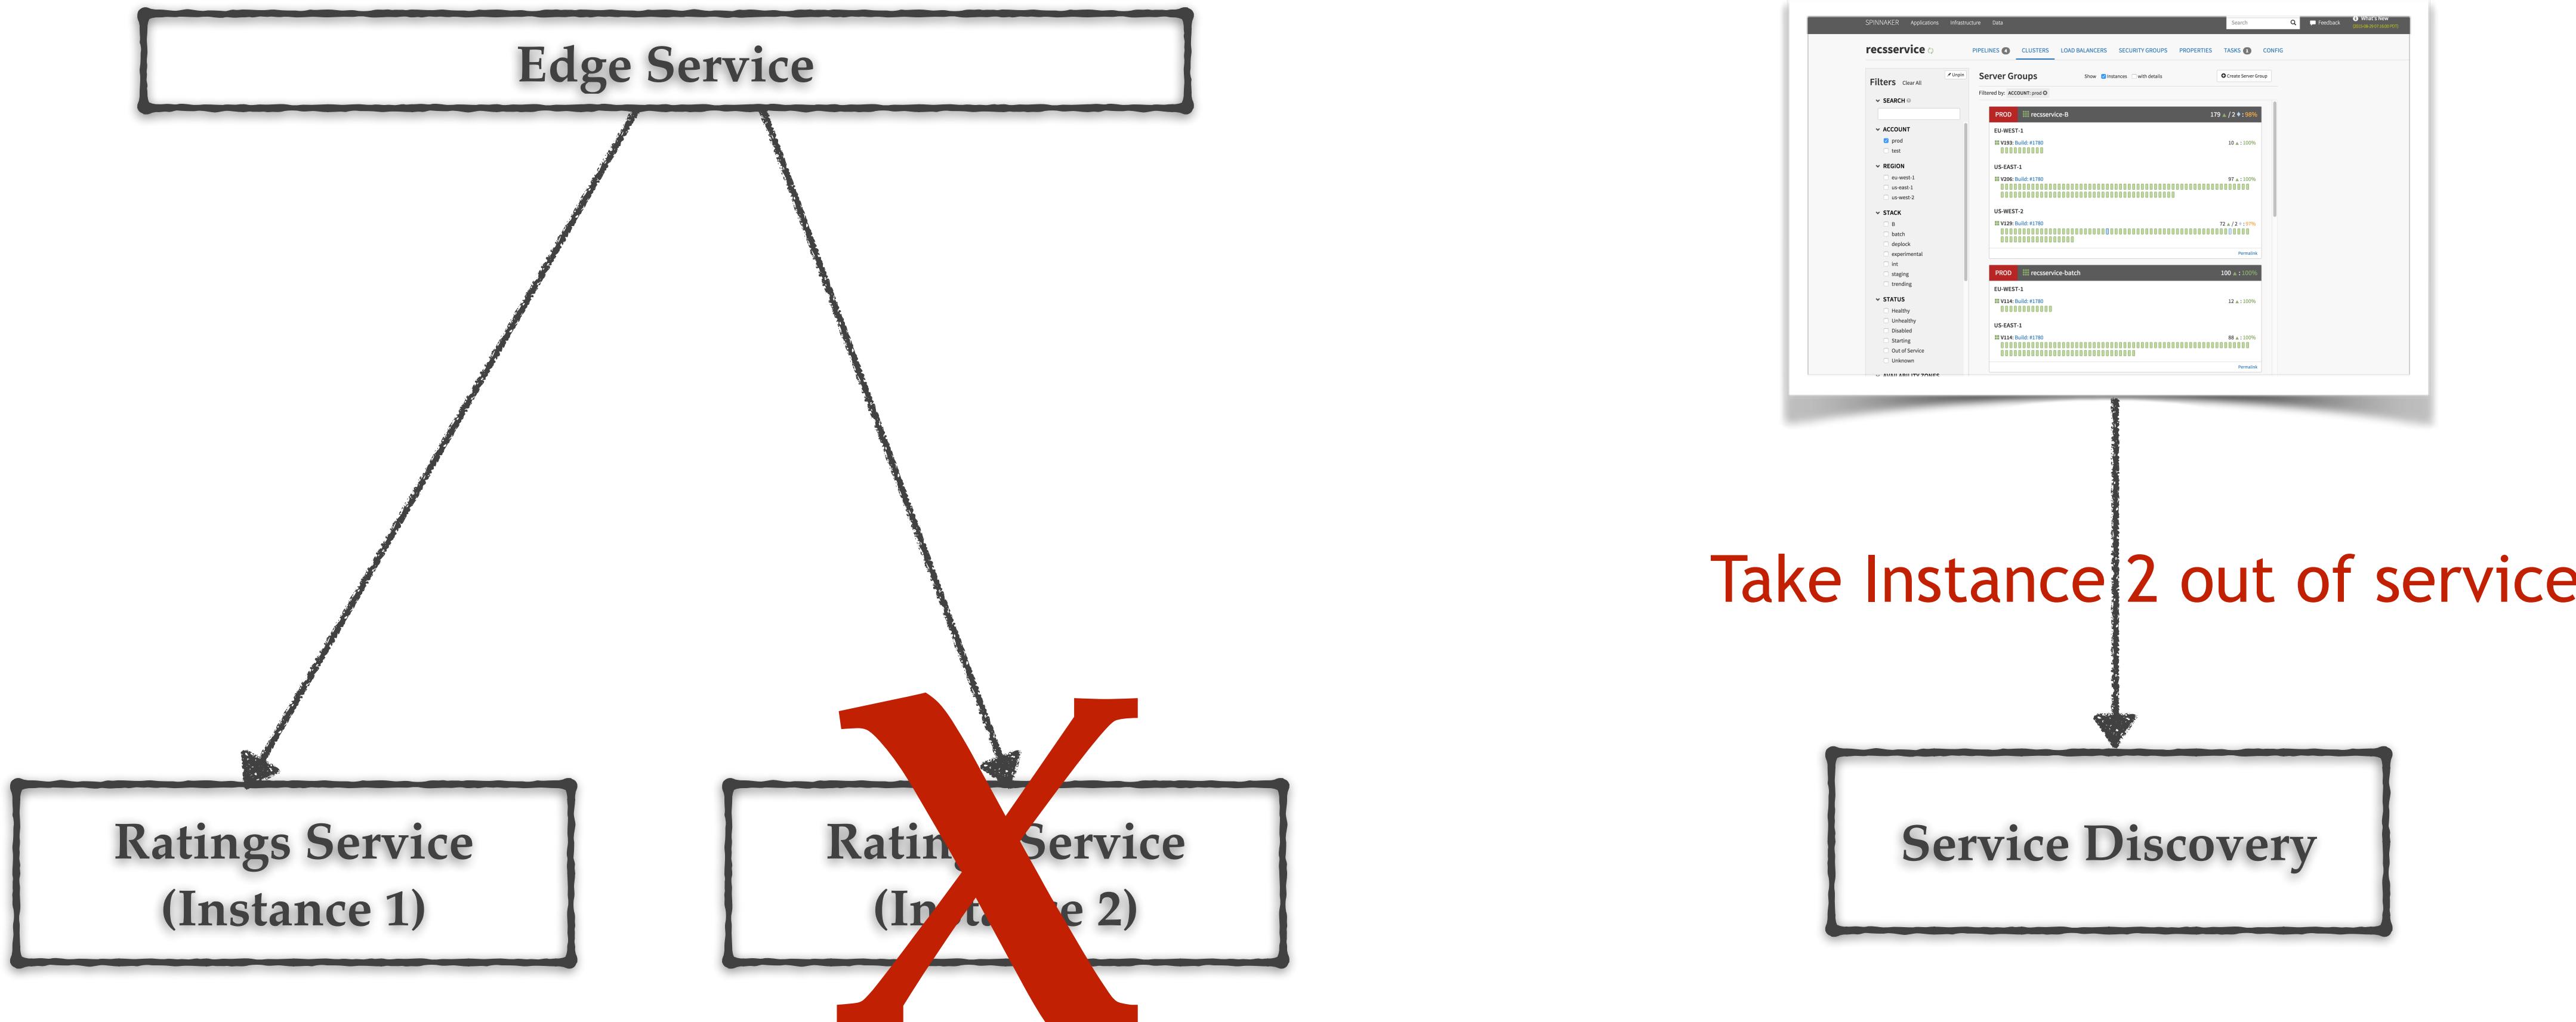
Node status override



Node status override



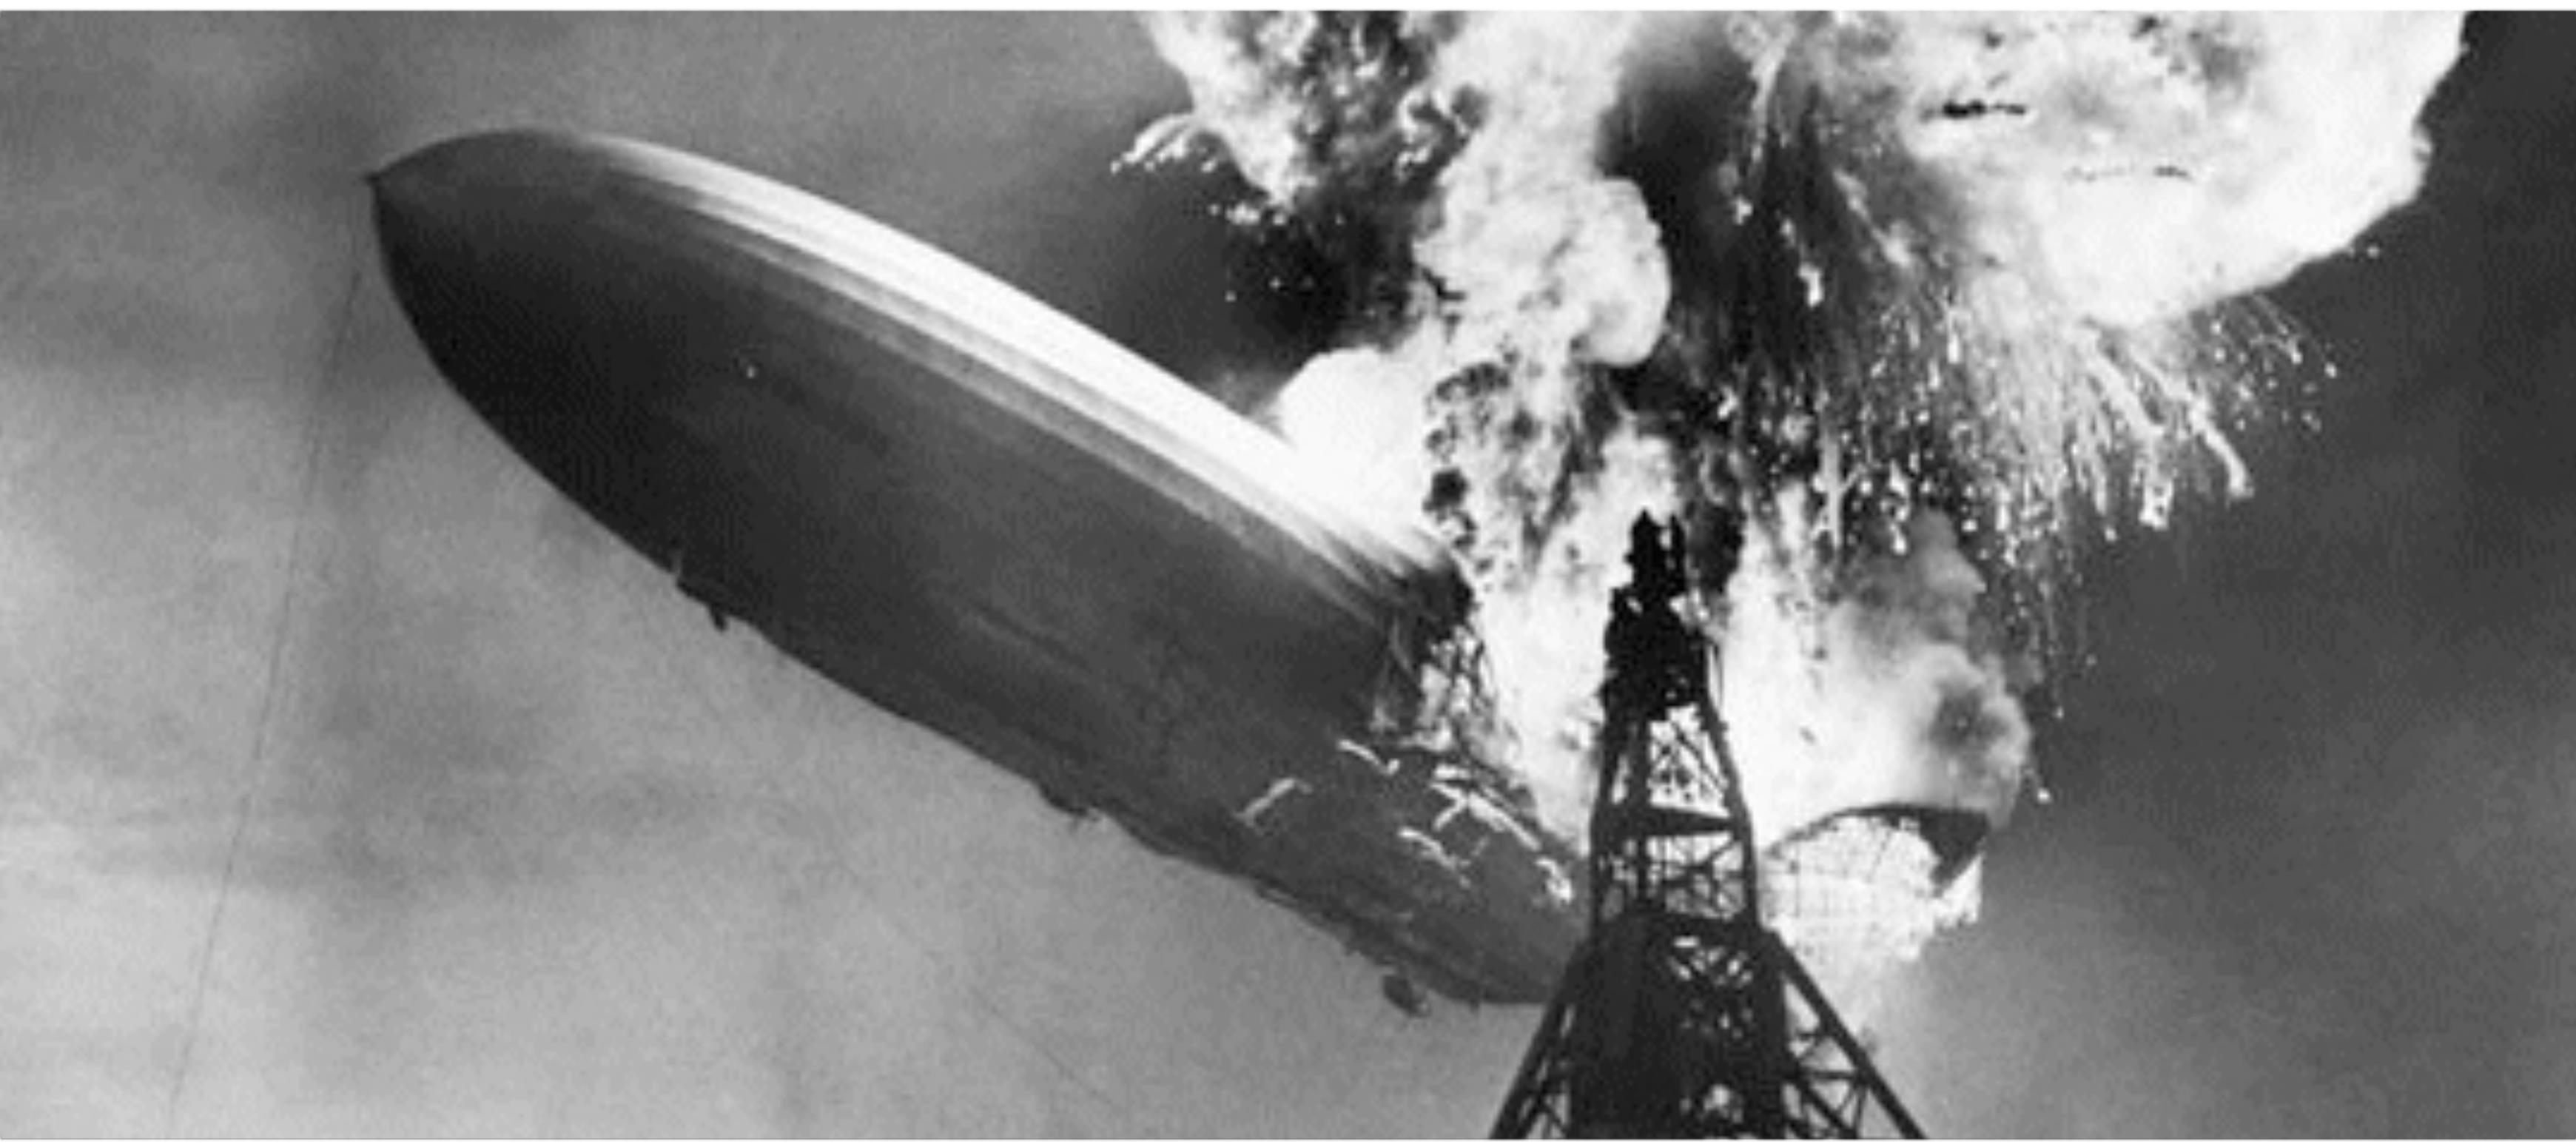
Node status override



Nodes can be isolated for debugging.

Nodes can be started in isolation.

Service discovery controls visibility of nodes but does not guarantee availability.



Failures

What if?

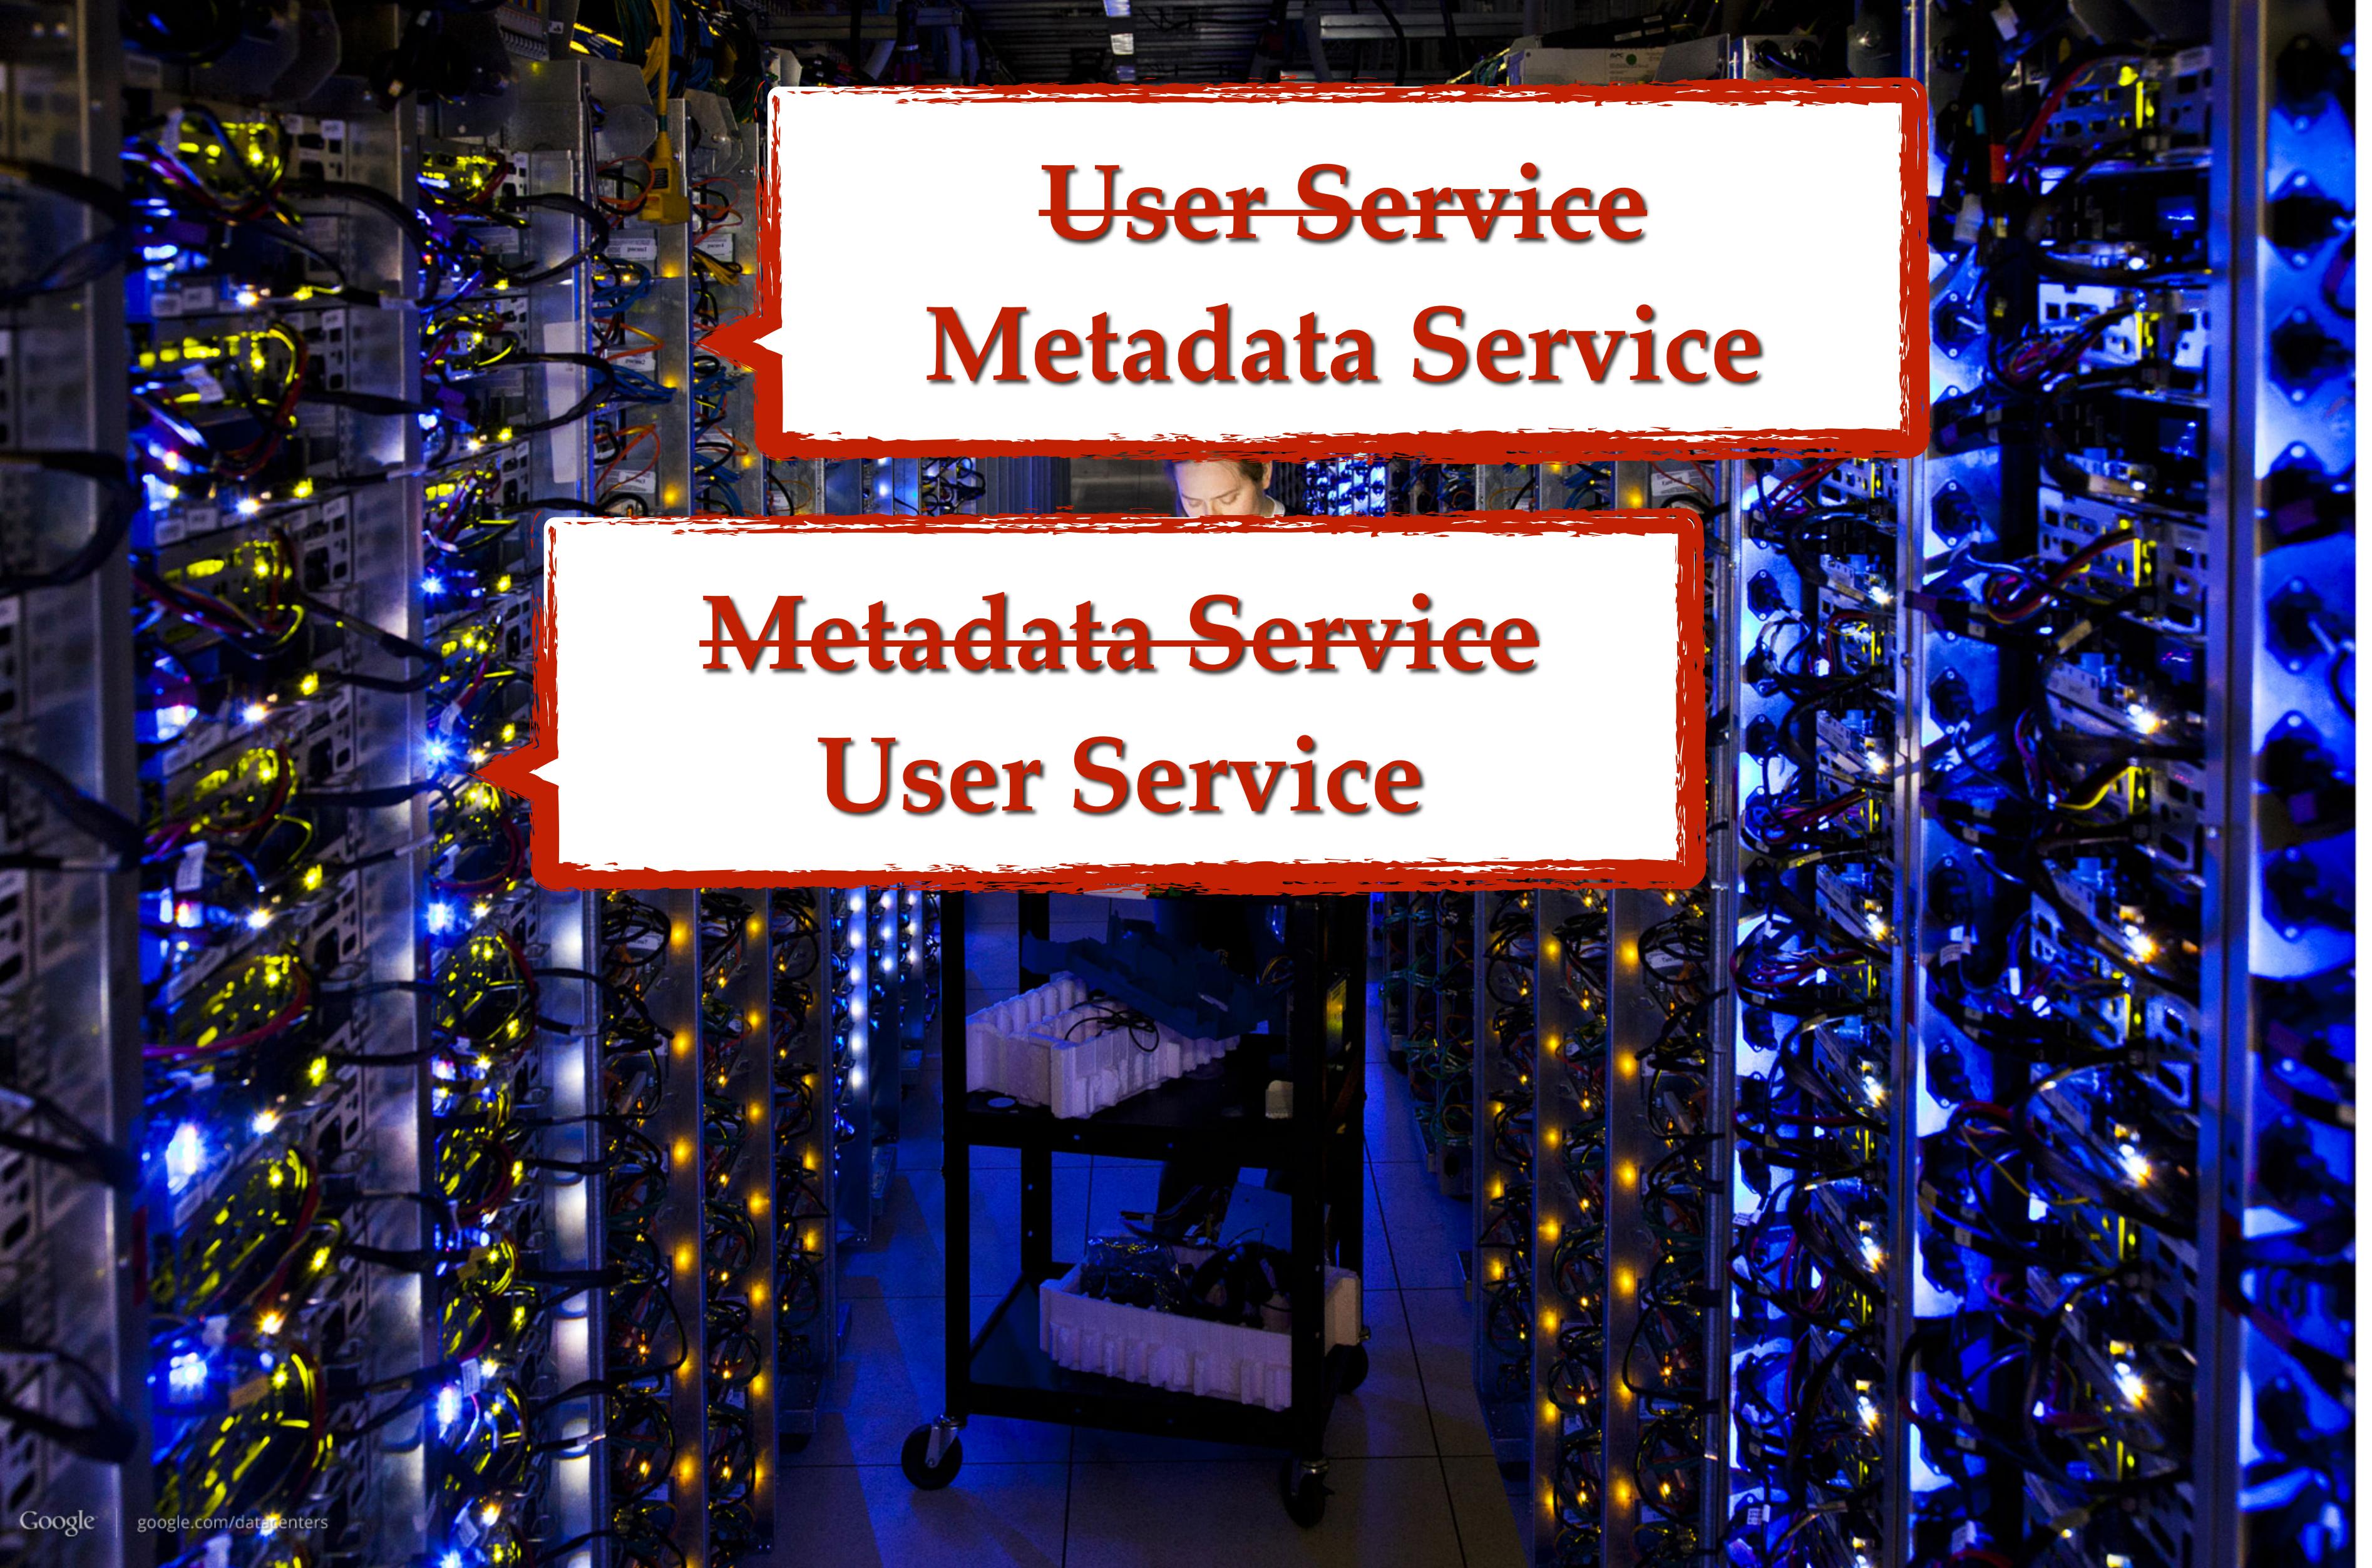
What if ...

Service discovery is unavailable?

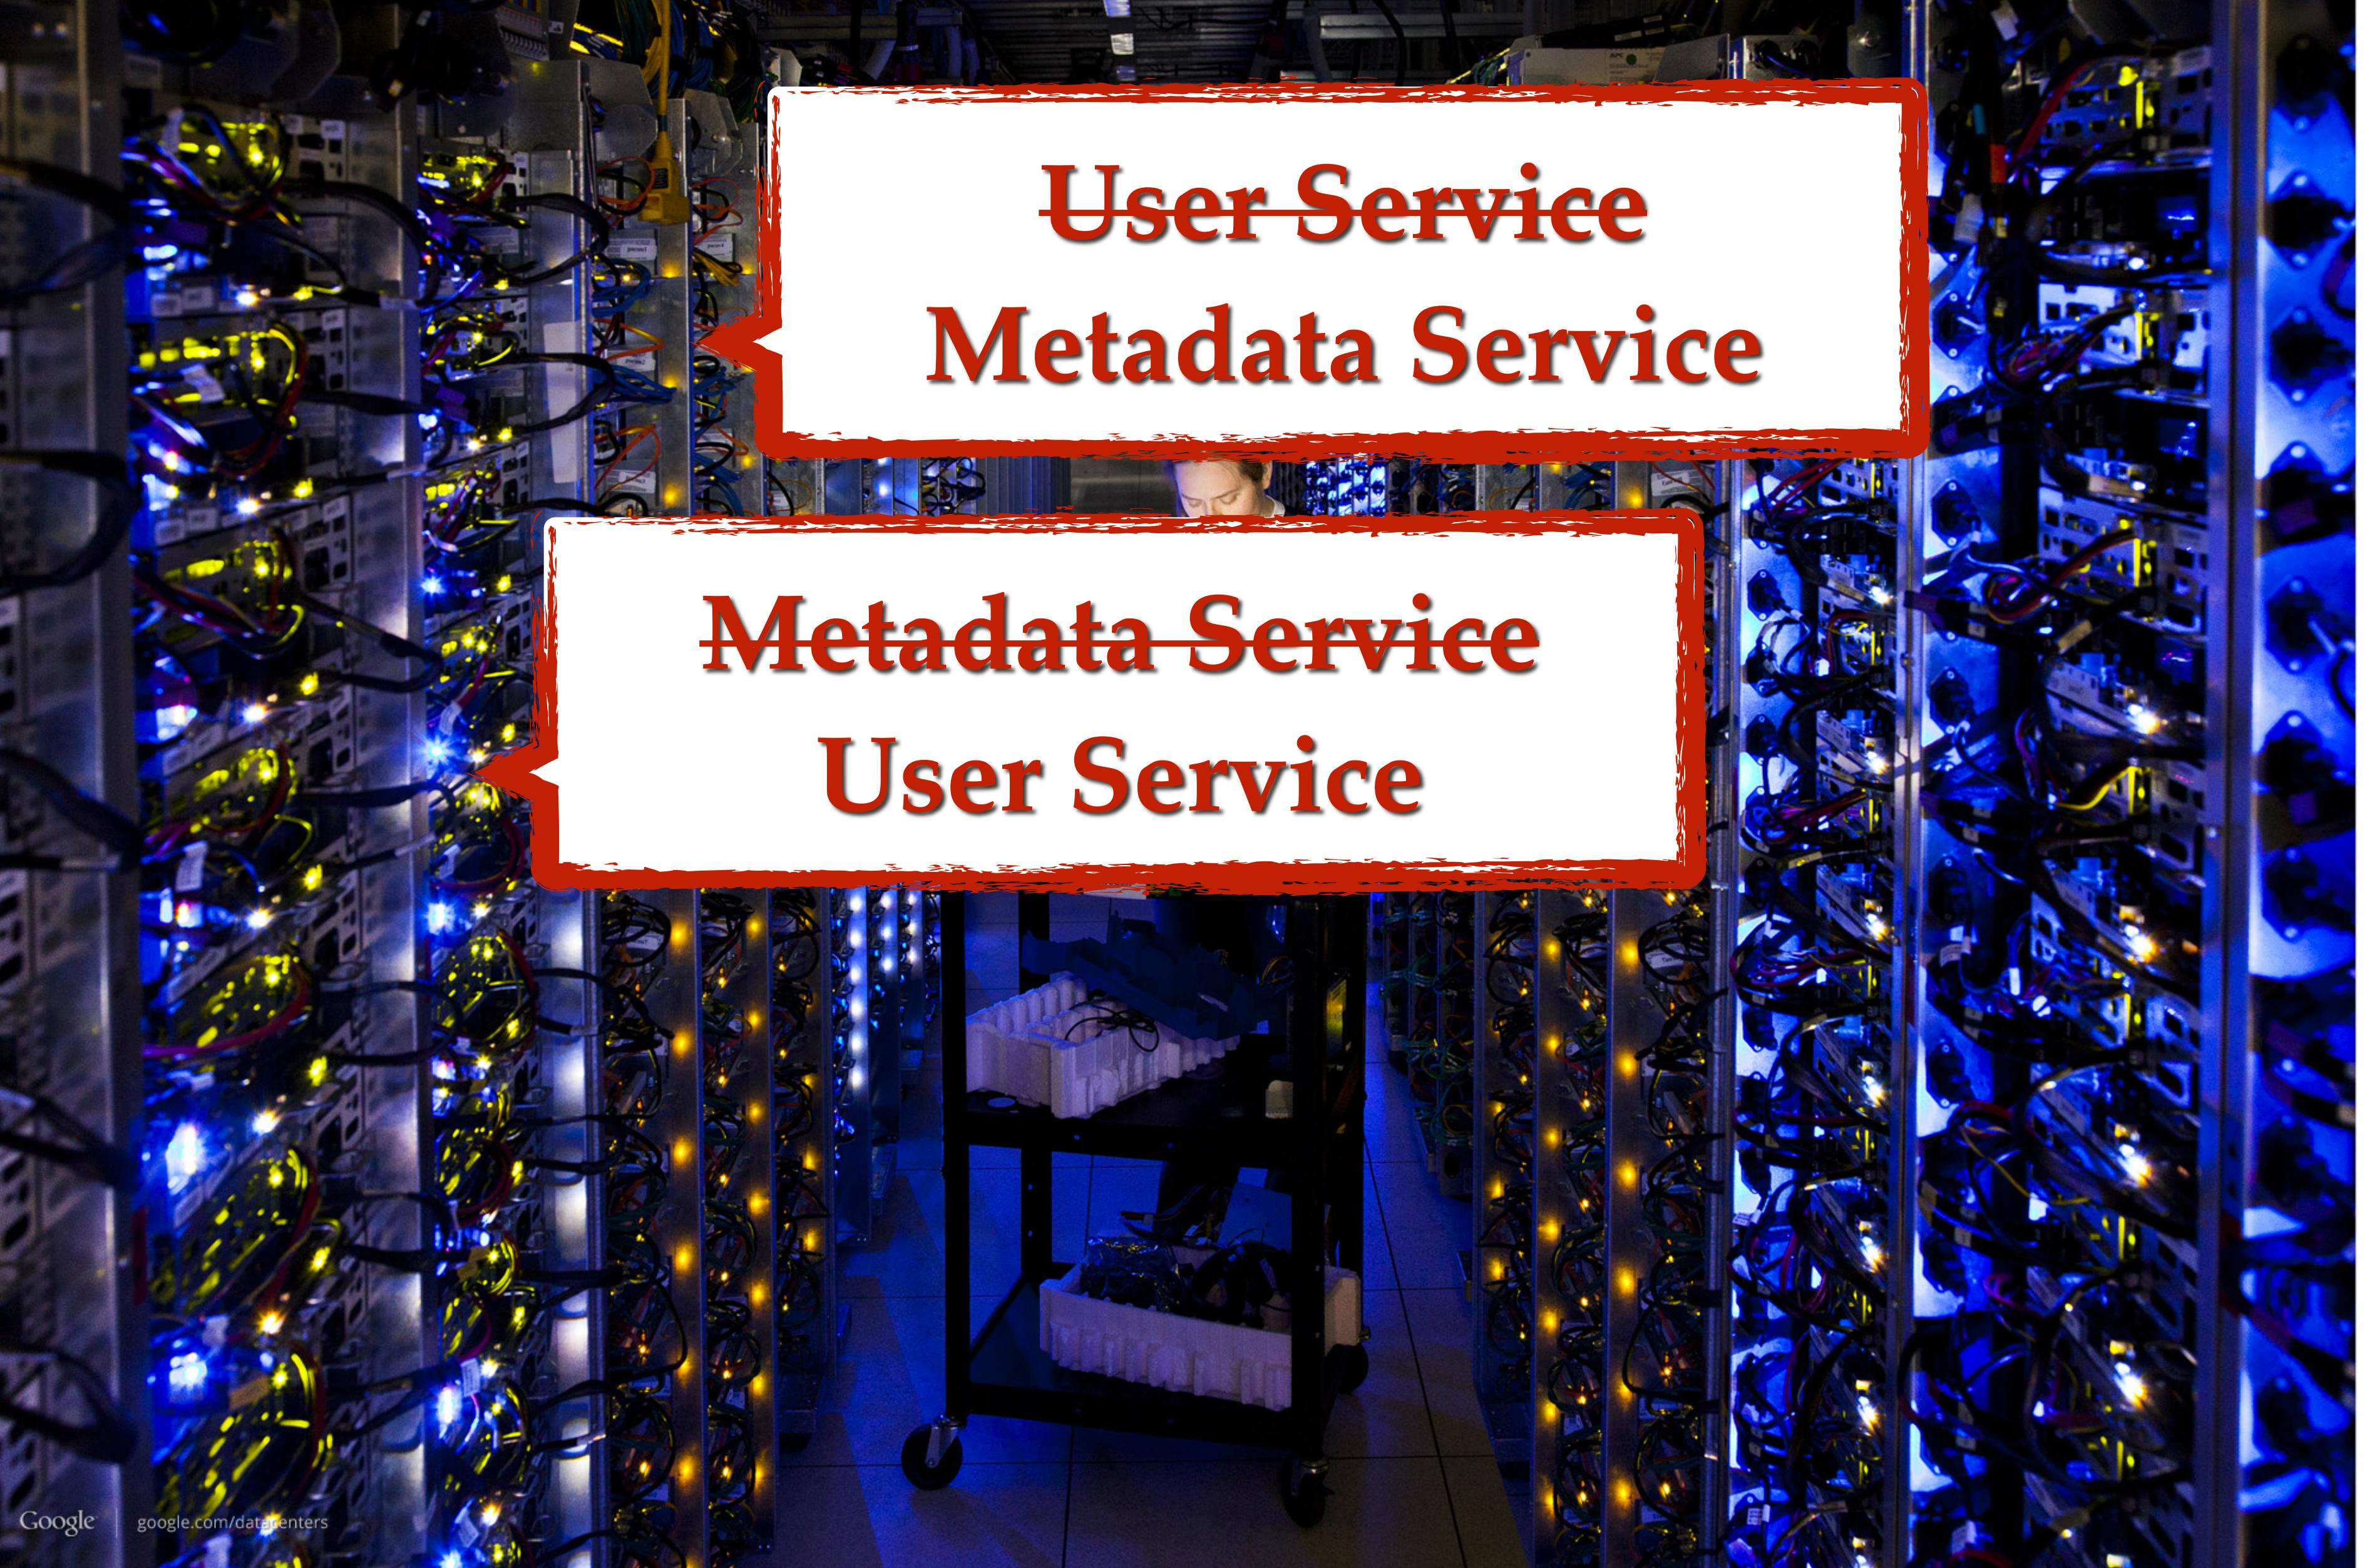
Service discovery controls visibility of nodes but does not guarantee availability.

All you lose is the visibility of new nodes.

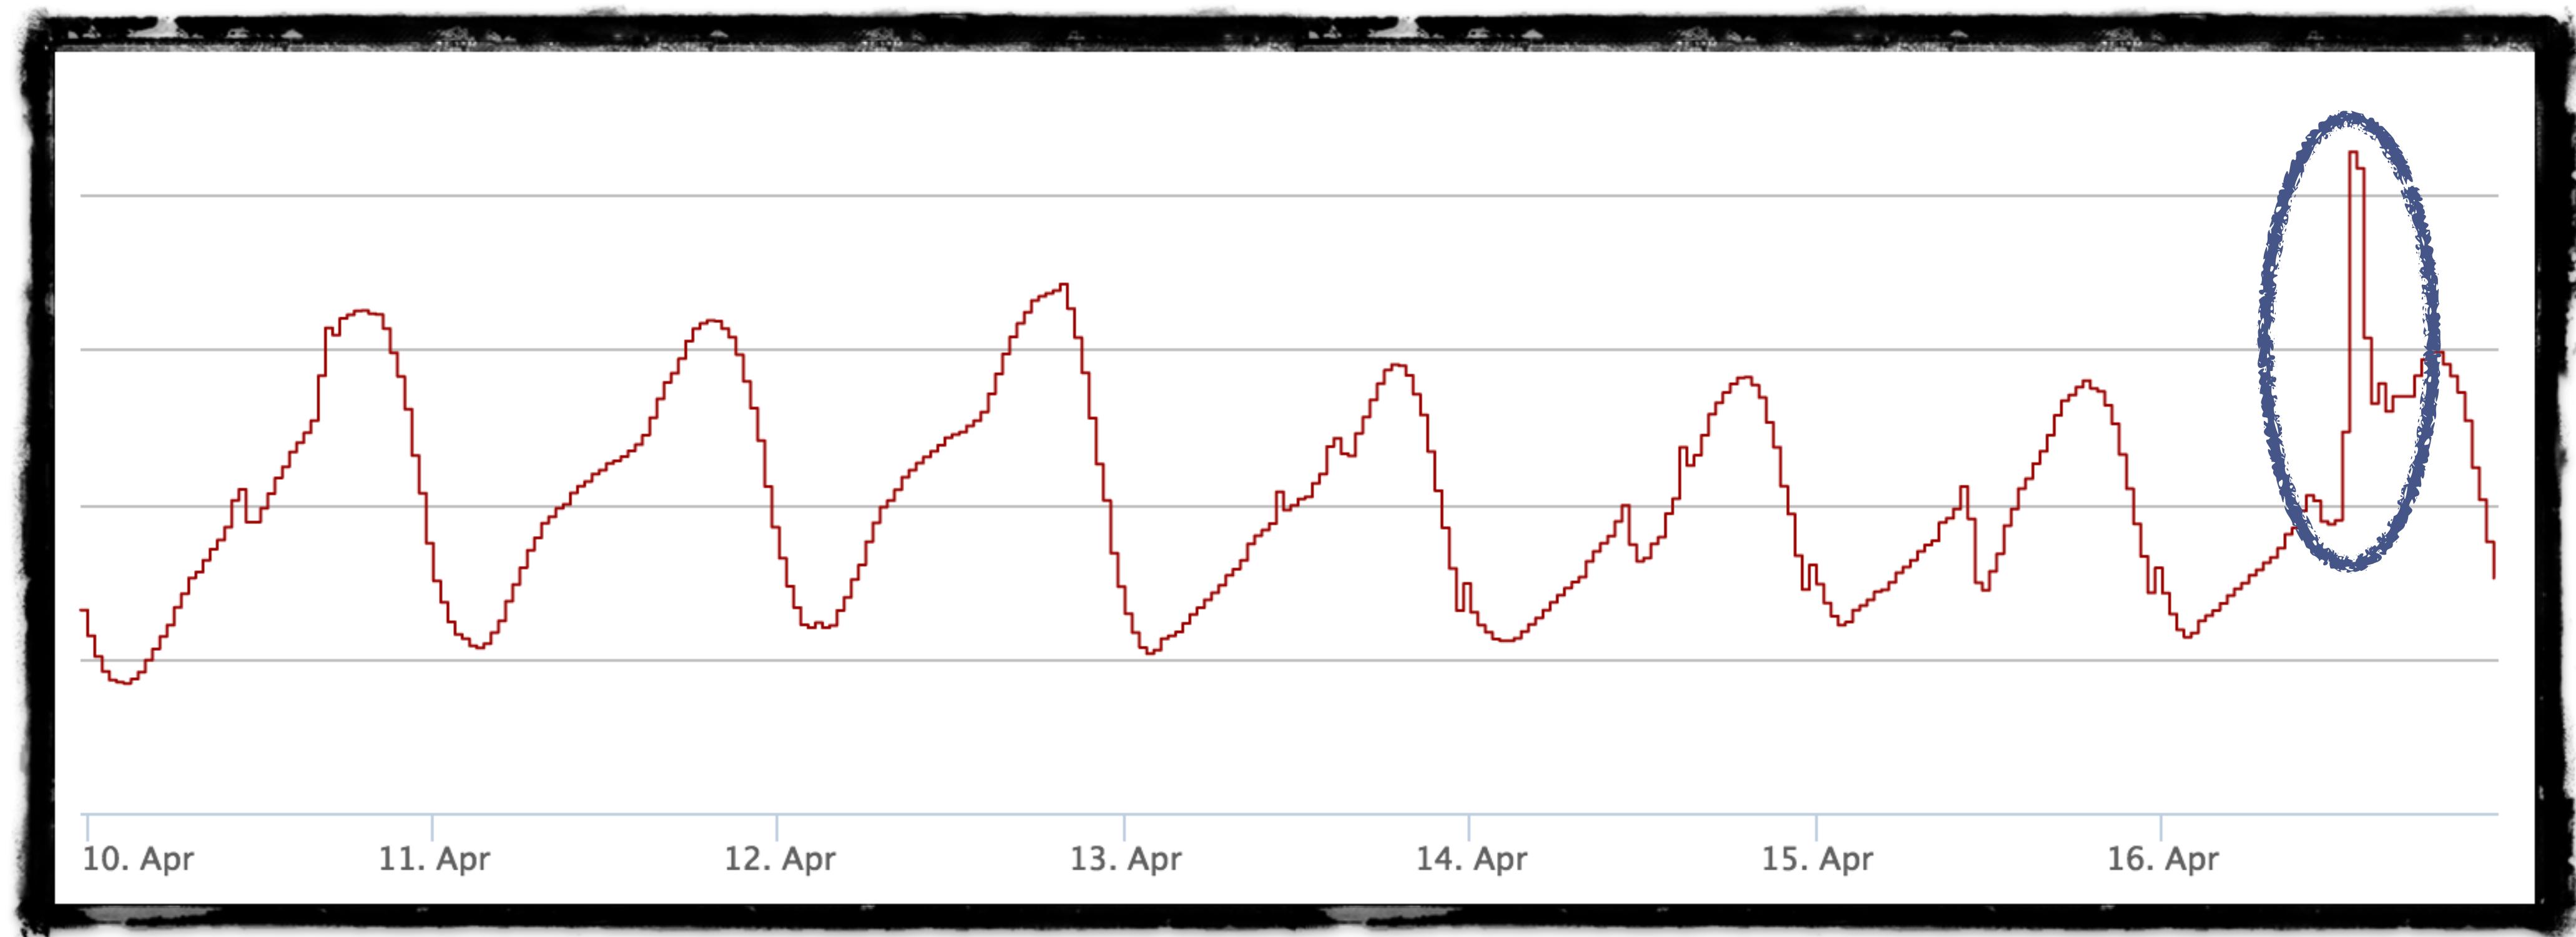




User Service
Metadata Service



Metadata Service
User Service



Netflix 'down': Users complain video streaming site is 'not working'



HE HE HE



High availability is a strong
requirement.

CAP theorem



Visibility

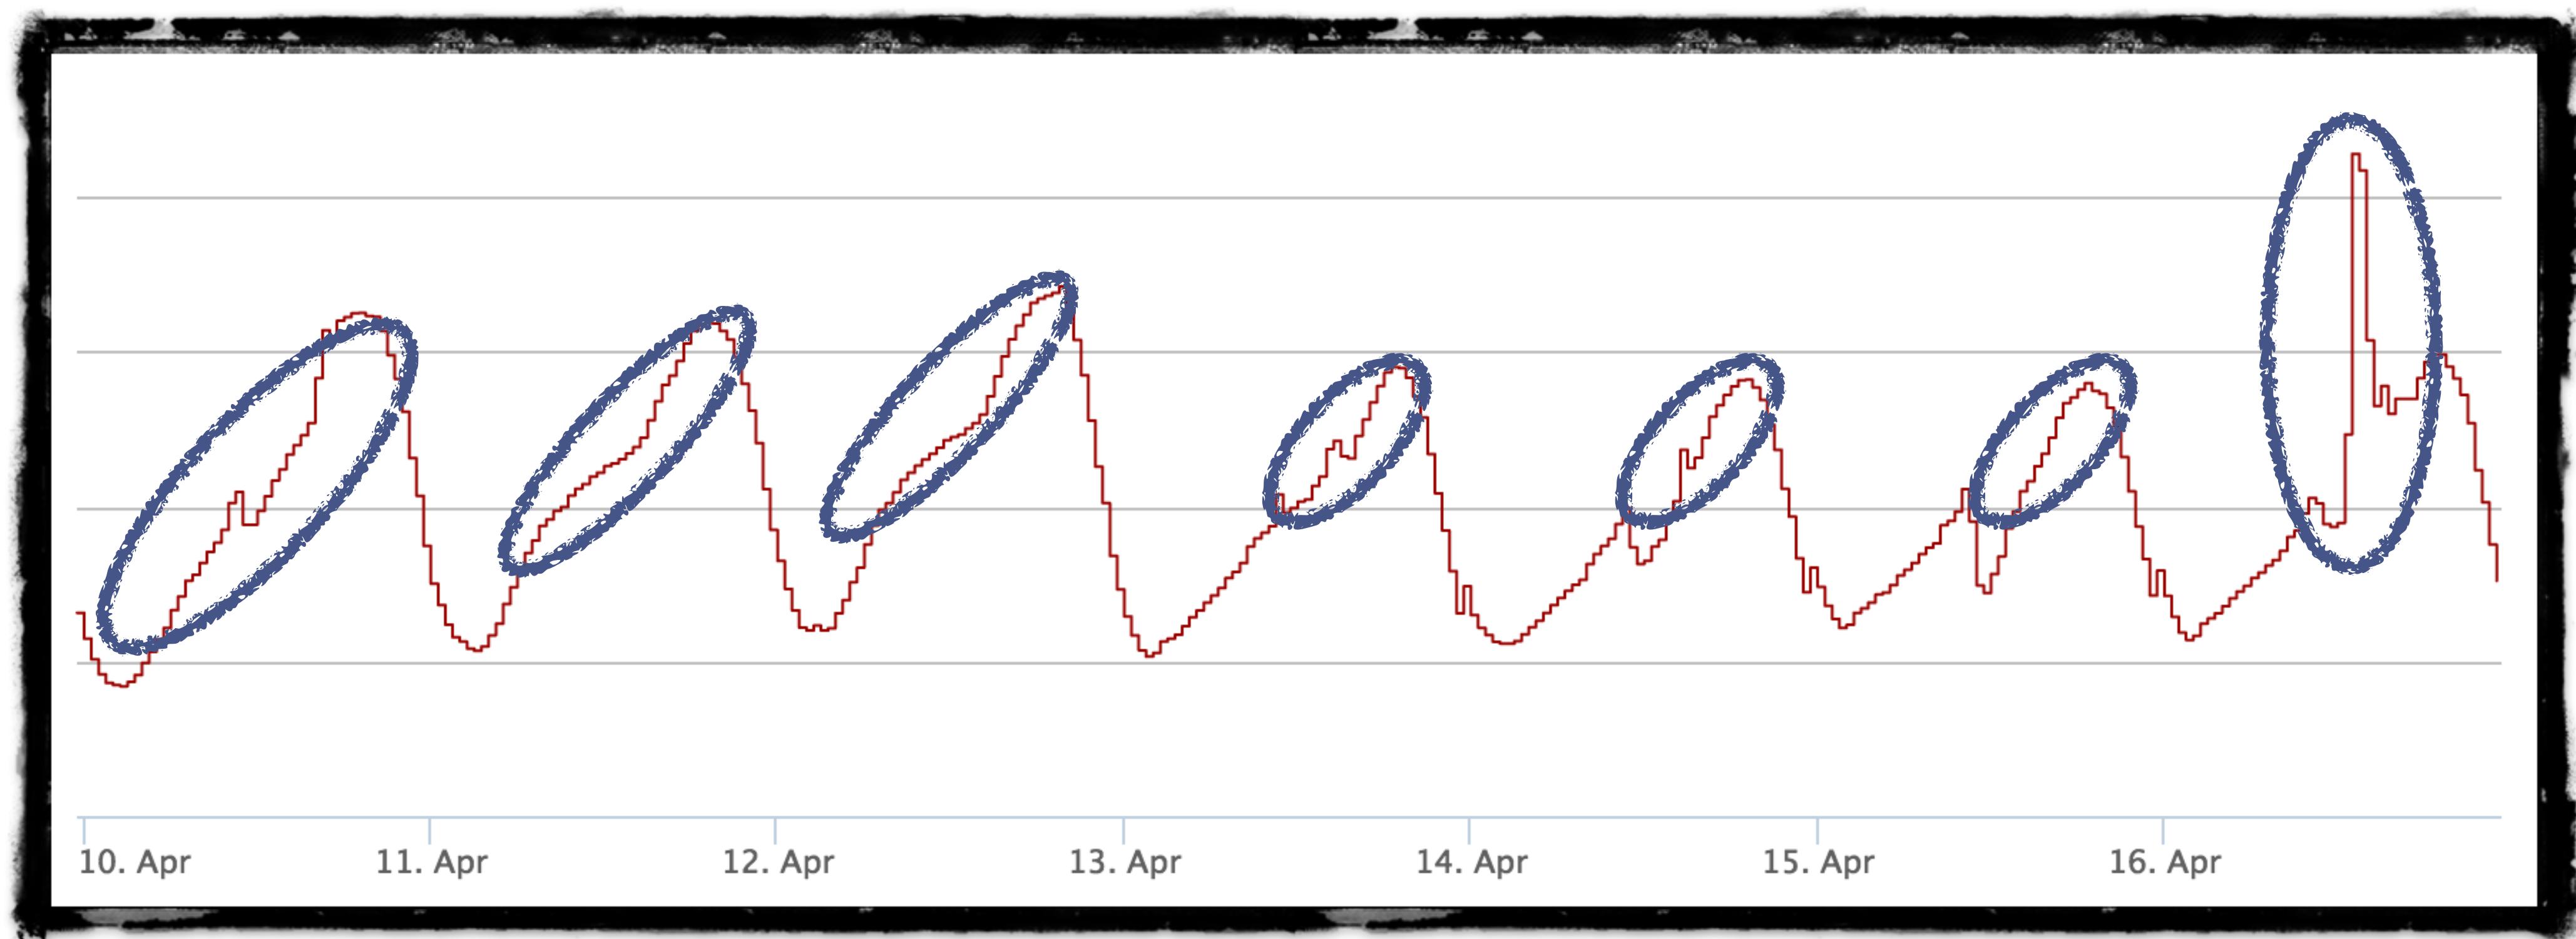
Consistency

Availability

Visibility

Consistency

A



Worst time

Visibility

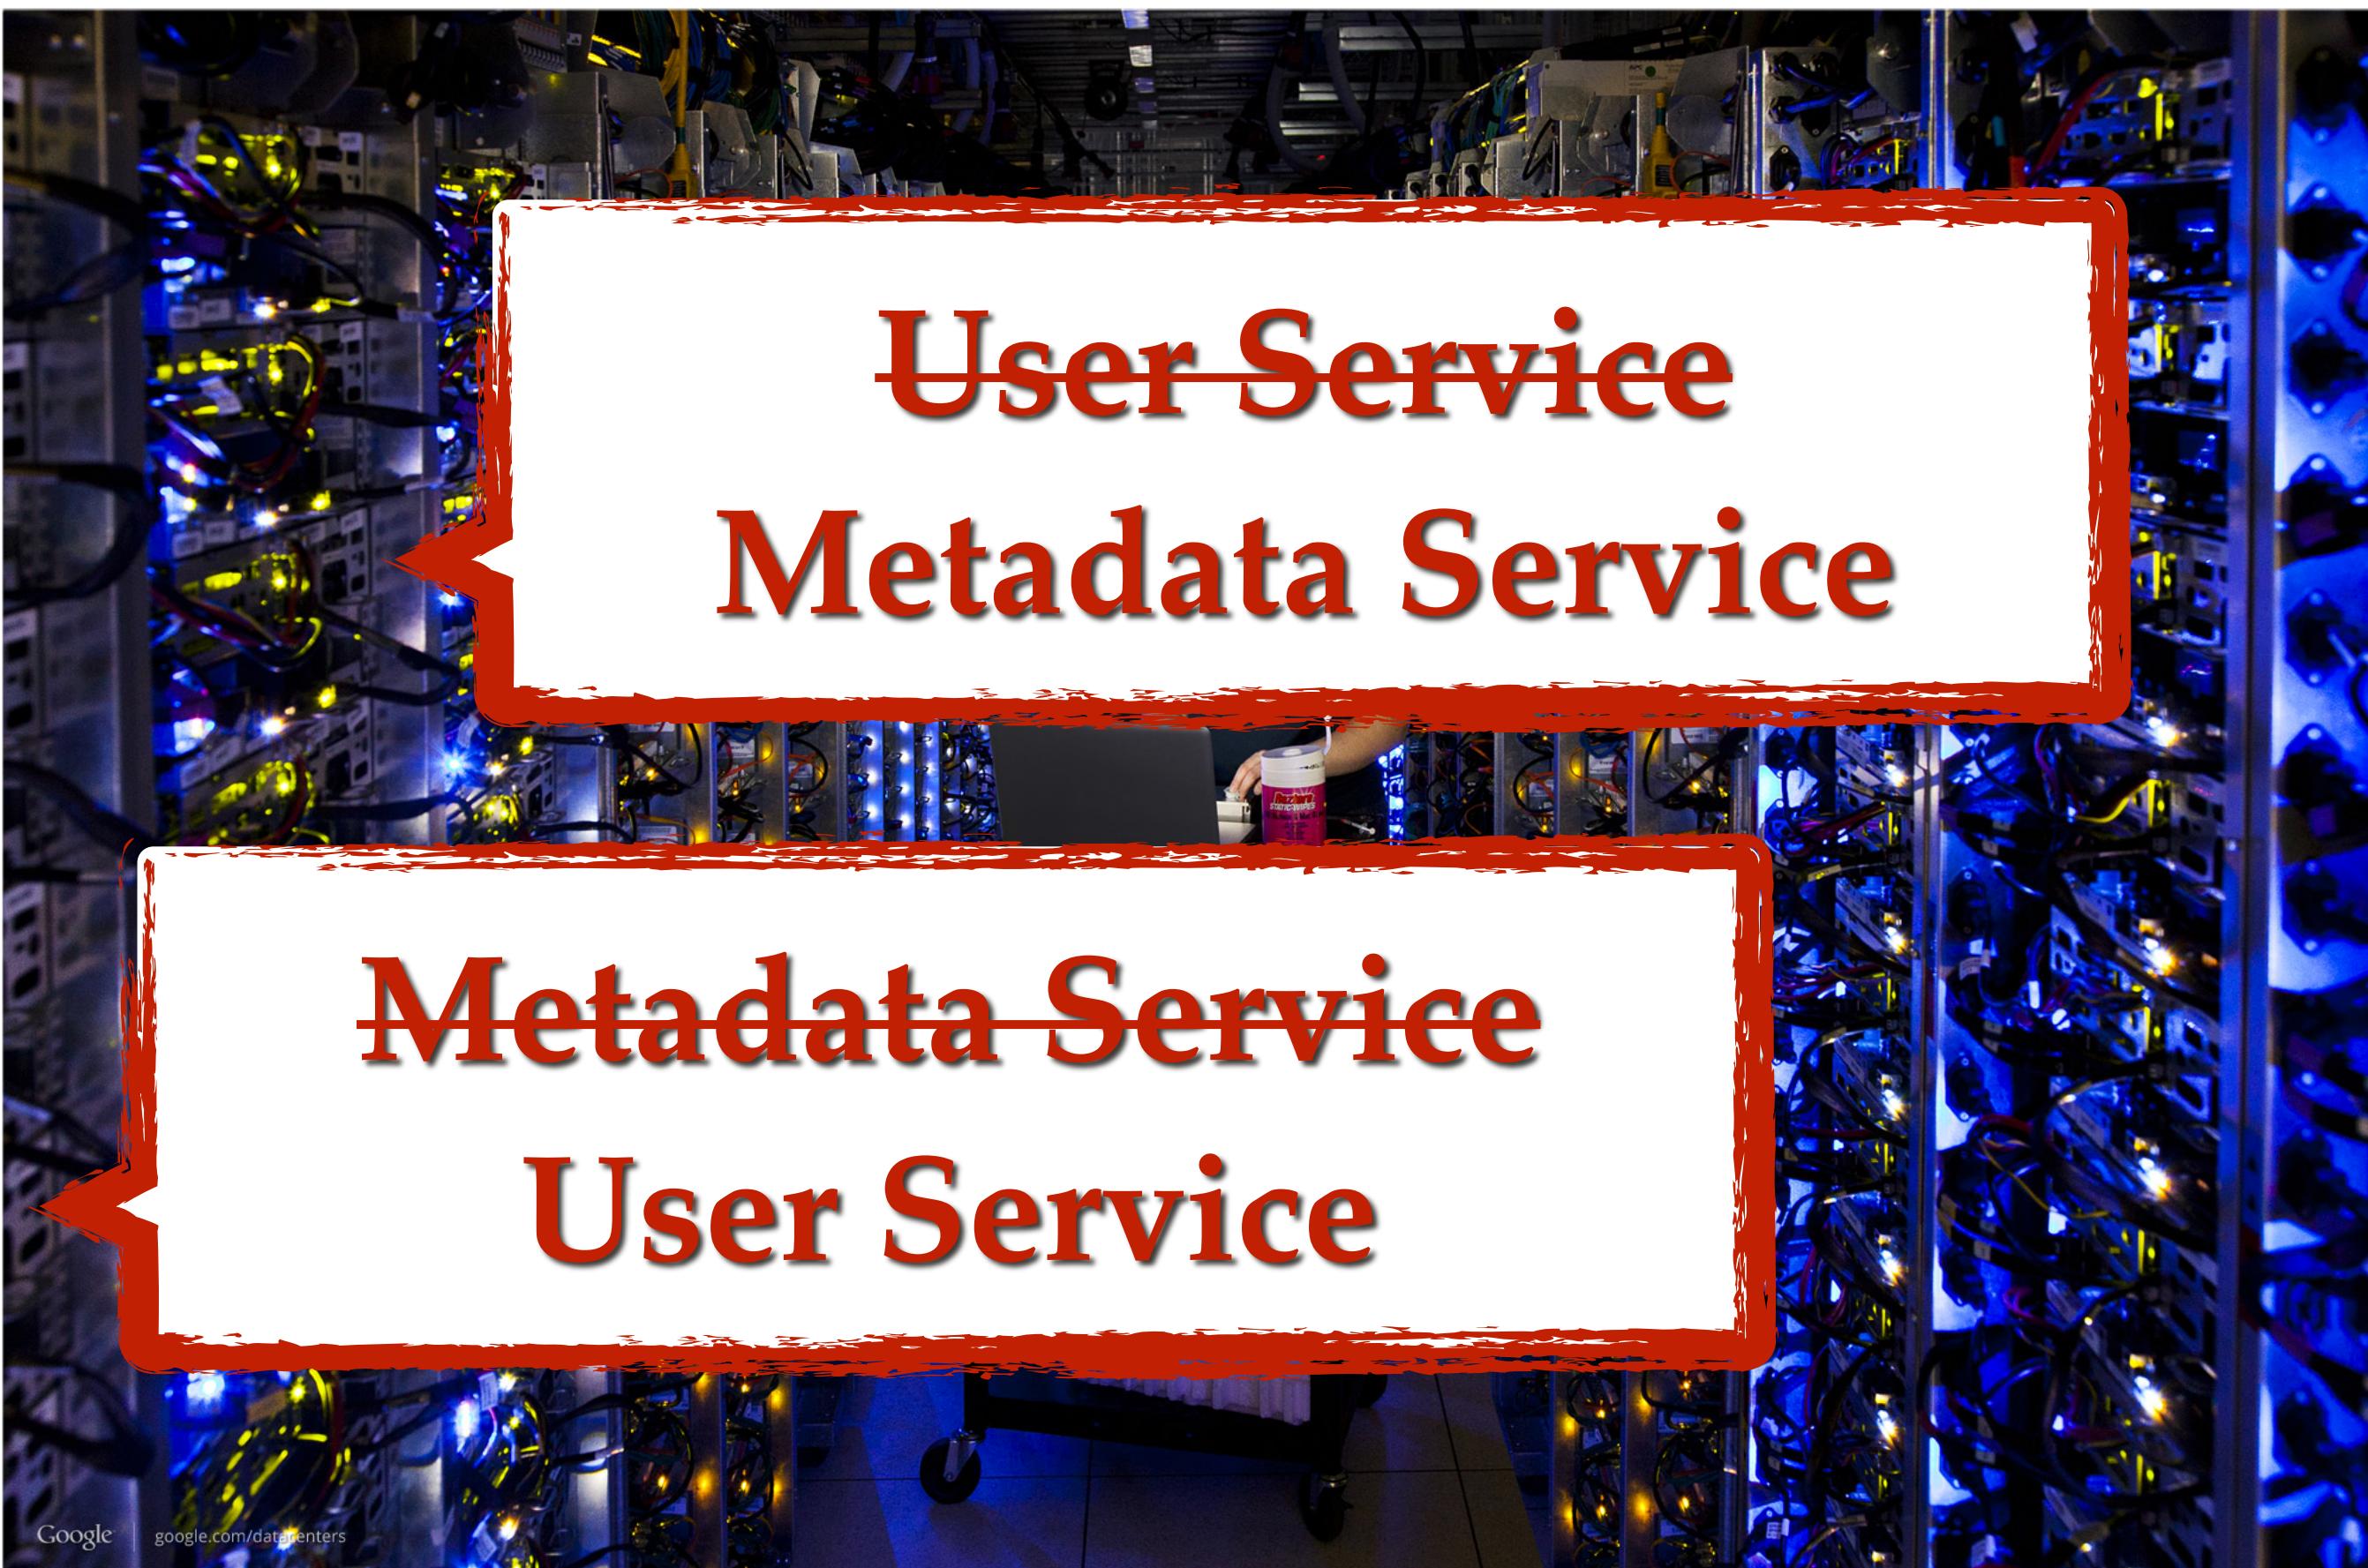
Consistency

Availability

Visibility

C

A**vailability**



Worst case

A persistent, verified connection
based interaction solves this issue.

Edge Service

Connect

Recommendations Service

Edge Service

Connect

Verify

Are you a Recs service instance?

Recommendations Service

Edge Service

Connect

Verify

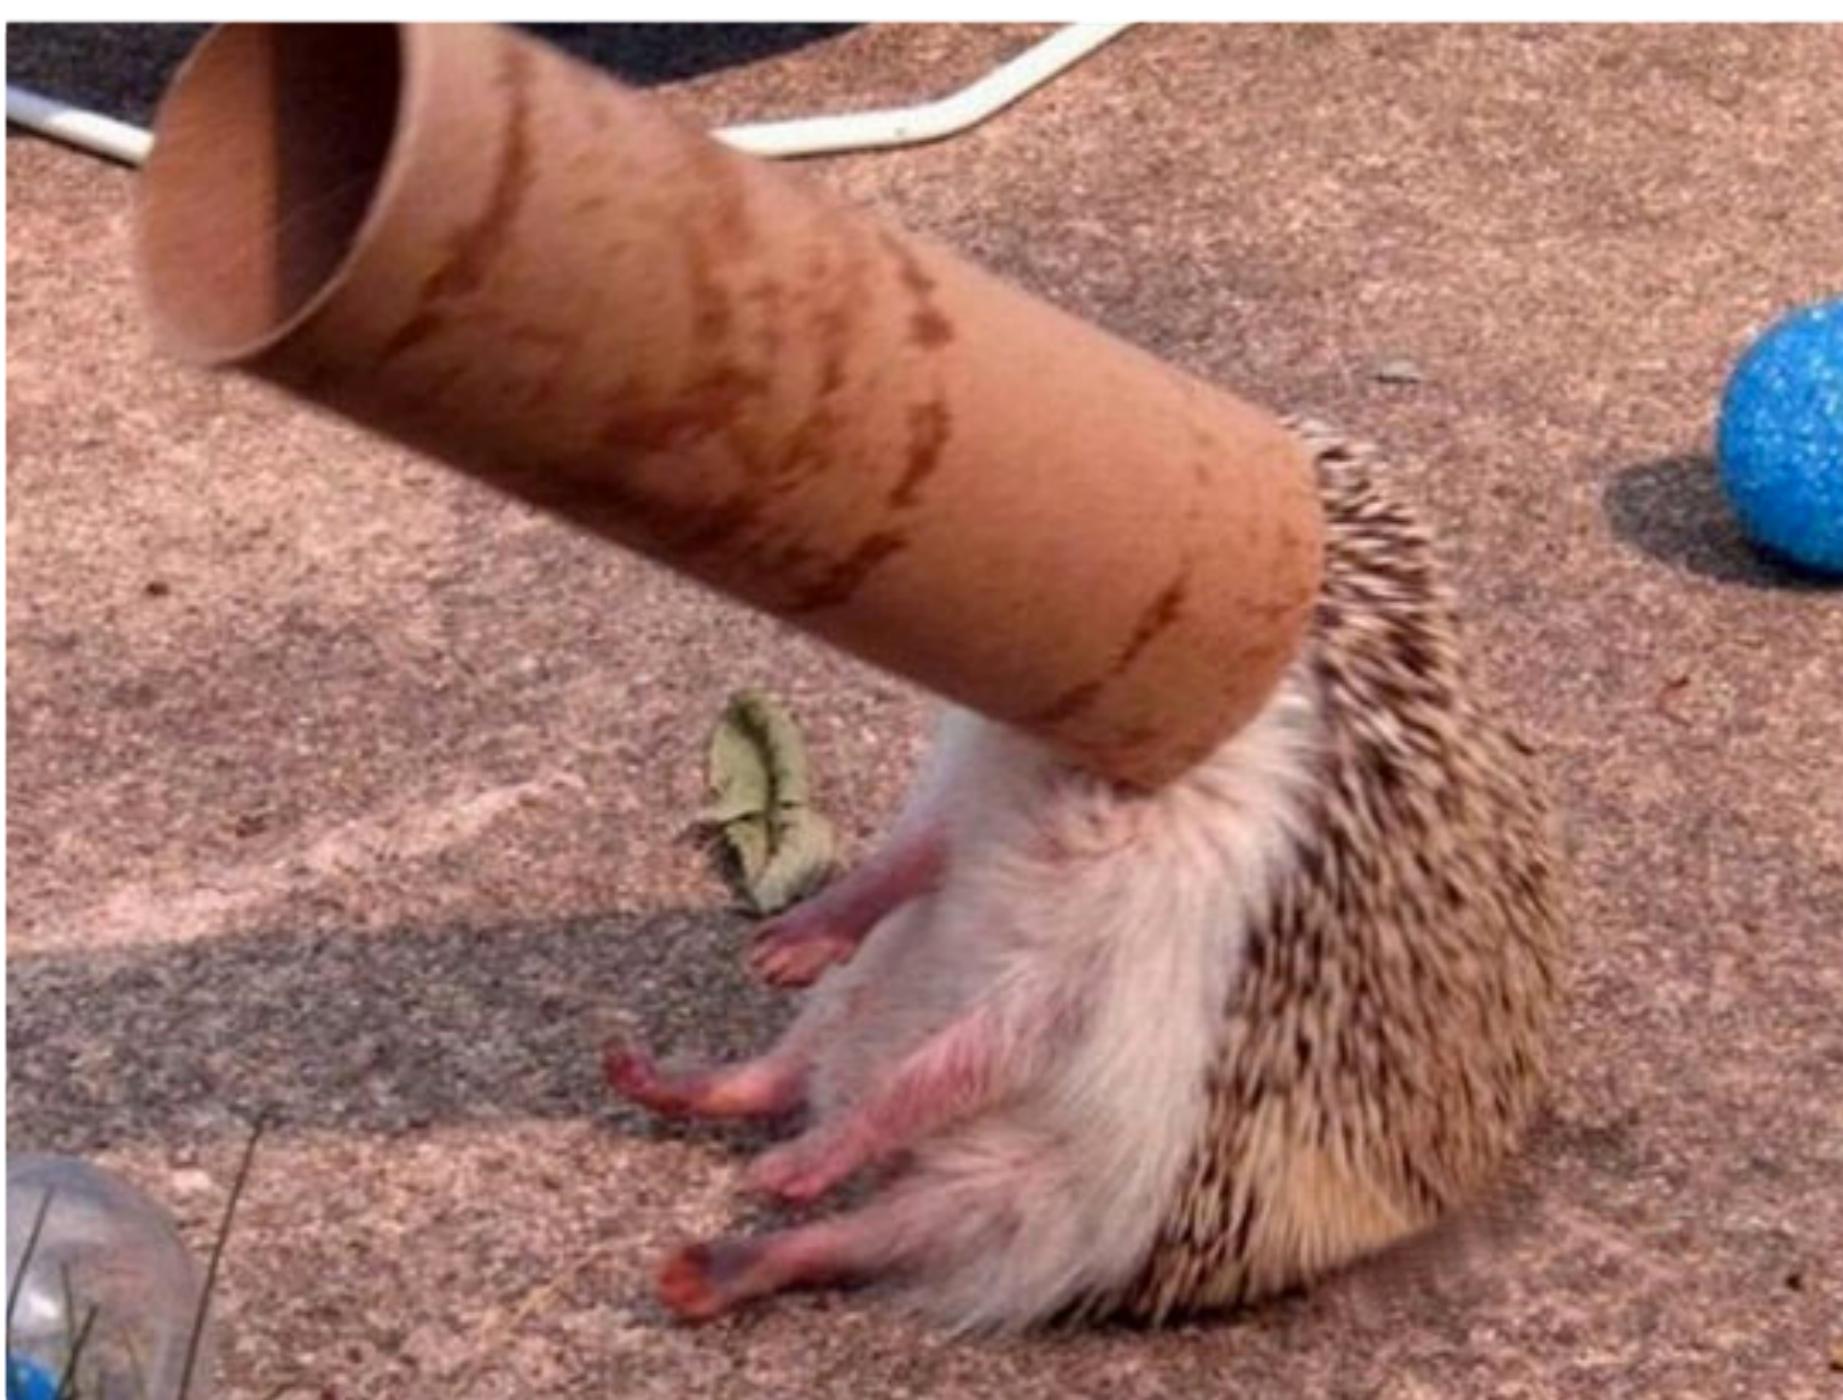
Are you a Recs service instance?

Requests

-
-
-

Recommendations Service

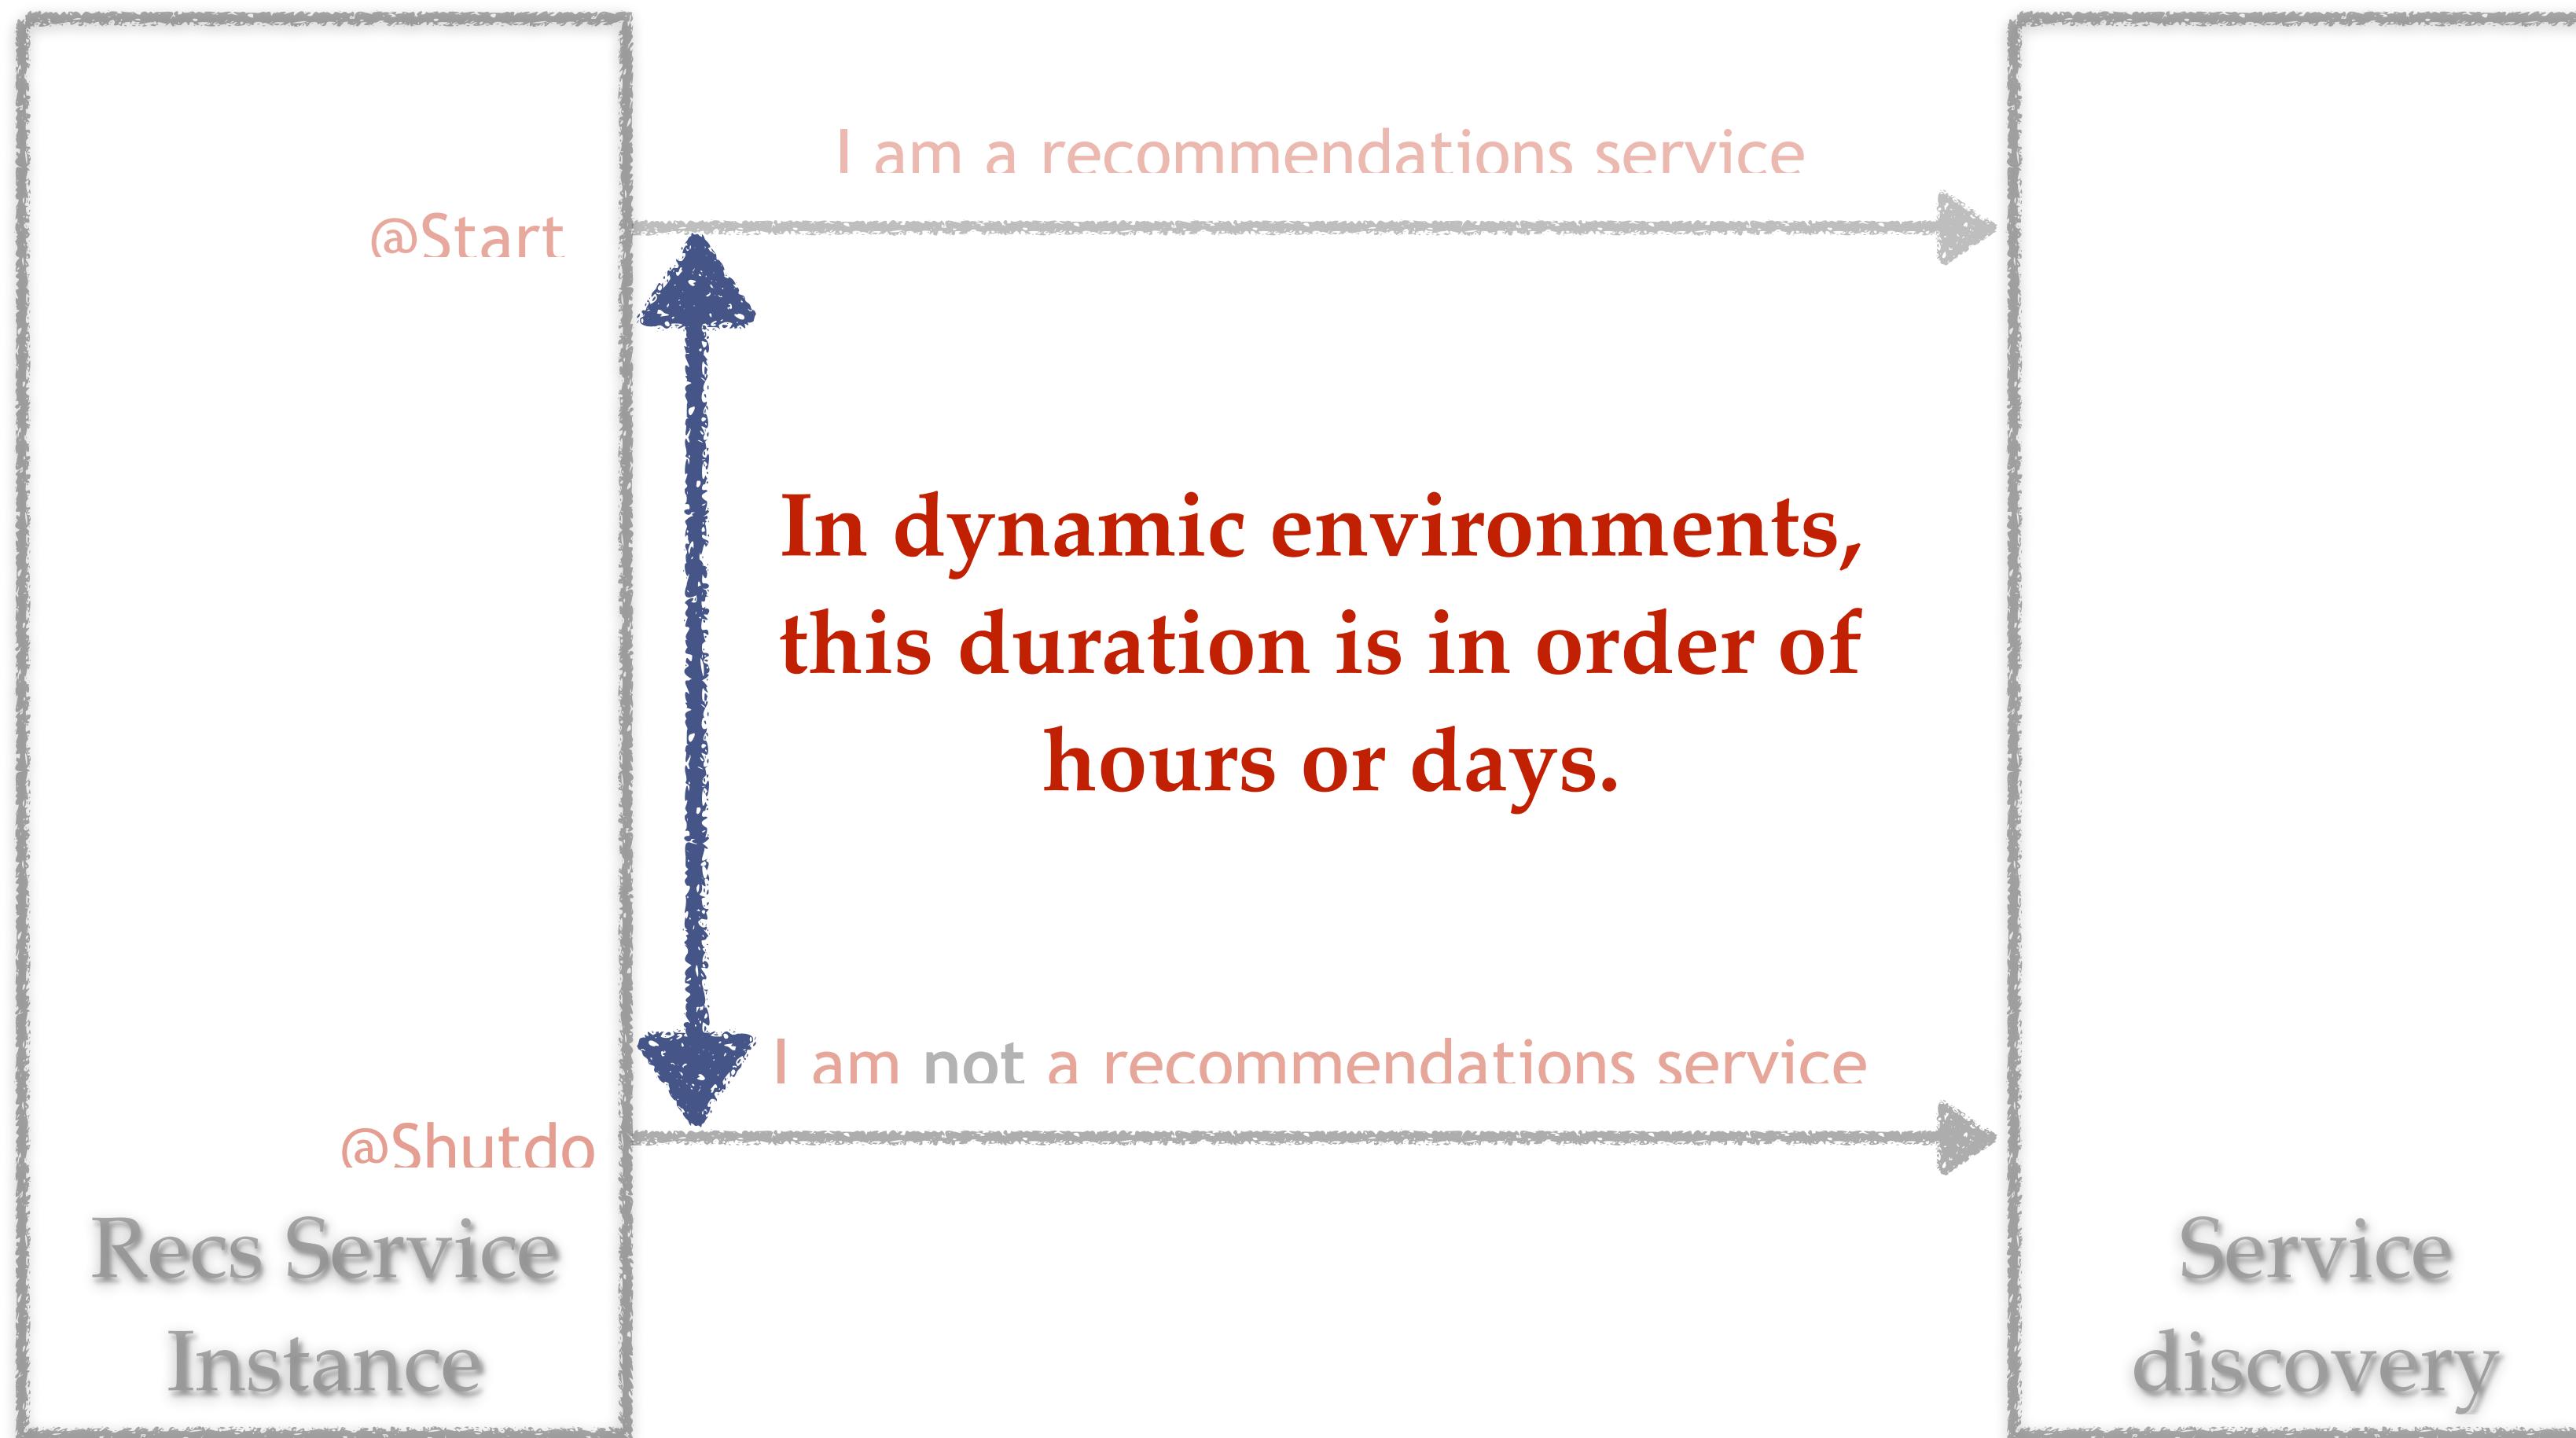
In most cases choosing availability over consistency for Service Discovery is the correct thing to do.



Durability

**Service discovery data is
ephemeral and re-generatable.**

Ephemeral



Re-generatable data

Service Discovery data is the **instance information**.
It is always available to the instance.

node-categorization
re-generatable

availability

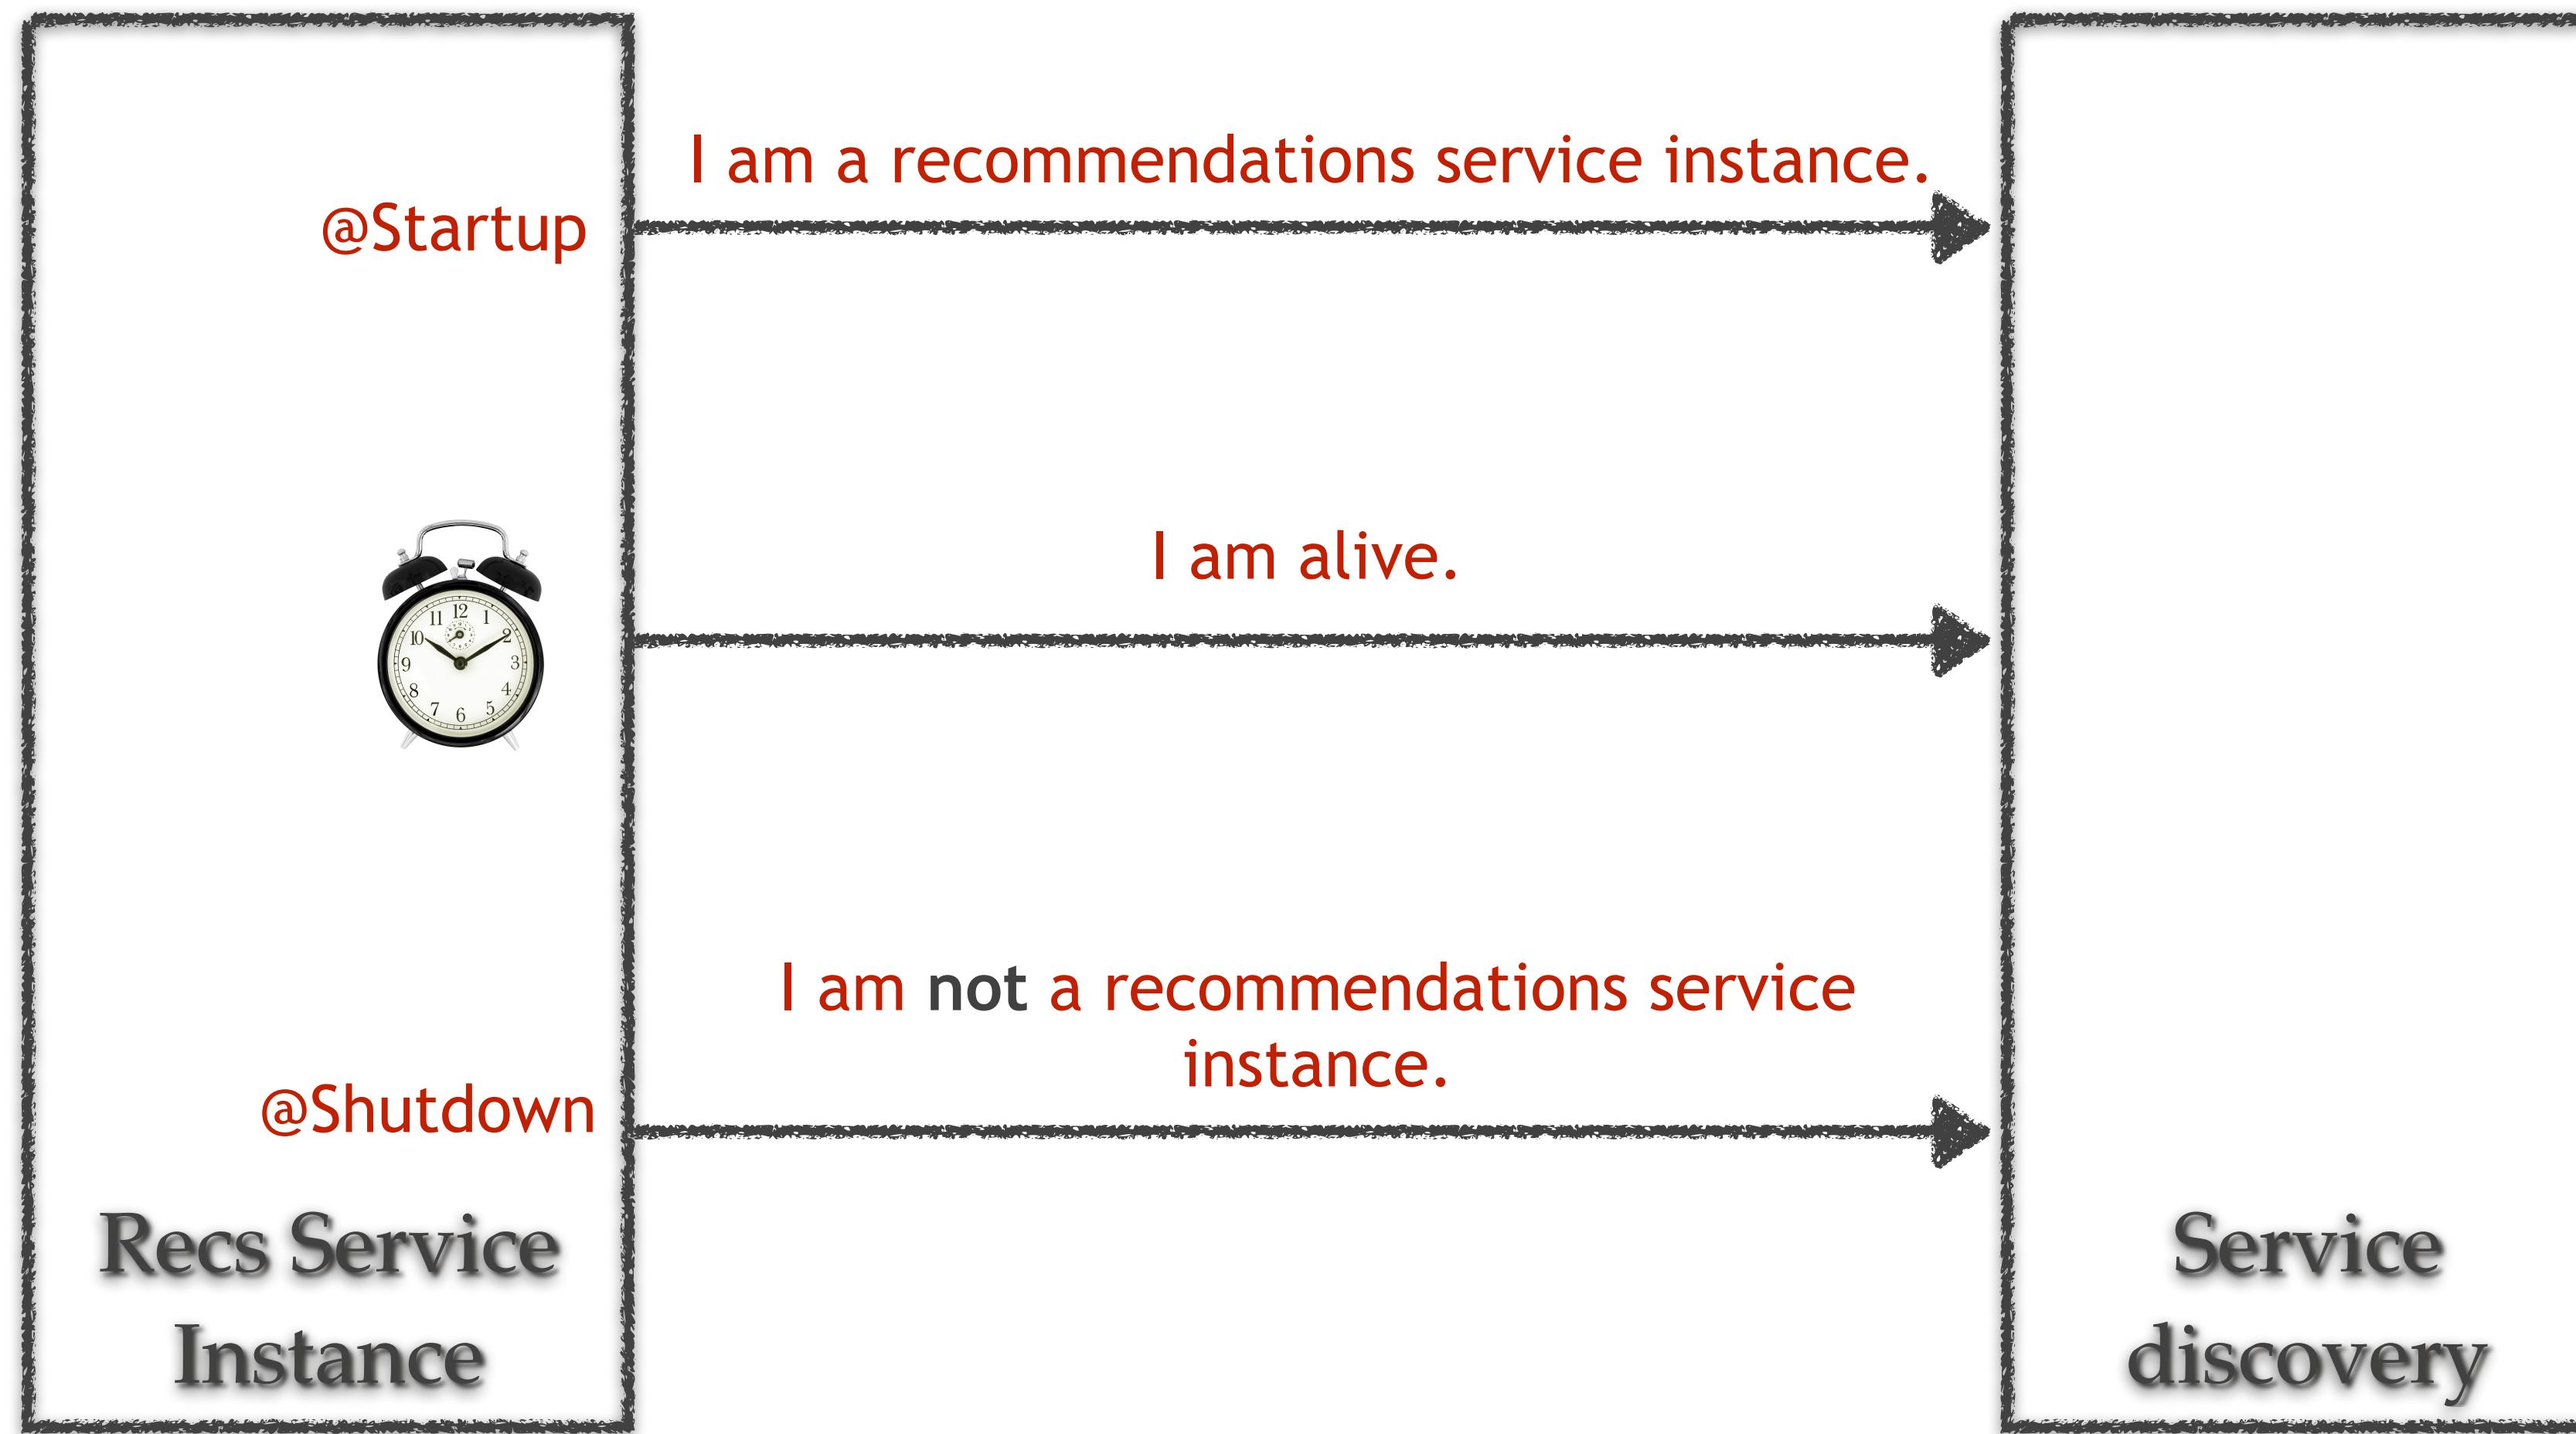
heartbeats ephemeral

node-availability-out-of-scope



Can we create a
co-ordination free service discovery?

Stateful Client-Server interaction



Stateful Client-Server interaction

Needs a stateful protocol.



Nice Benefits

Causal ordering

Nice Benefits

Causal ordering

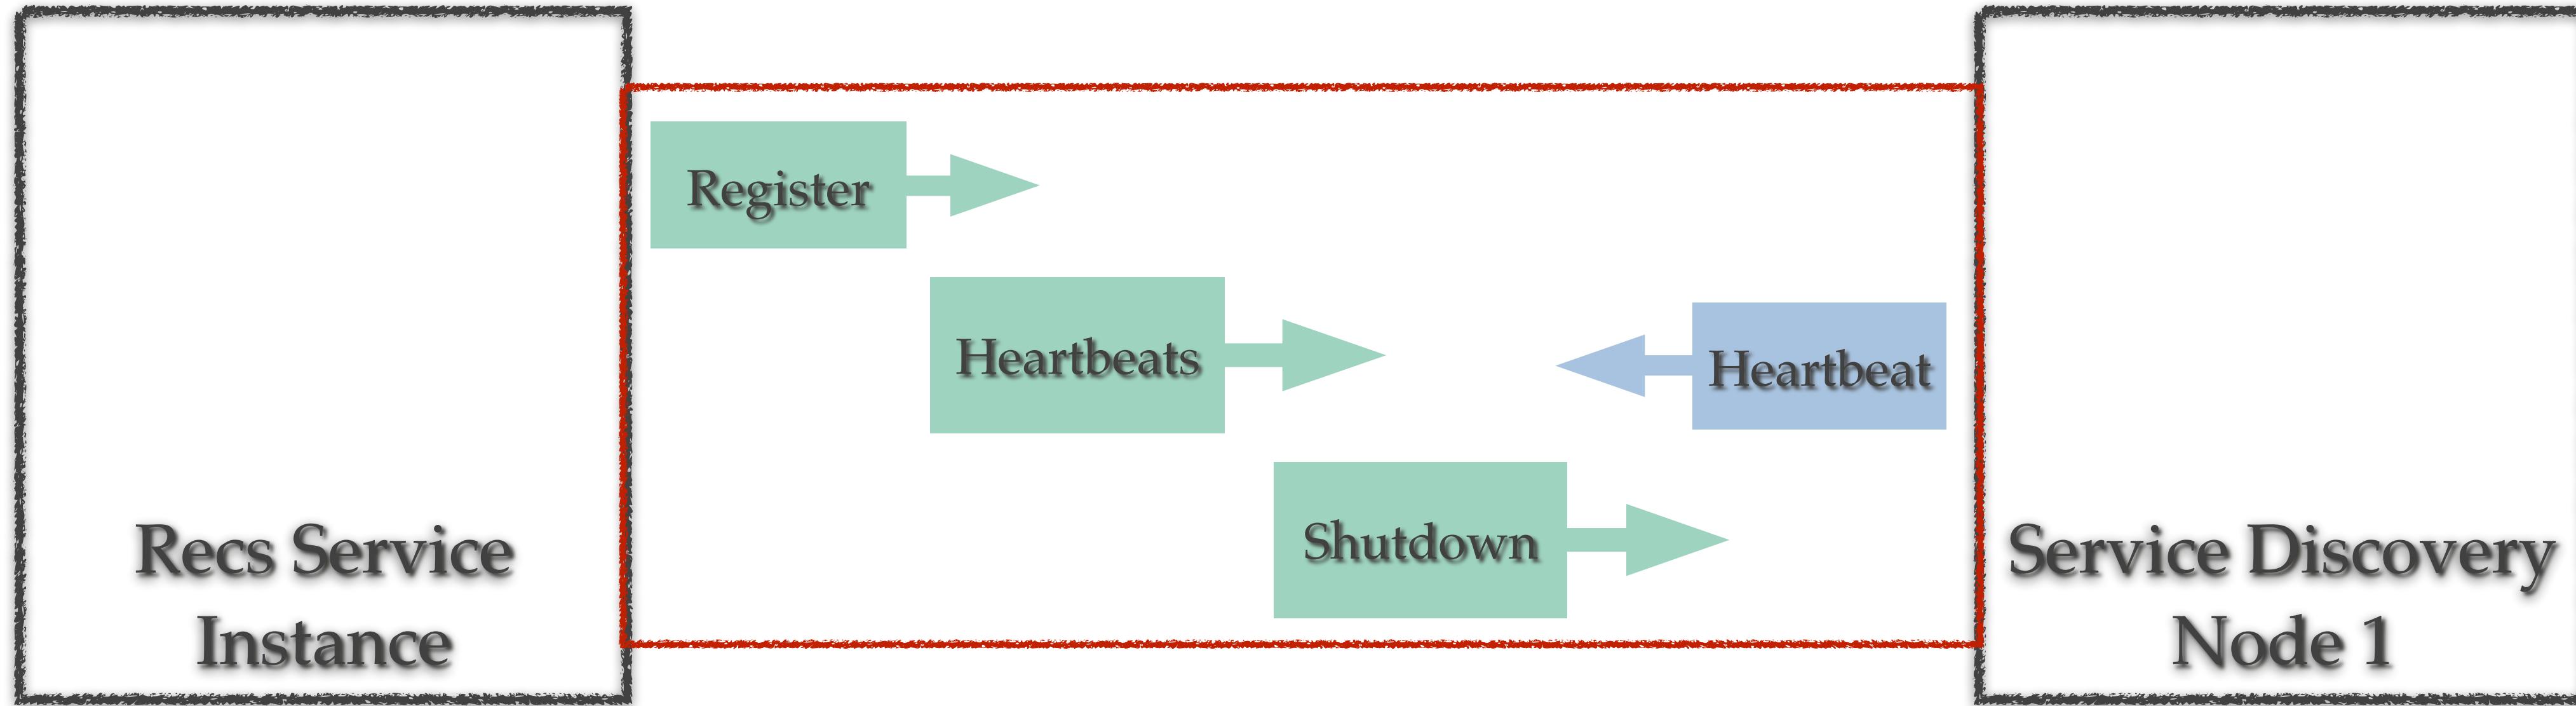
Natural lifecycle (connection)

Nice Benefits

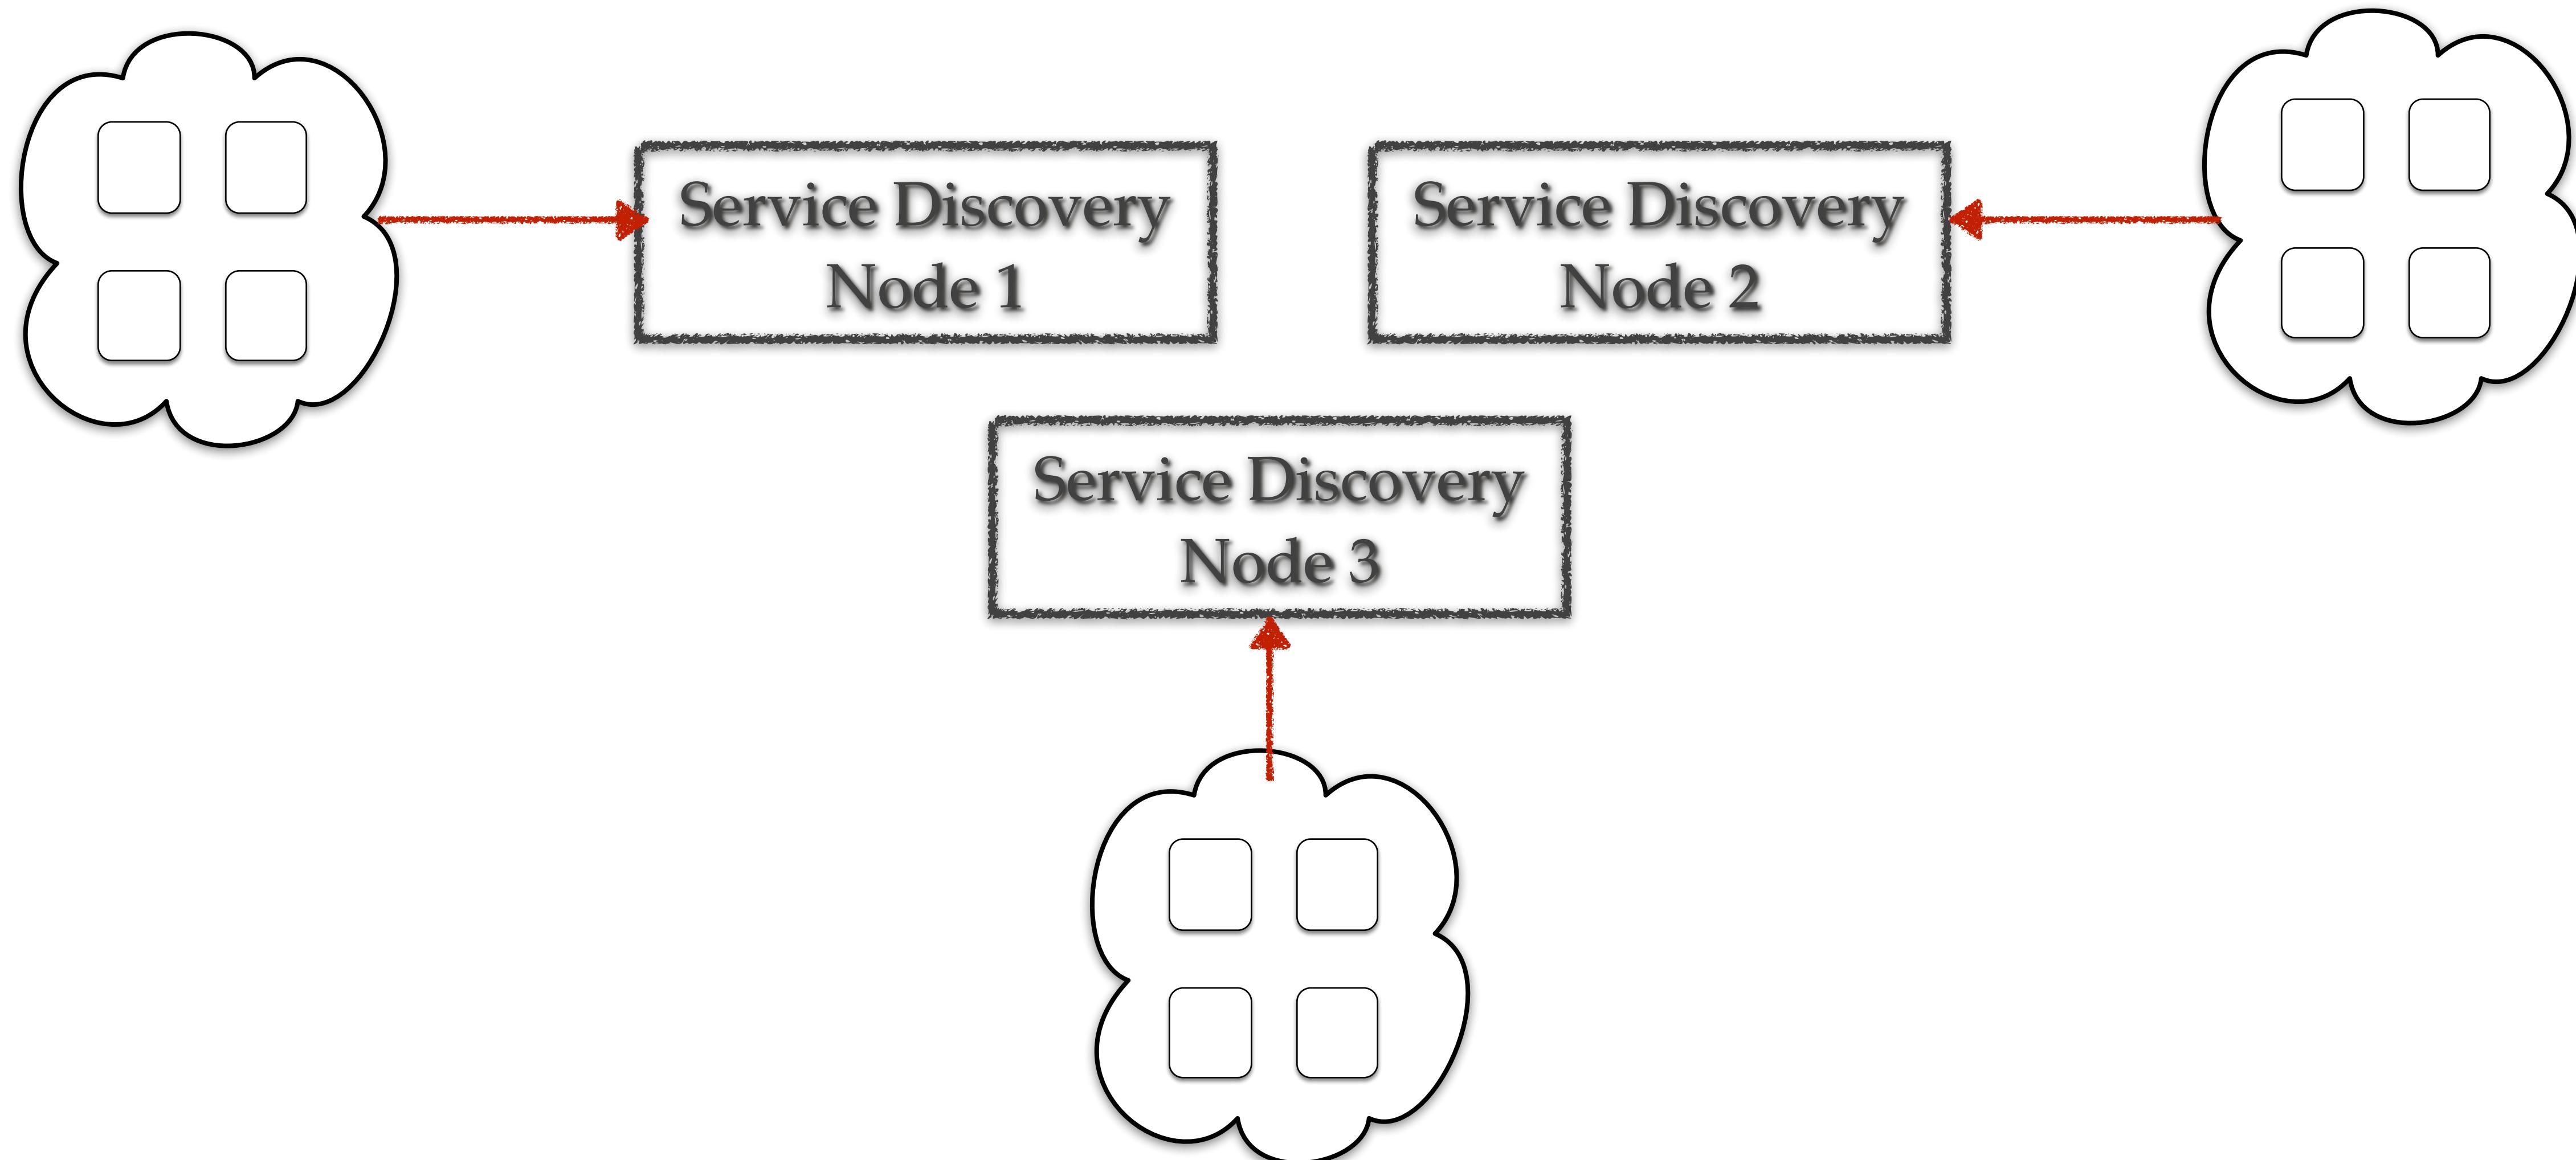
Causal ordering

Natural lifecycle (connection)

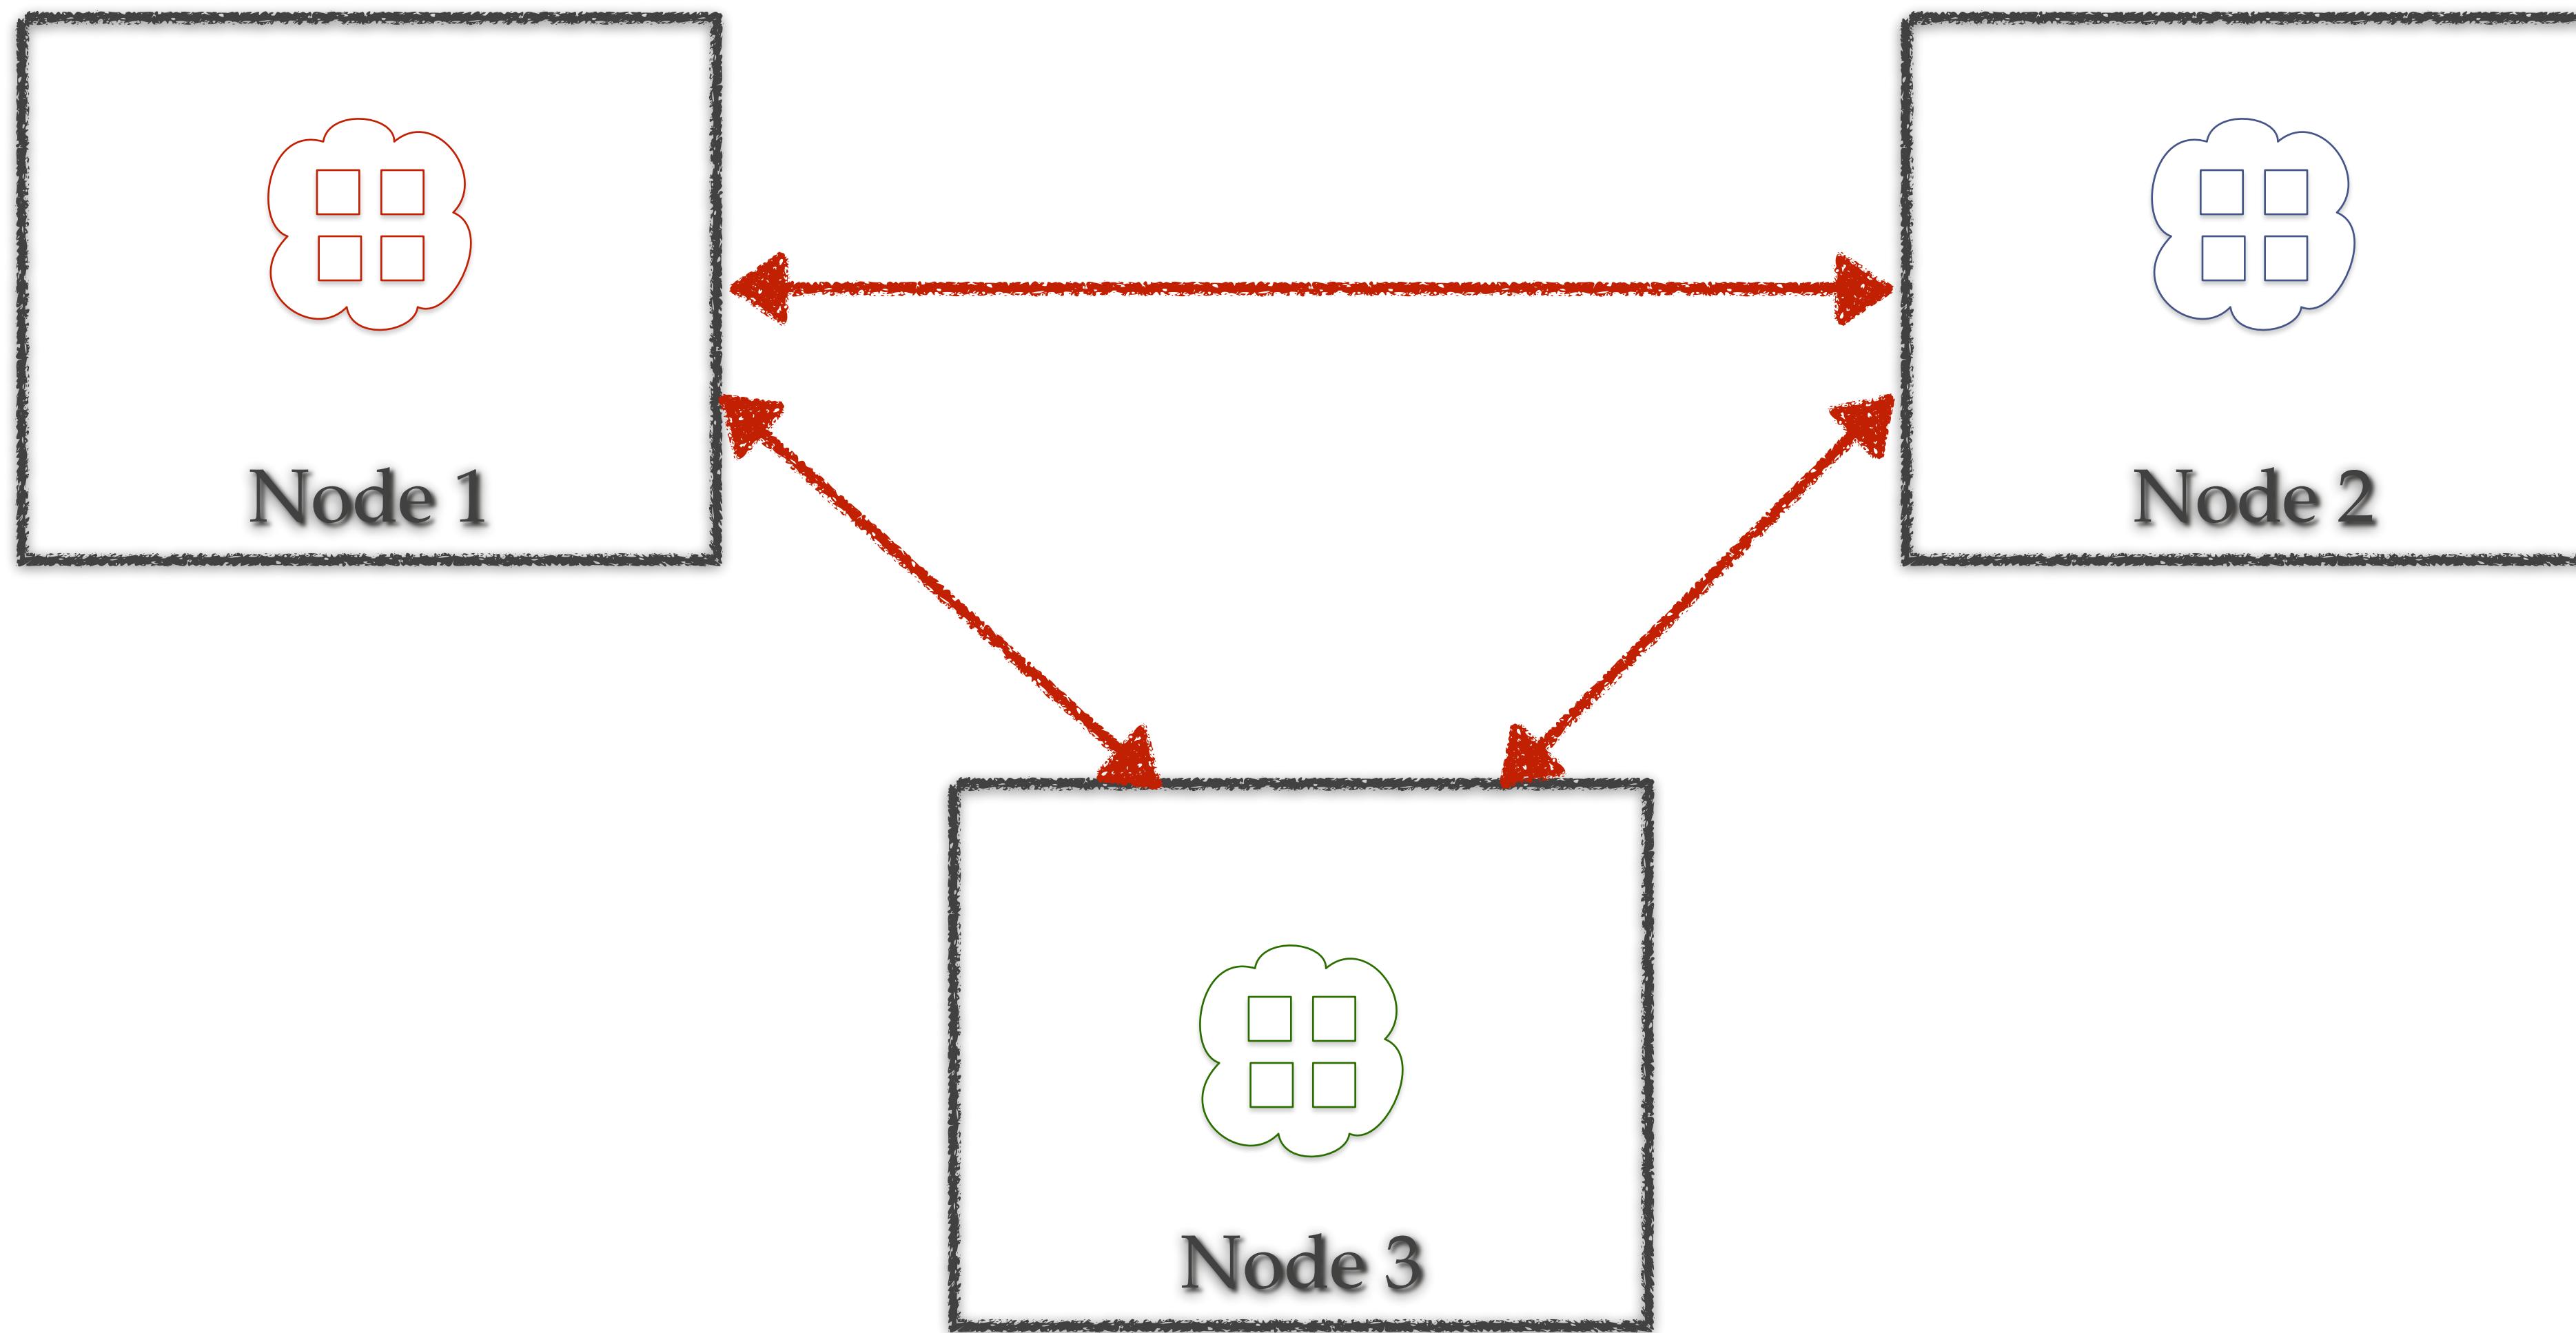
Ability to send data “diff”.



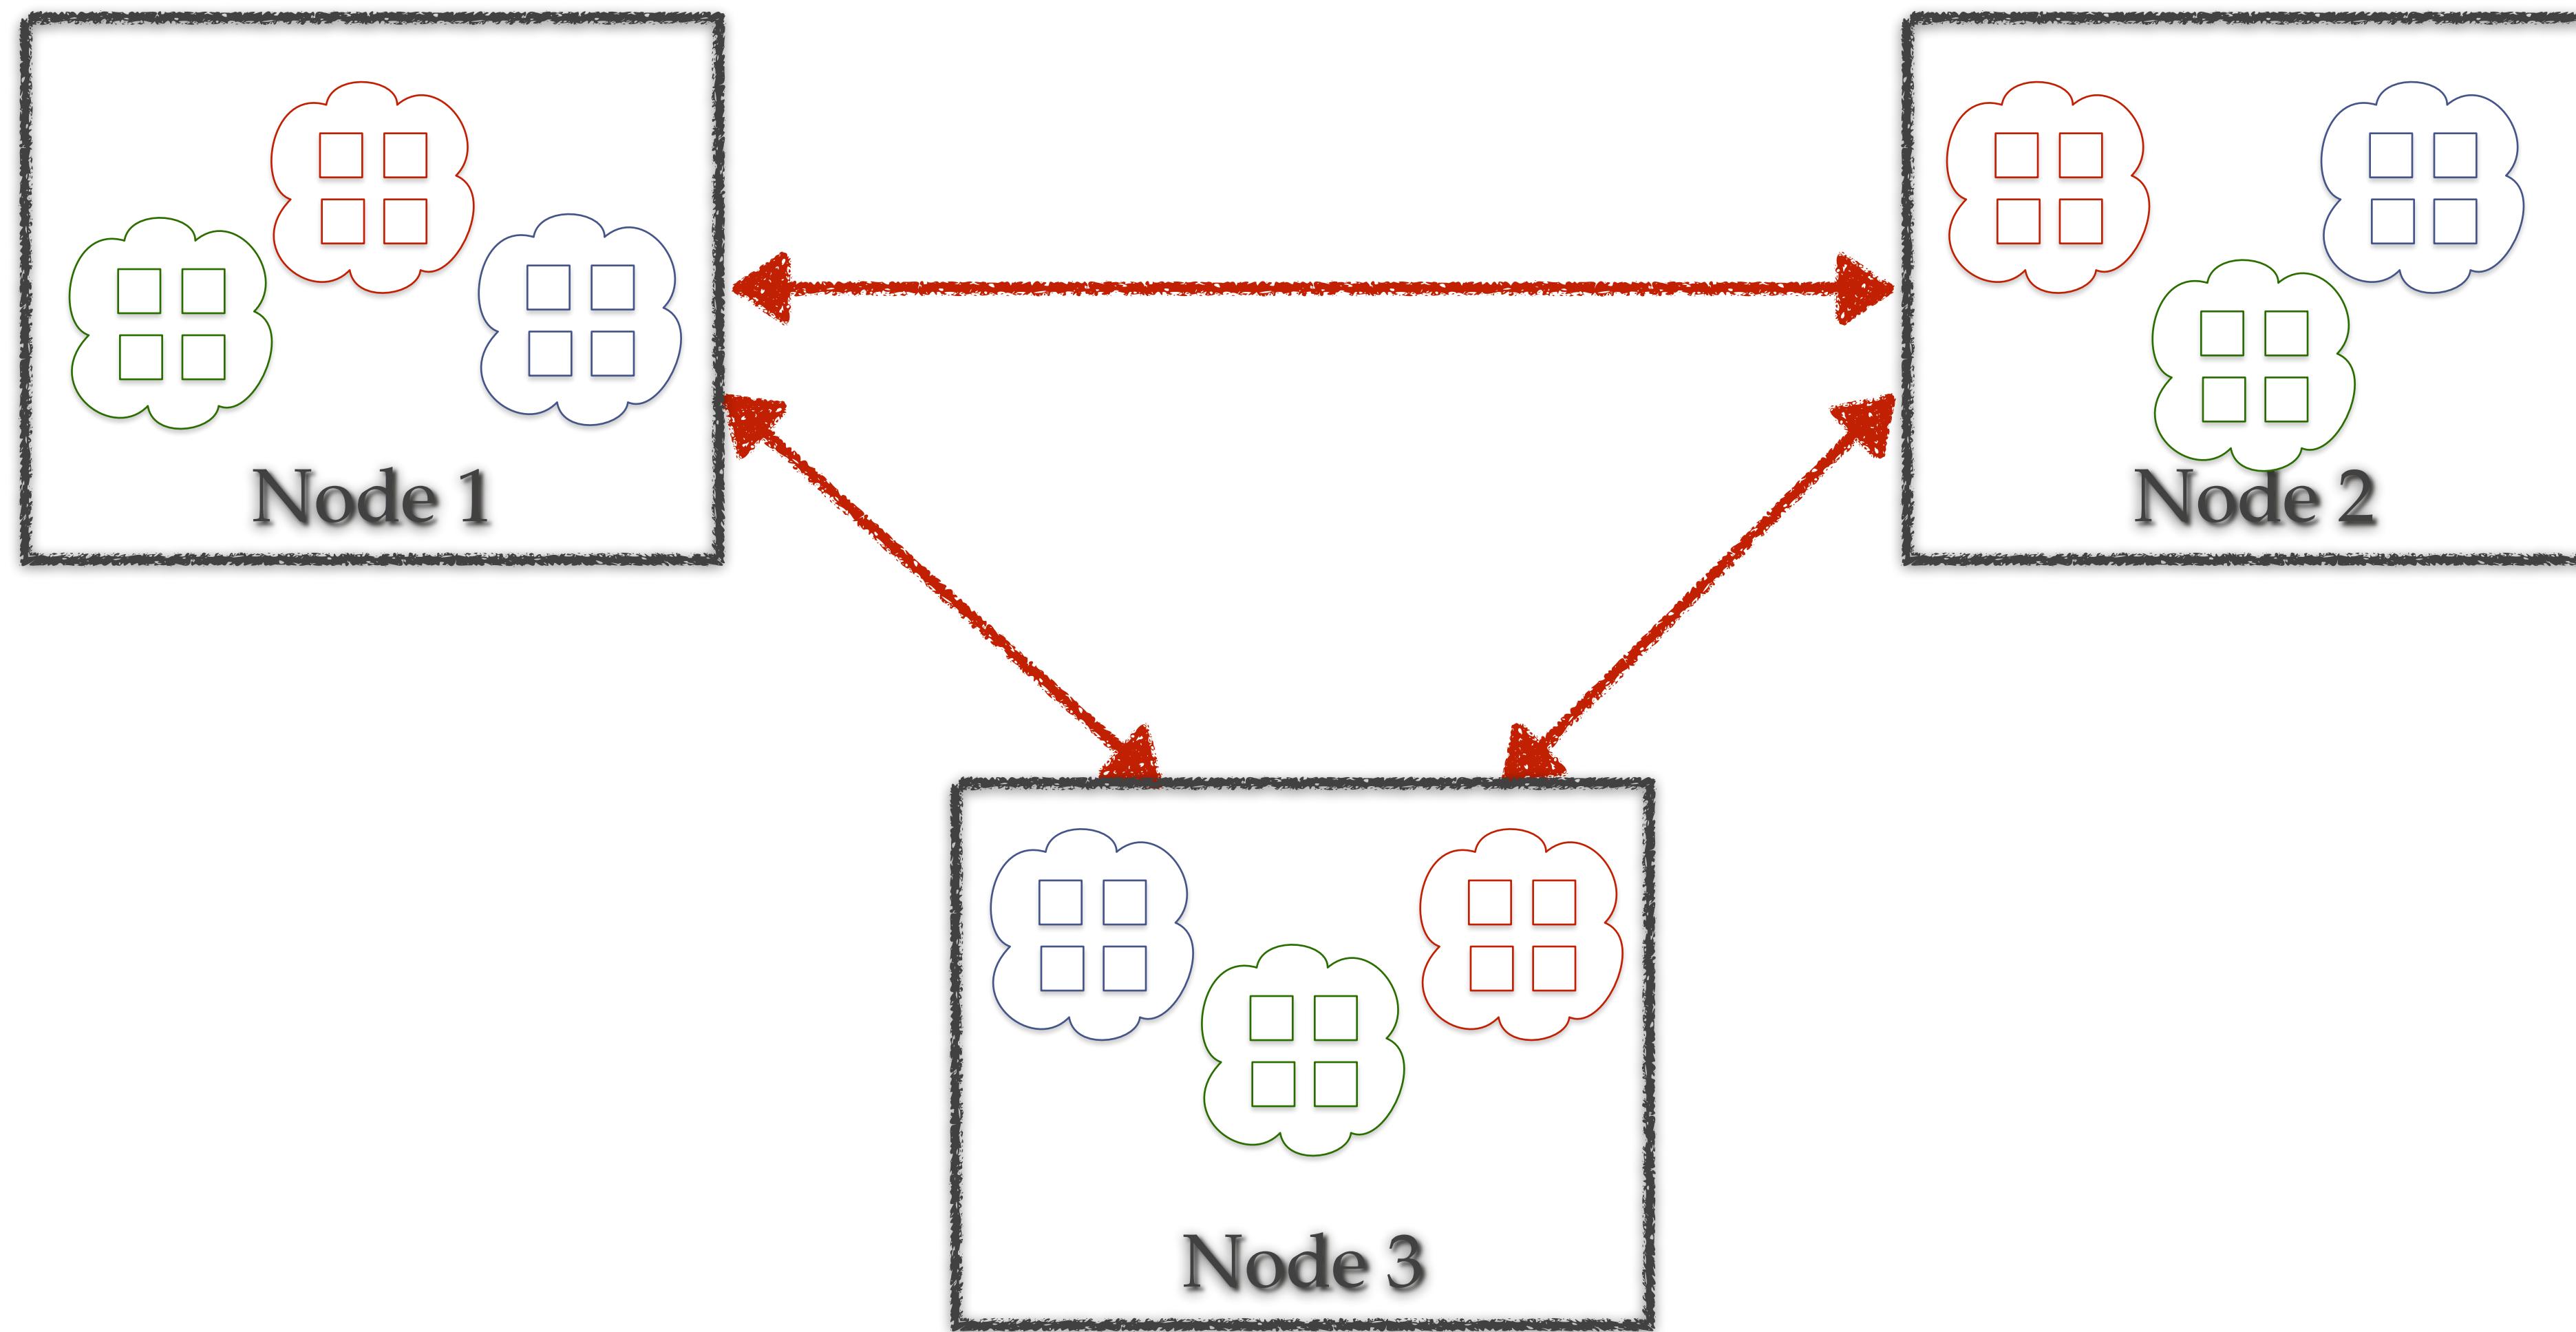
Bigger Picture

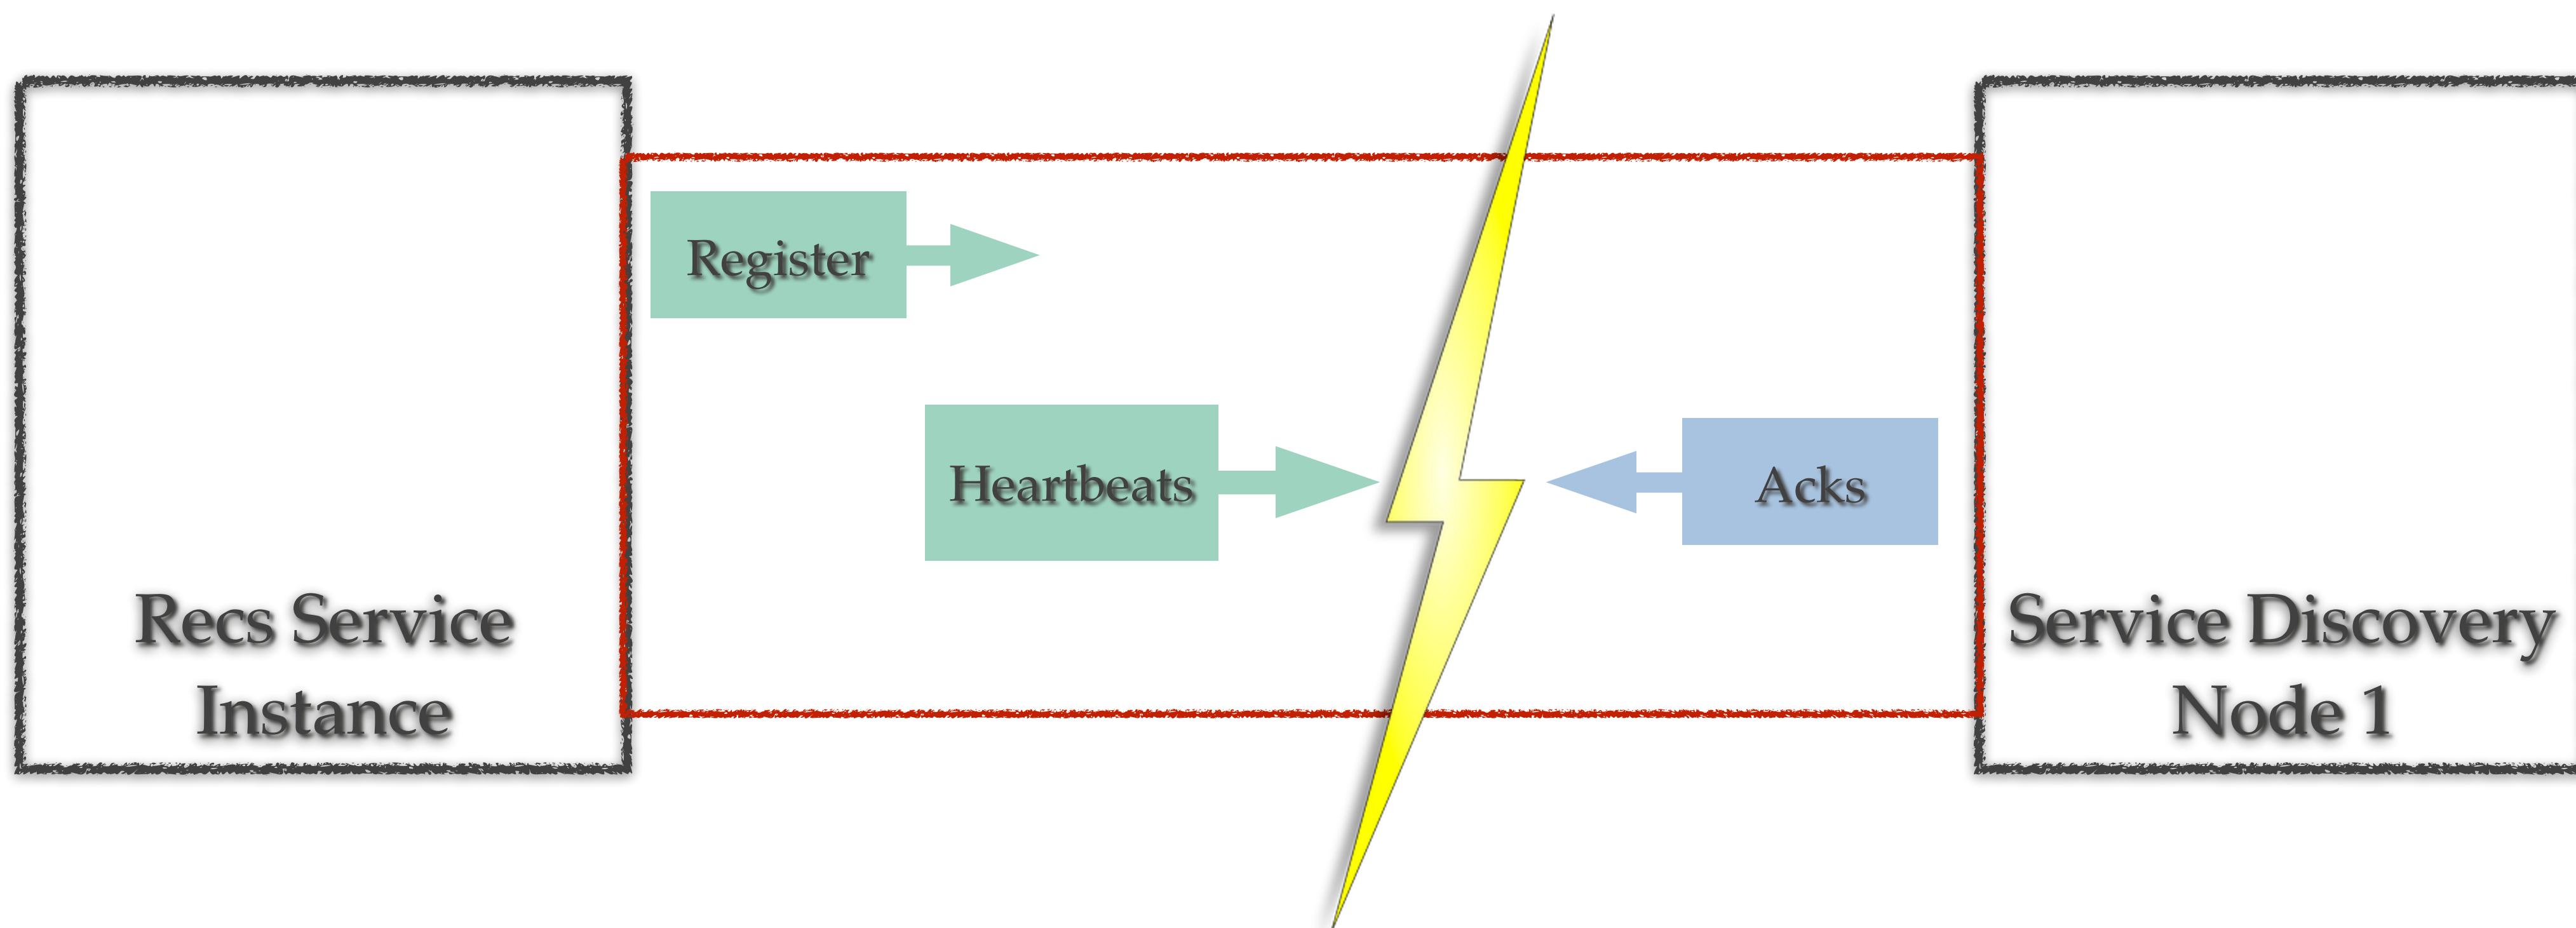


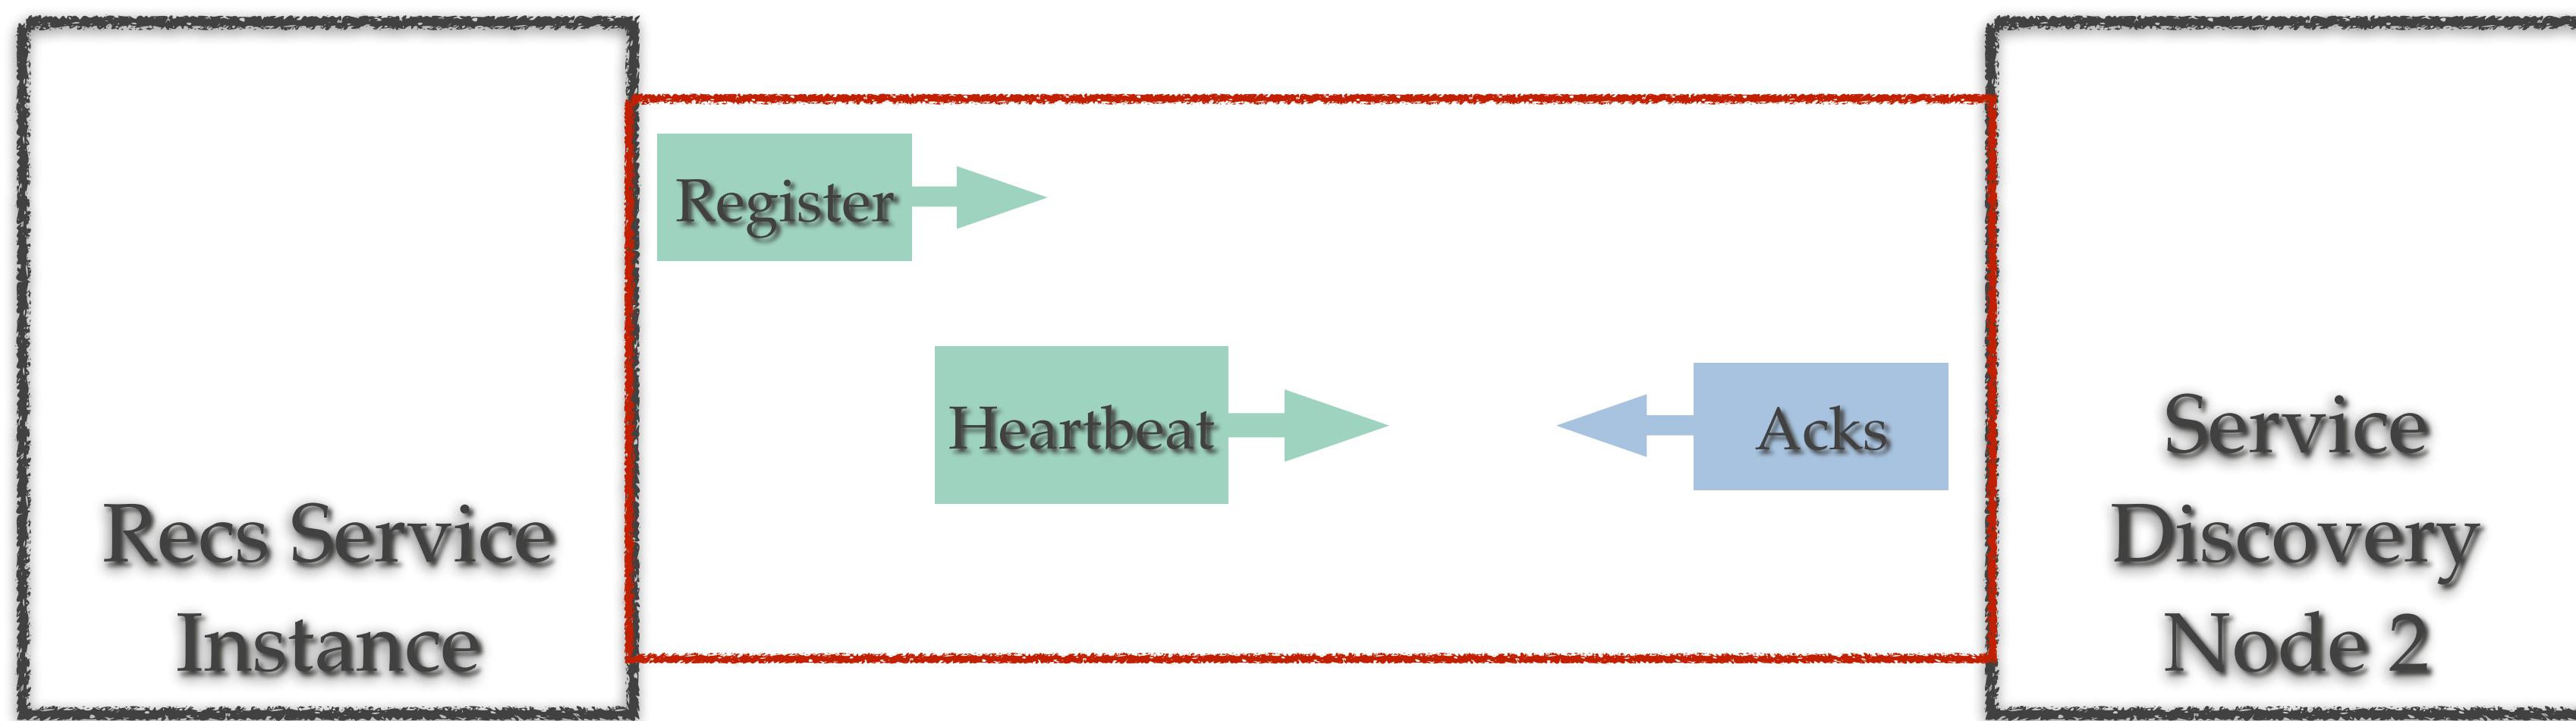
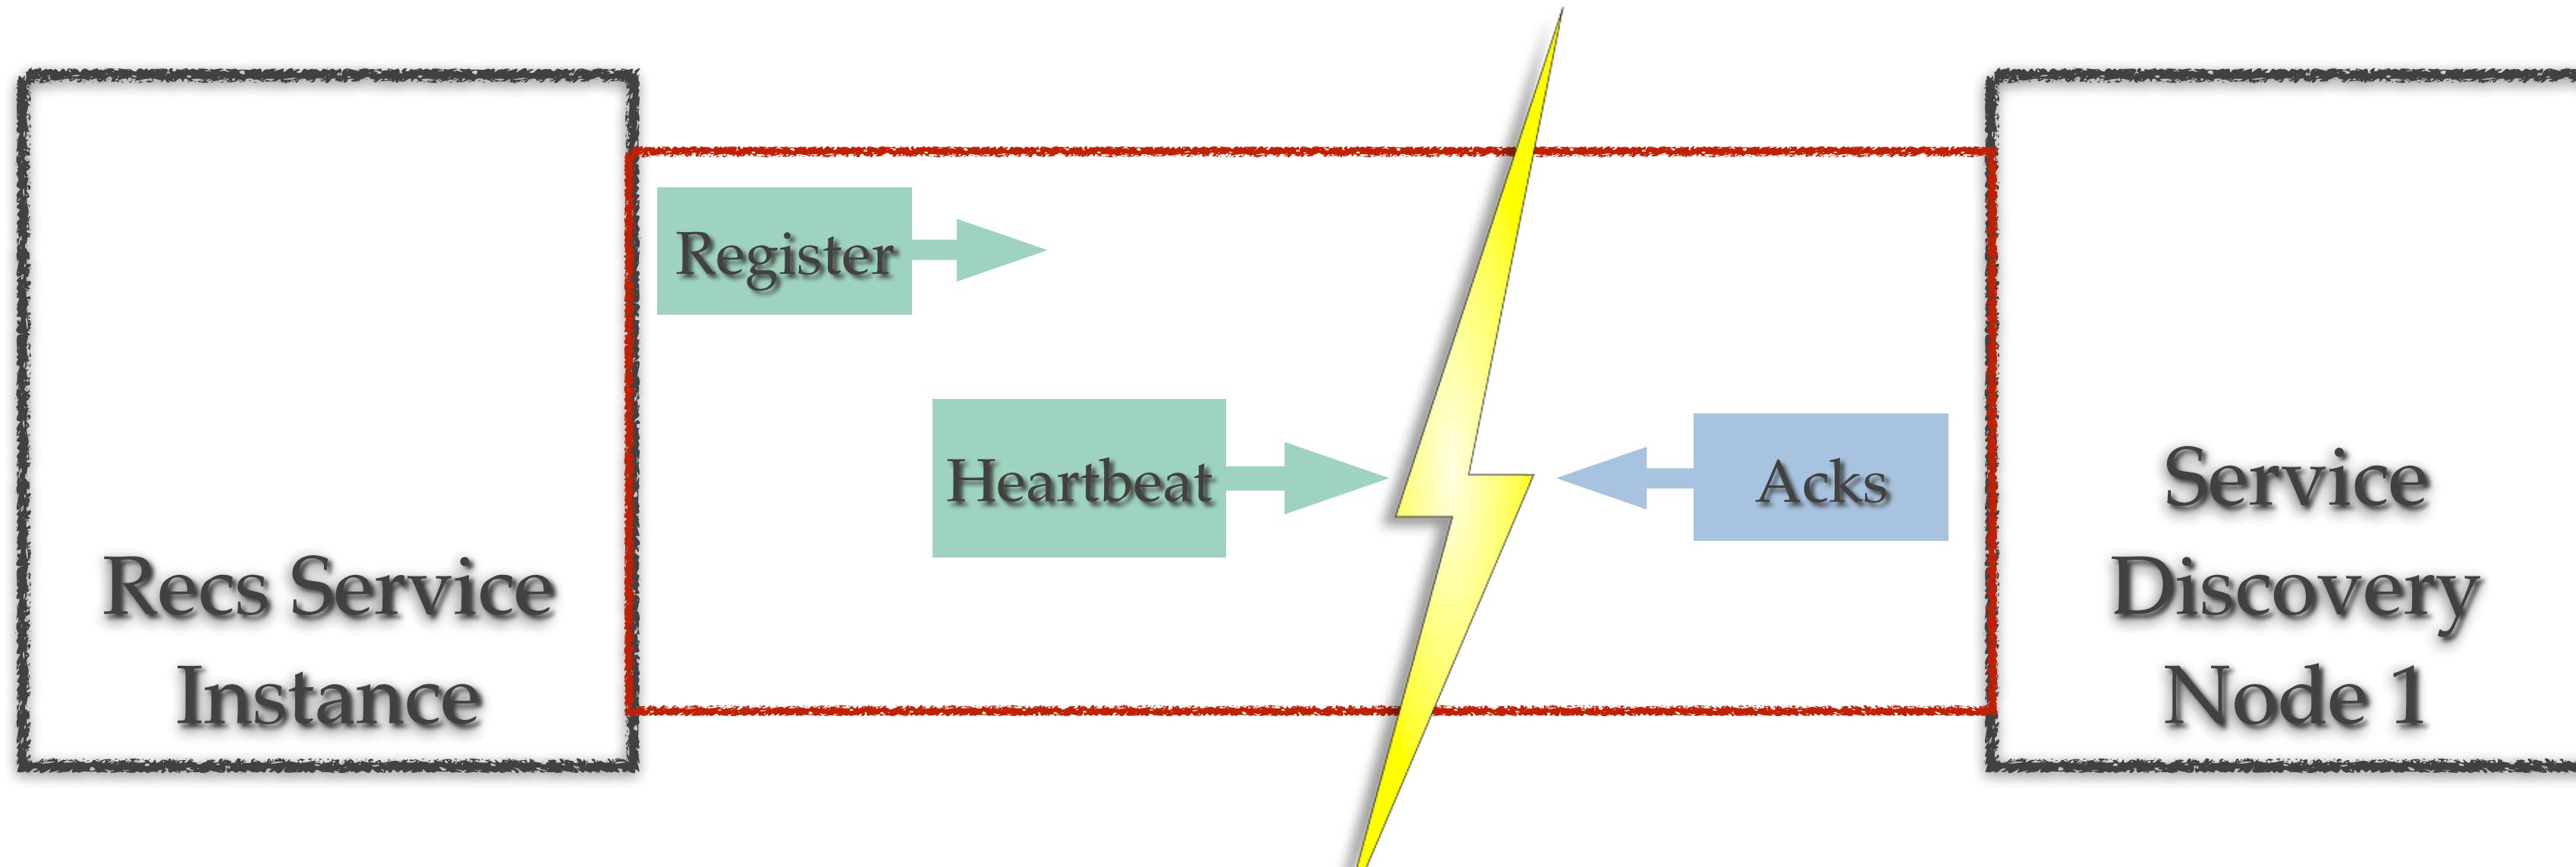
Replication

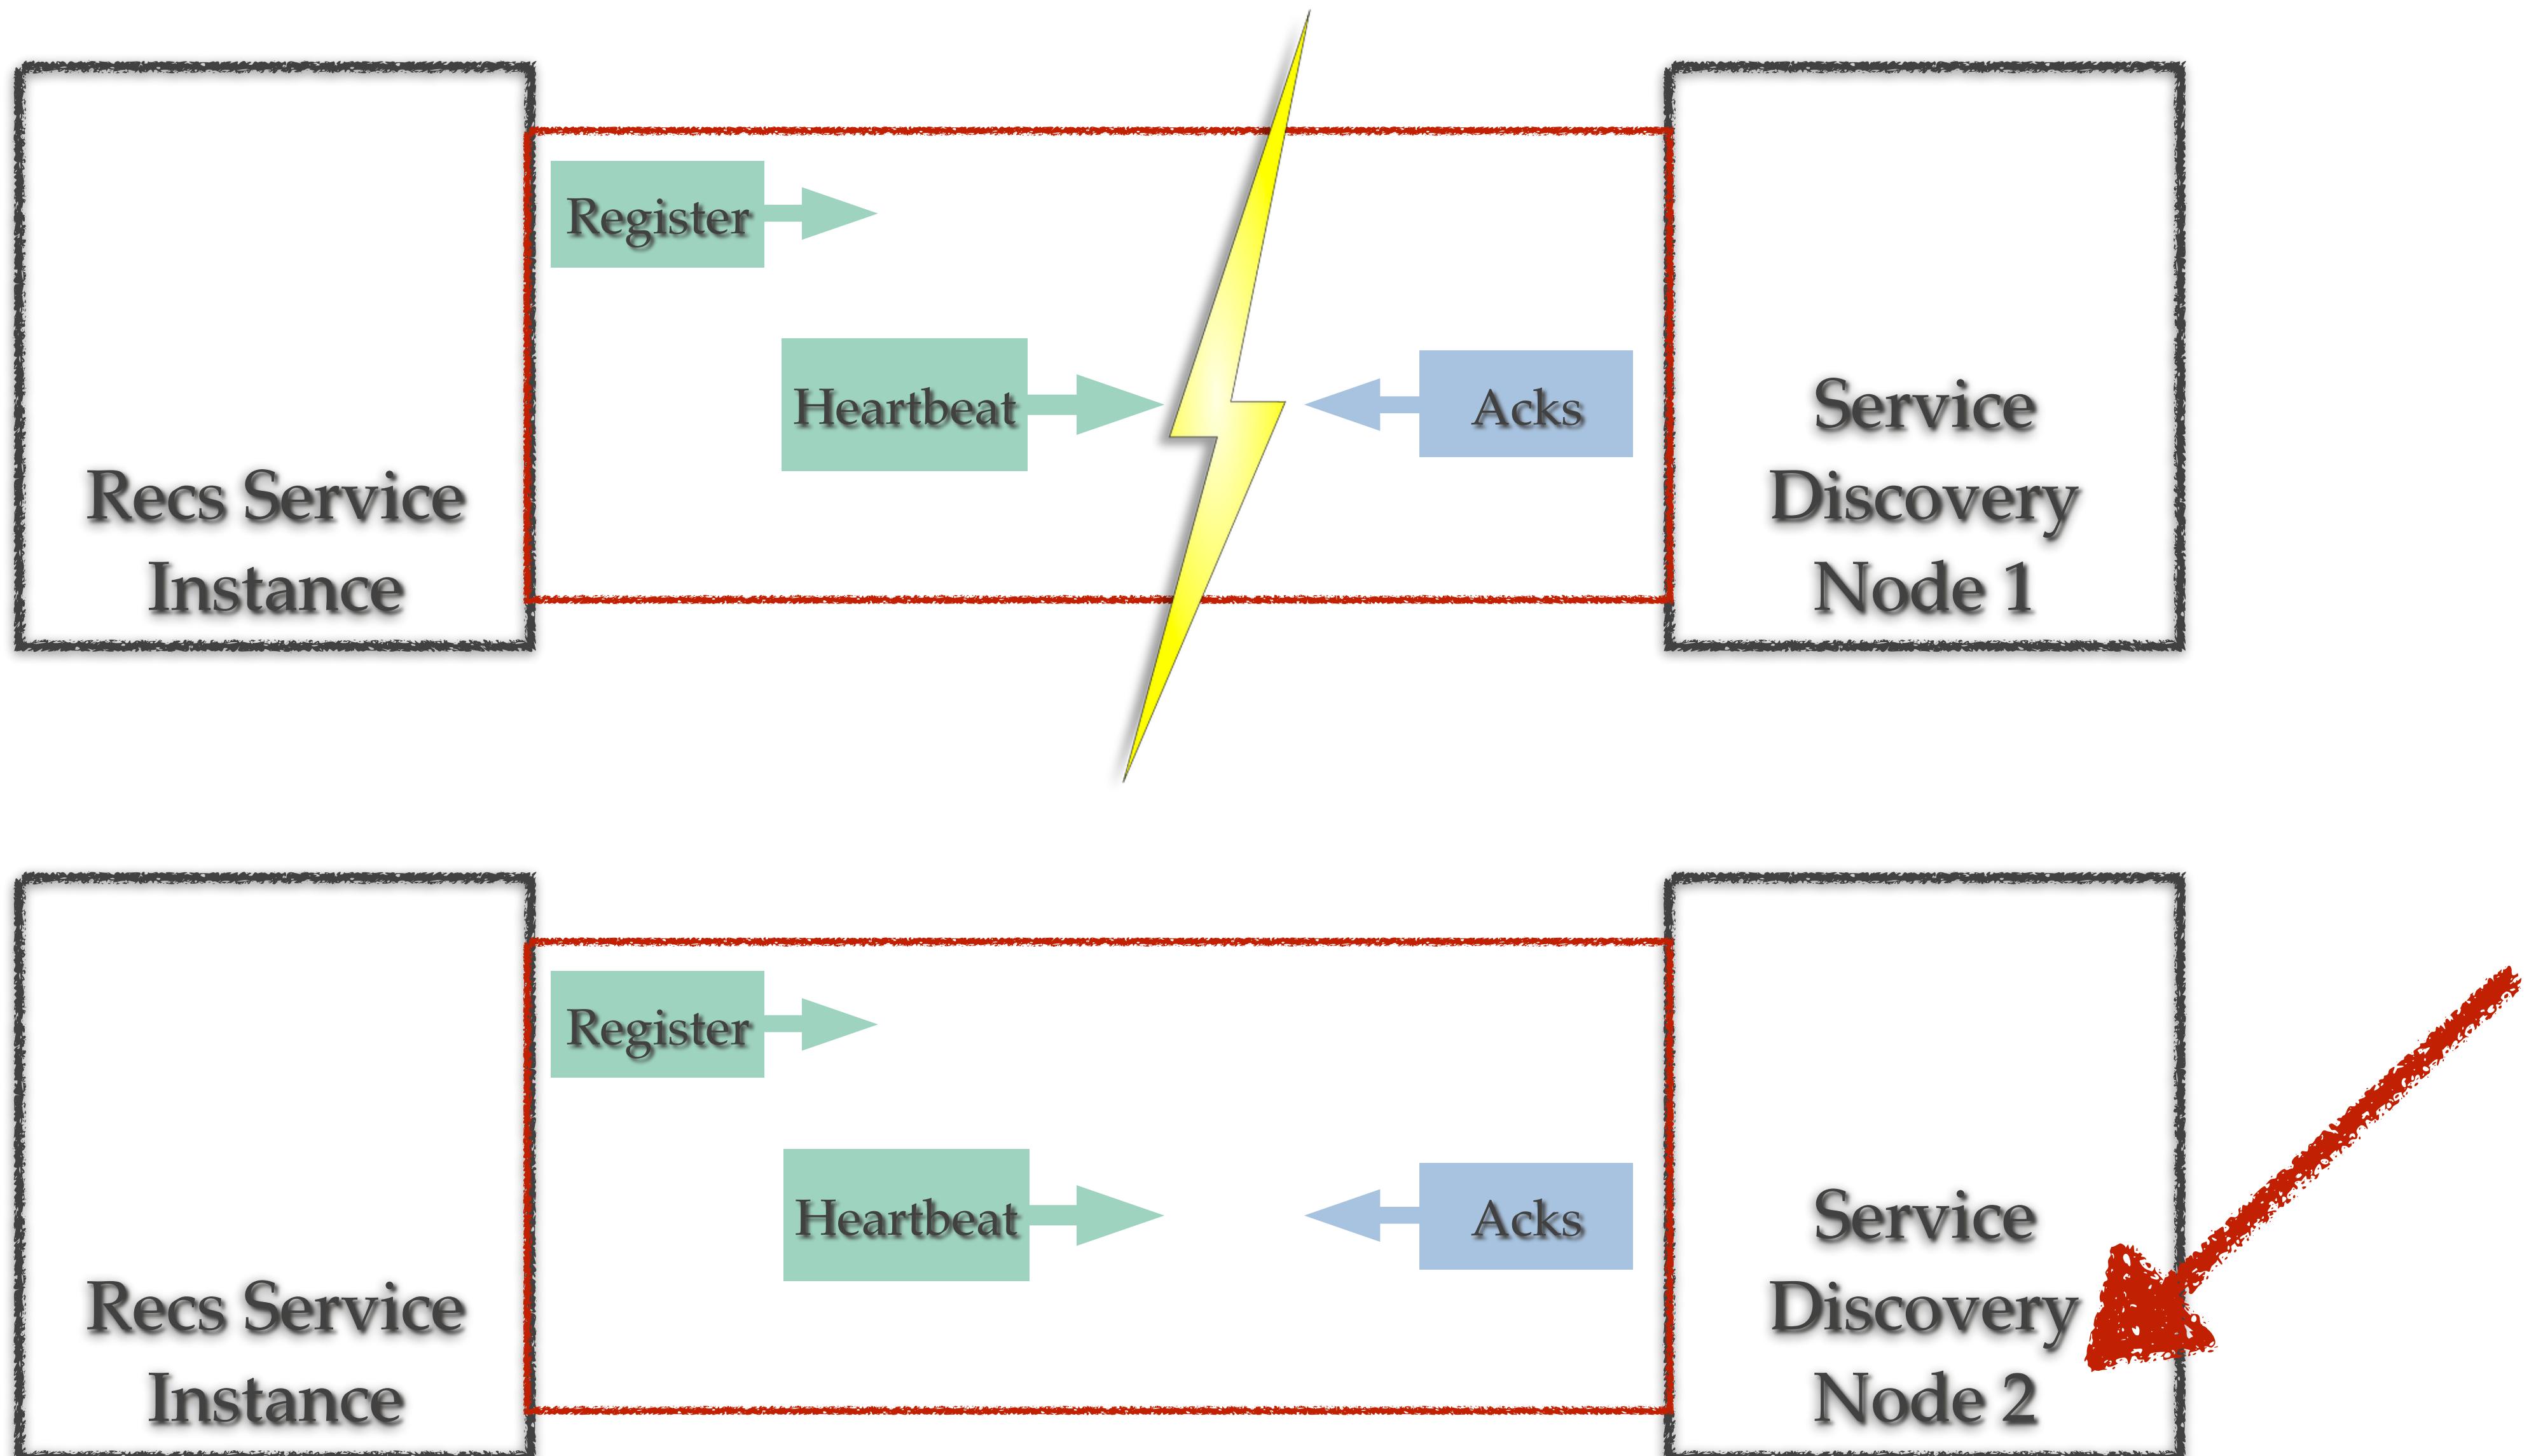


Replication



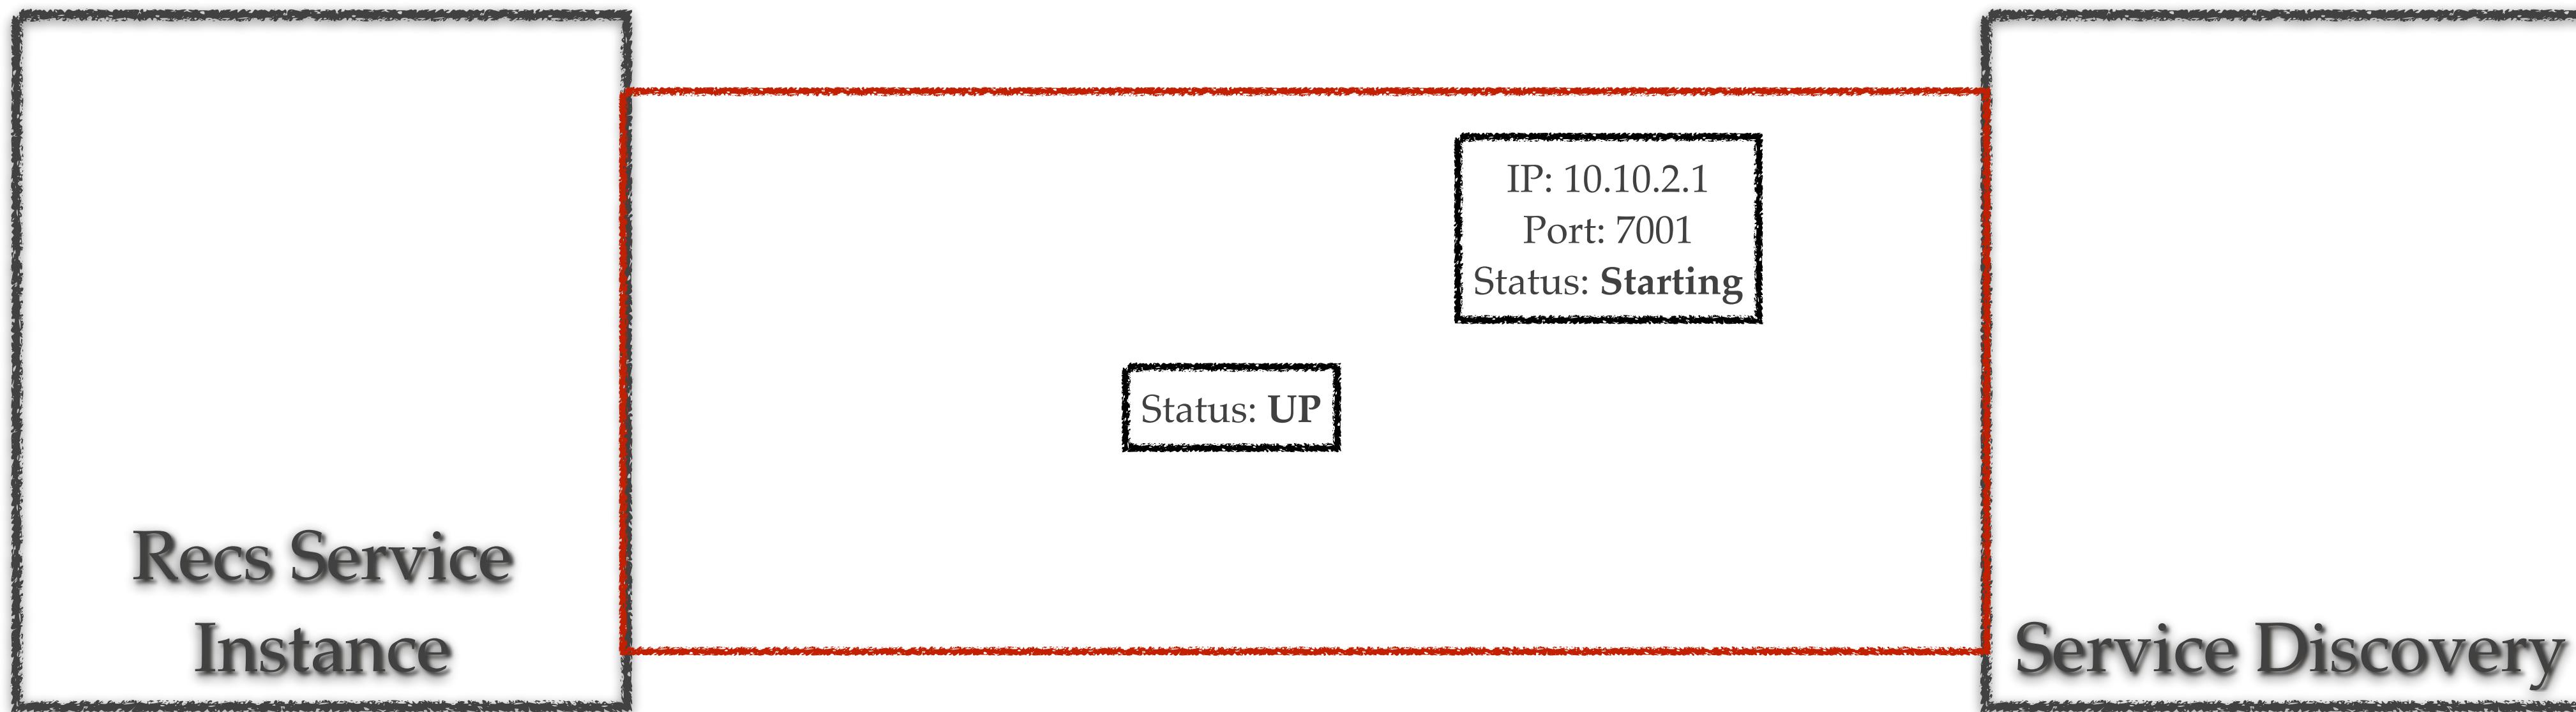




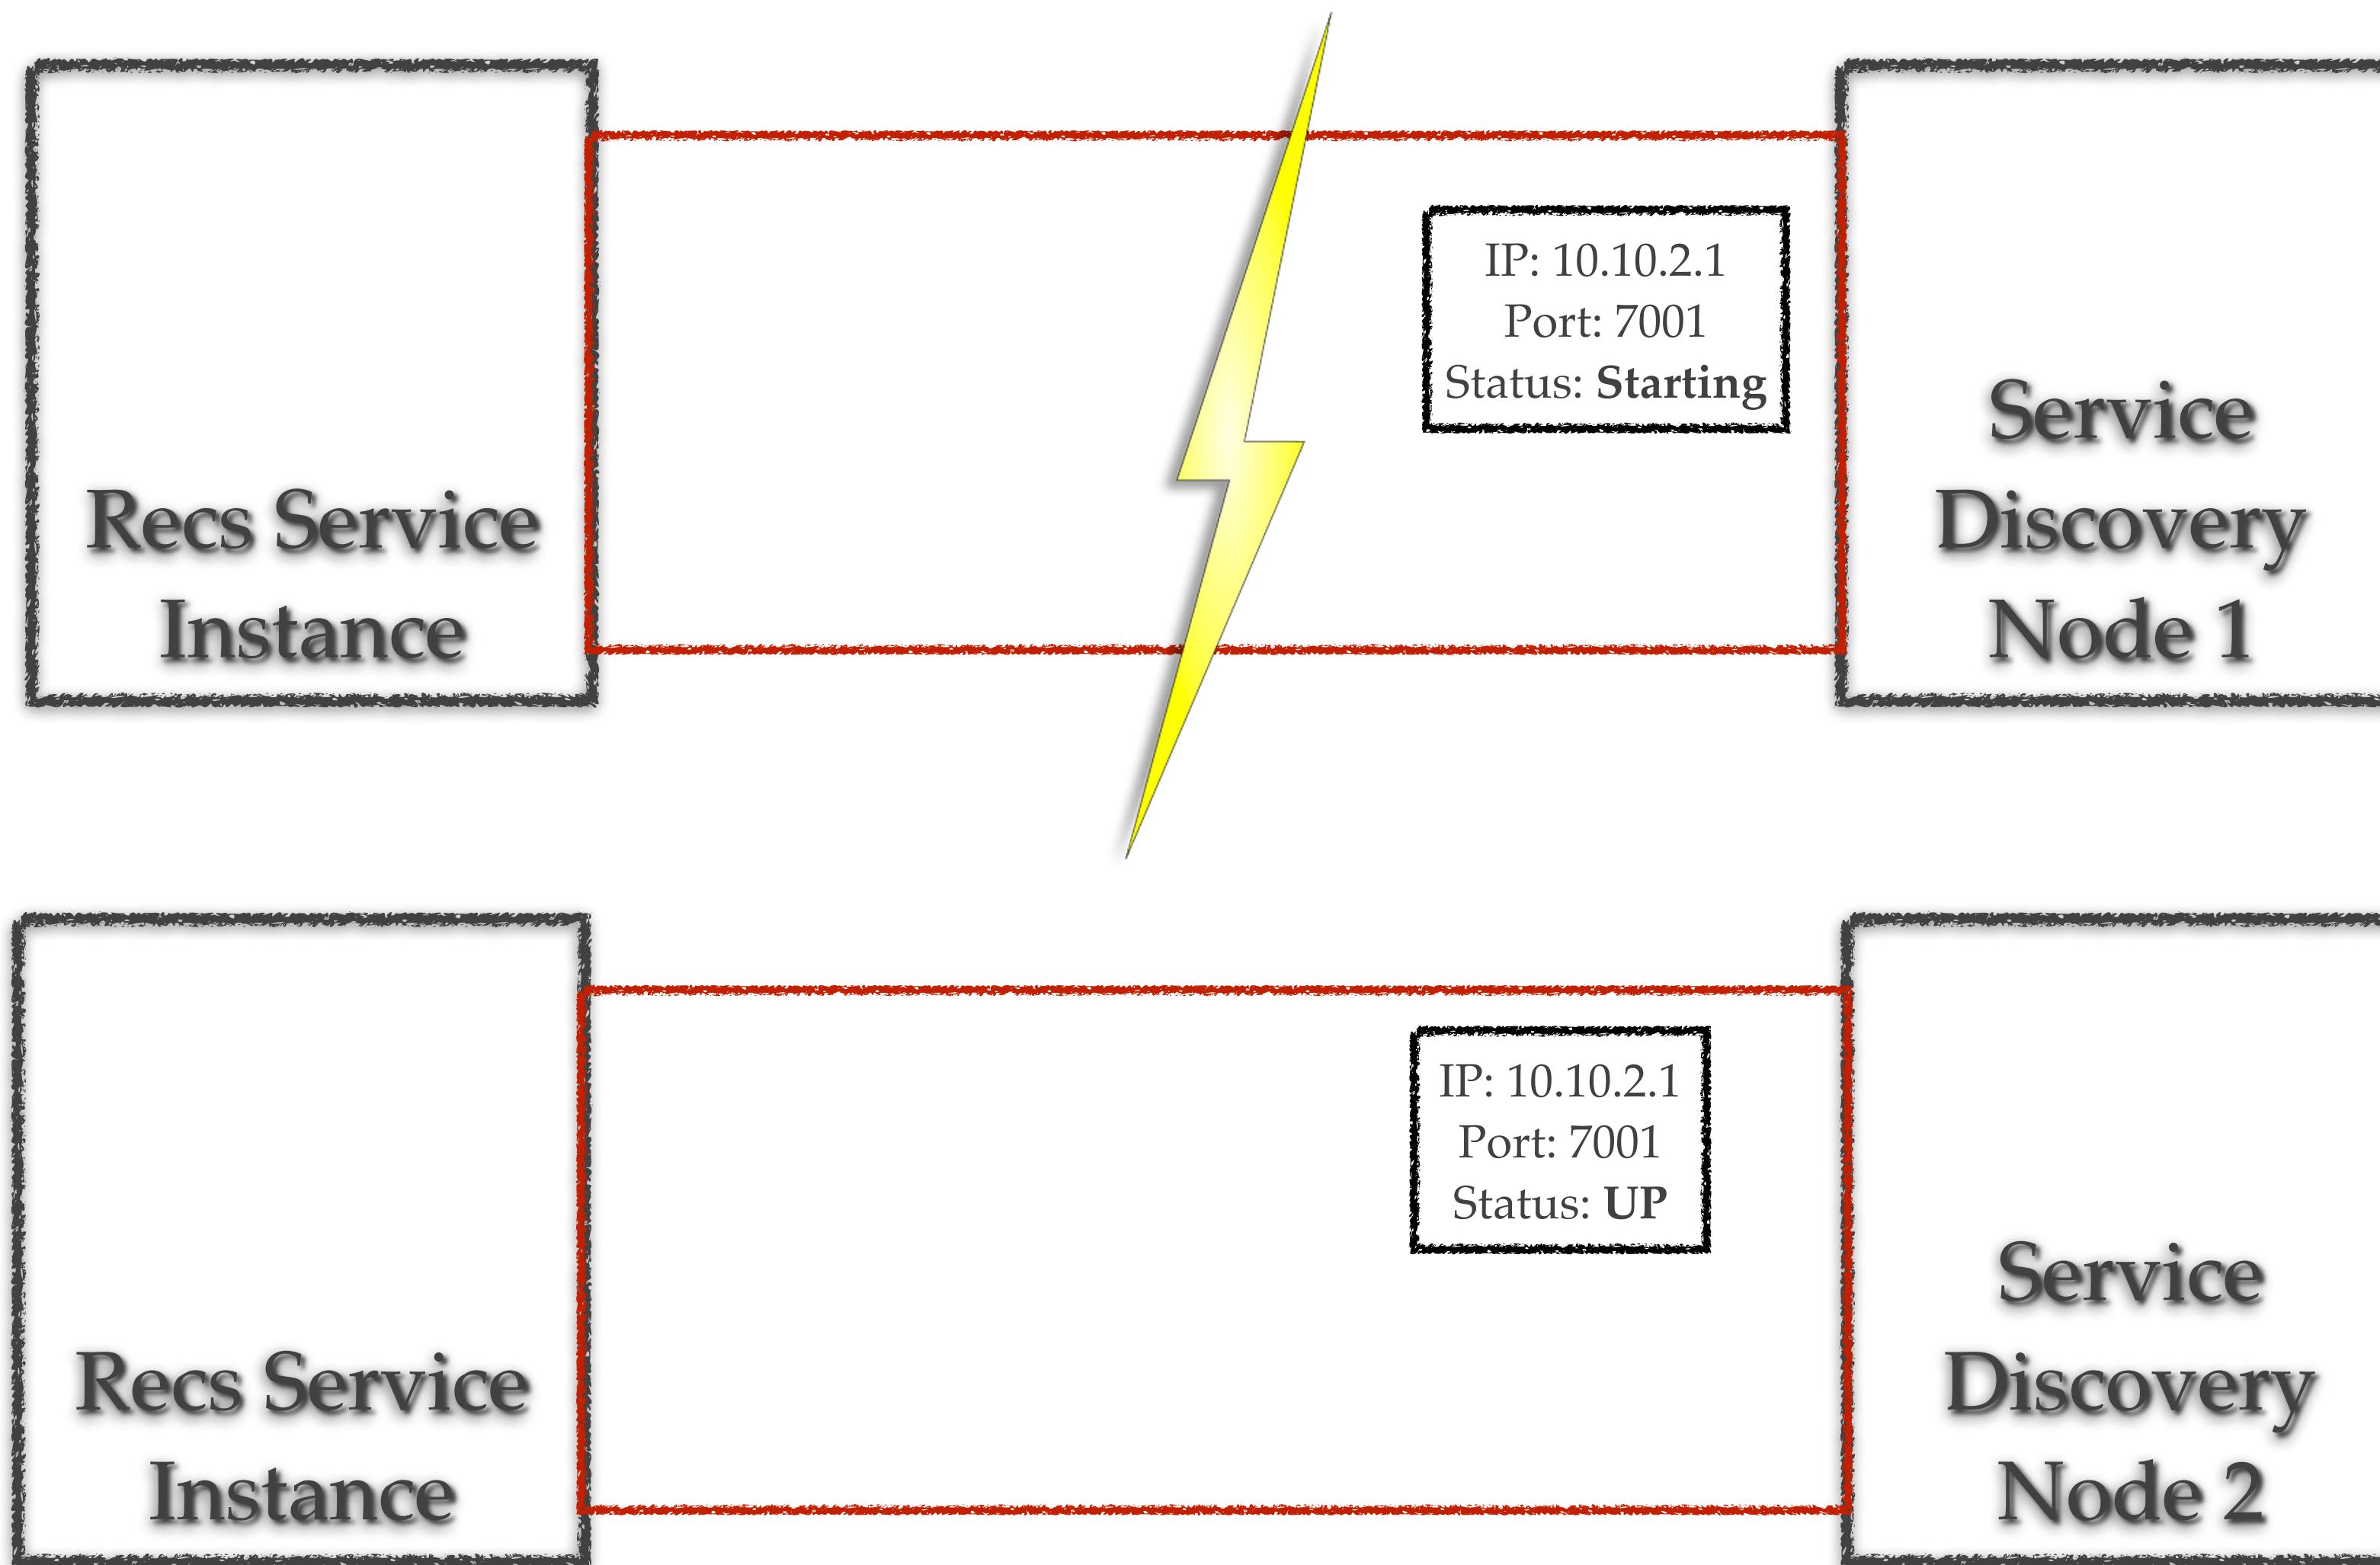


Data conflicts

Conflicts



Conflicts



Conflicts

IP: 10.10.2.1
Port: 7001
Status: **Starting**

From Service Discovery Node 1

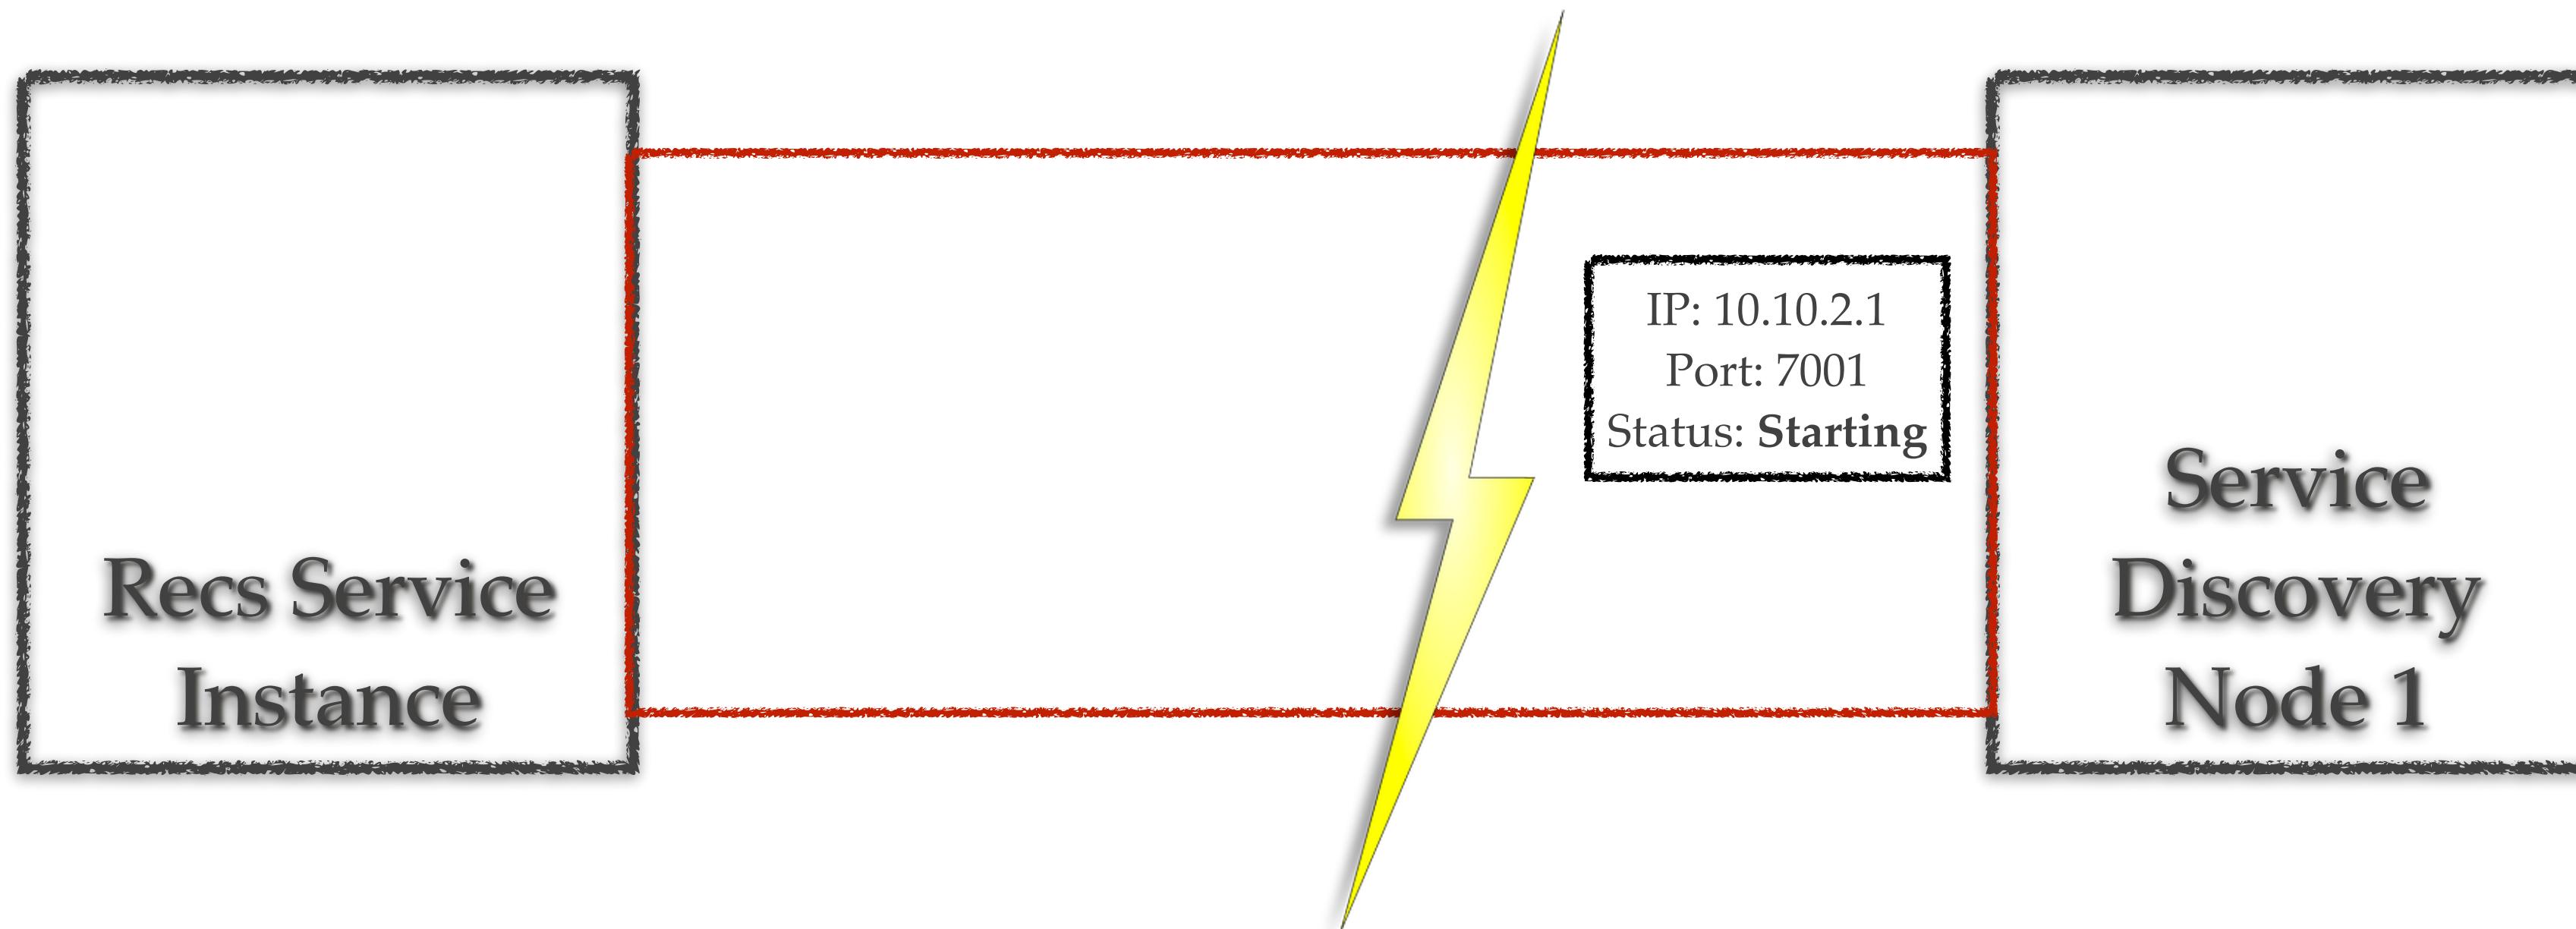
IP: 10.10.2.1
Port: 7001
Status: **UP**

From Service Discovery Node 2

Service Discovery Node 3

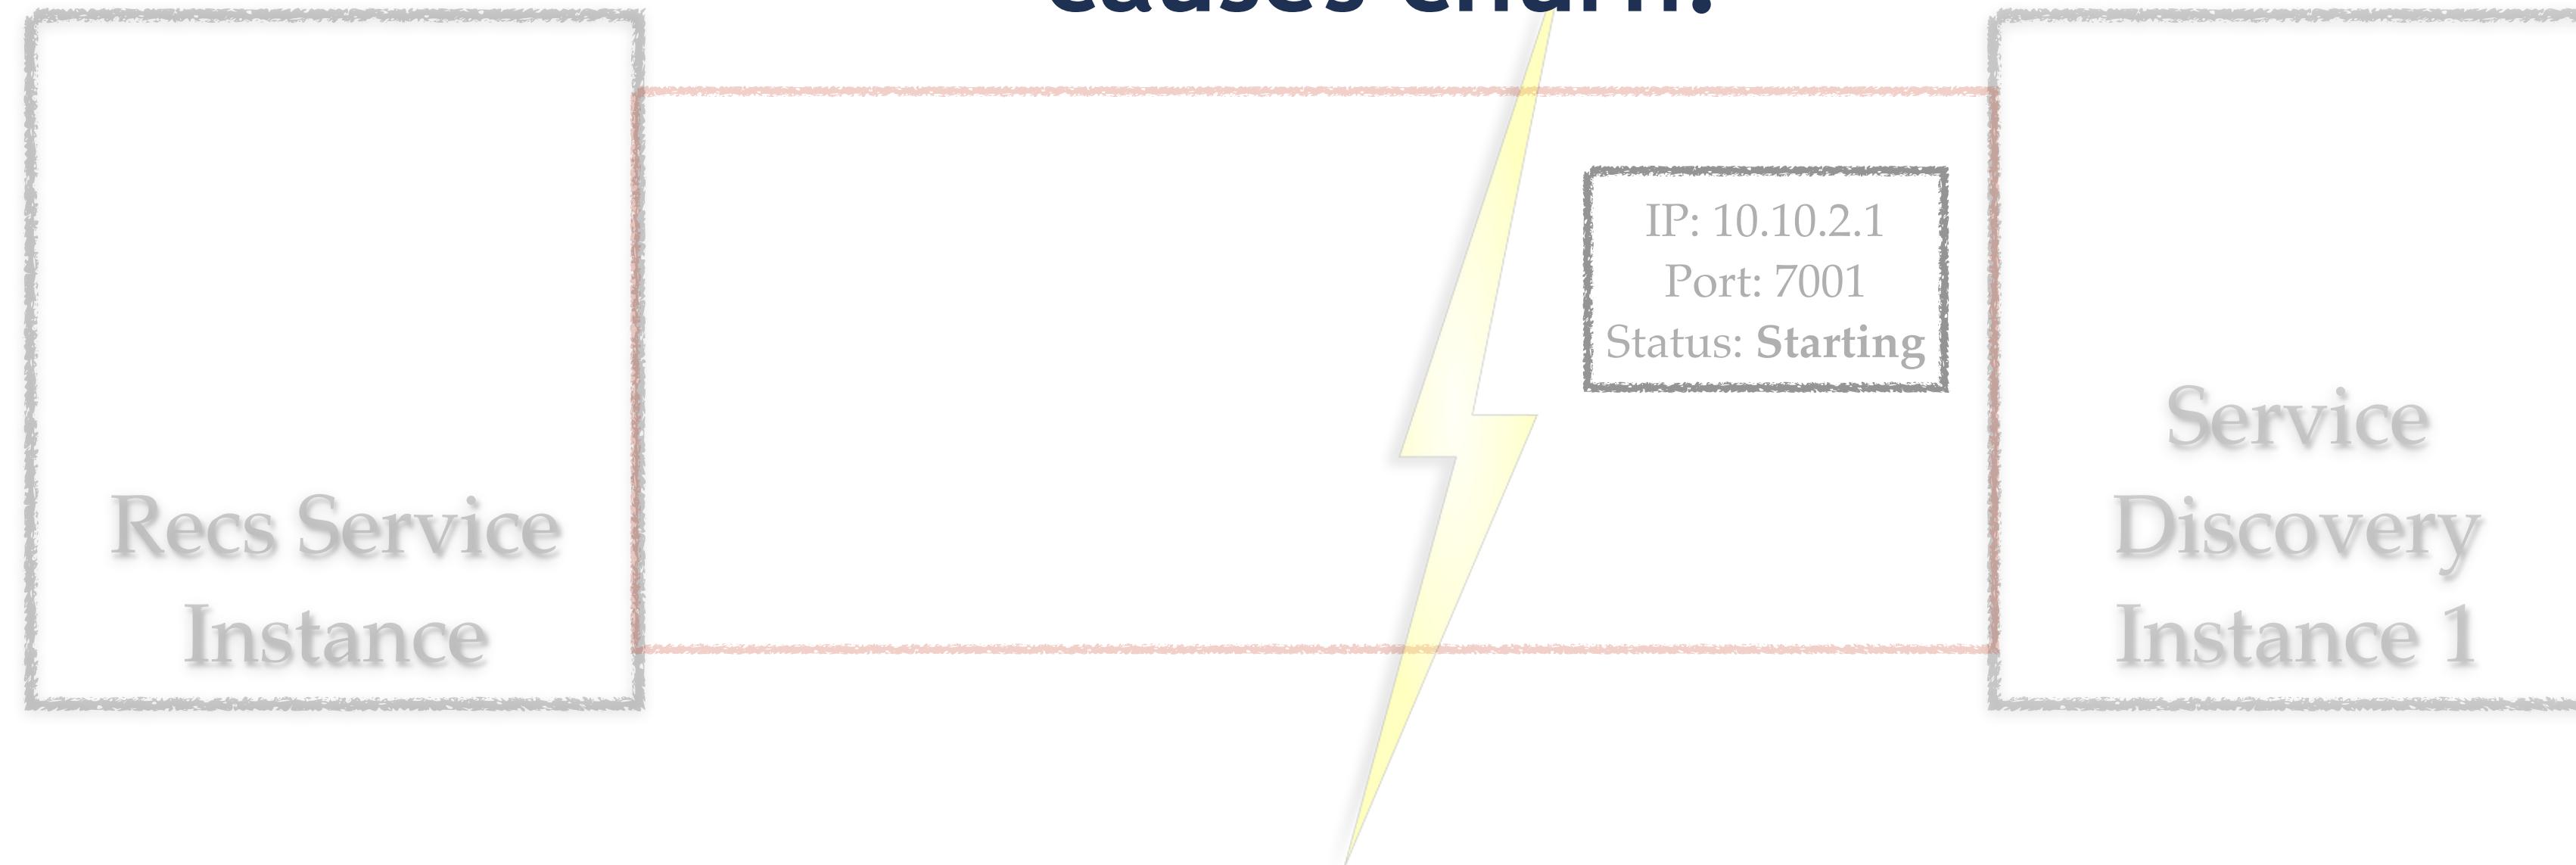
Conflict resolution

What should be the state of Recs Service Instance data now?



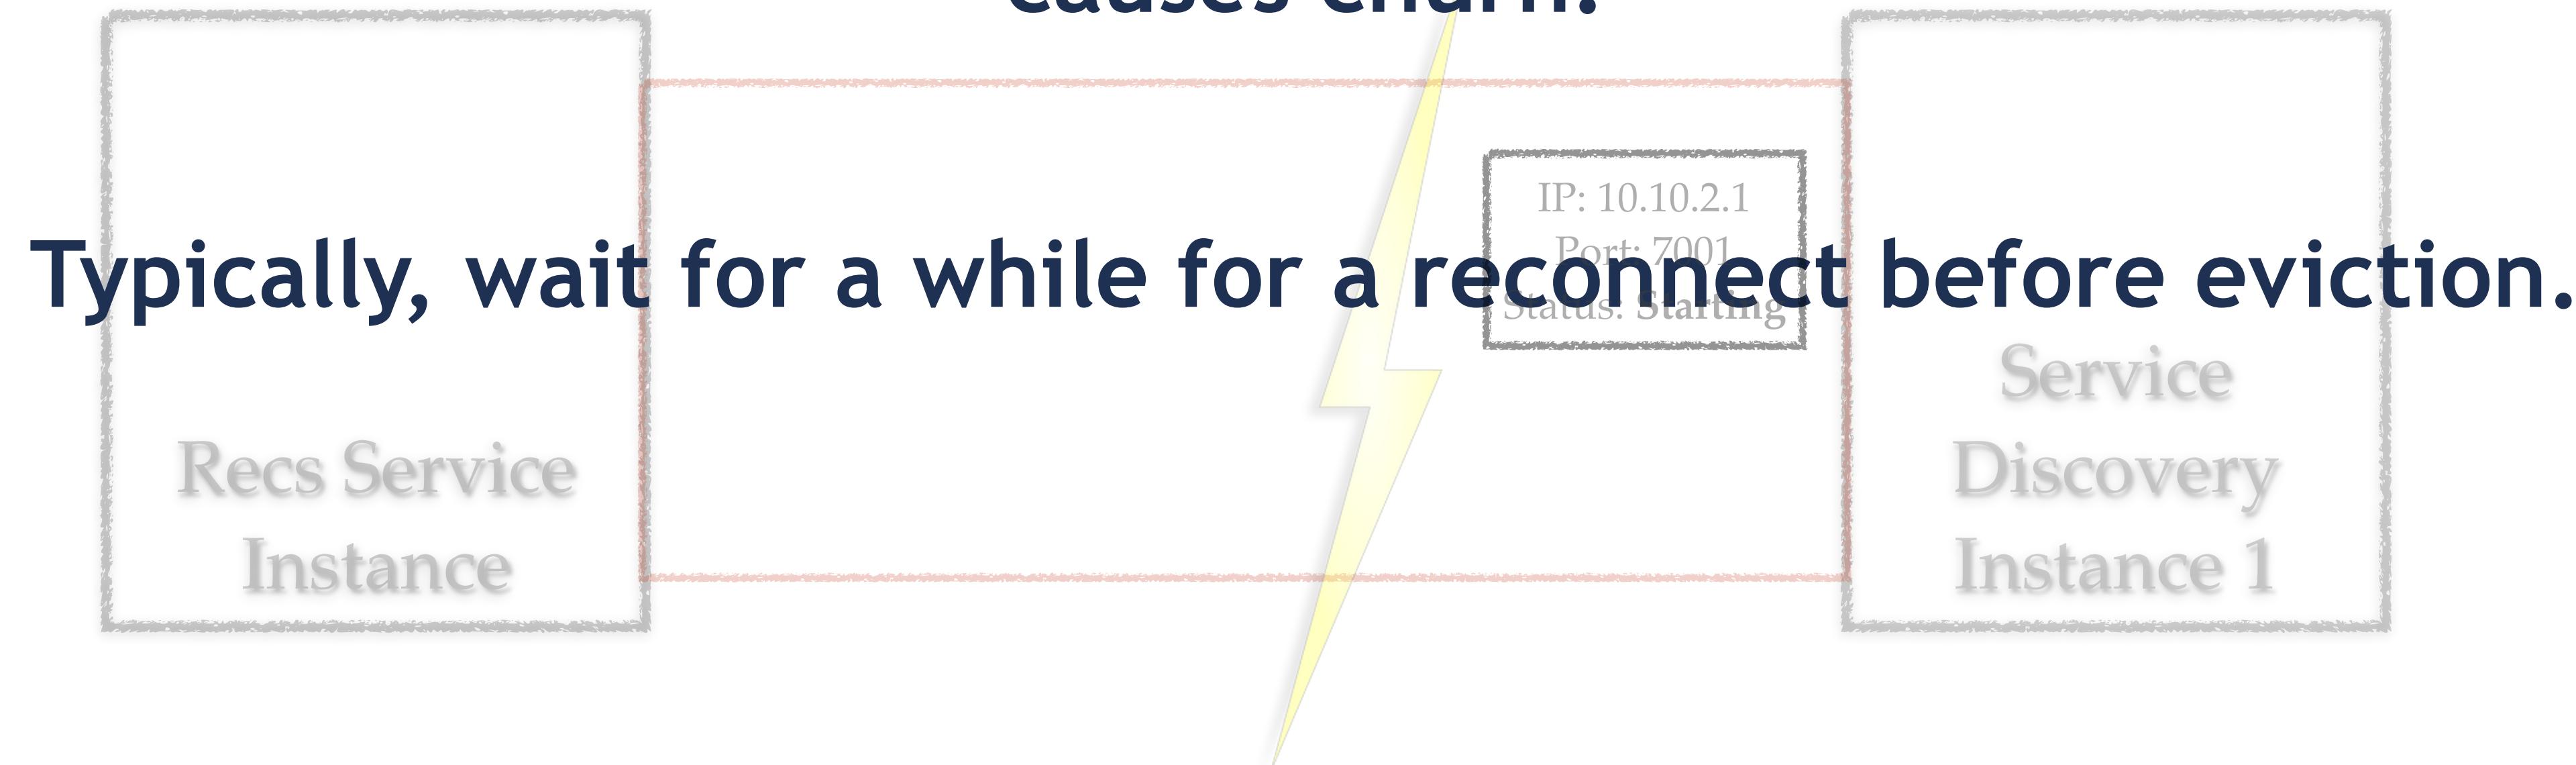
What should be the state of Recs Service Instance data now?

Show tolerance towards broken connections, evicting an instance too early causes churn.



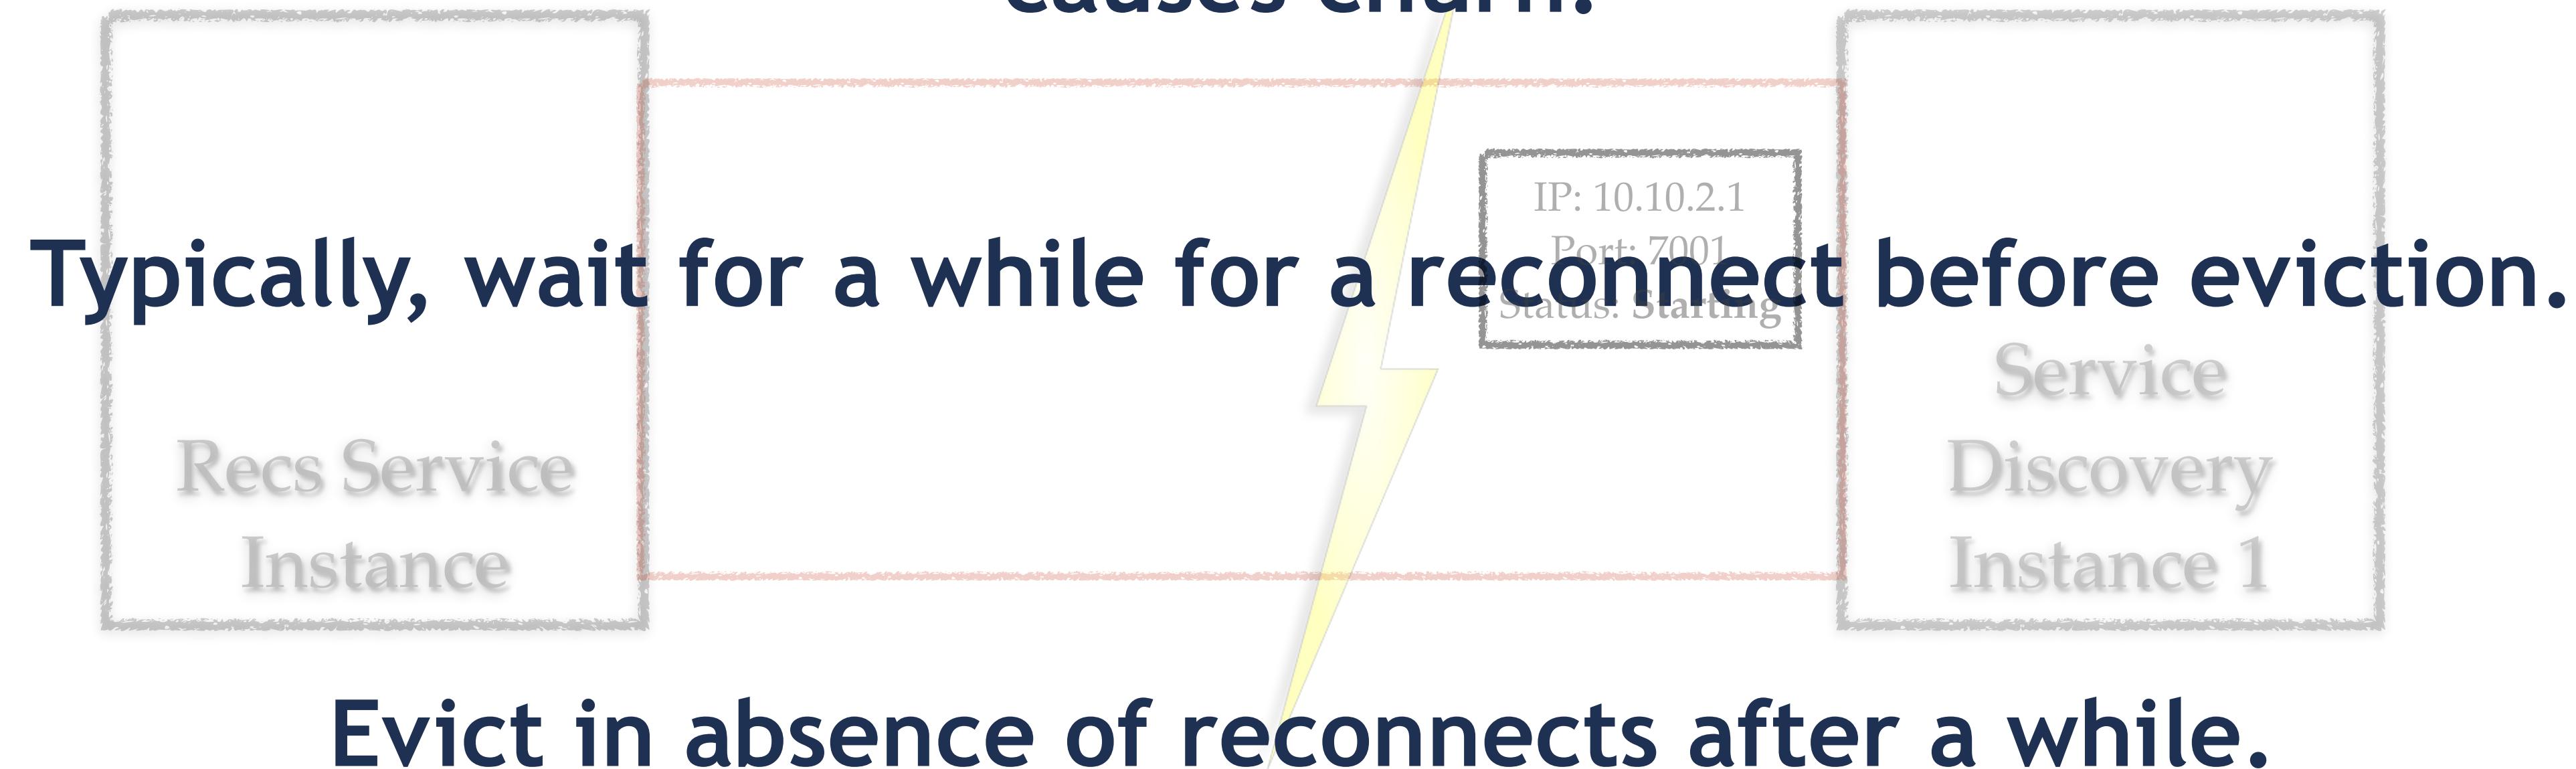
What should be the state of Recs Service Instance data now?

Show tolerance towards broken connections, evicting an instance too early causes churn.



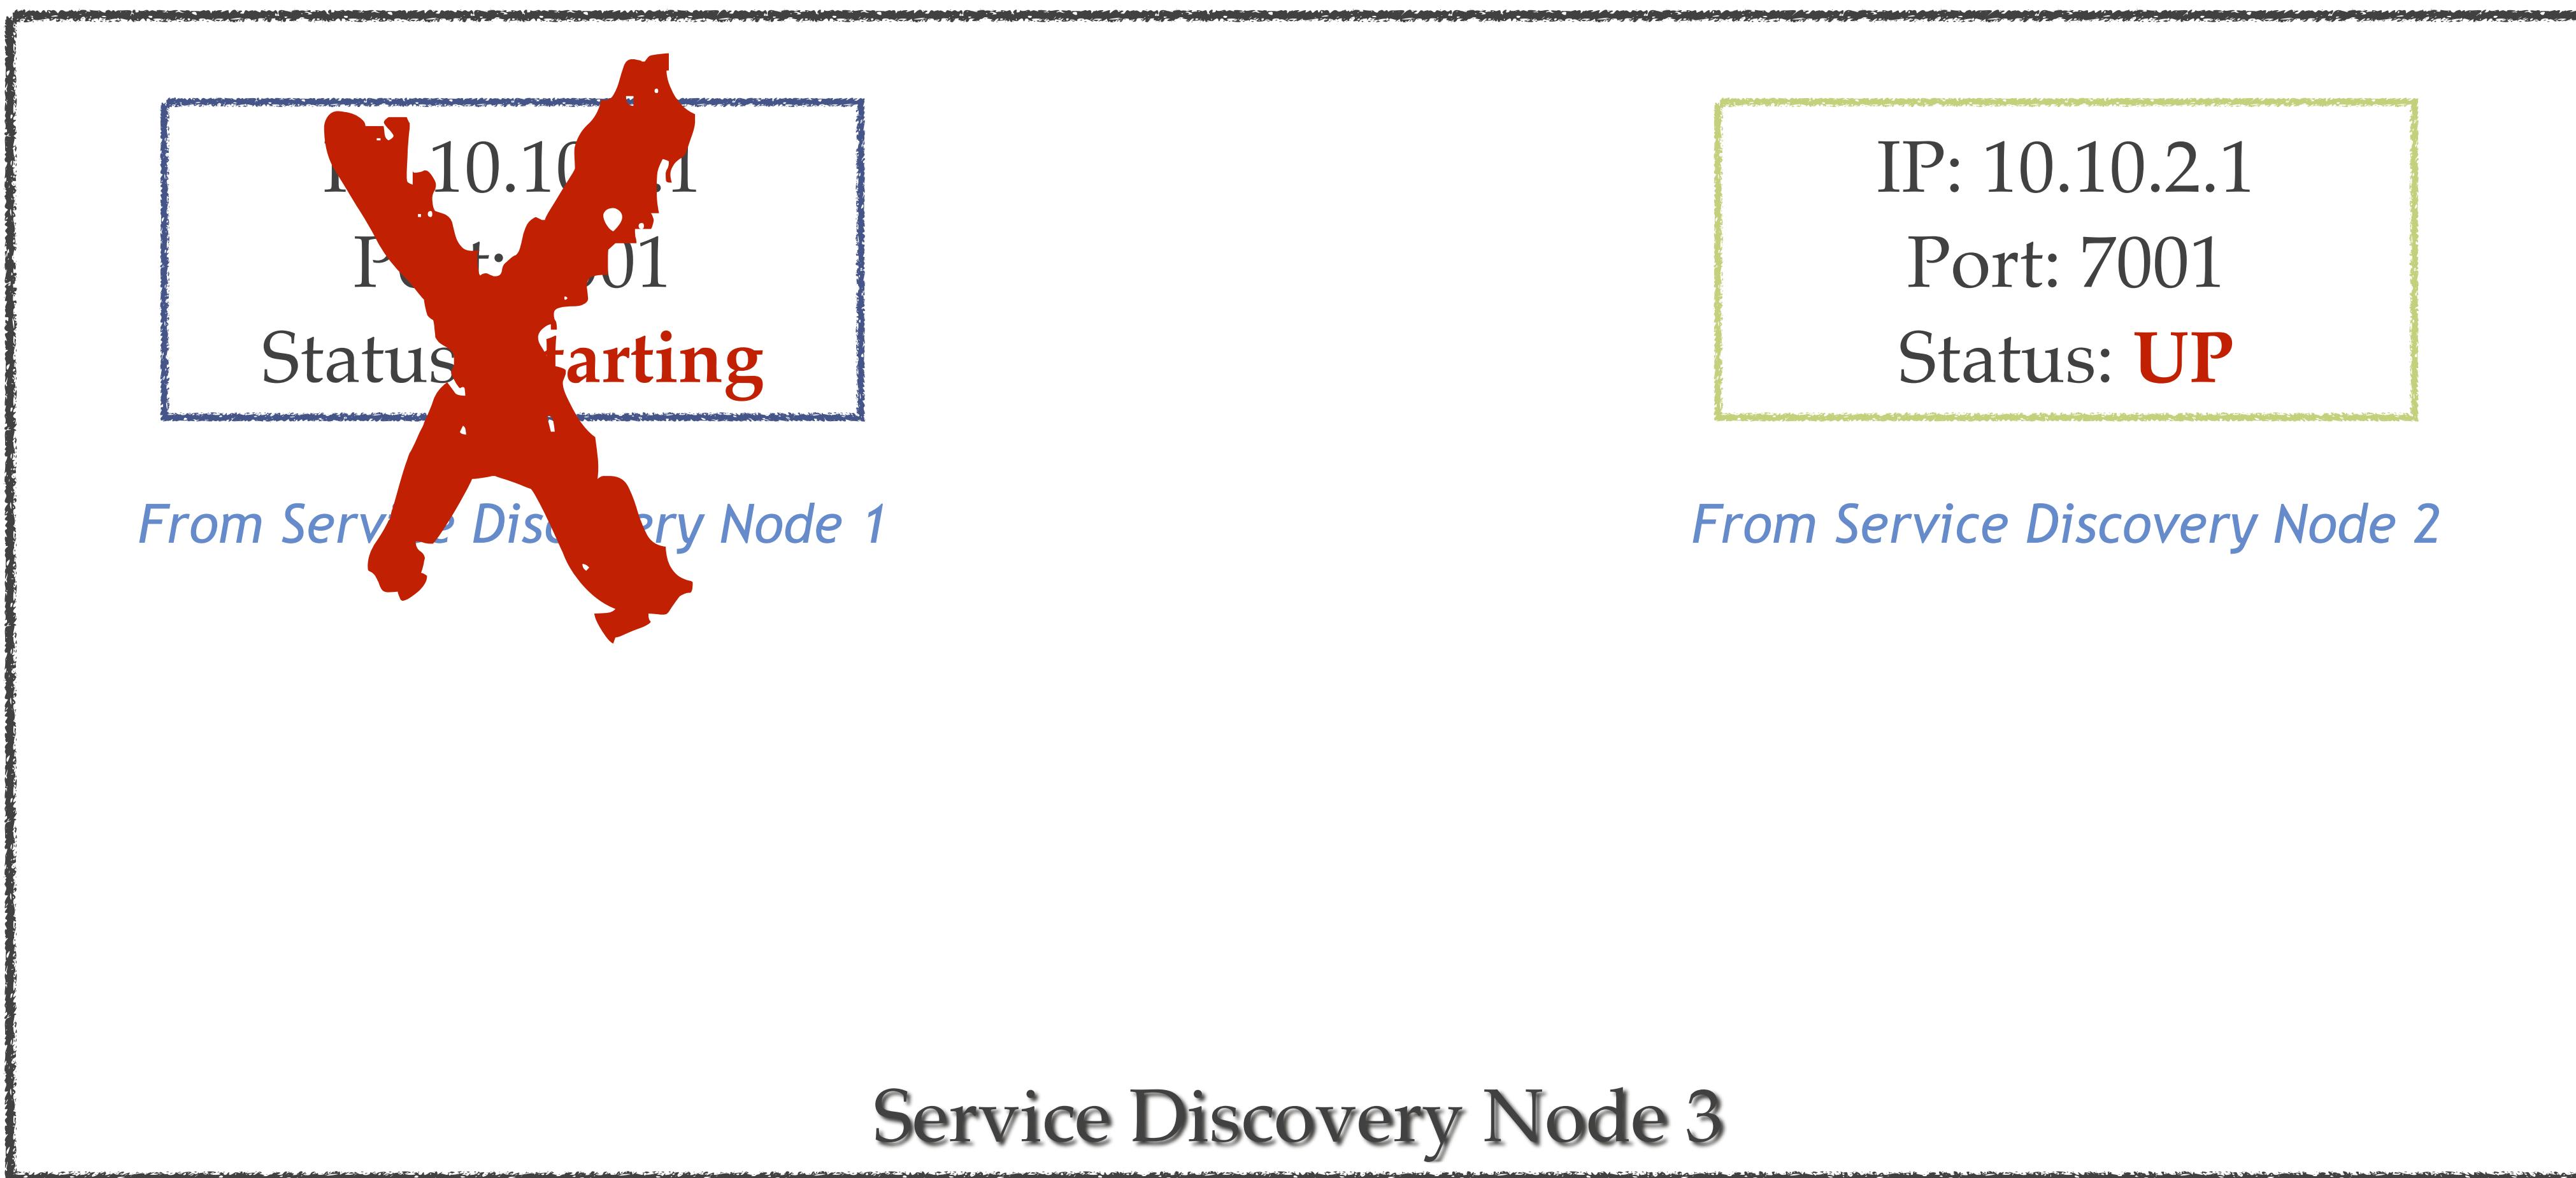
What should be the state of Recs Service Instance data now?

Show tolerance towards broken connections, evicting an instance too early causes churn.



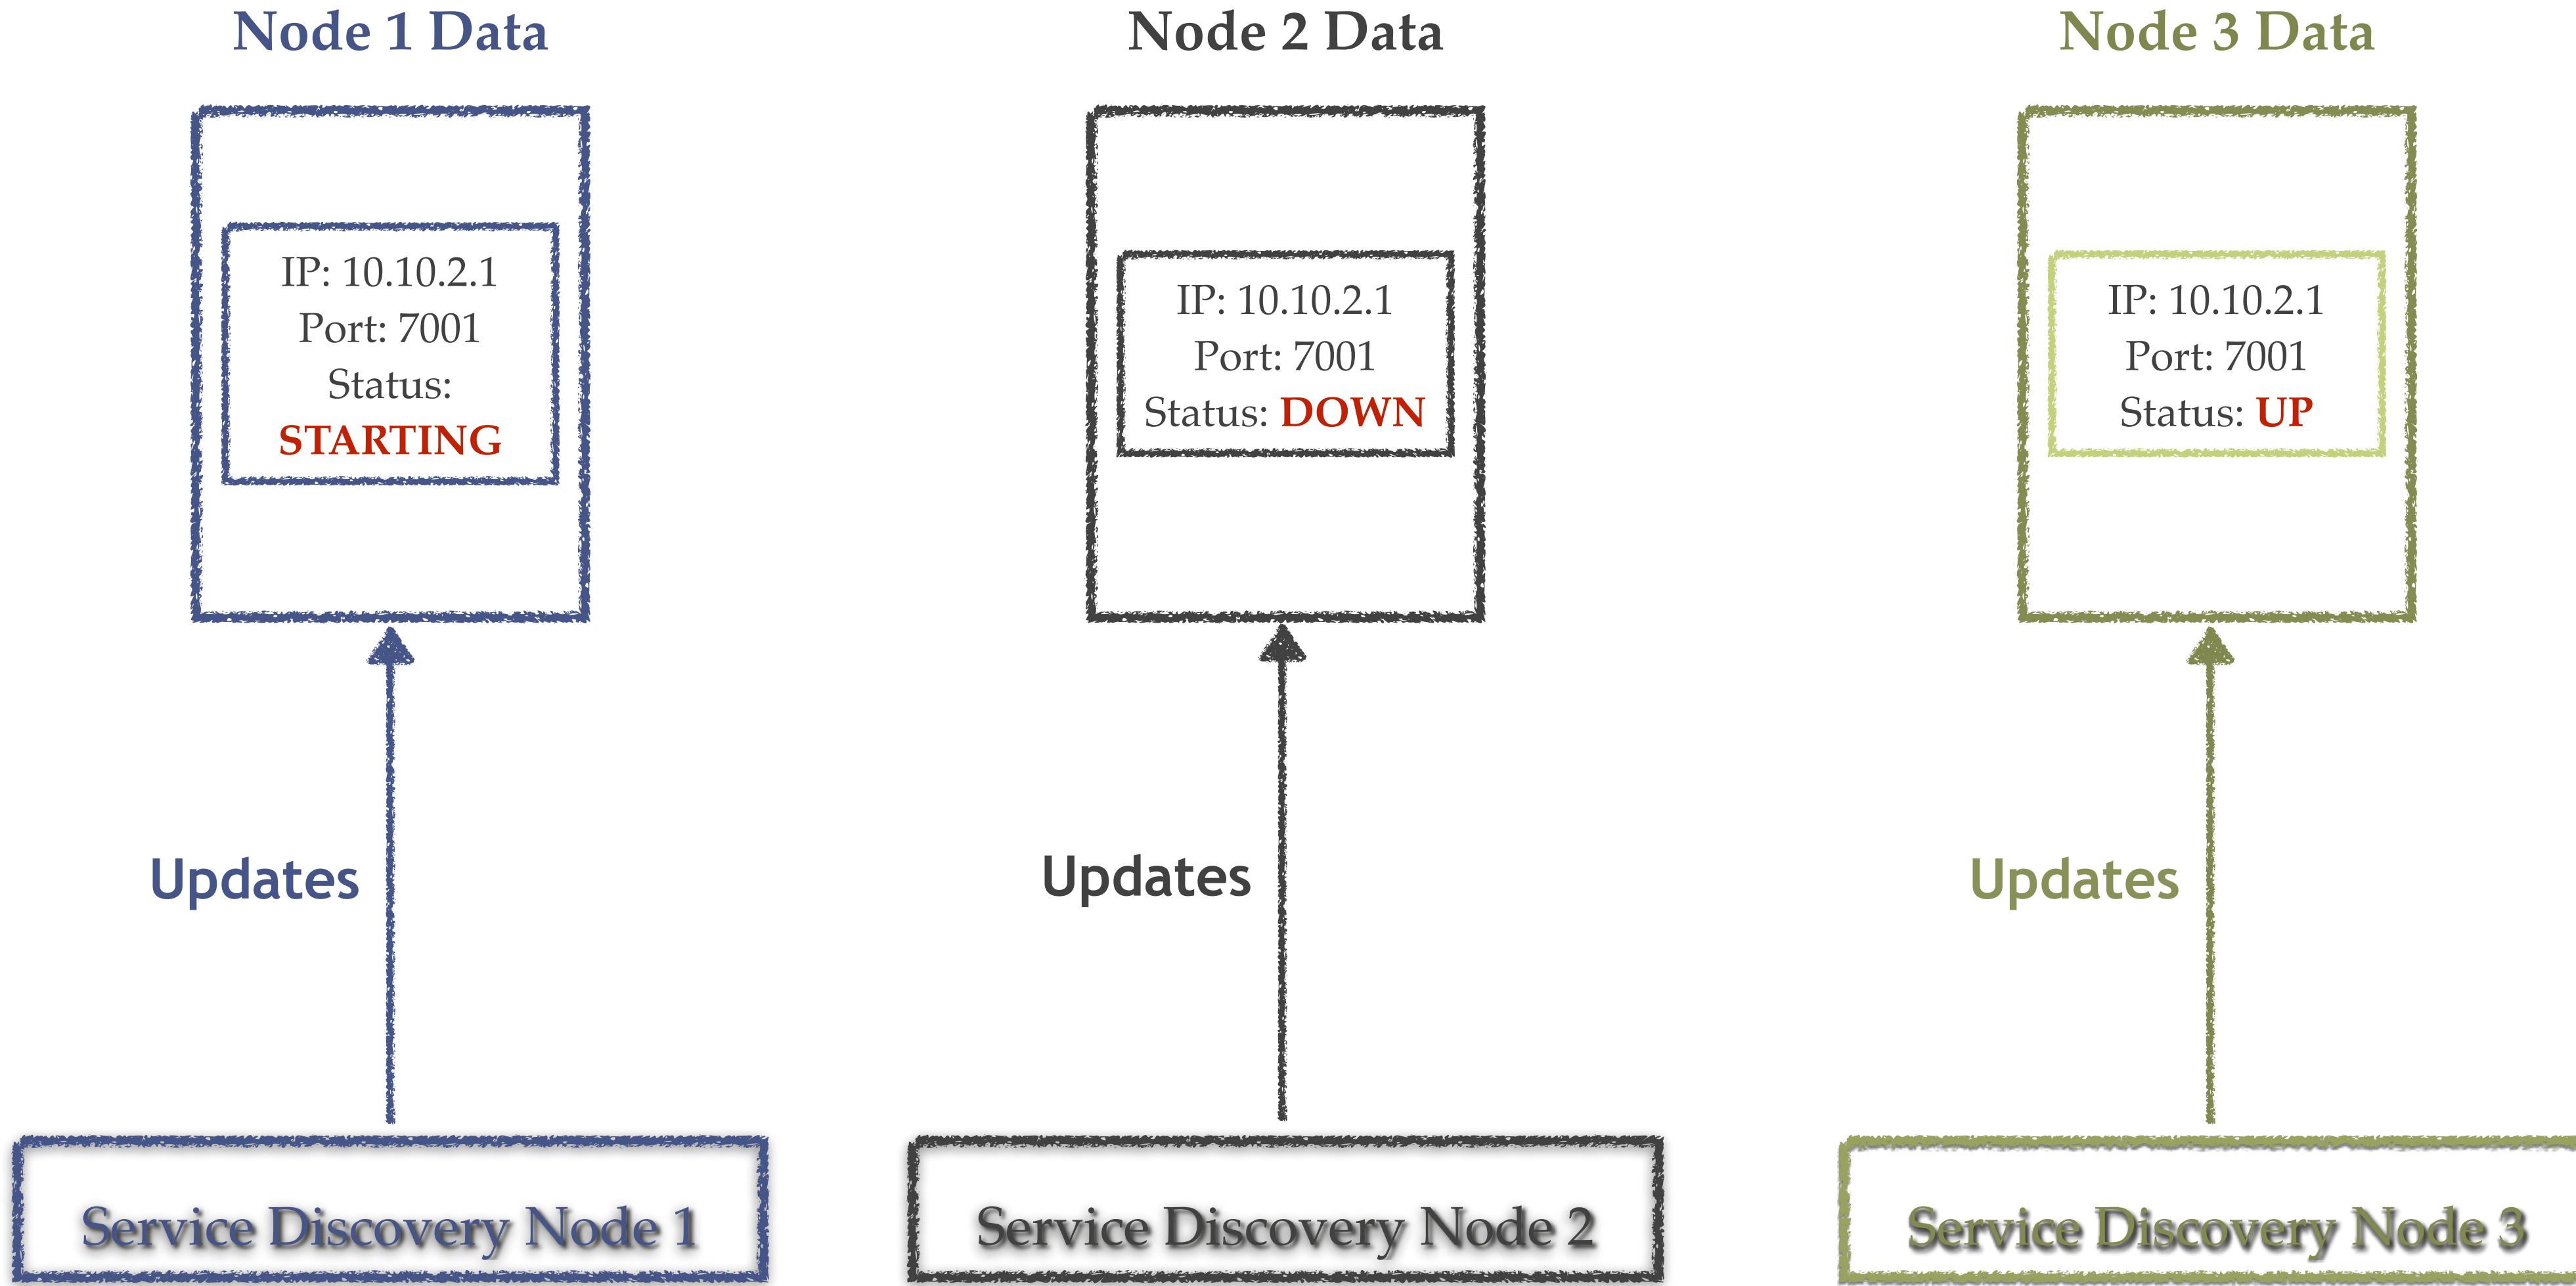
Well behaved clients at steady state connects
to a single Service Discovery node.

Conflicts

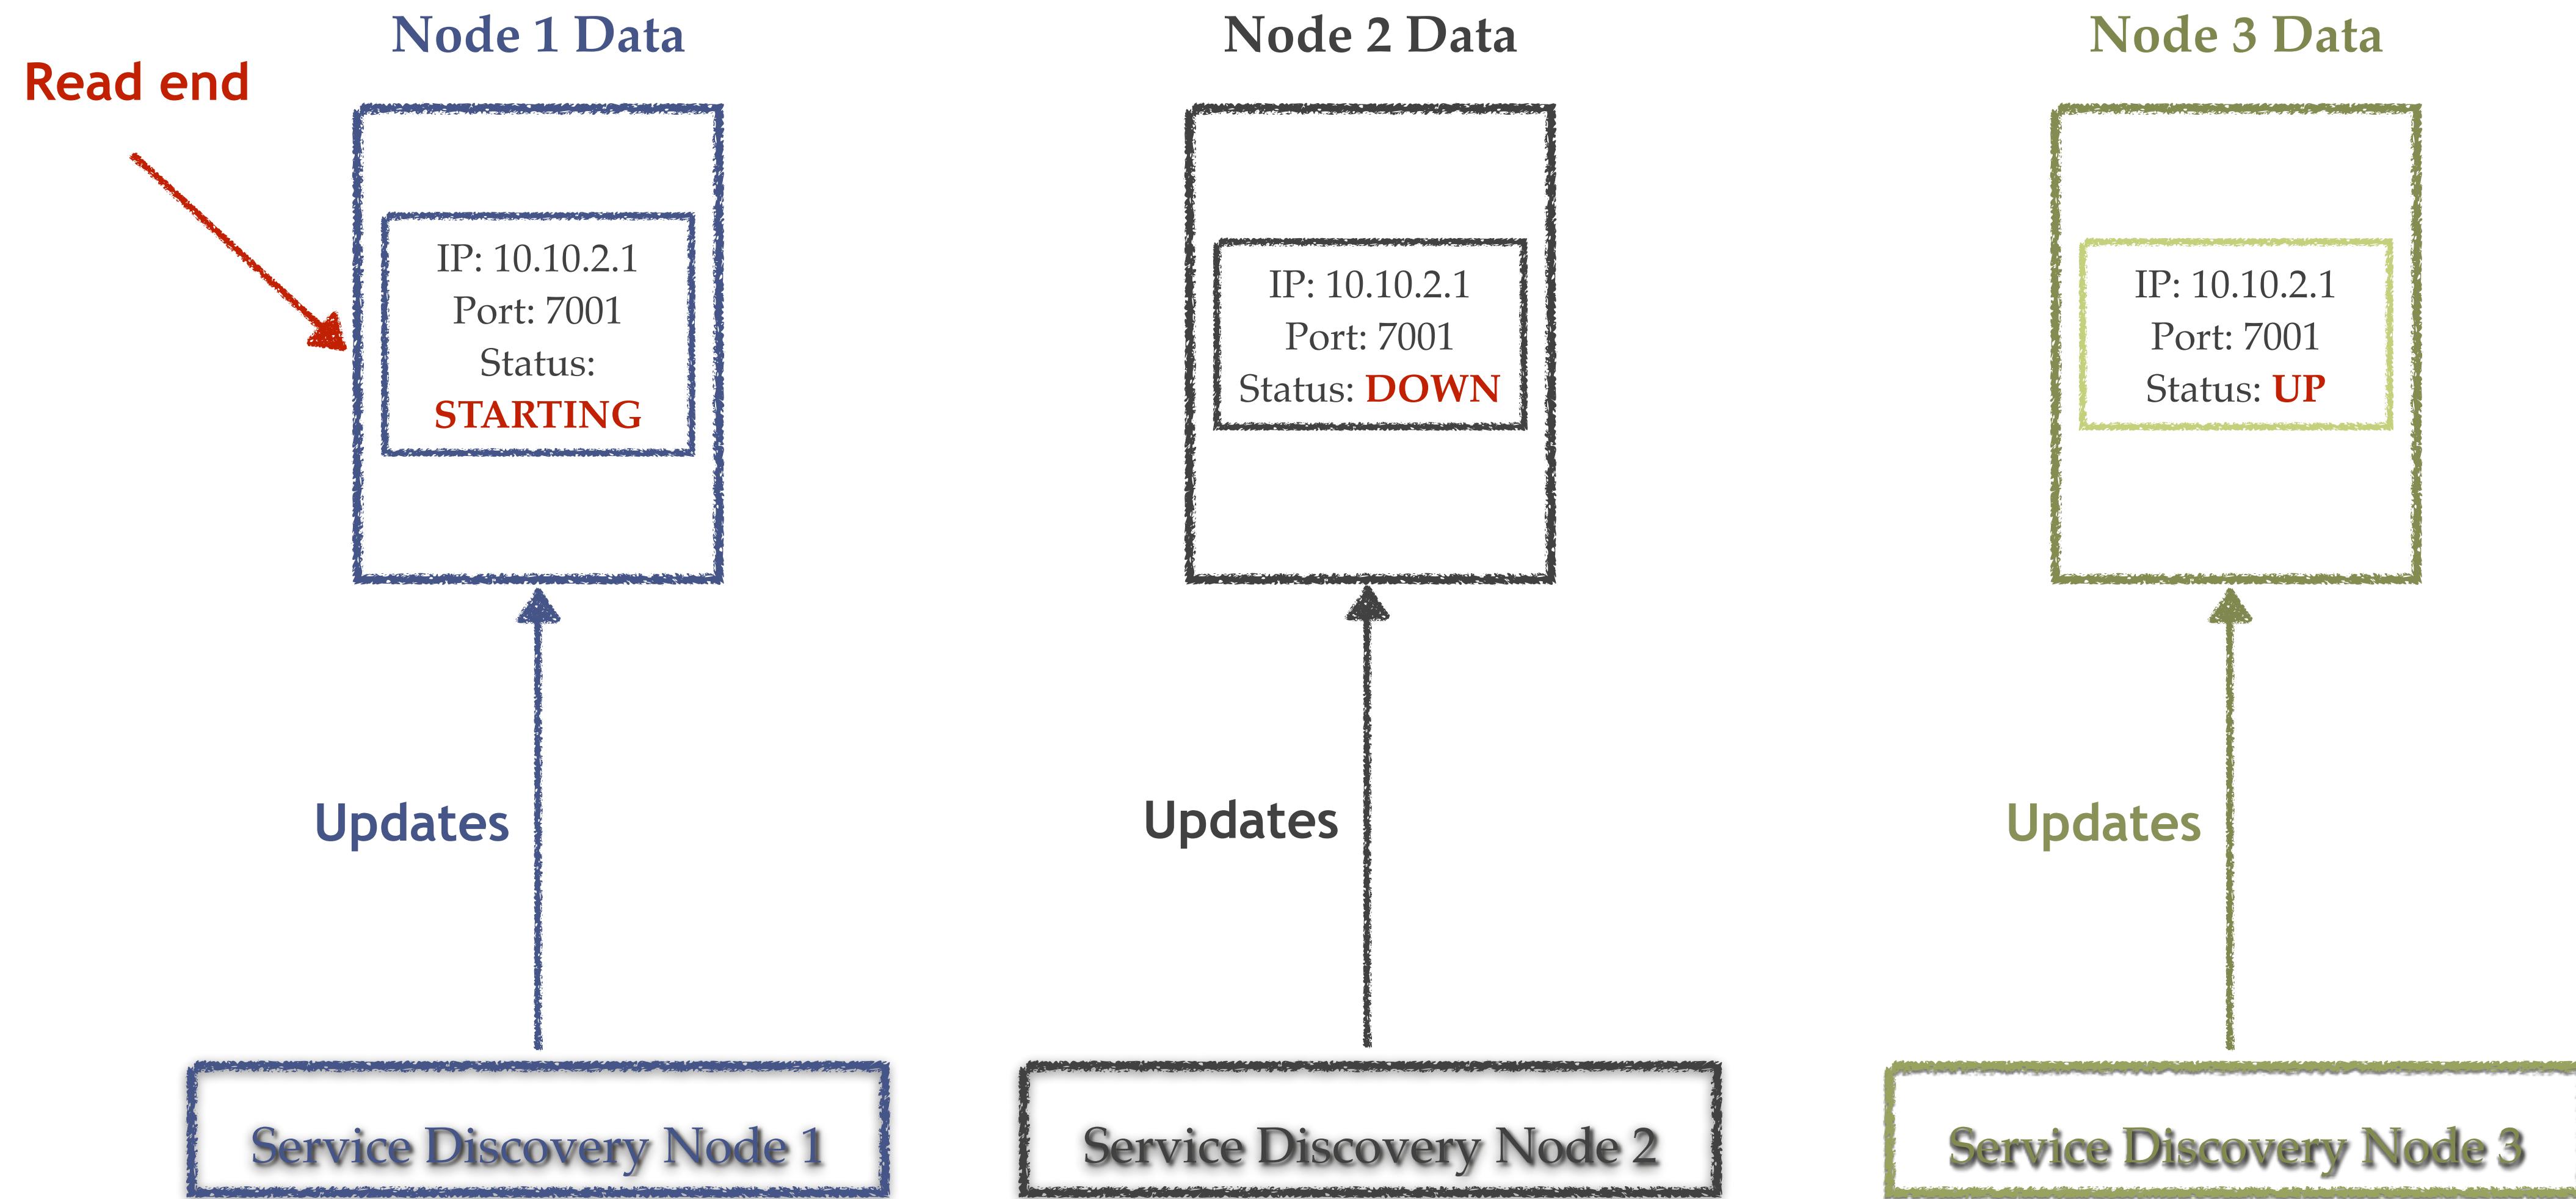


Data conflicts are resolved naturally
for service discovery.

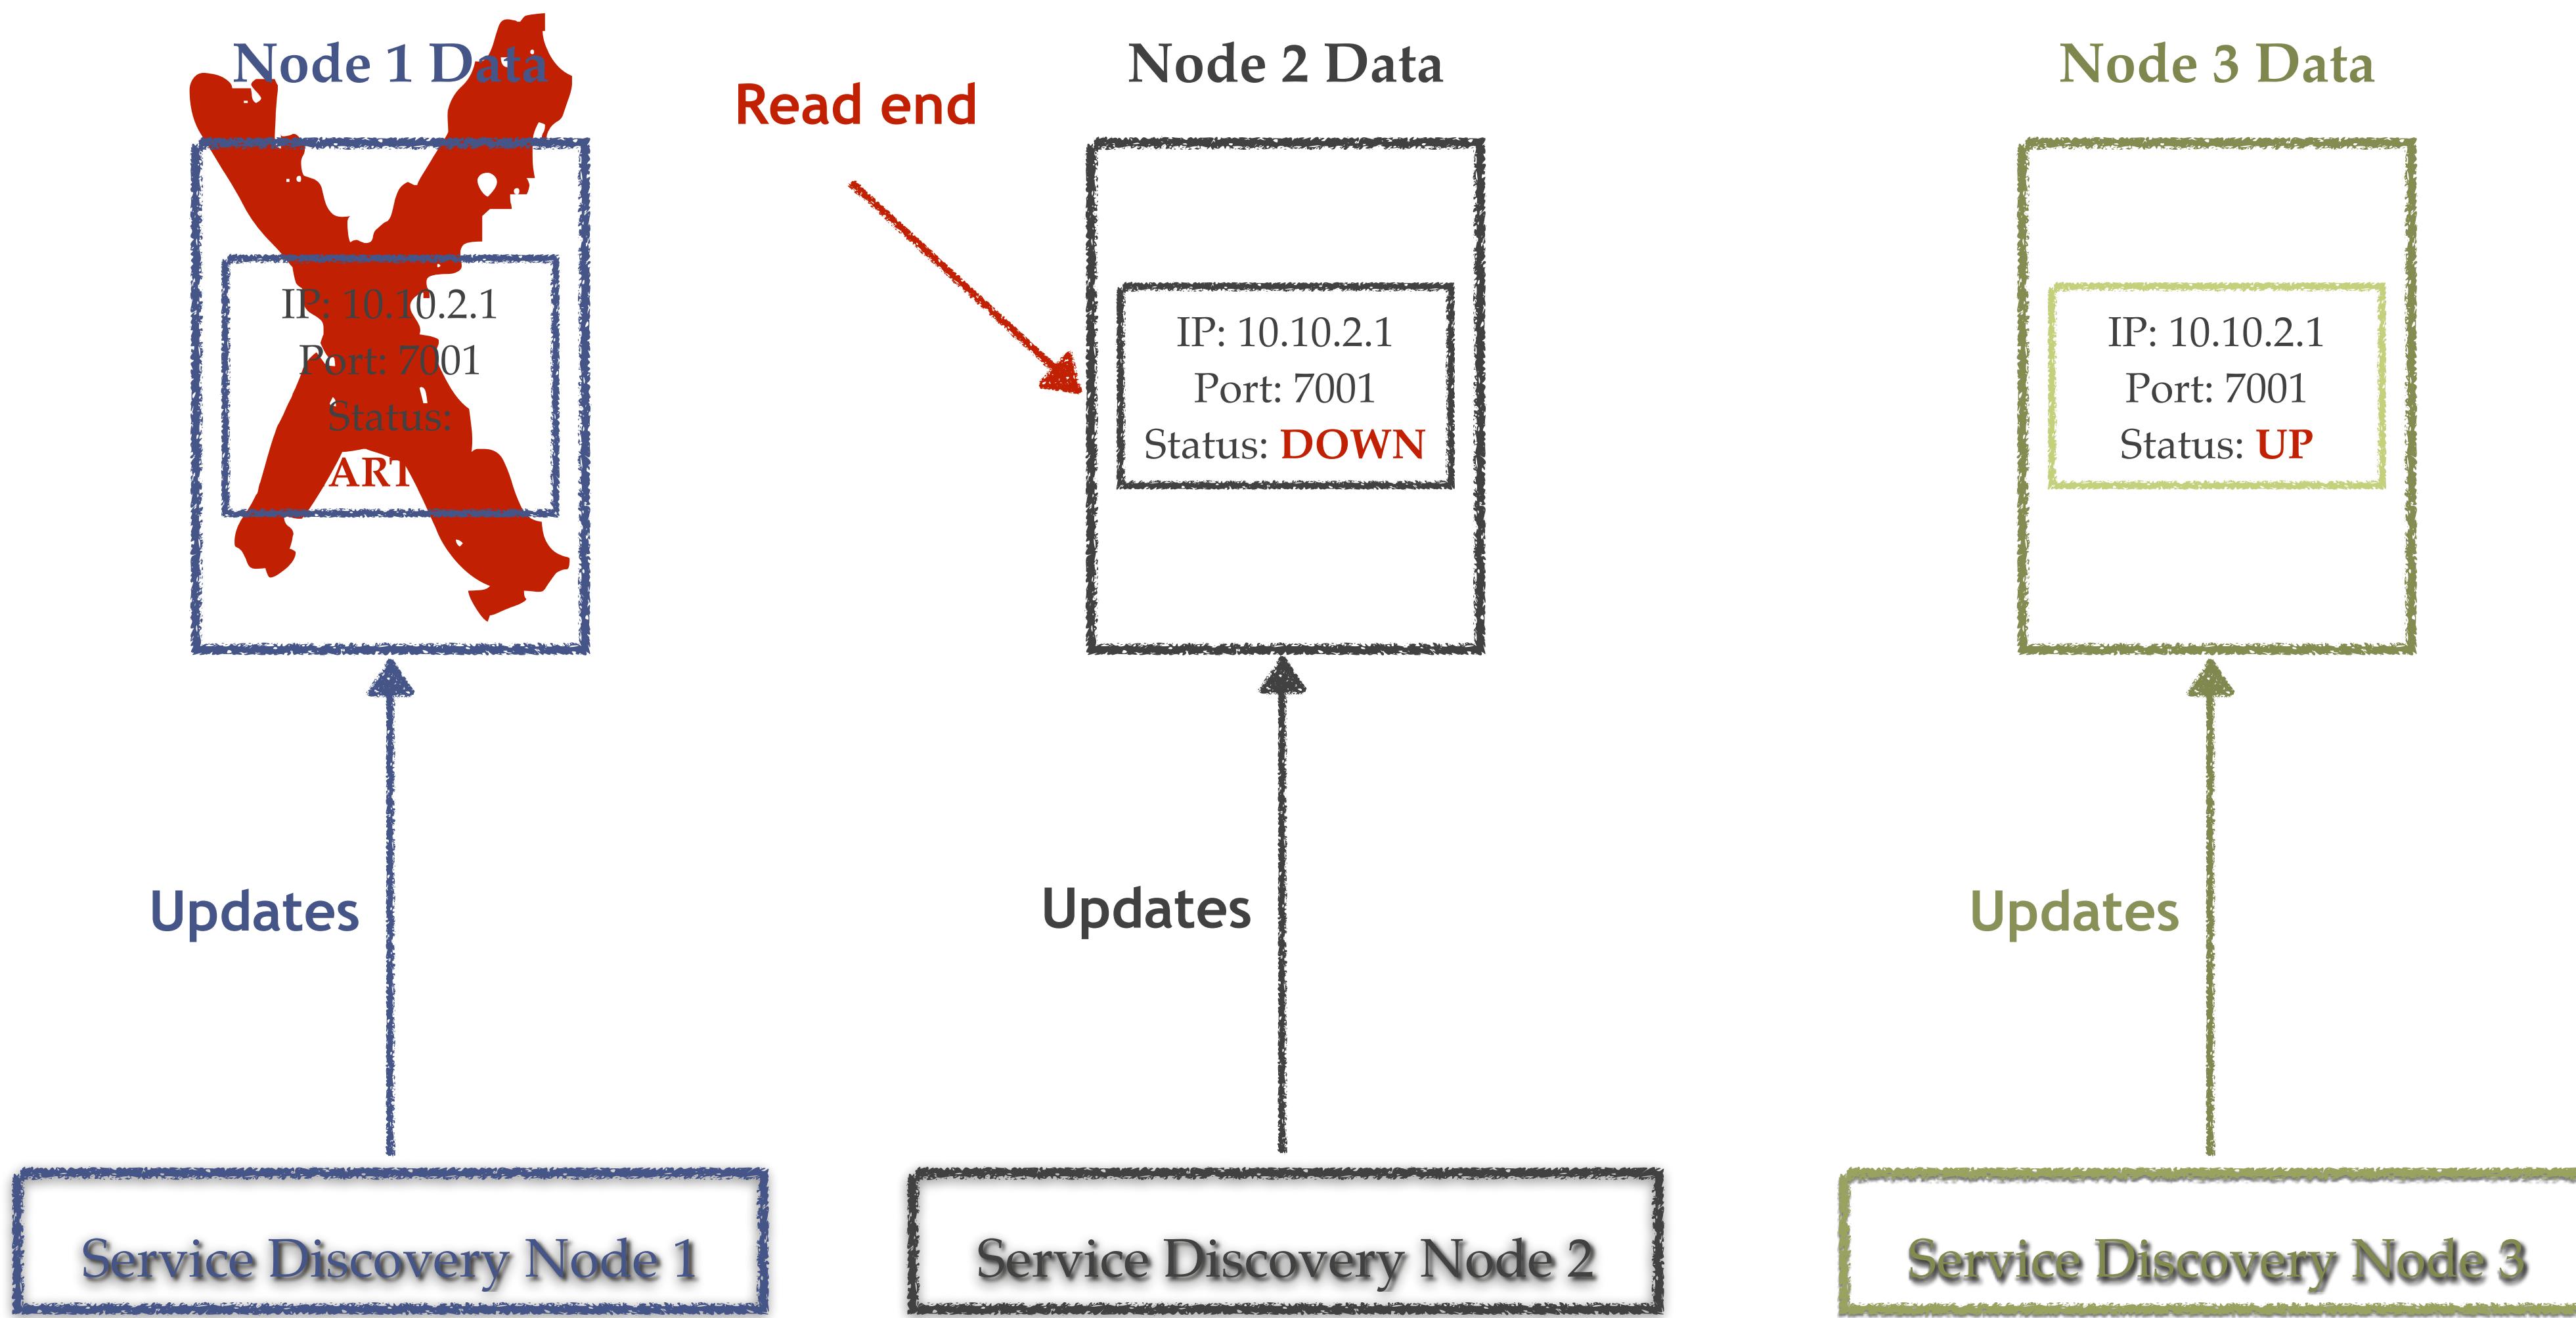
Tolerate temporary conflicts



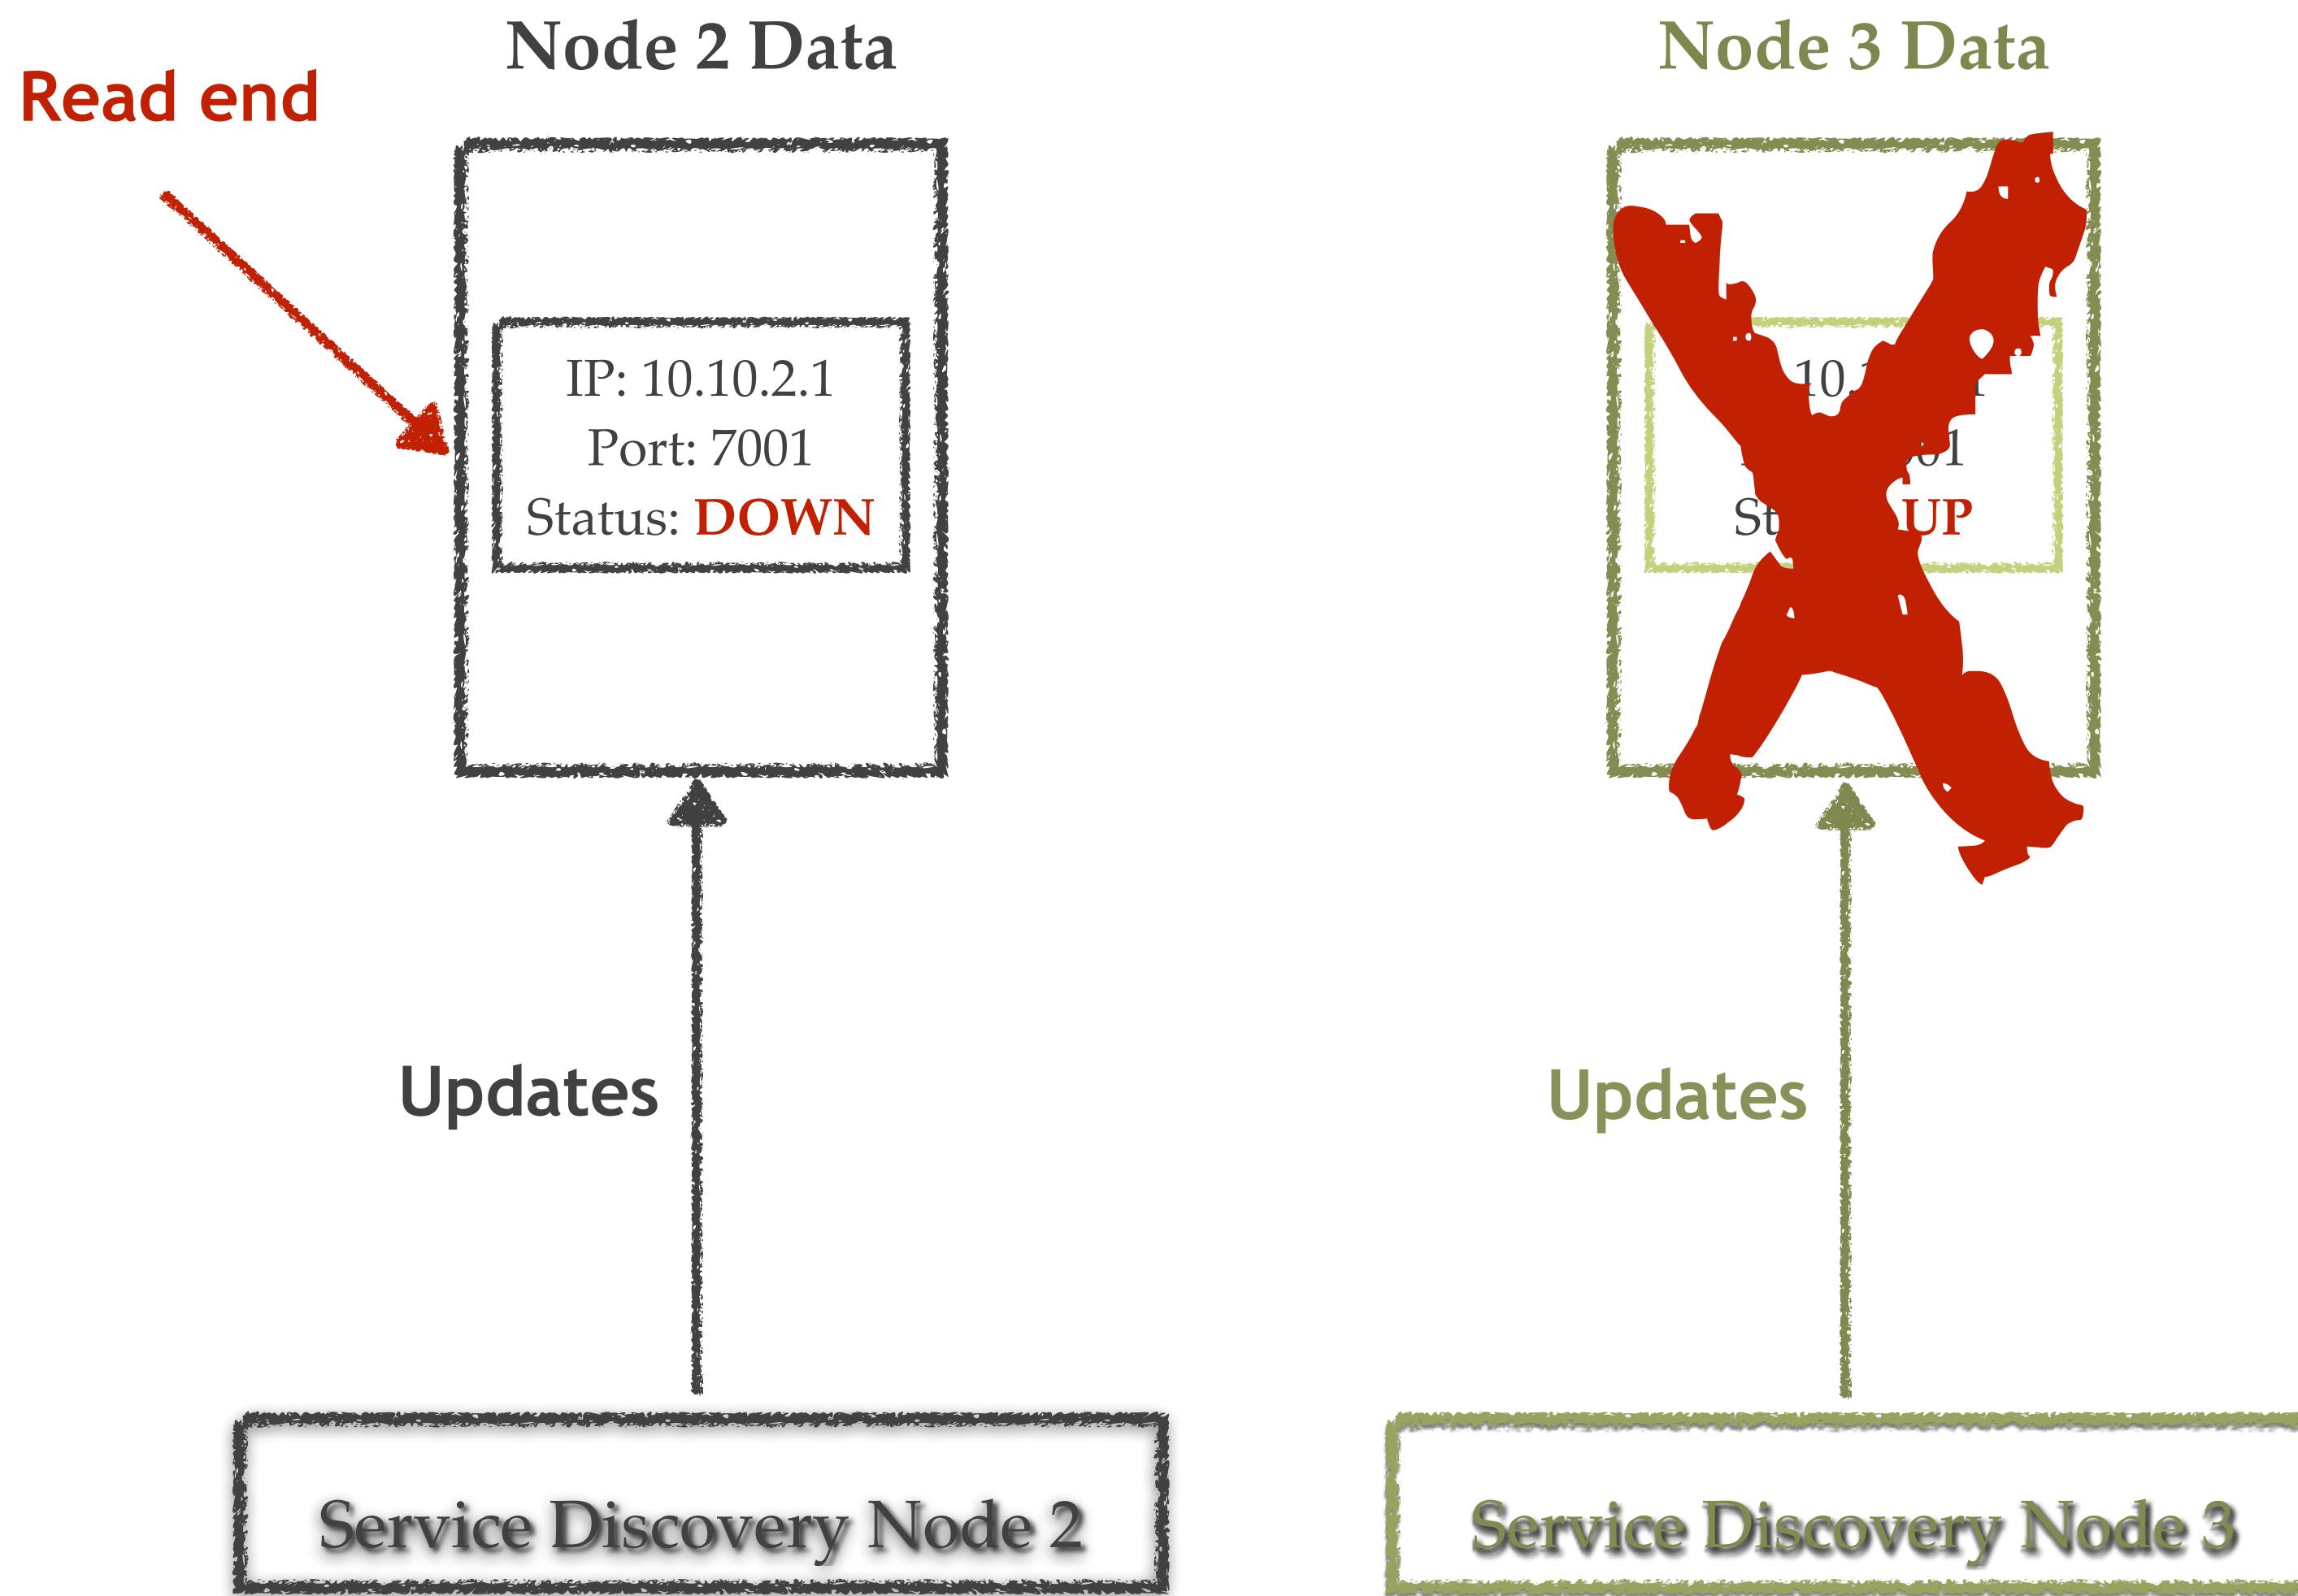
Read from a version (till it is gone)



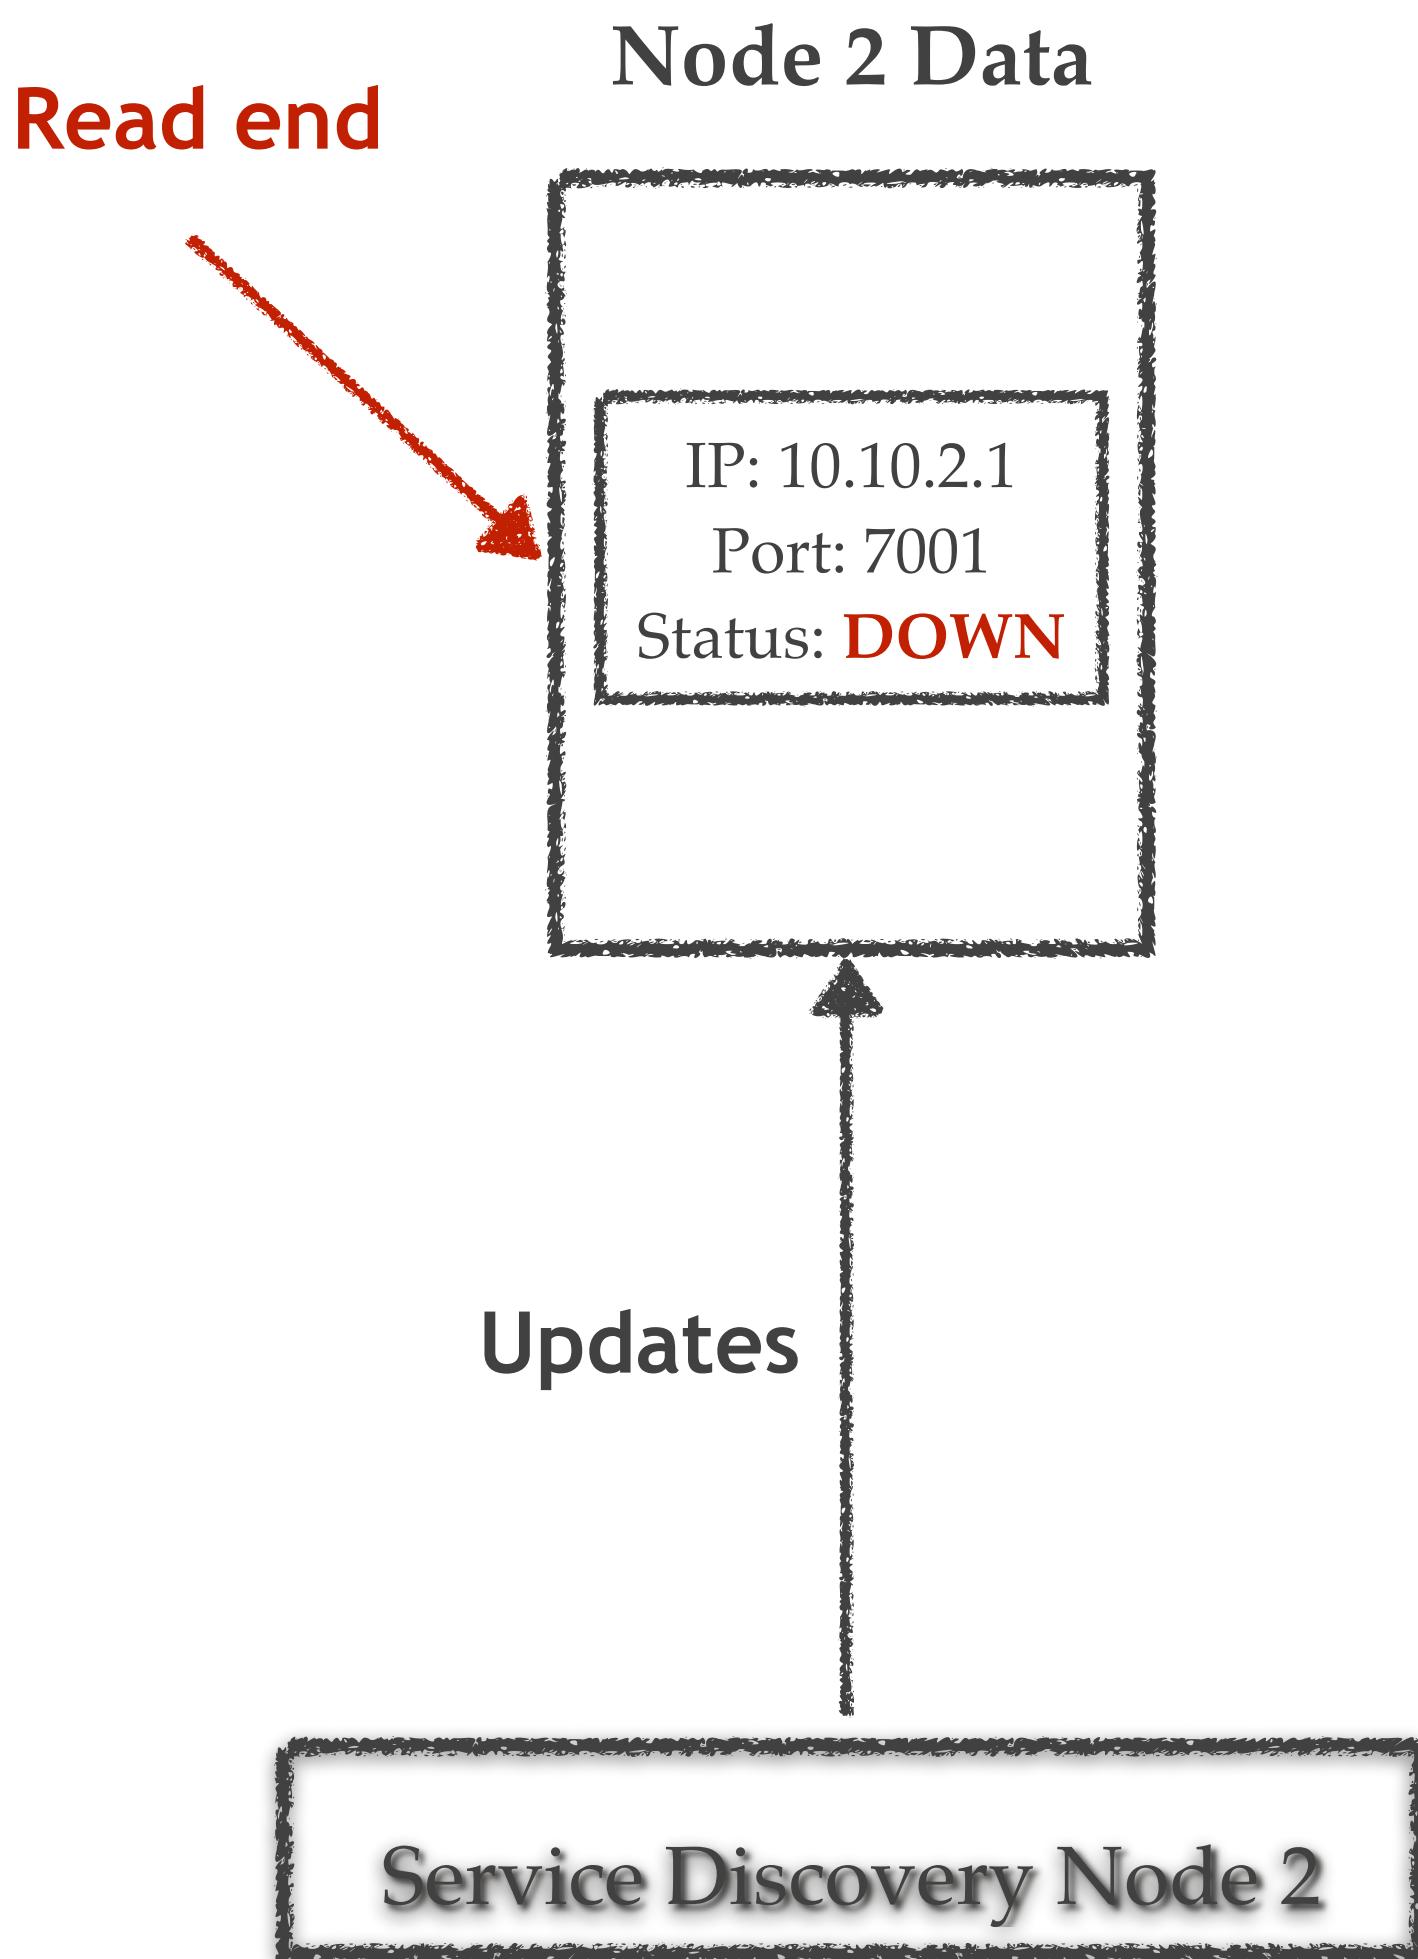
Read from a version (till it is gone)



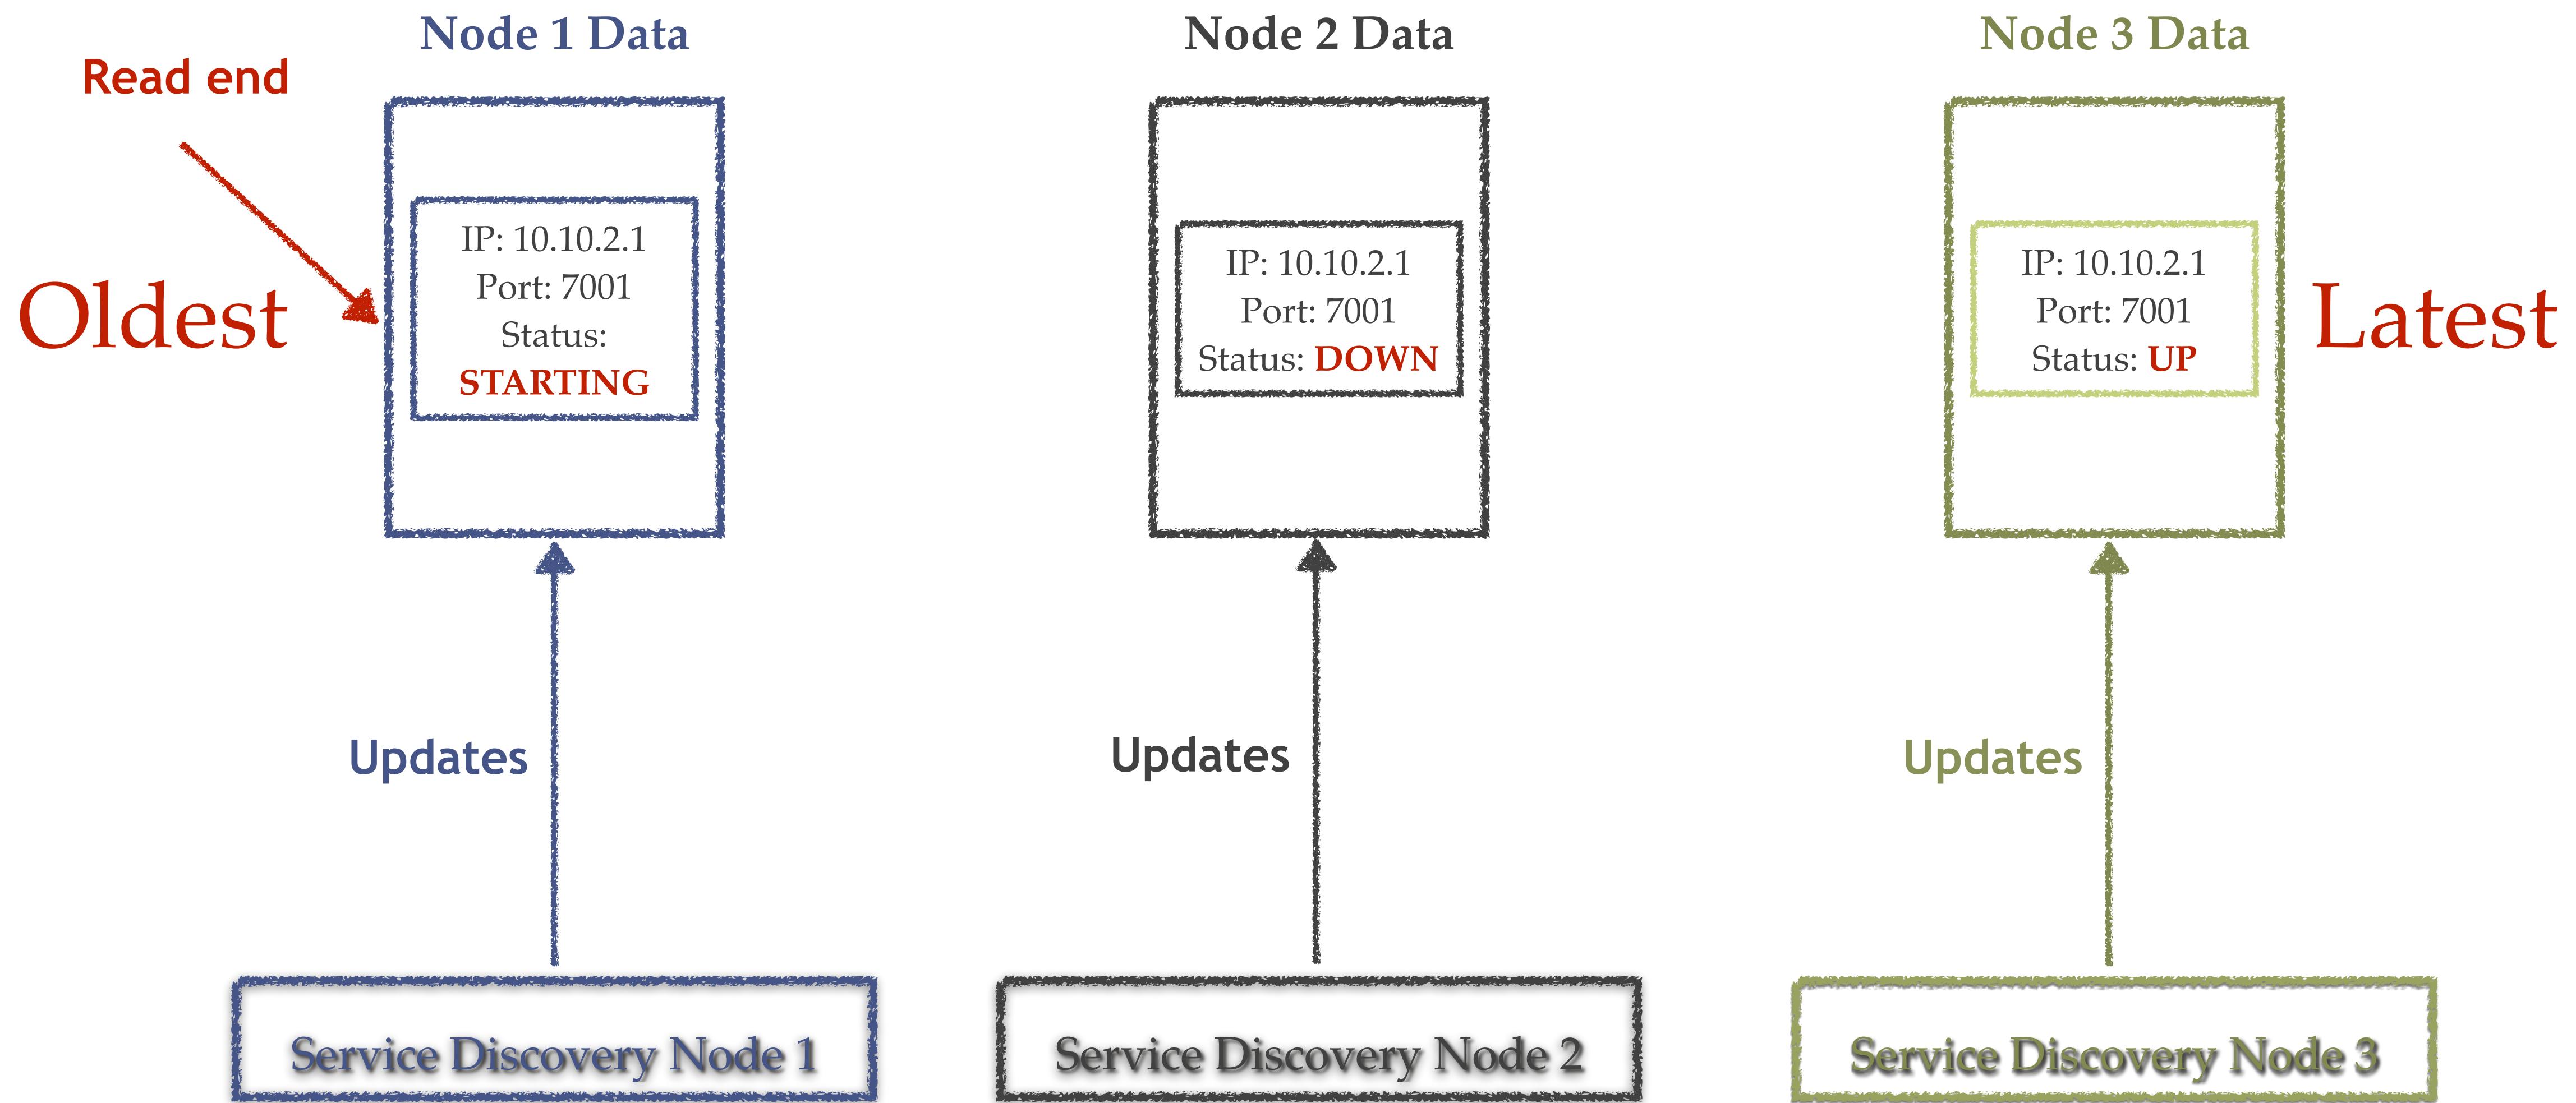
Read from a version (till it is gone)



Steady State



Time to converge (worst case)



Time to converge (worst case)

Time to evict stale copies. (*Constant*)

+

Time to replicate from the owner node.

Time to evict stale copy

Heartbeat interval

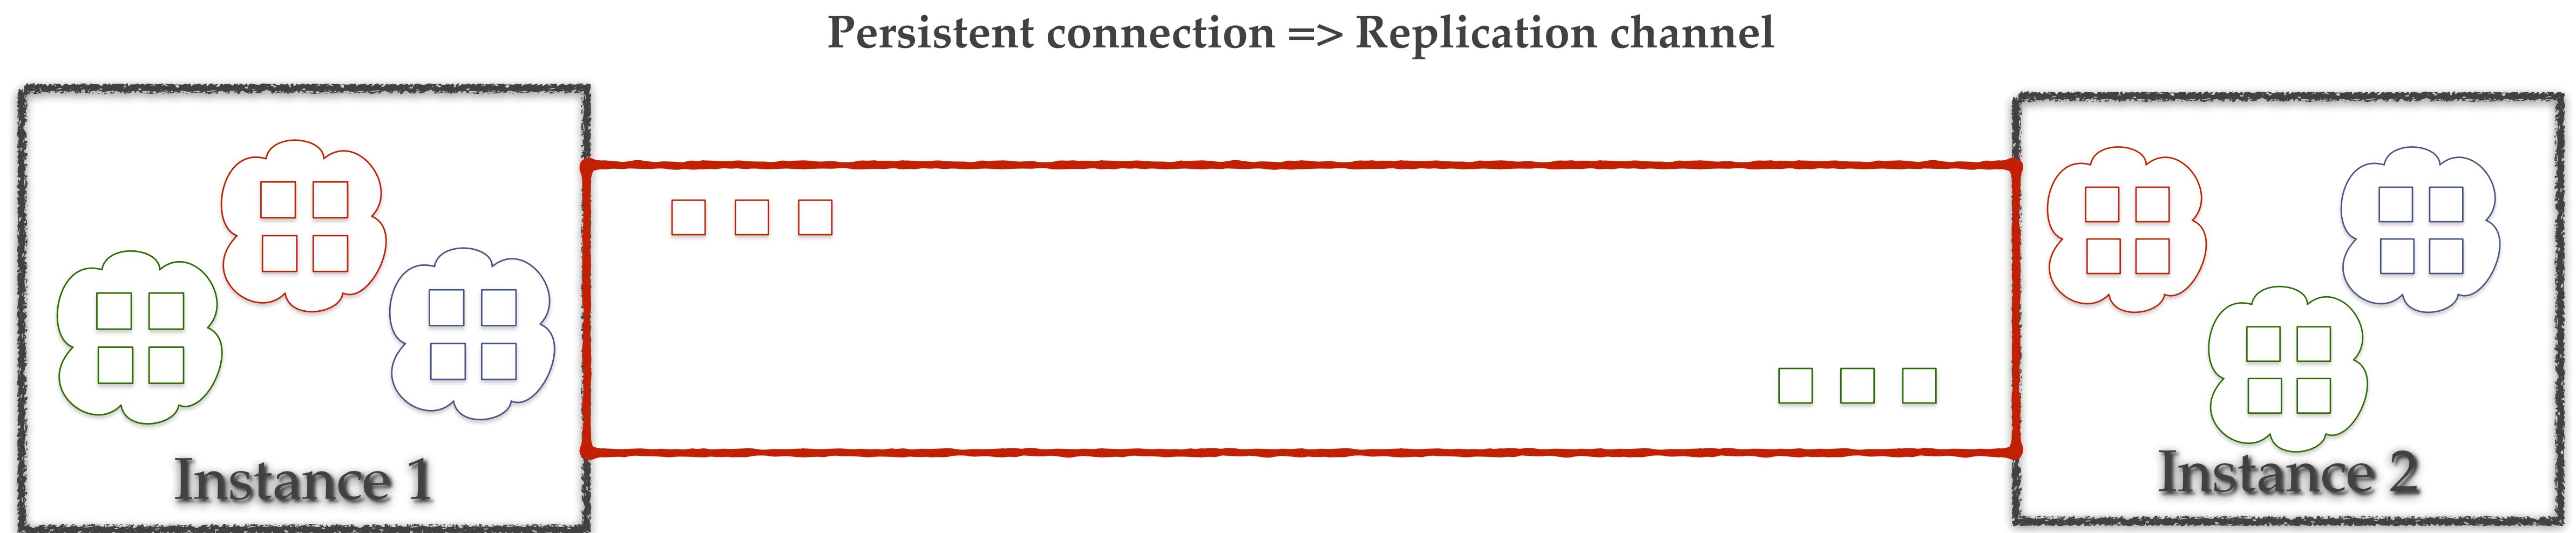
*

Number of tolerated missing heartbeats

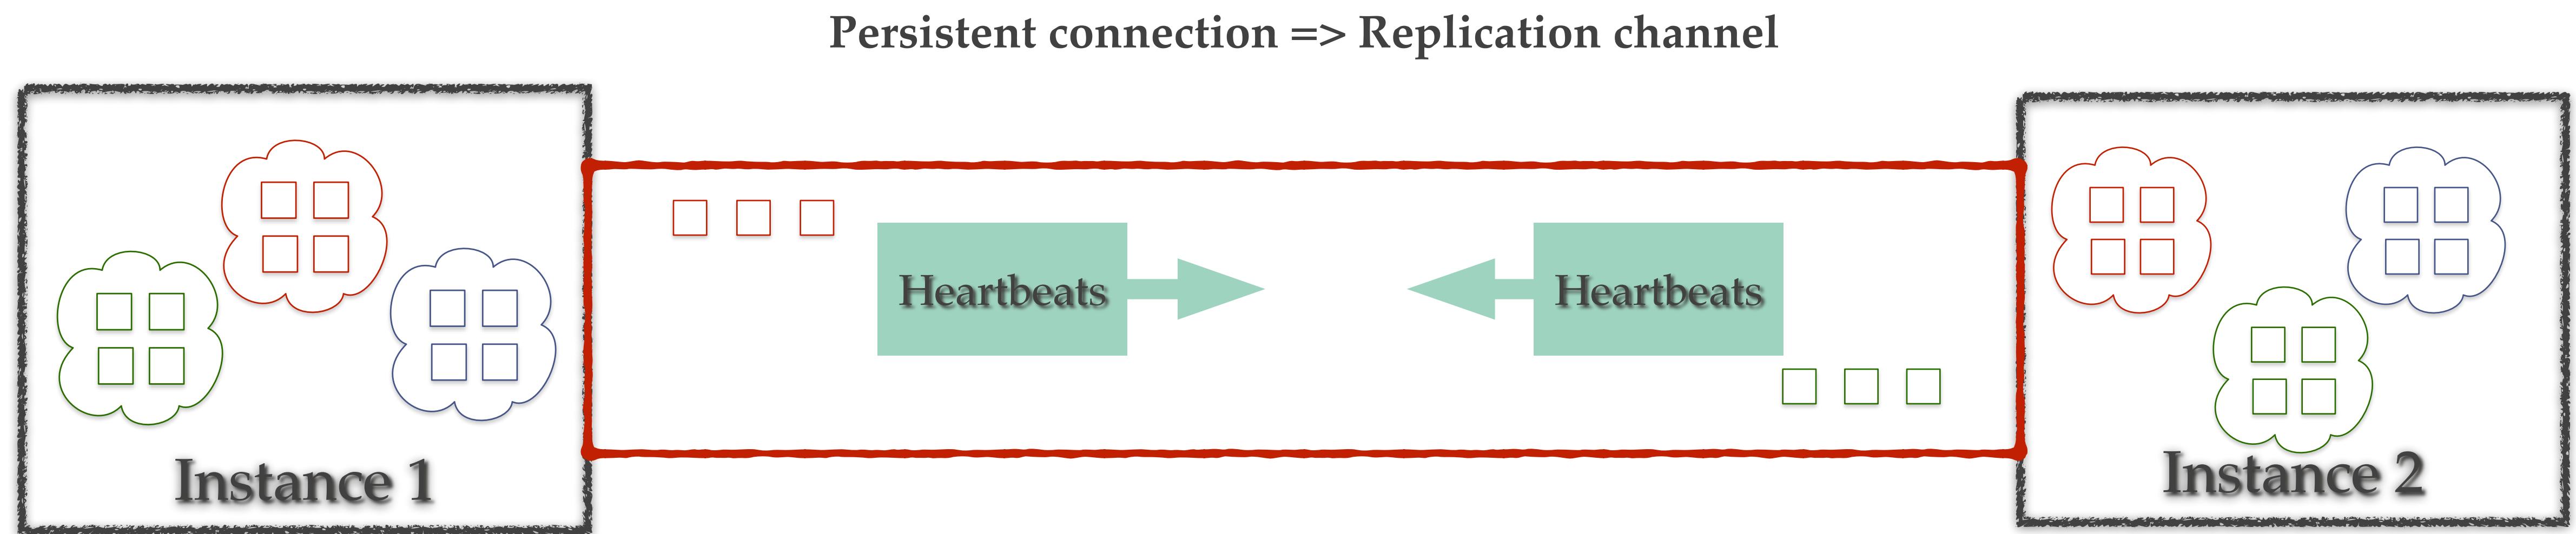
Time to replicate from the owner node

Theoretically unbounded

Adding bounds



Adding bounds



Time to replicate from the owner node

Somewhat bounded

Heartbeat interval

*

Number of tolerated missing heartbeats

Time to converge (worst case)

Cost of divergence?

Cost of divergence?

Instance data is hardly changing!

In most cases,
there isn't a divergence
and when it happens,
the impact is low.

Cost of divergence?

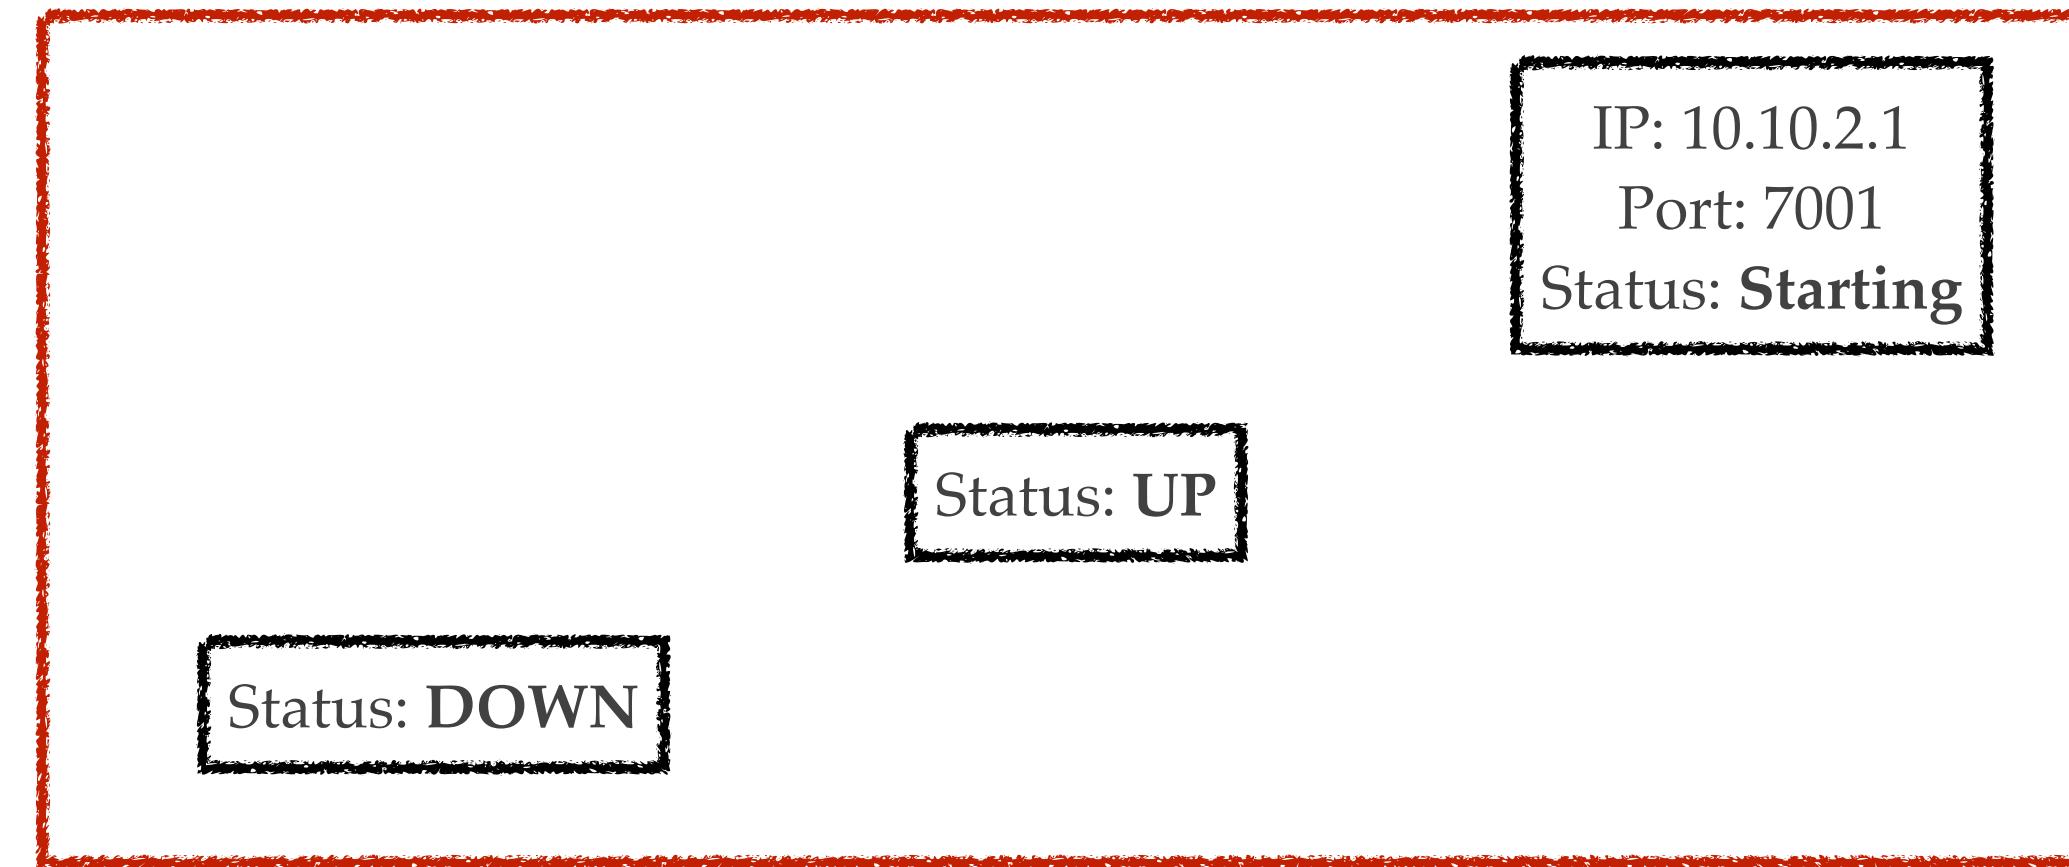
Service discovery controls visibility of nodes but does not guarantee availability.



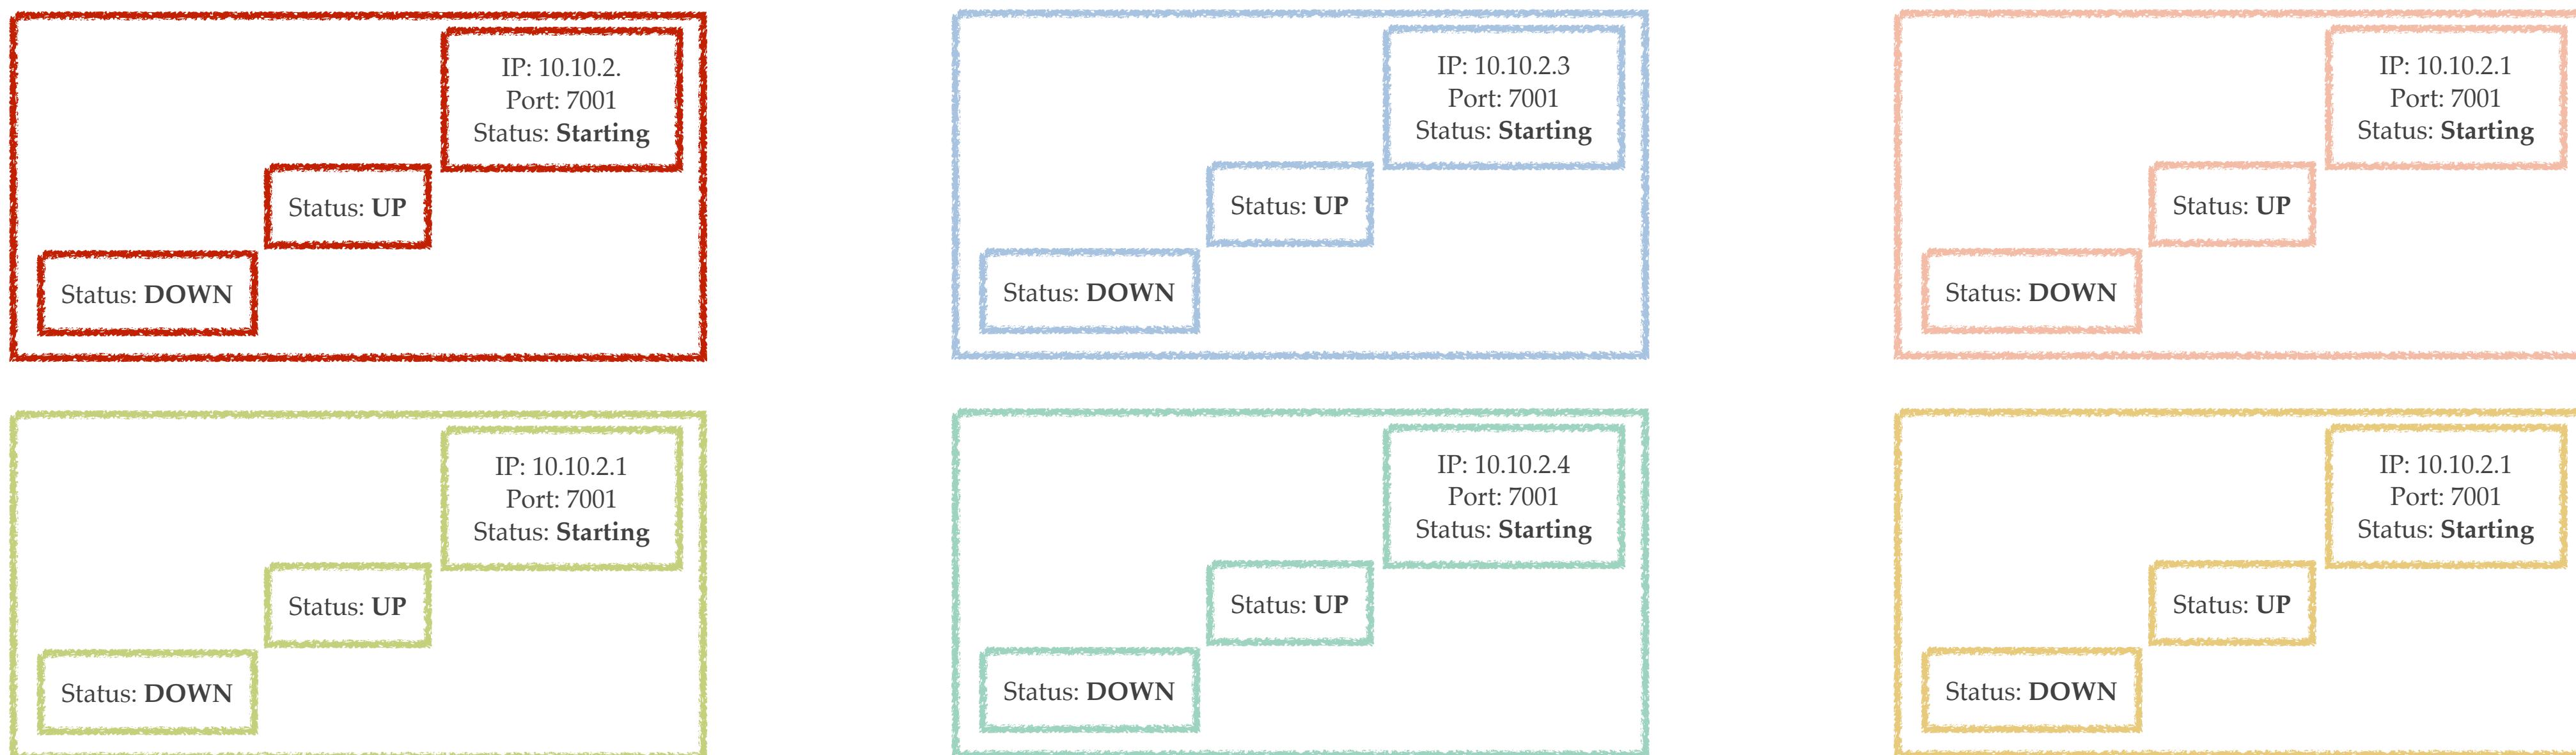
Reads

How to implement reads on service discovery data?

Service discovery data is an ordered stream.

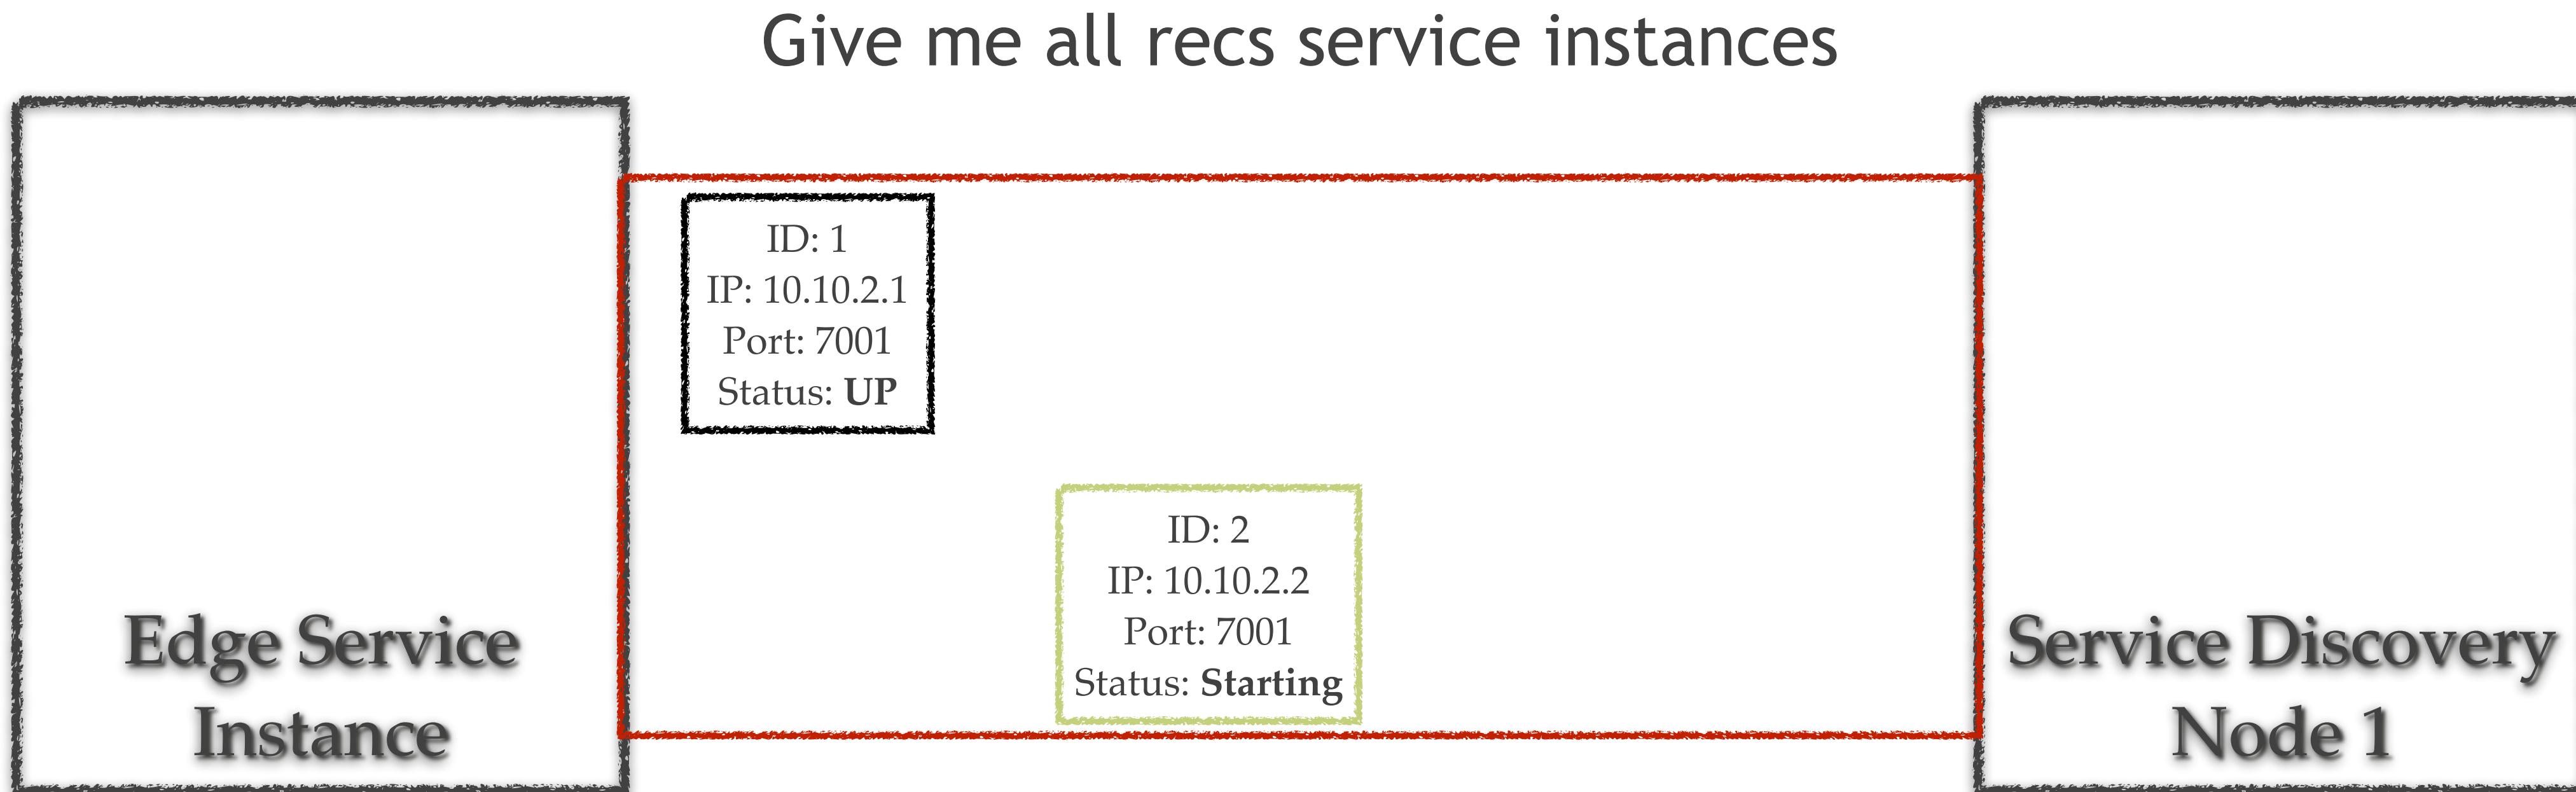


Service discovery data is a ~~ordered stream~~ merged ordered stream

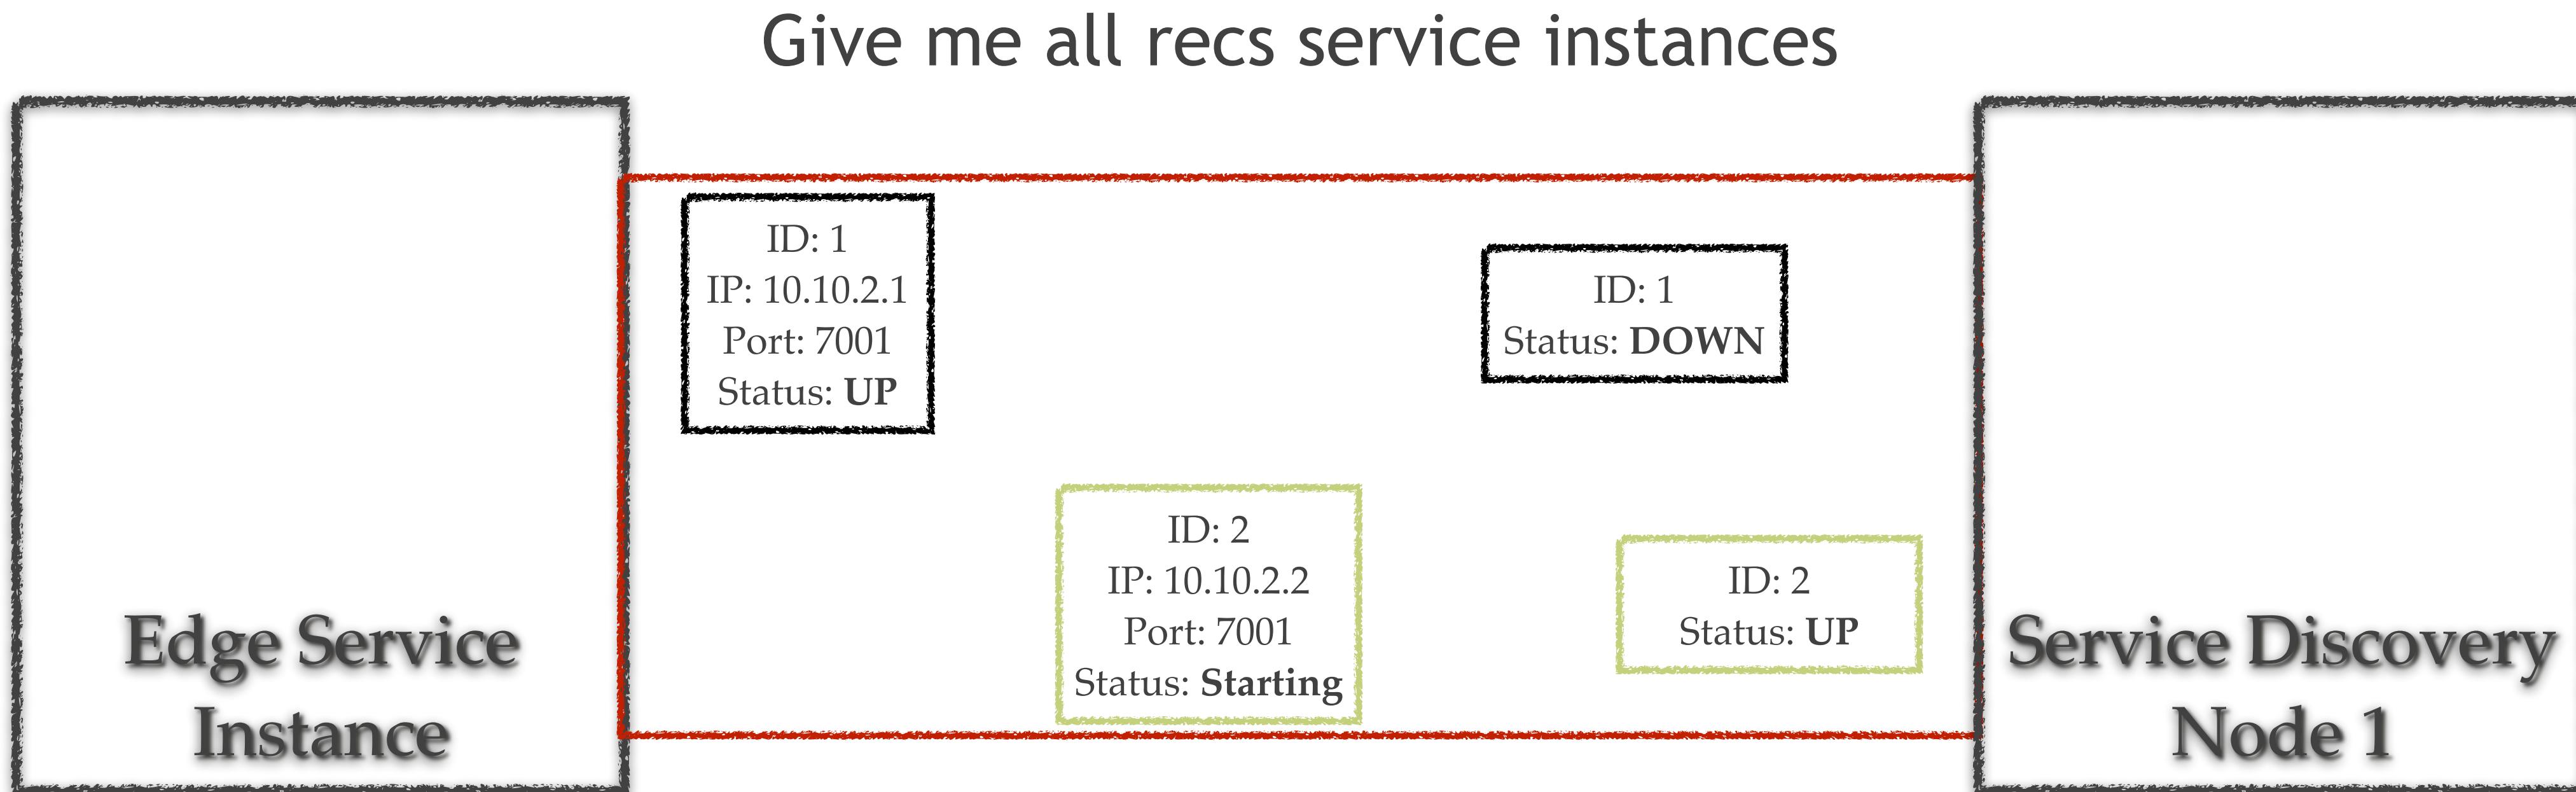


Data as a stream (Lookup)

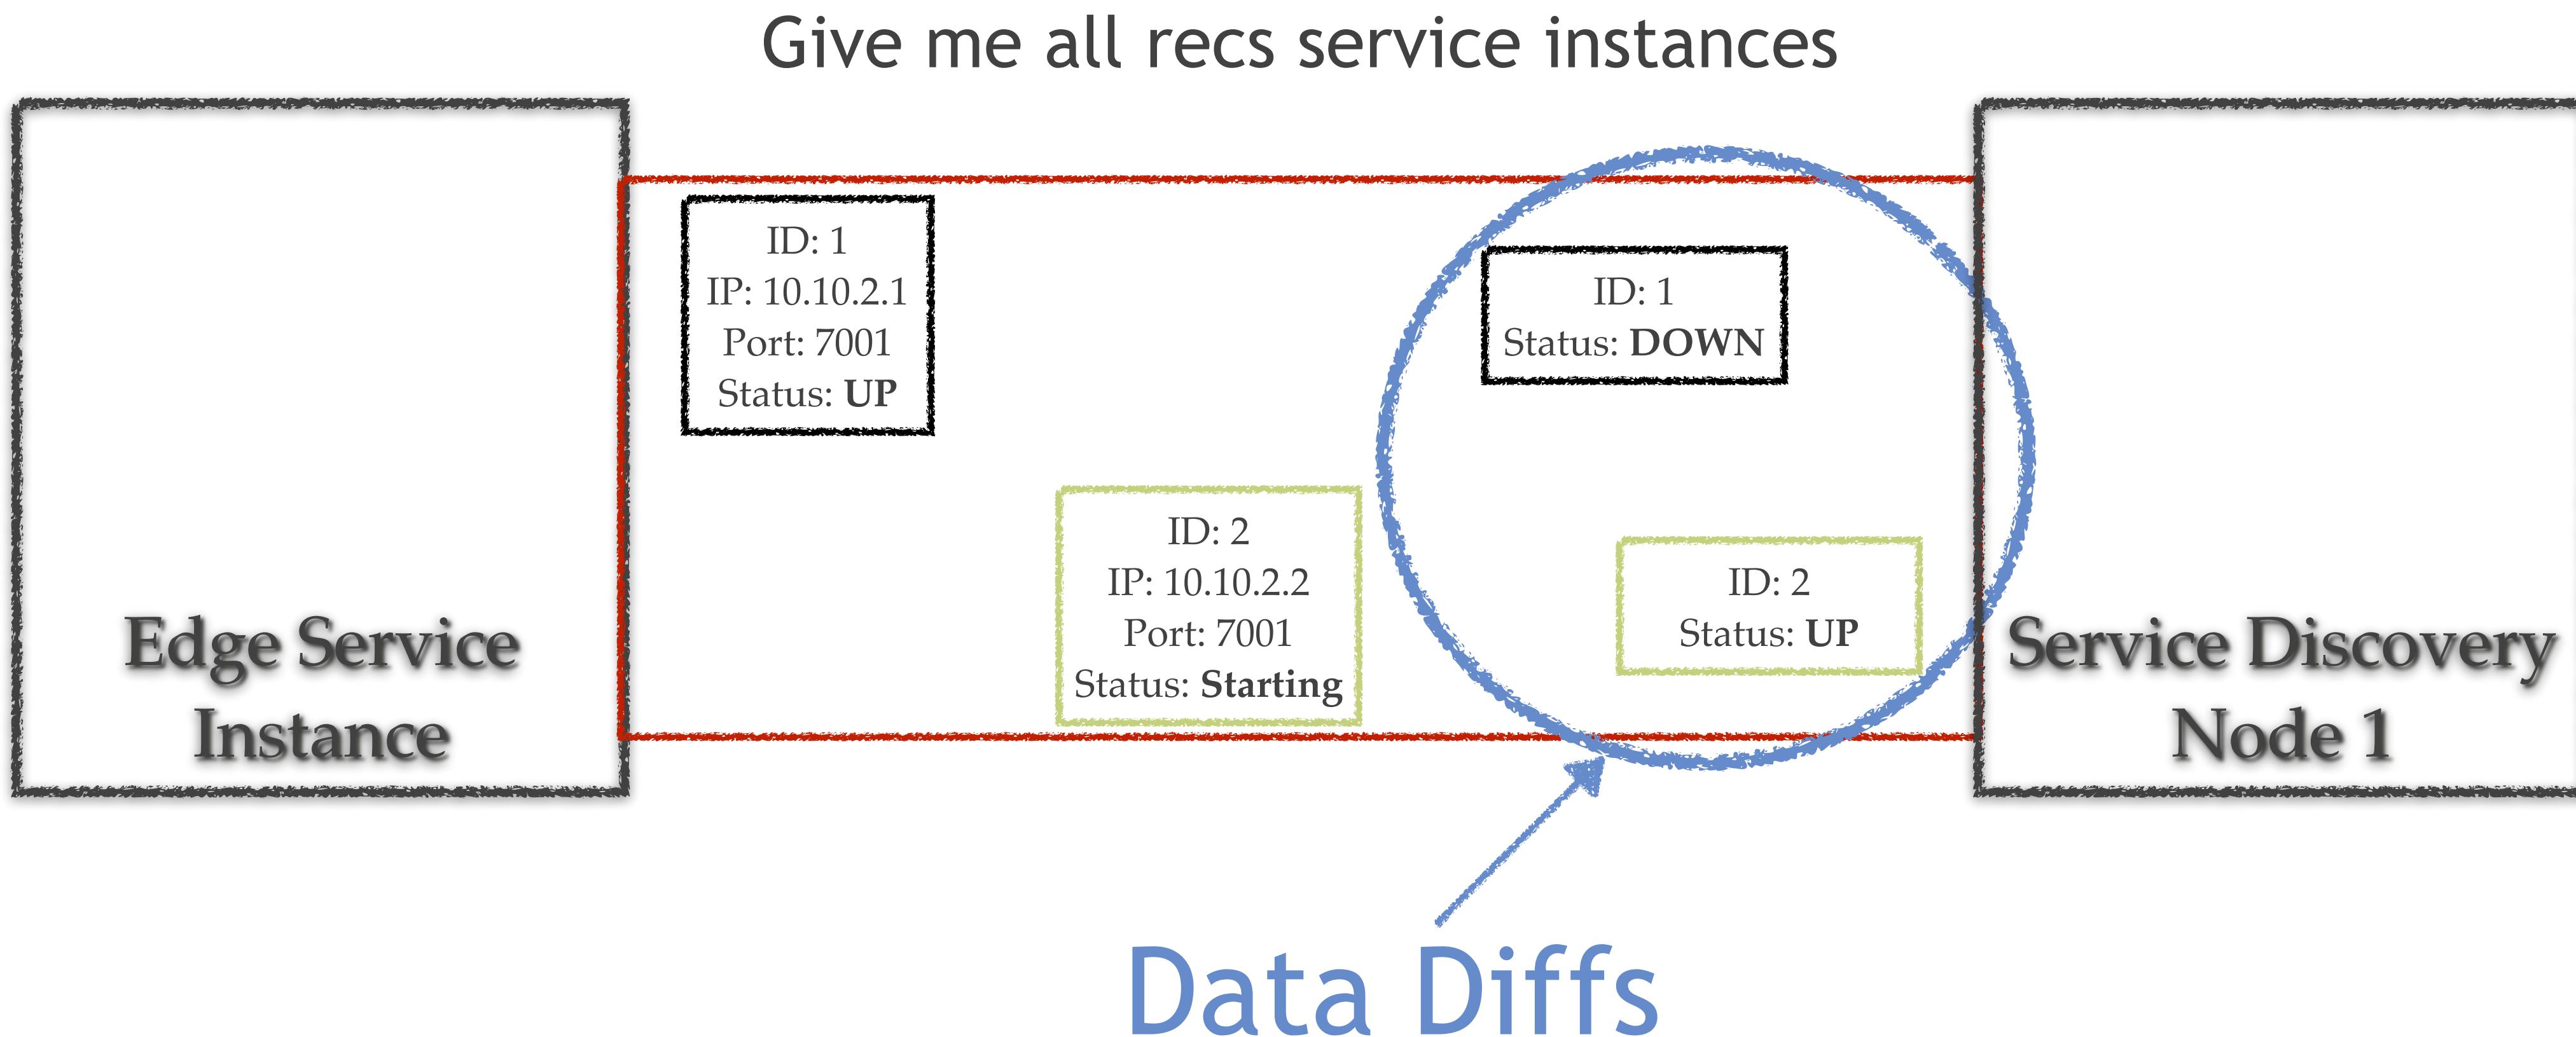
Data as a stream (Lookup)



Data as a stream (Lookup)



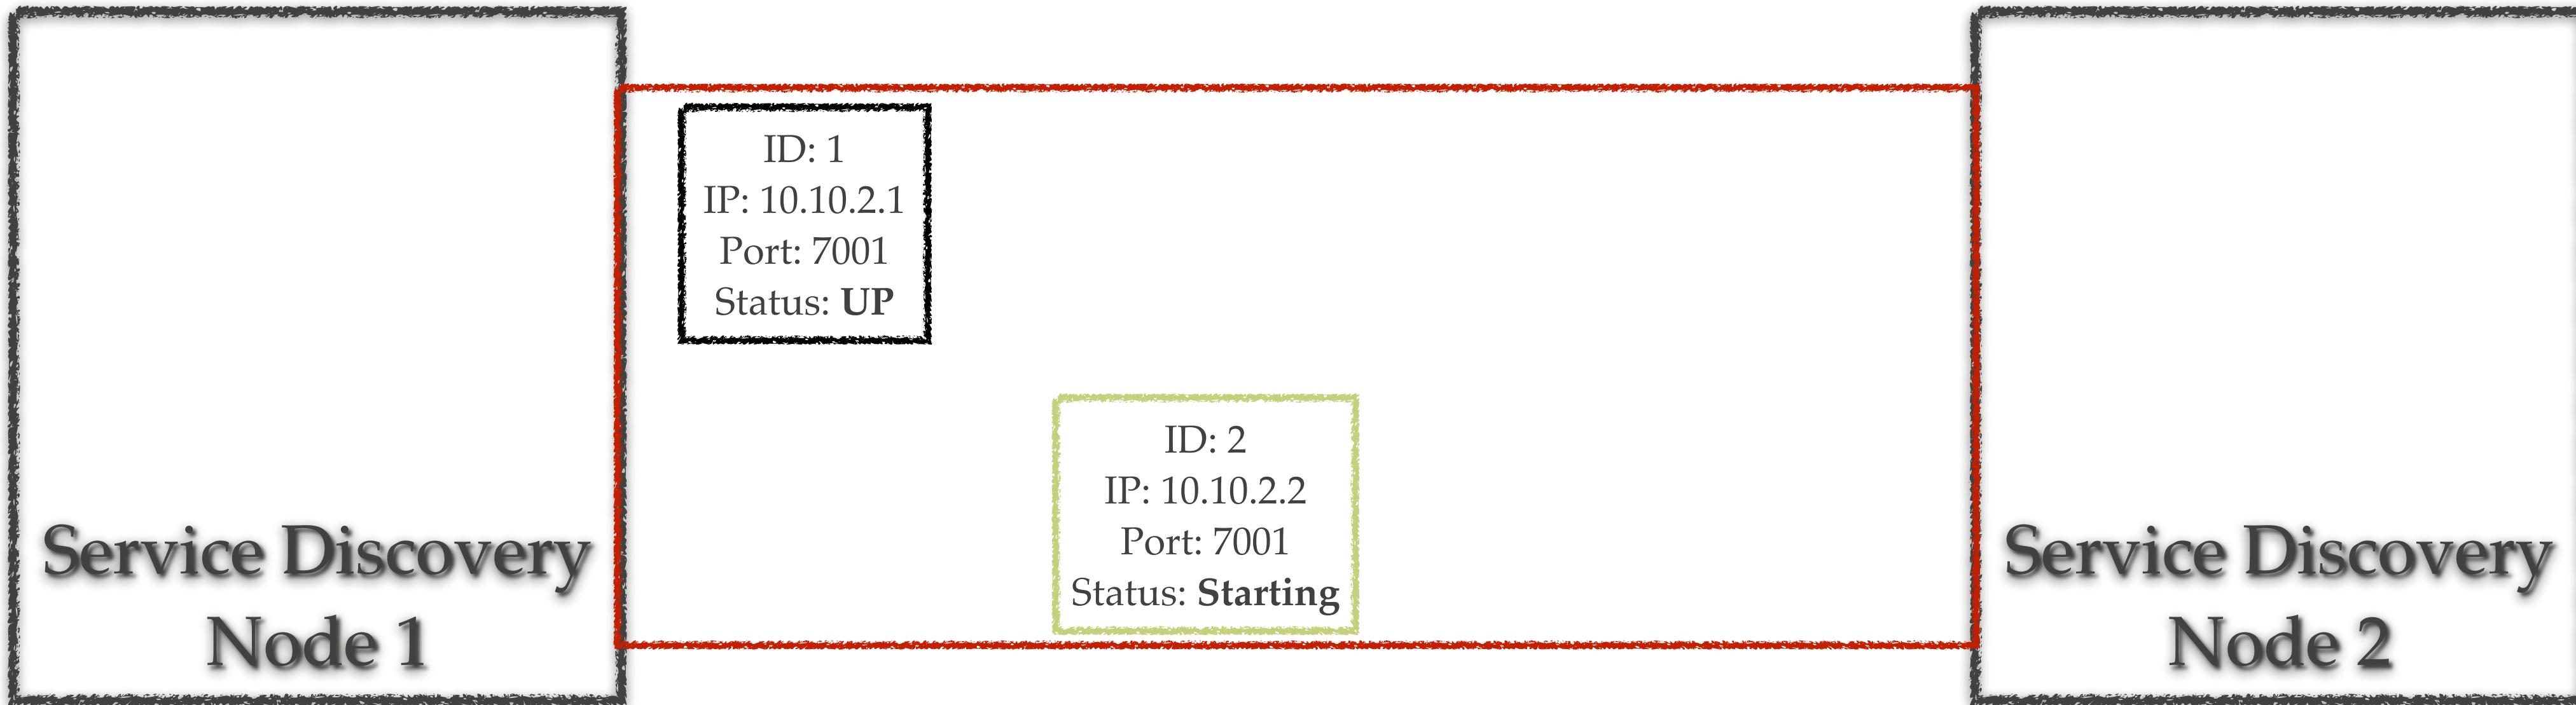
Data as a stream (Lookup)



Data as a stream (Replication)

Data as a stream (Replication)

Give me all “non-replicated” instances



Data as a stream (Replication)



KEEP CALM



AND EMBRACE STREAMS

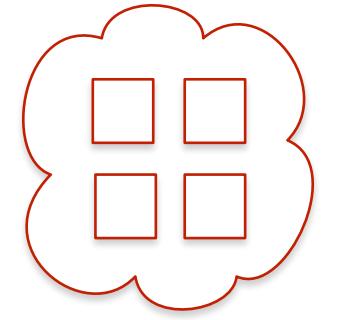
www.meme-generator.net

Learnt the hard-way

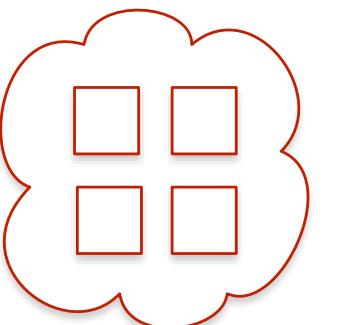


i've been to
the dark side
they lied about
the cookies

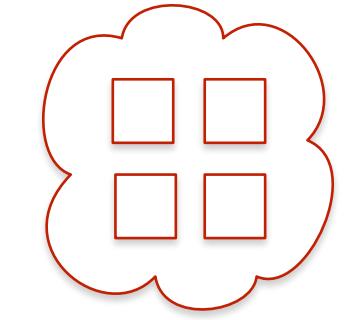
Network partitions gone wild!



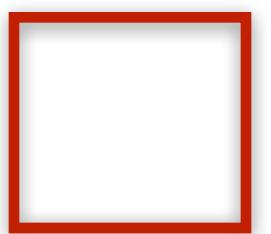
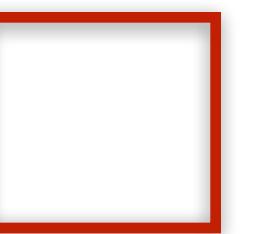
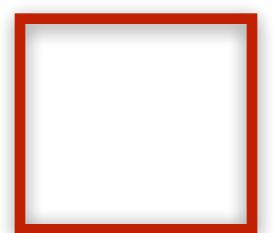
**Service Discovery
Node 2**

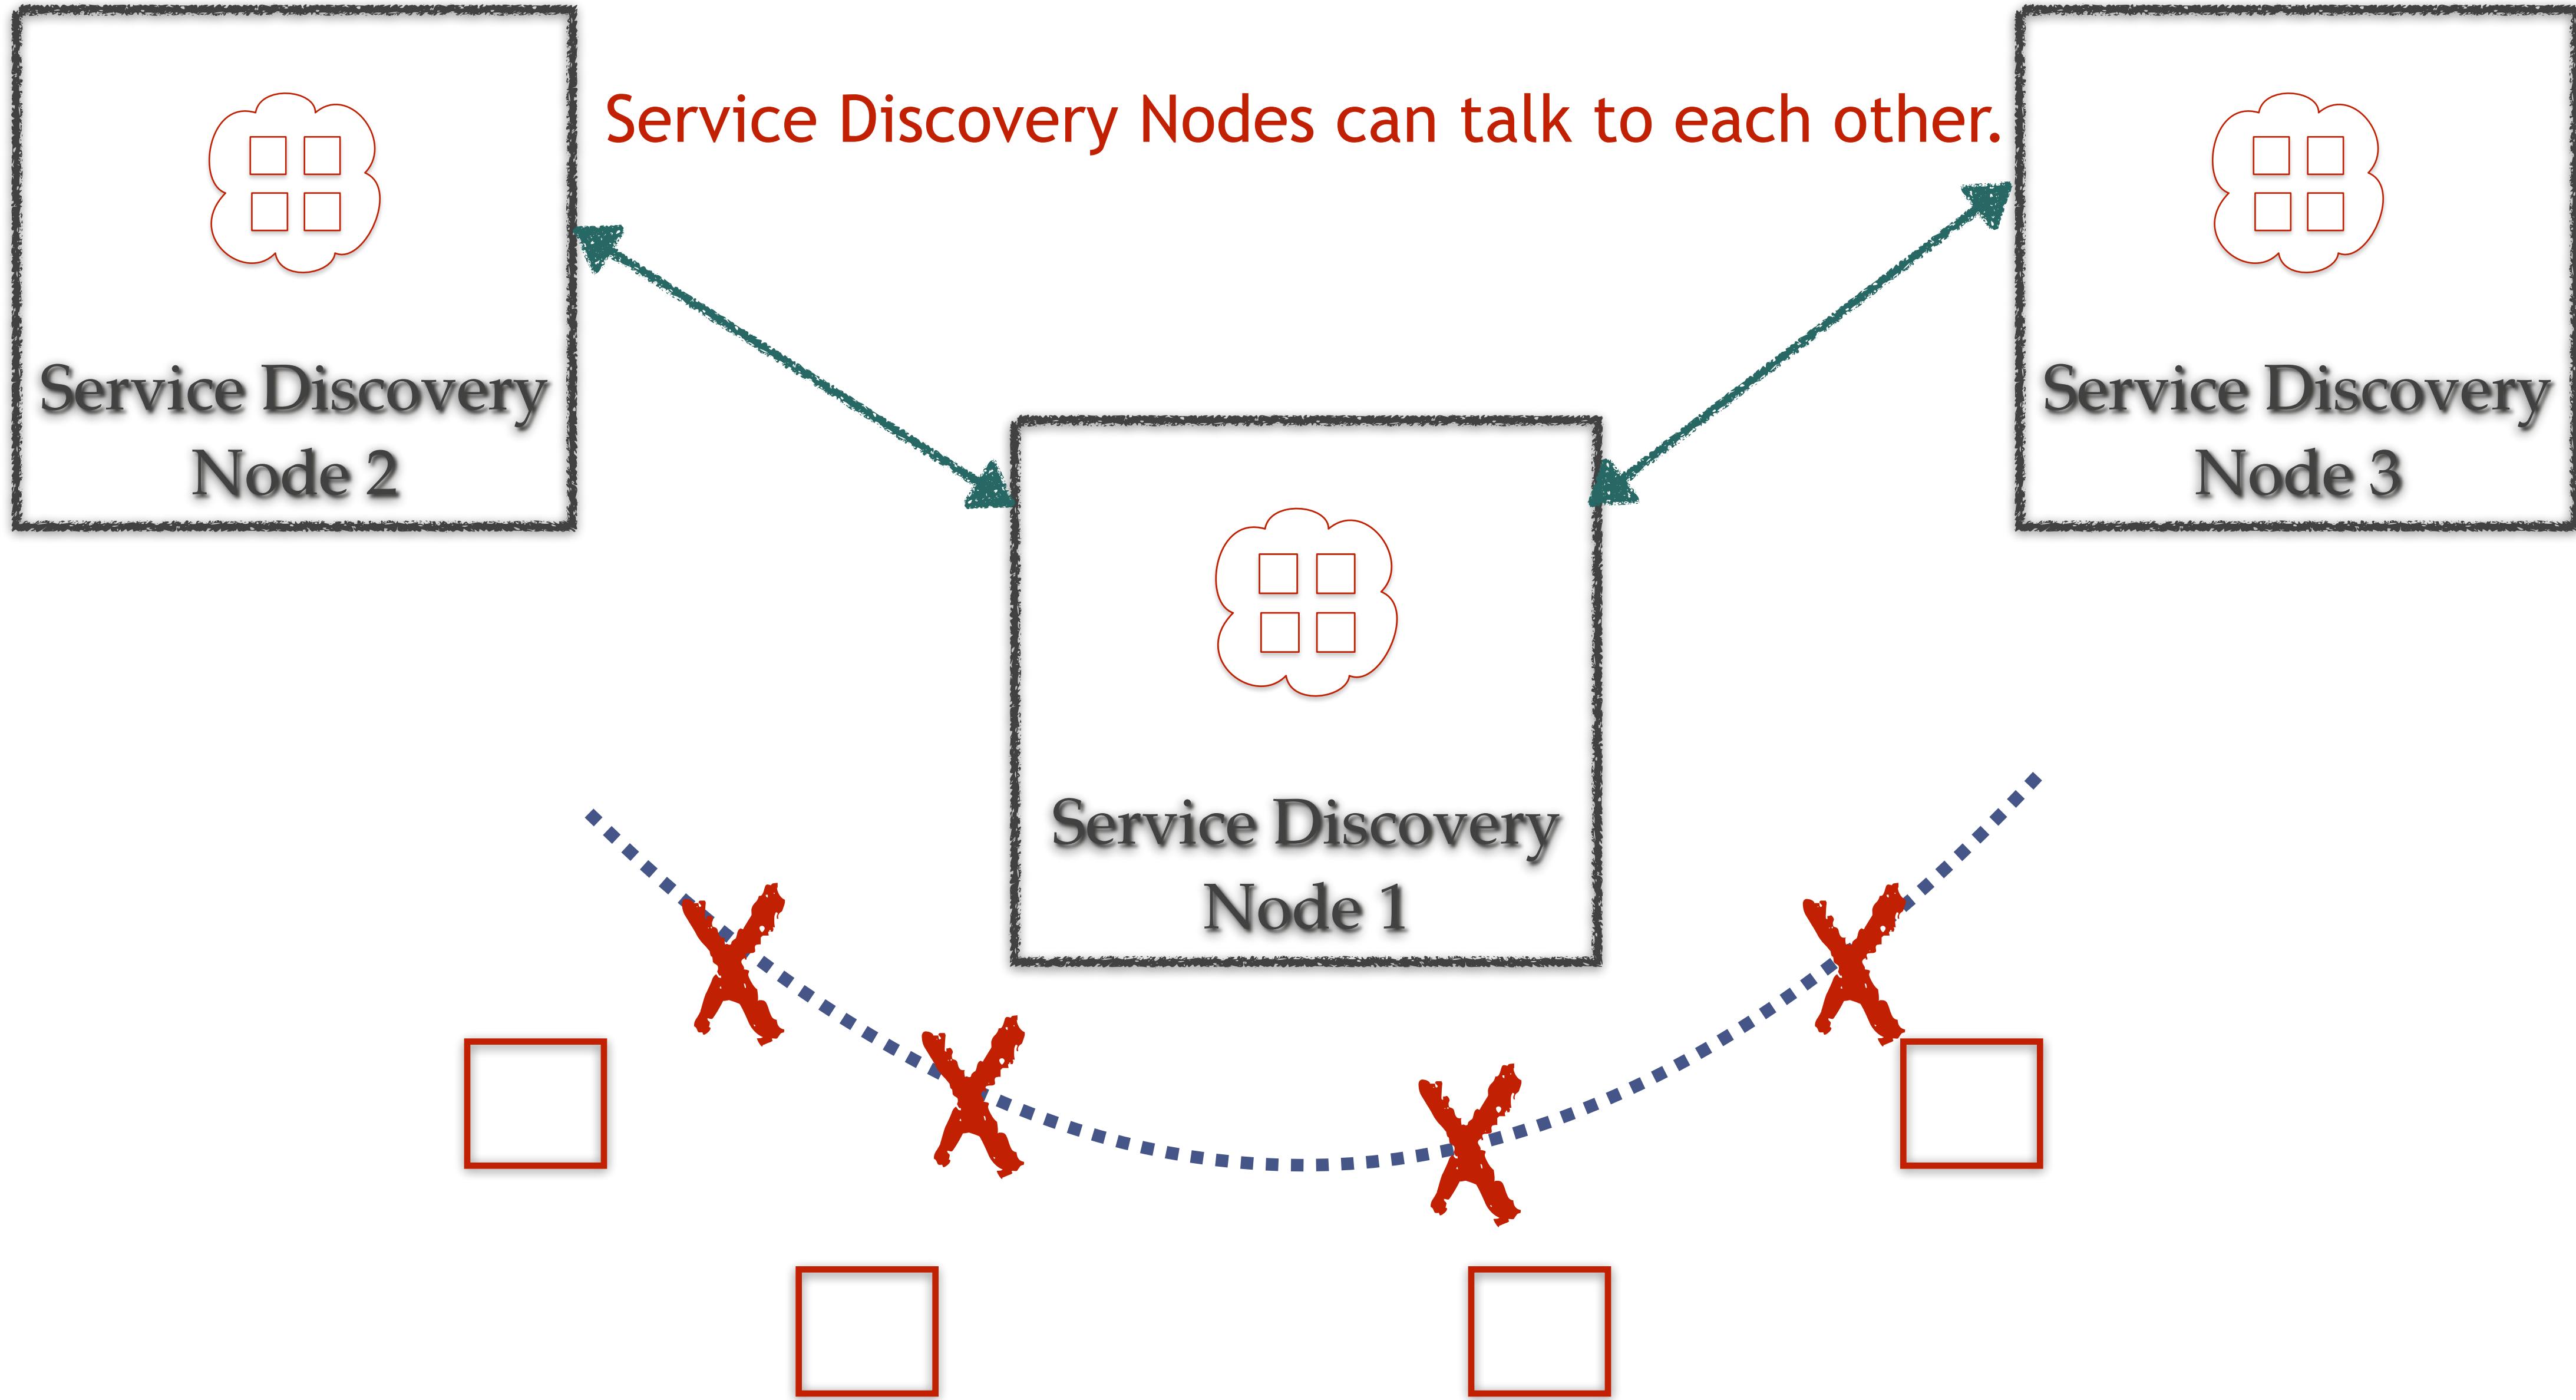


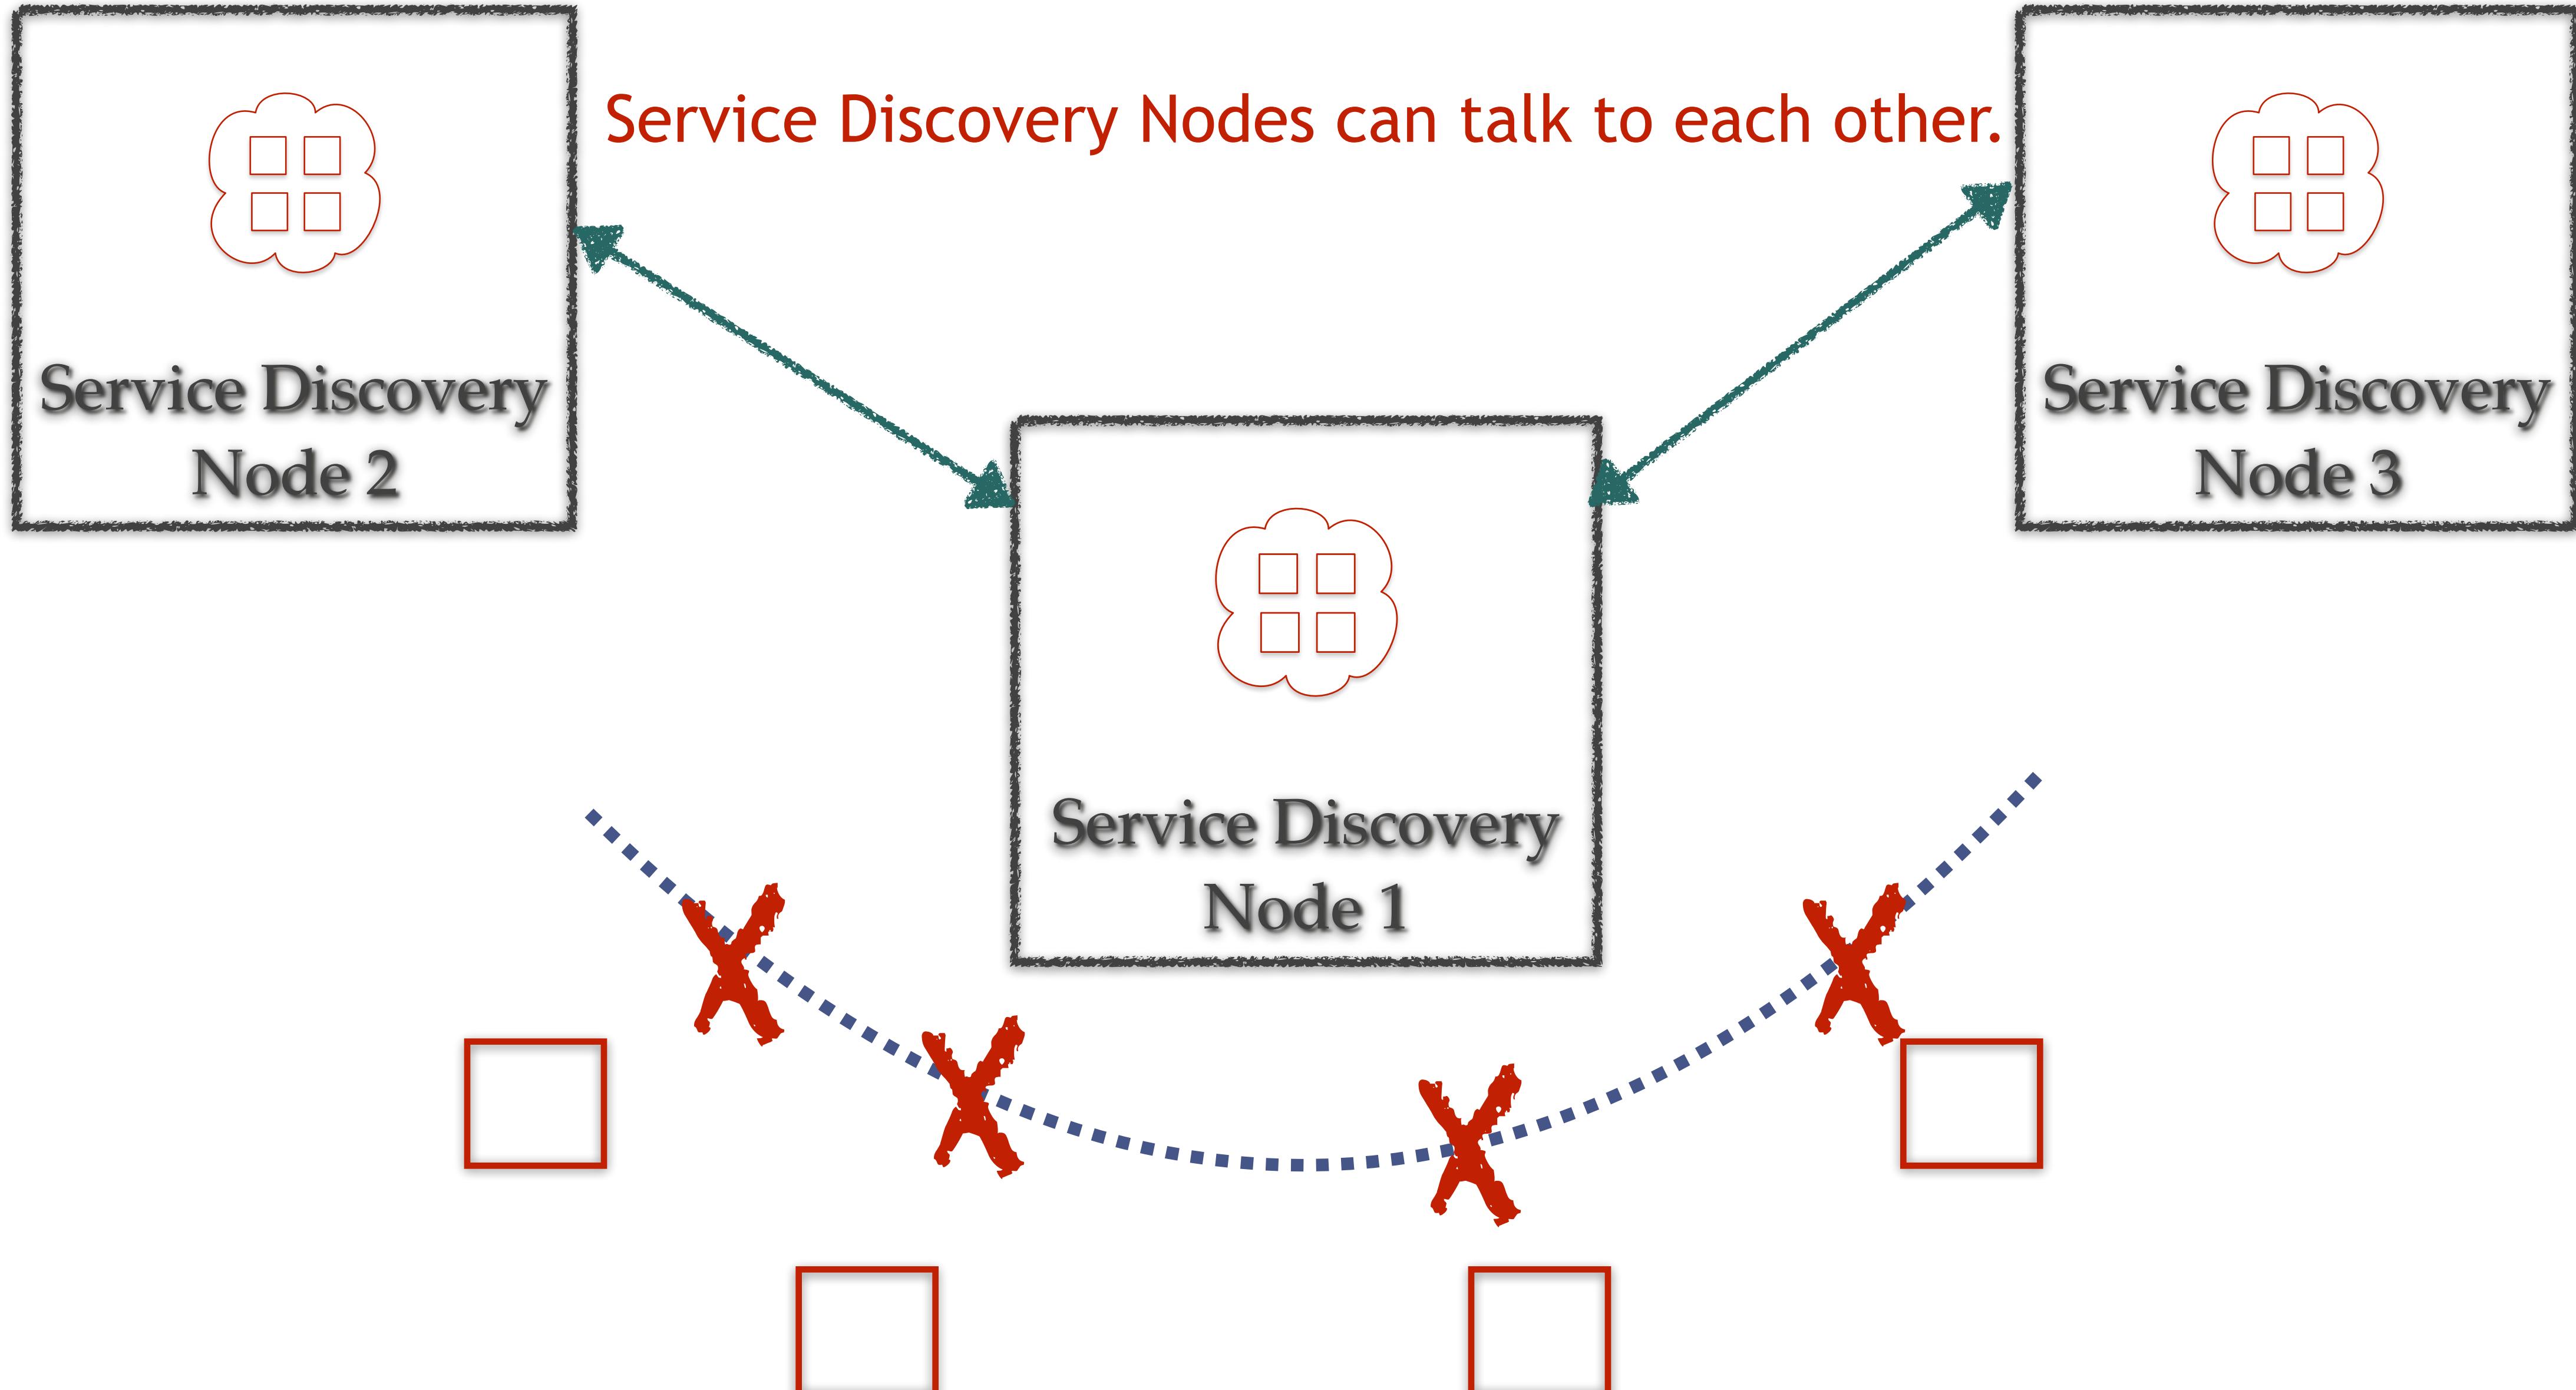
**Service Discovery
Node 1**



**Service Discovery
Node 3**







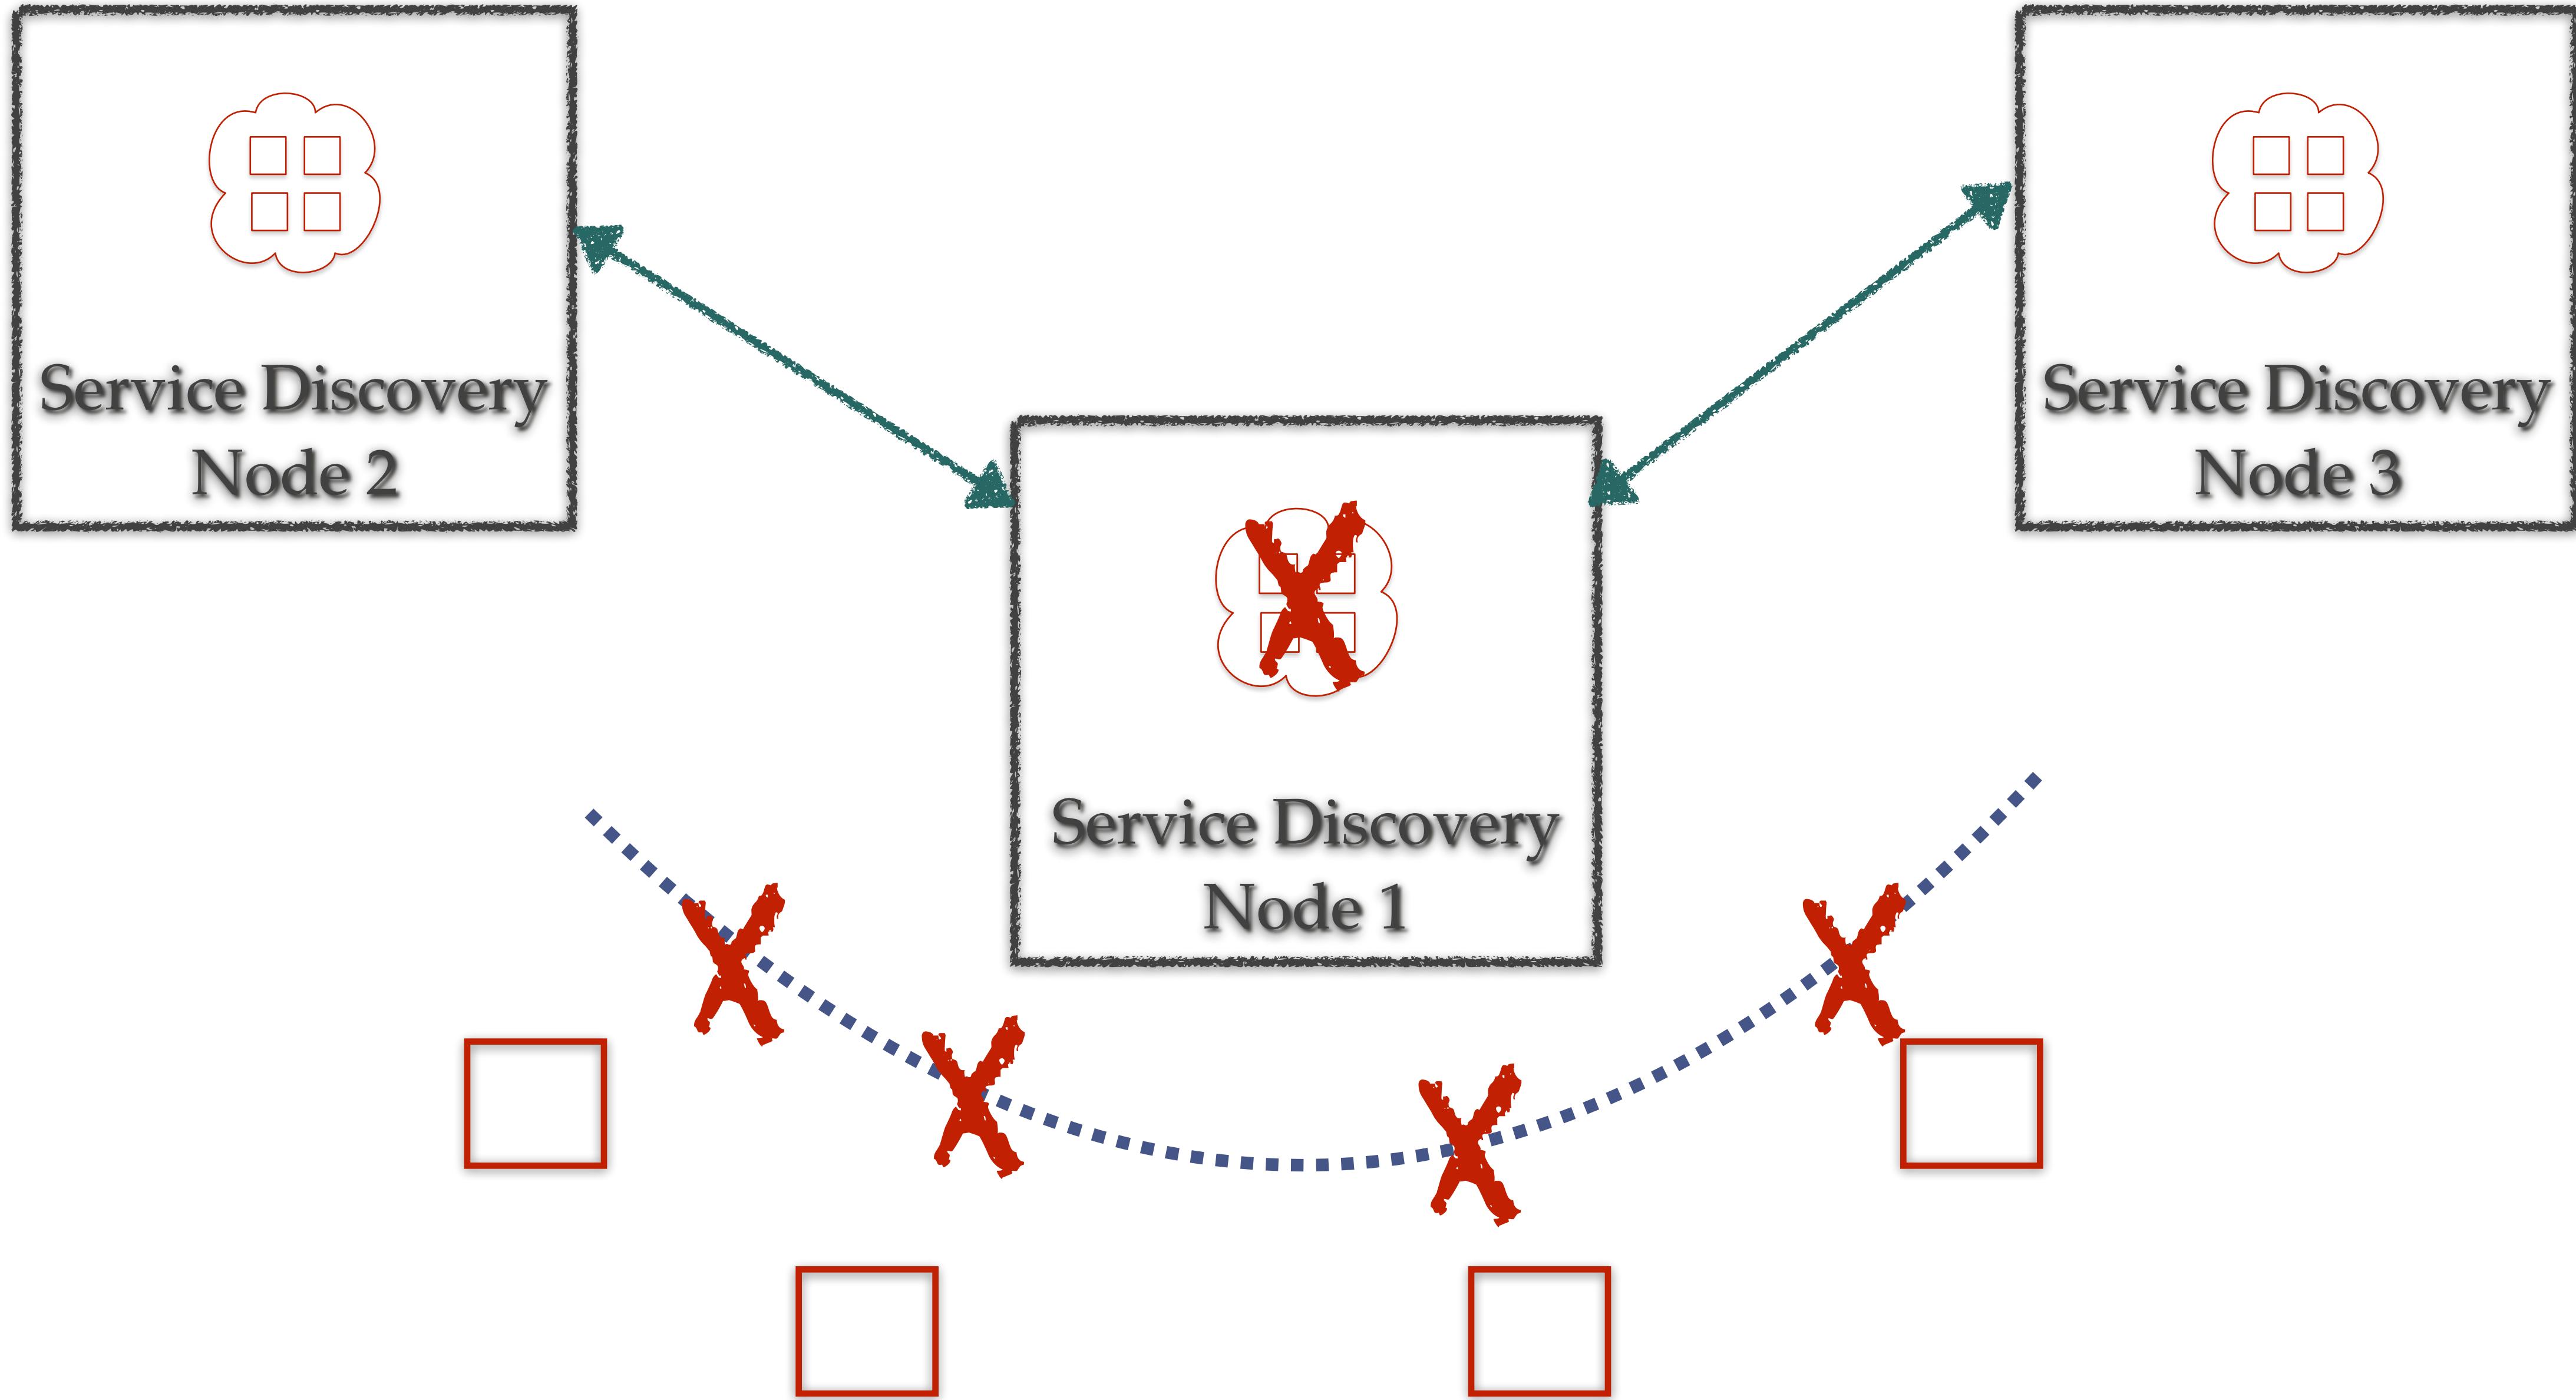
Service Discovery Nodes can talk to each other.

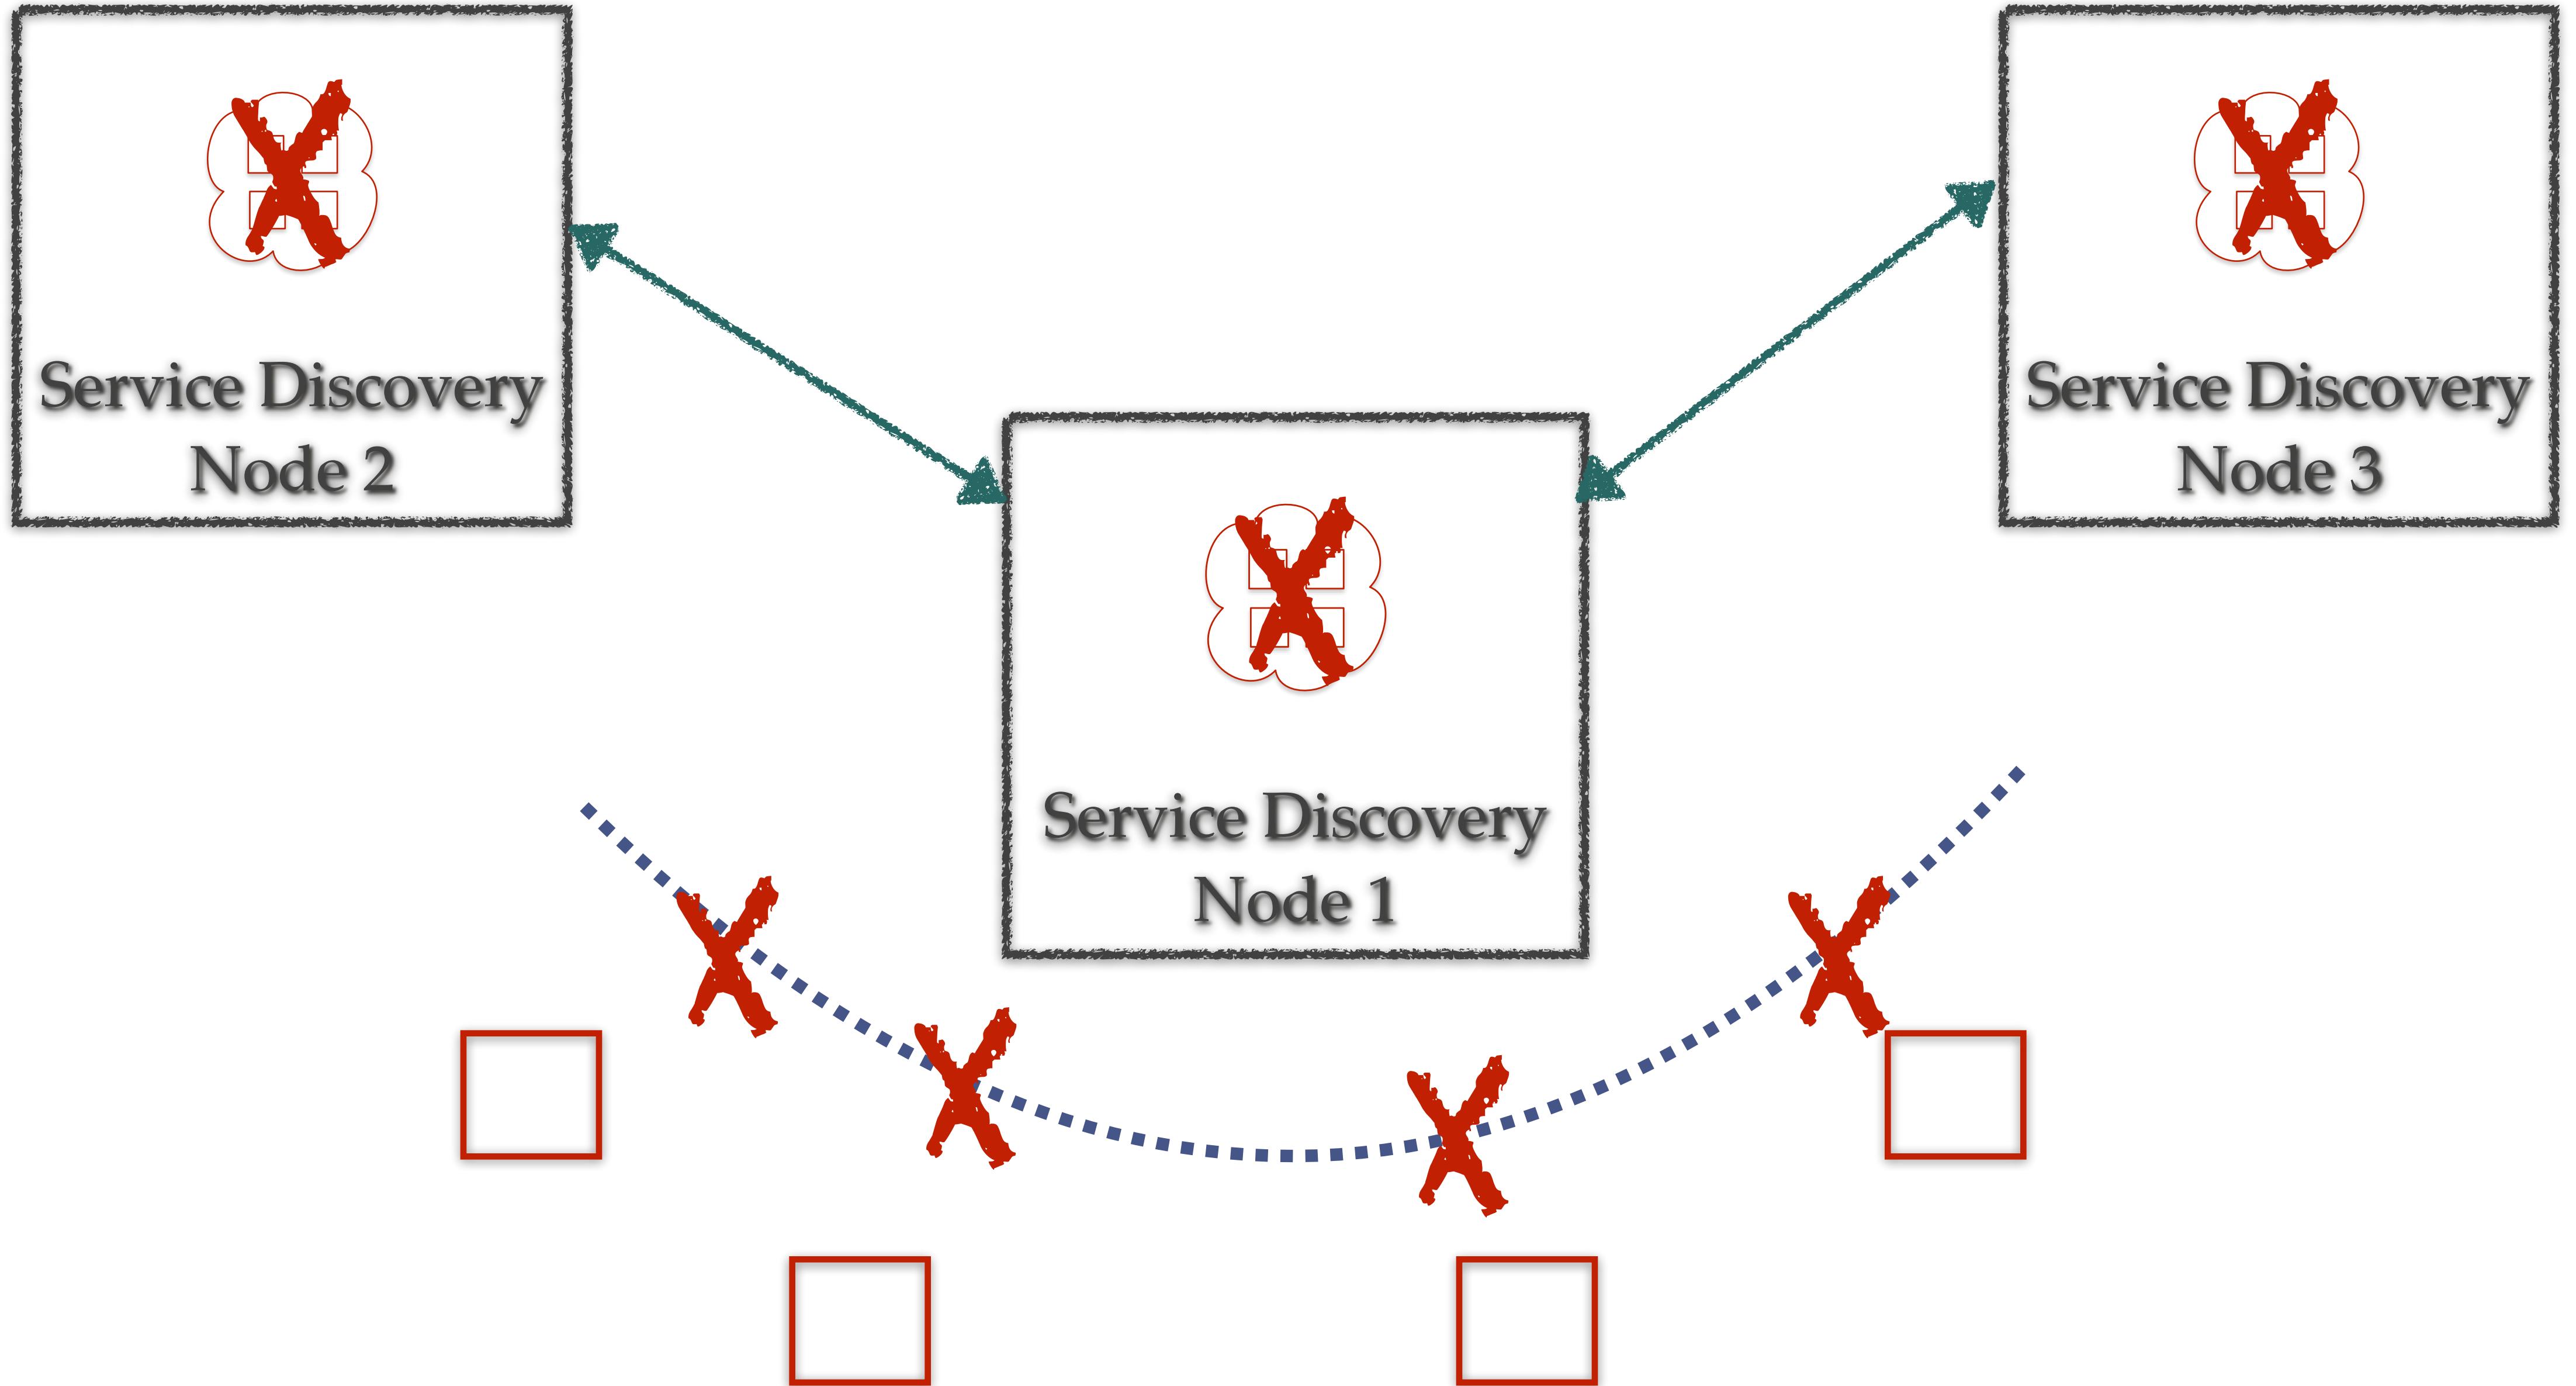
Service Discovery
Node 2

Service Discovery
Node 1

Service Discovery
Node 3

No outside instance can talk to a node.





One fine day

... you lost most of your service instances ...

One fine day
... you lost most of your service instances ...

because ...
one node of service discovery was
partitioned

Taming evictions
a.k.a
Self preservation

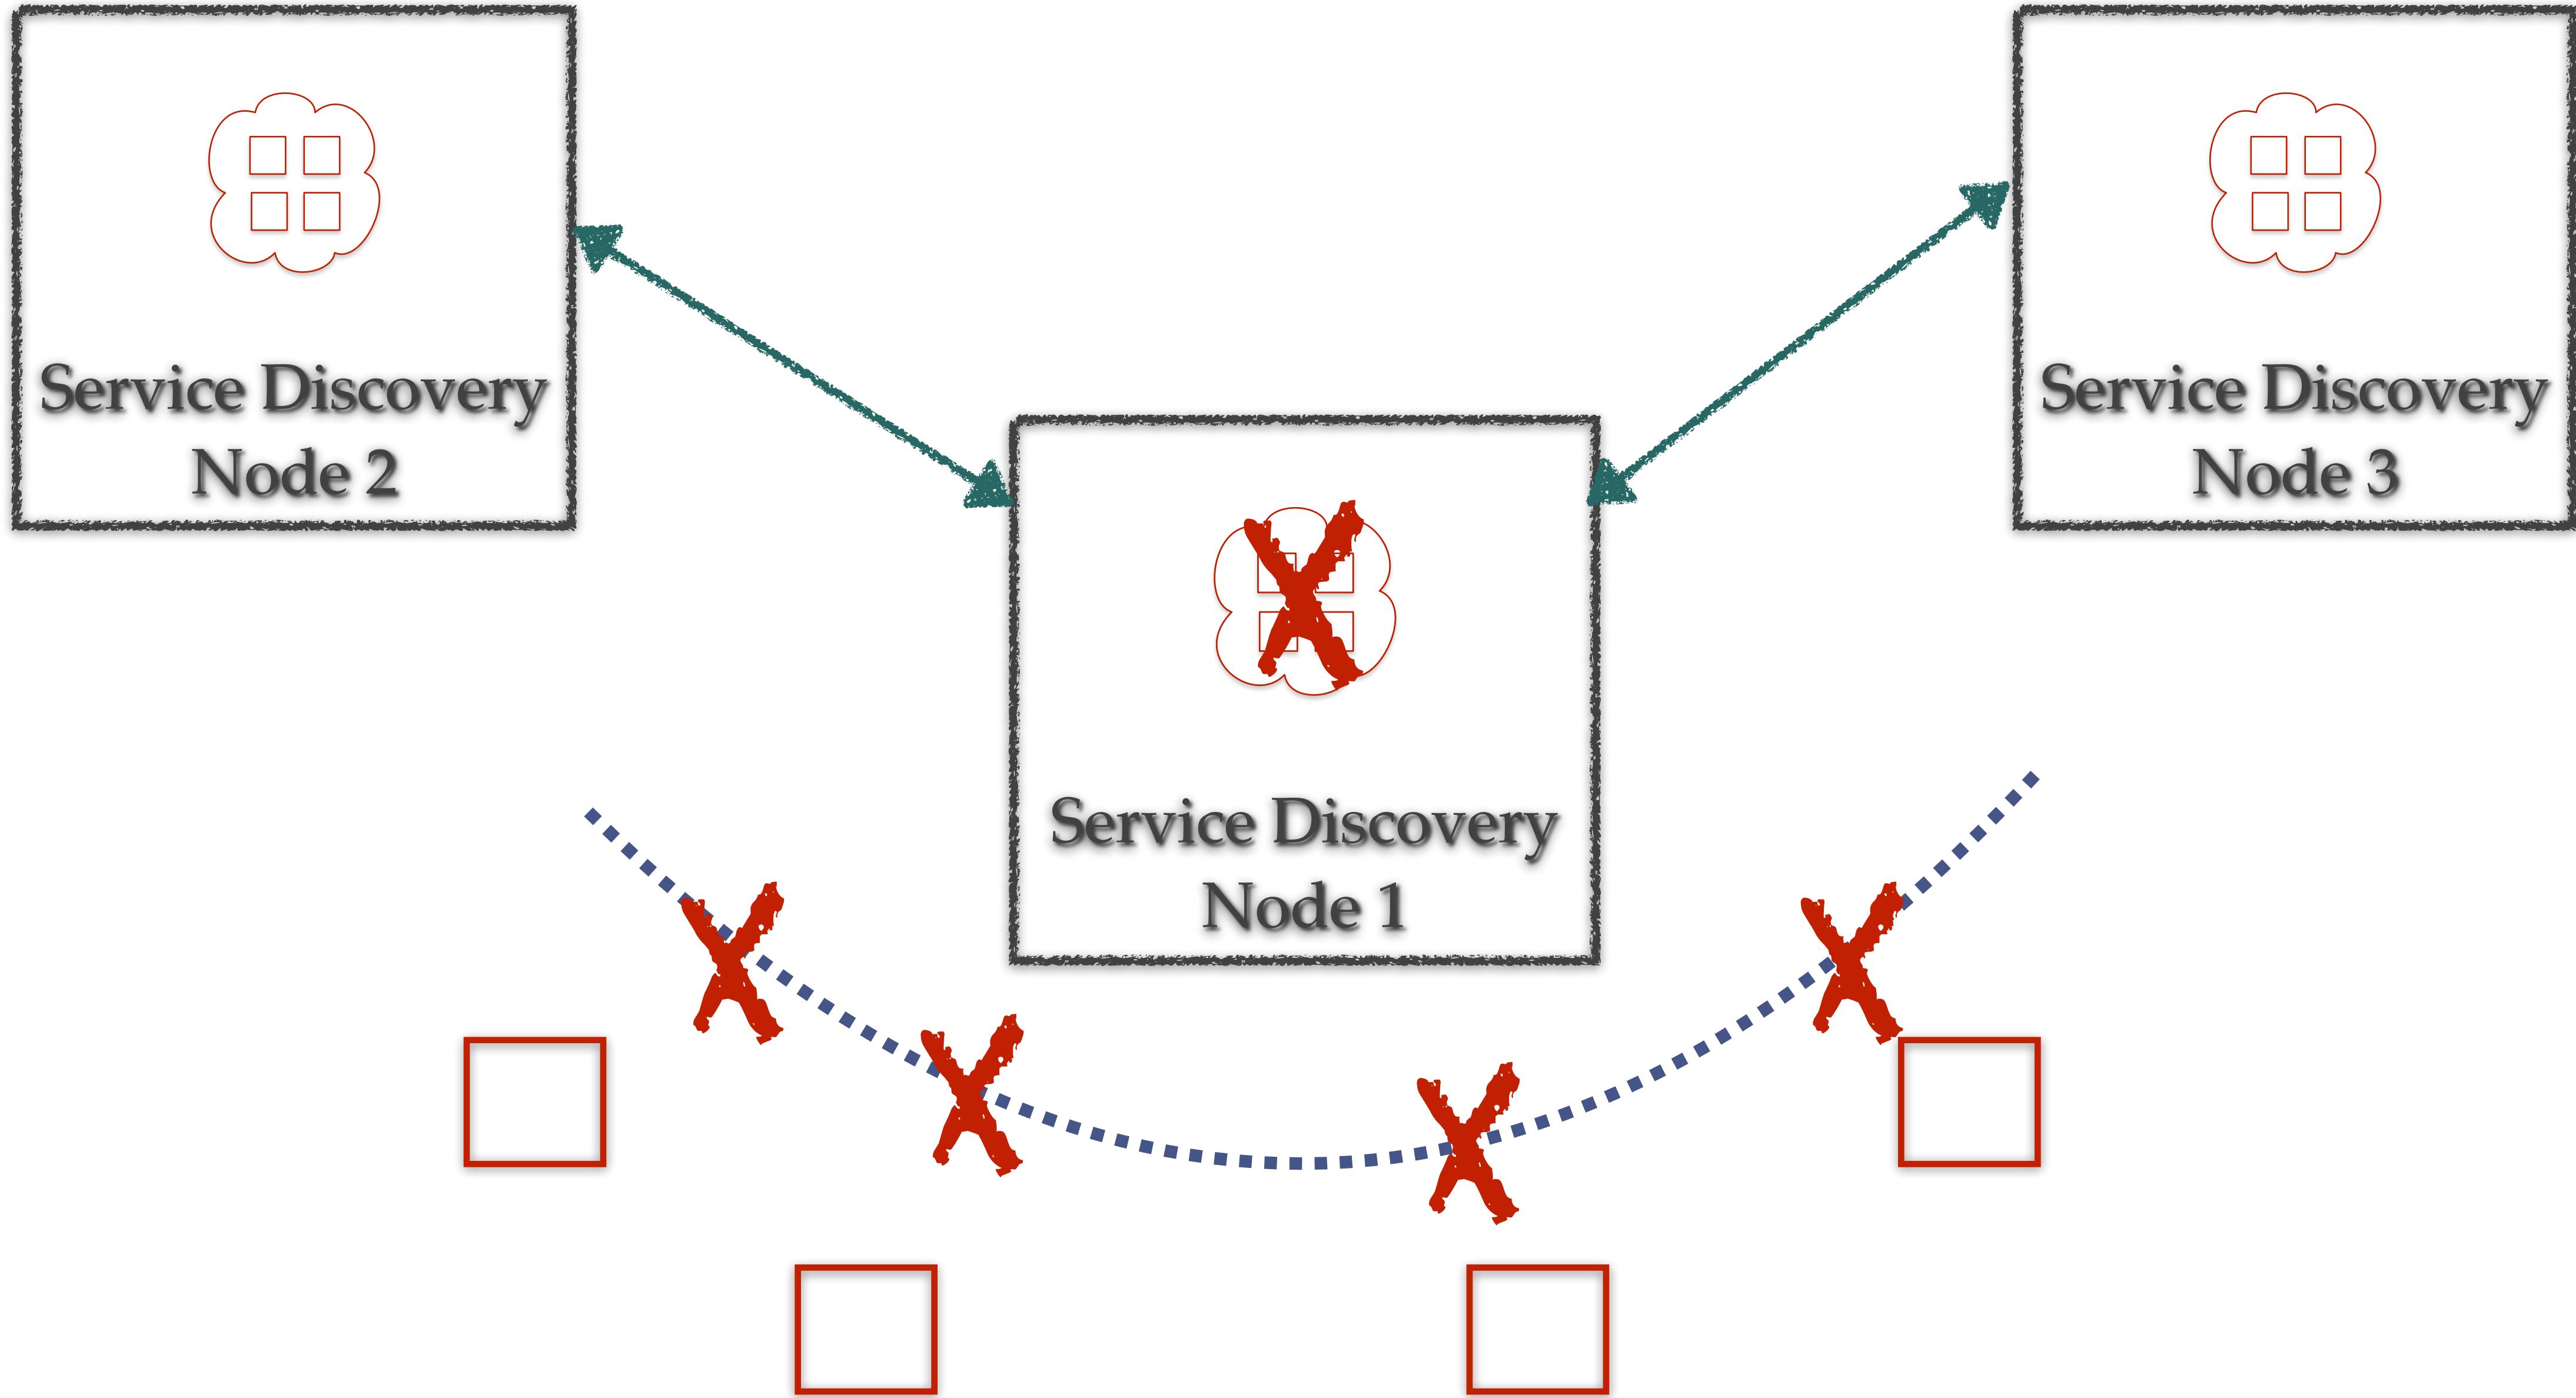
All instances “generally” do not
vanish

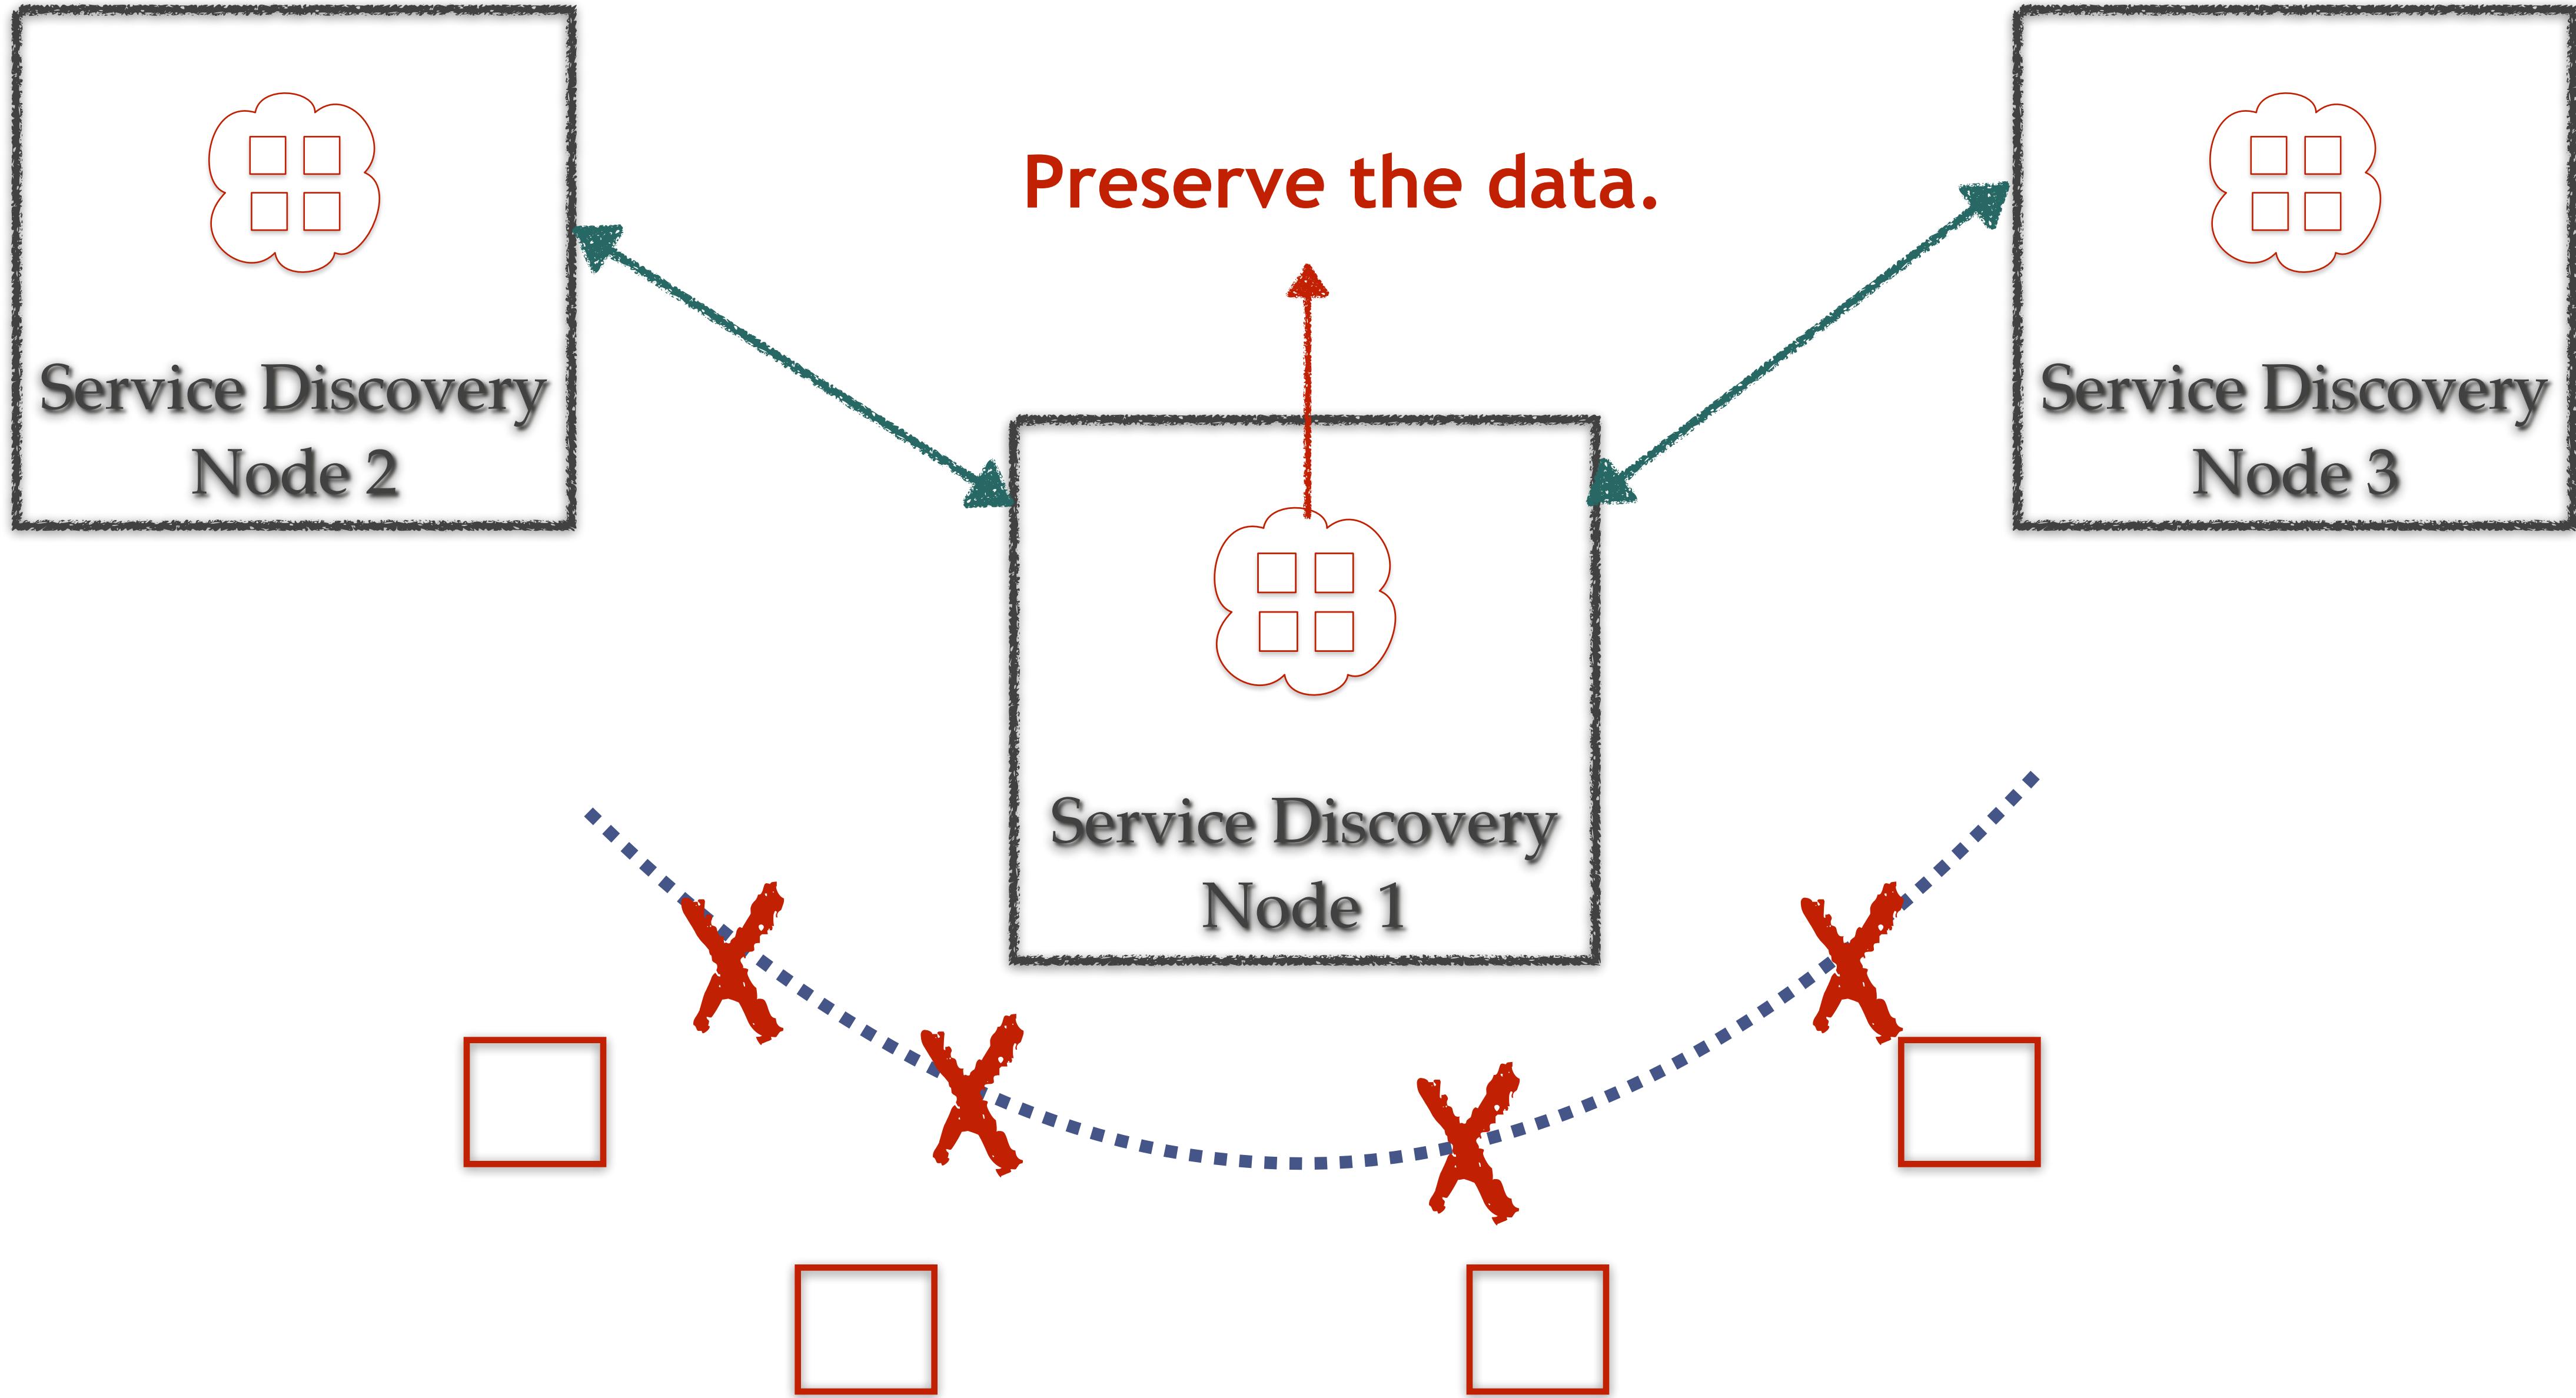
All instances “generally” do not
vanish

If they do, let local decisions prevail

Do NOT evict

If $> X\%$ of instances vanish





Clients talking to the “in-doubt” instance will detect failure, if any.

Total unavailability

a.k.a.

Apocalypse

Impact



Fatal

New nodes can not be started.

Impact

Fatal

Degraded

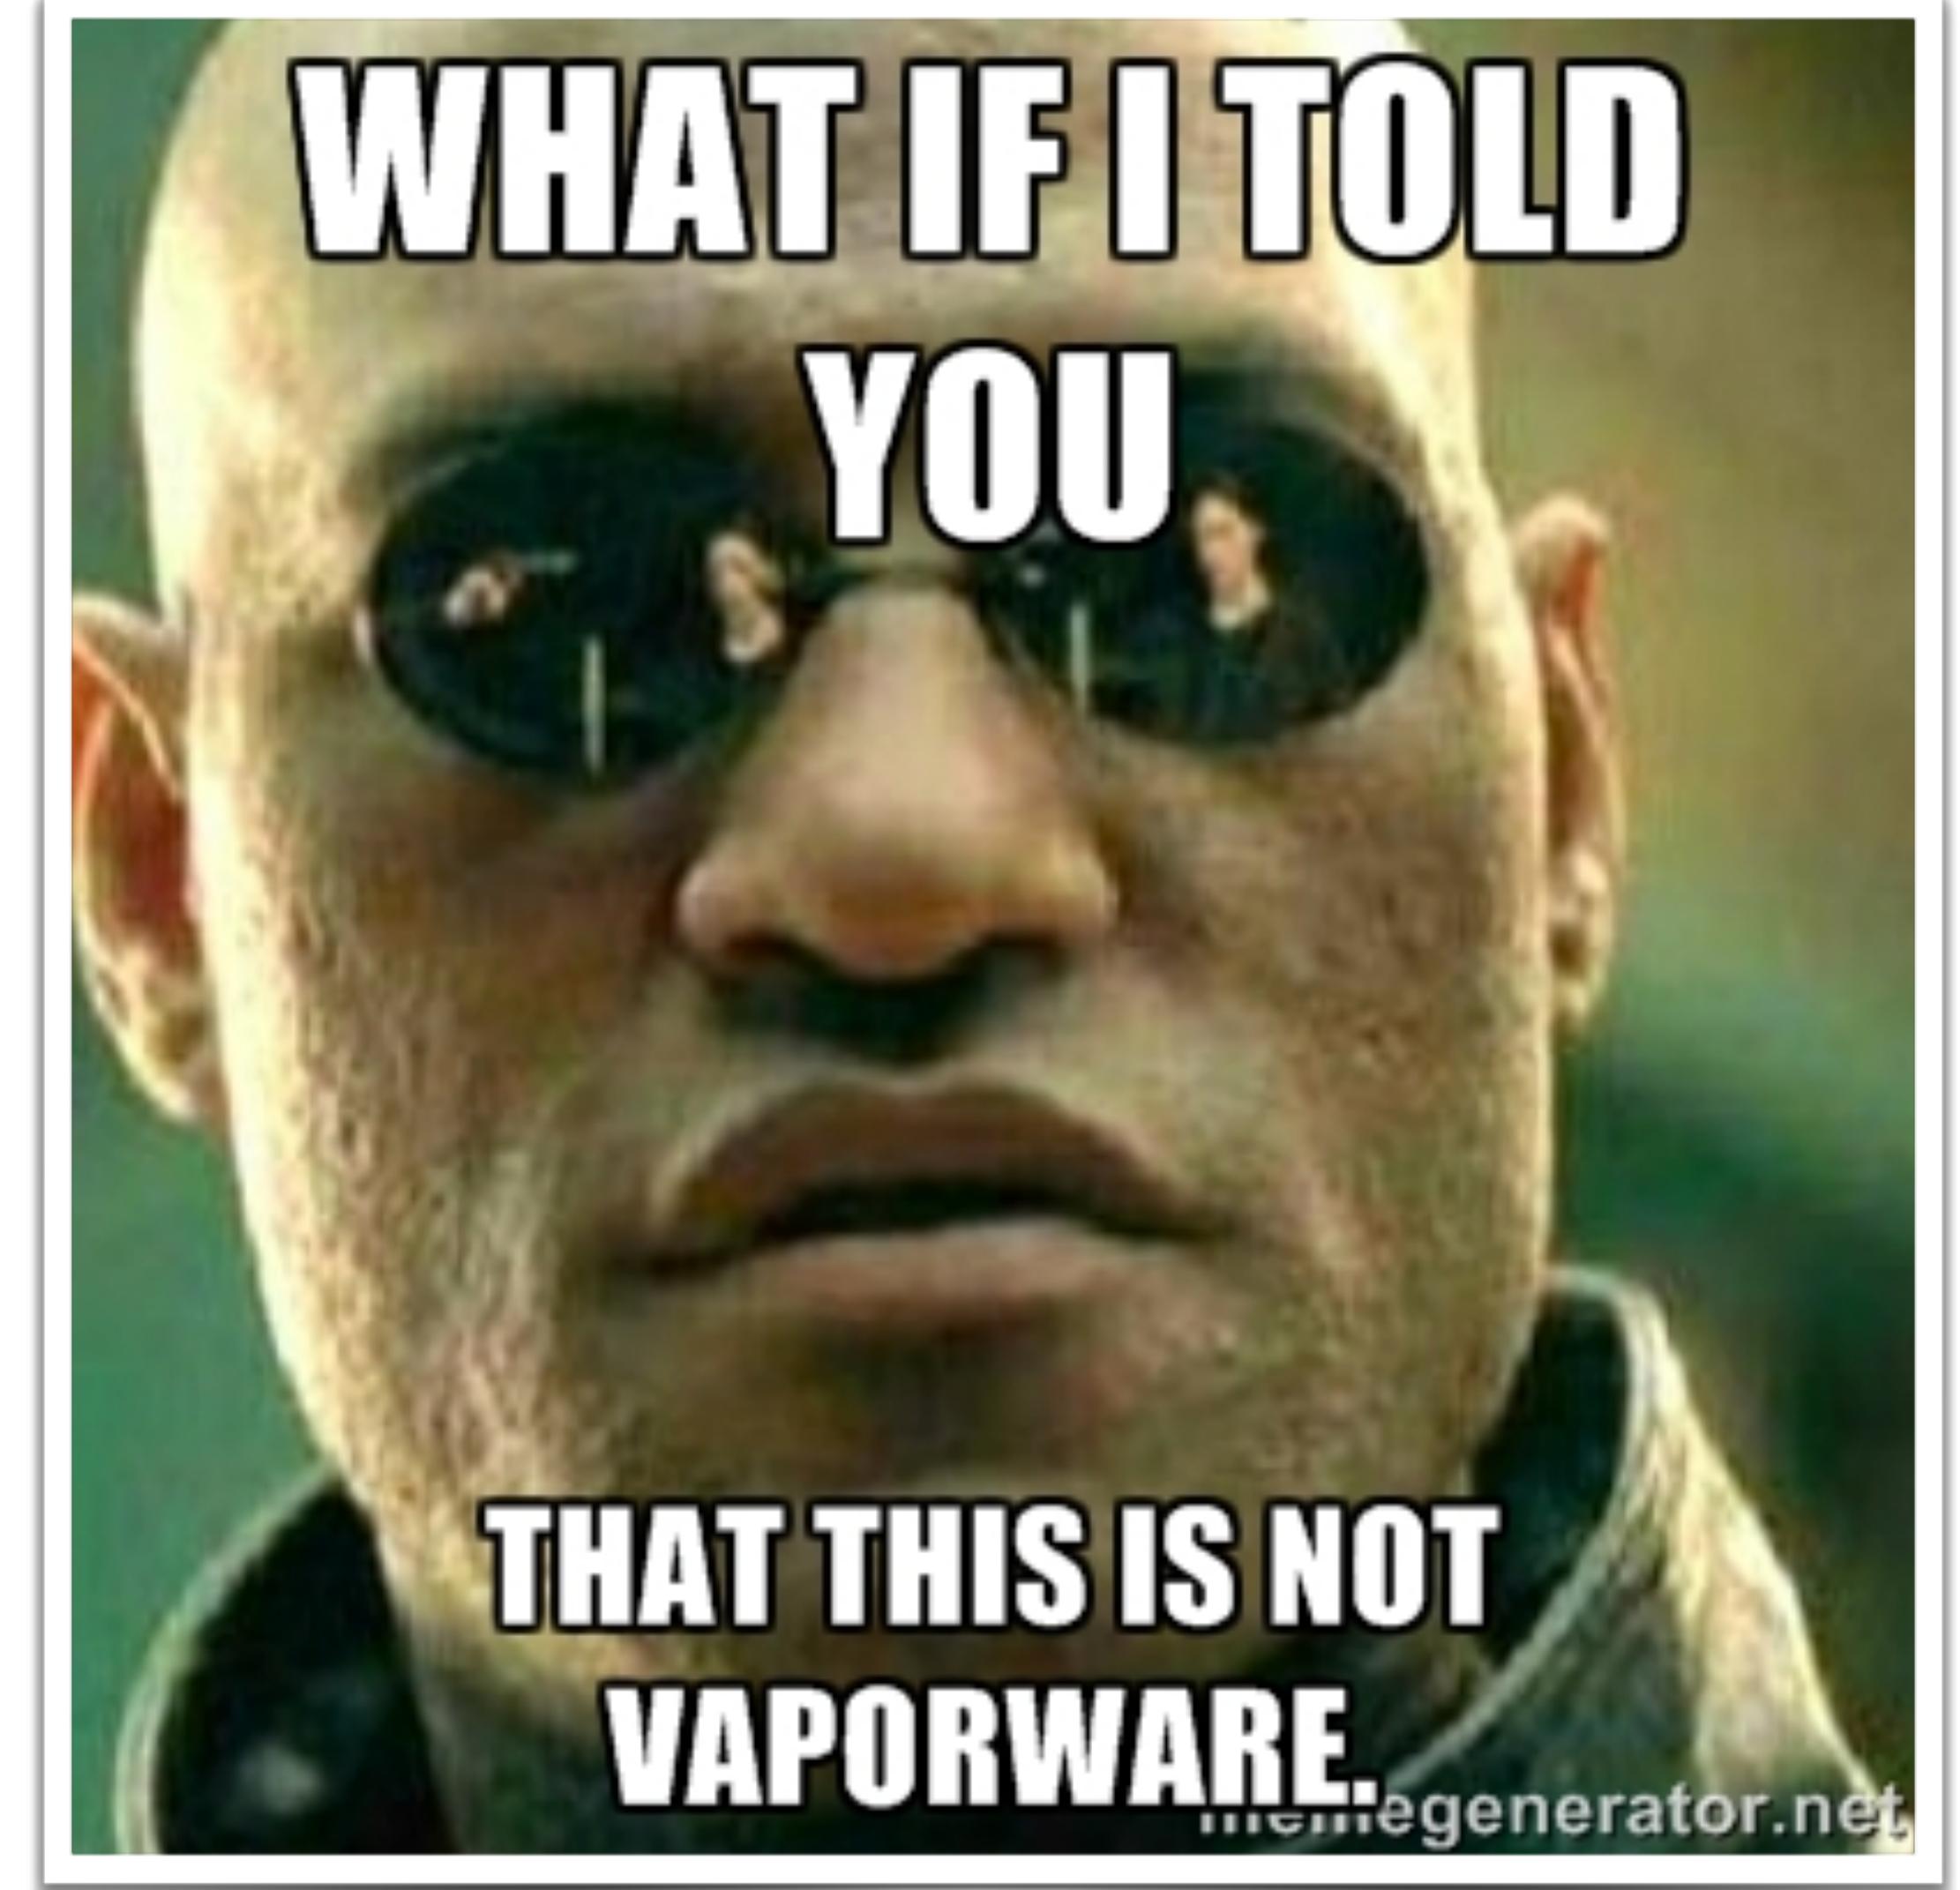
New nodes can not be started.

Existing nodes uses **cached** data.

NETFLIX | OSS

Netflix Eureka V2

<https://github.com/Netflix/eureka>



Don't make Service Discovery your Achilles heel.

Service discovery controls visibility of nodes but does **not** guarantee availability.



