



The Elizabeth H.  
and James S. McDonnell III

**McDONNELL  
GENOME INSTITUTE**  
at Washington University

## PTero

History, Architecture and Development  
of a scalable workflow system.

# What you're in for...

- **History and Background** (~15 min)  
Genome Modeling System -> Workflow -> Flow -> PTero
- **PTero Architecture** (~15 min)
  - Services
  - Service Interactions
  - Realistic Example
- **Development and Testing** (~5 min)
- **Deployment** (~2 min)
- **Live Demo?** (~3 min)



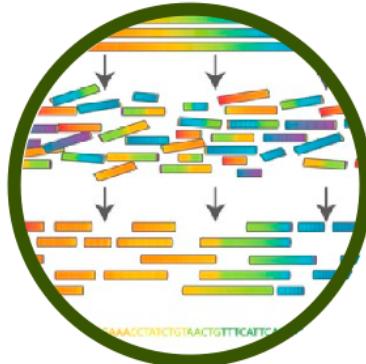
# History and Background

## Genetic Sequencing and Analysis at MGI

Sample Preparation



Alignment



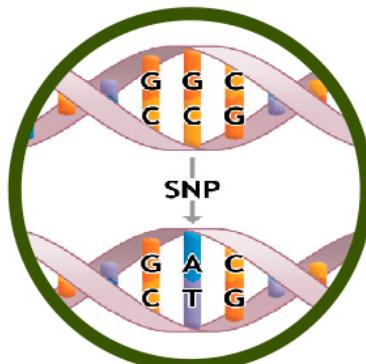
Annotation



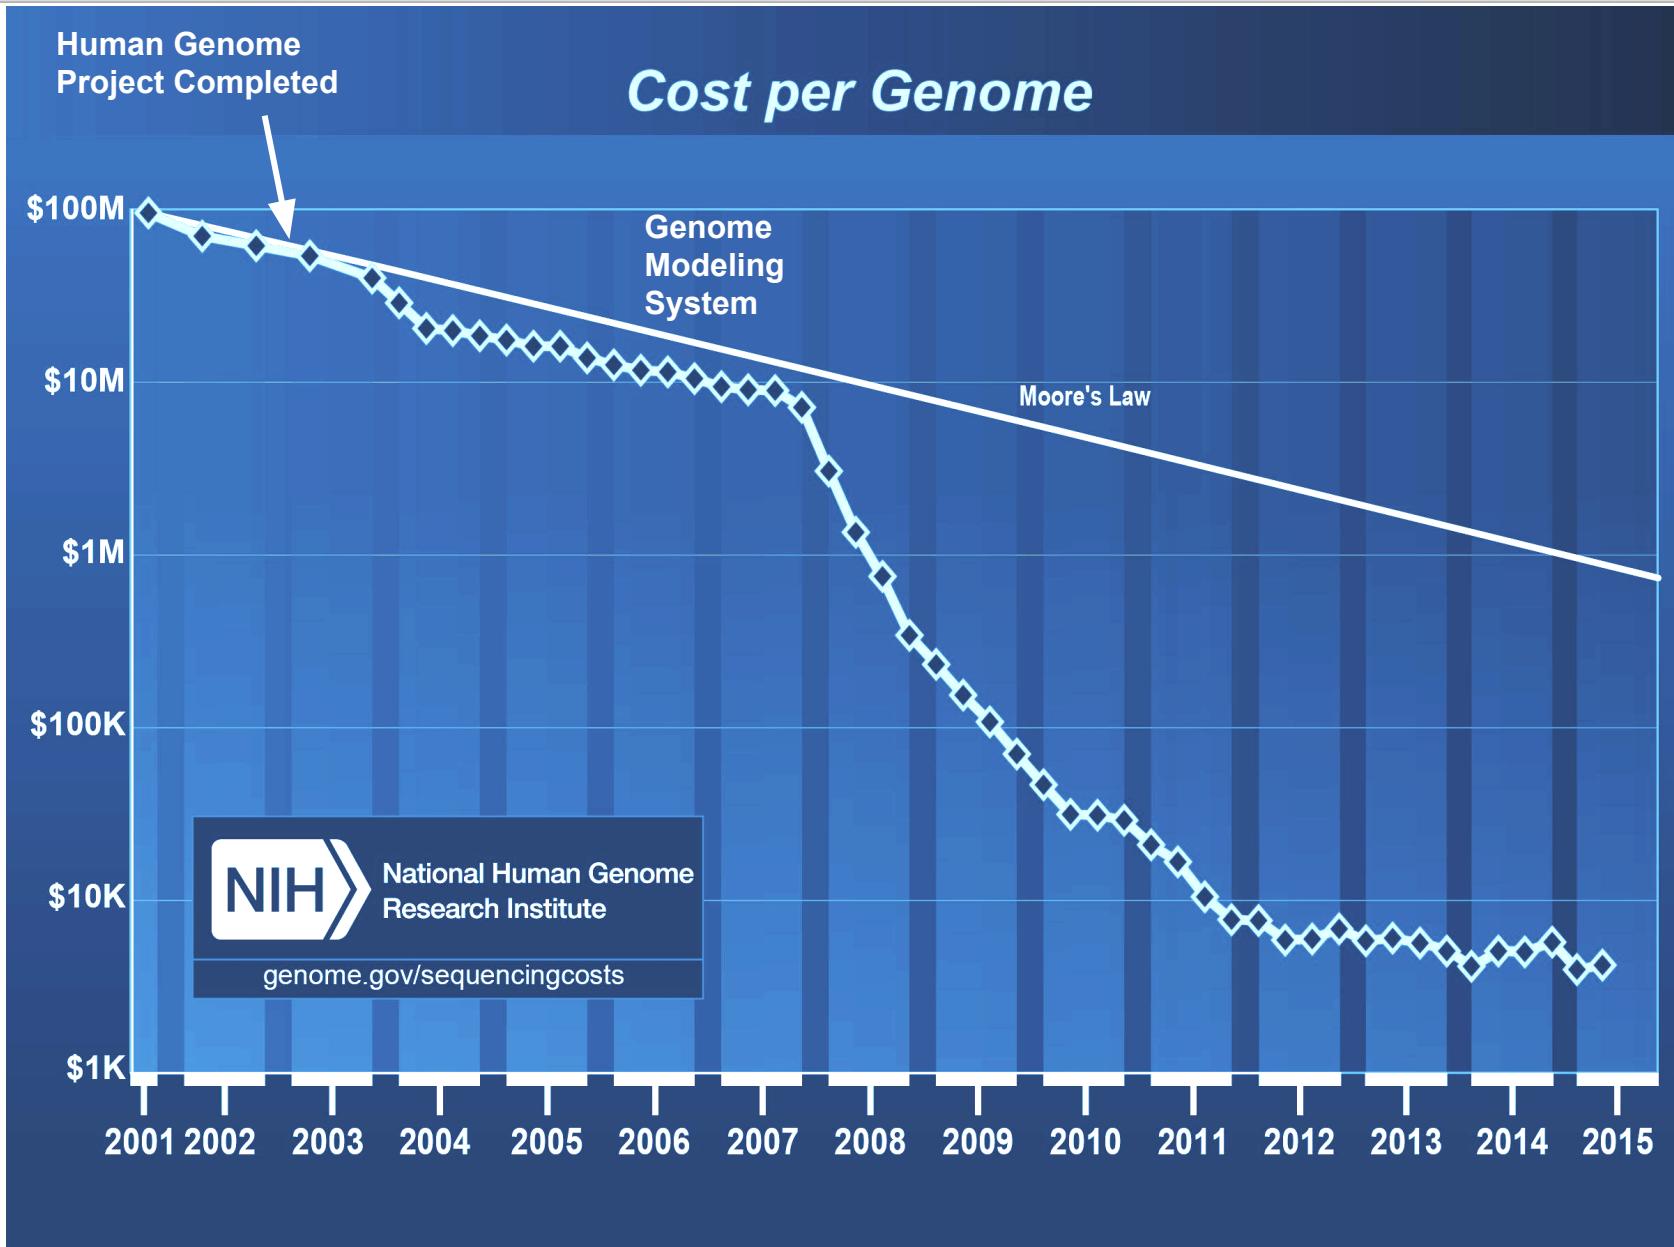
Sequencing



Variant Detection



# History and Background



# History and Background



## File Based Programs

- Long Running
- Noninteractive
- 3rd Party

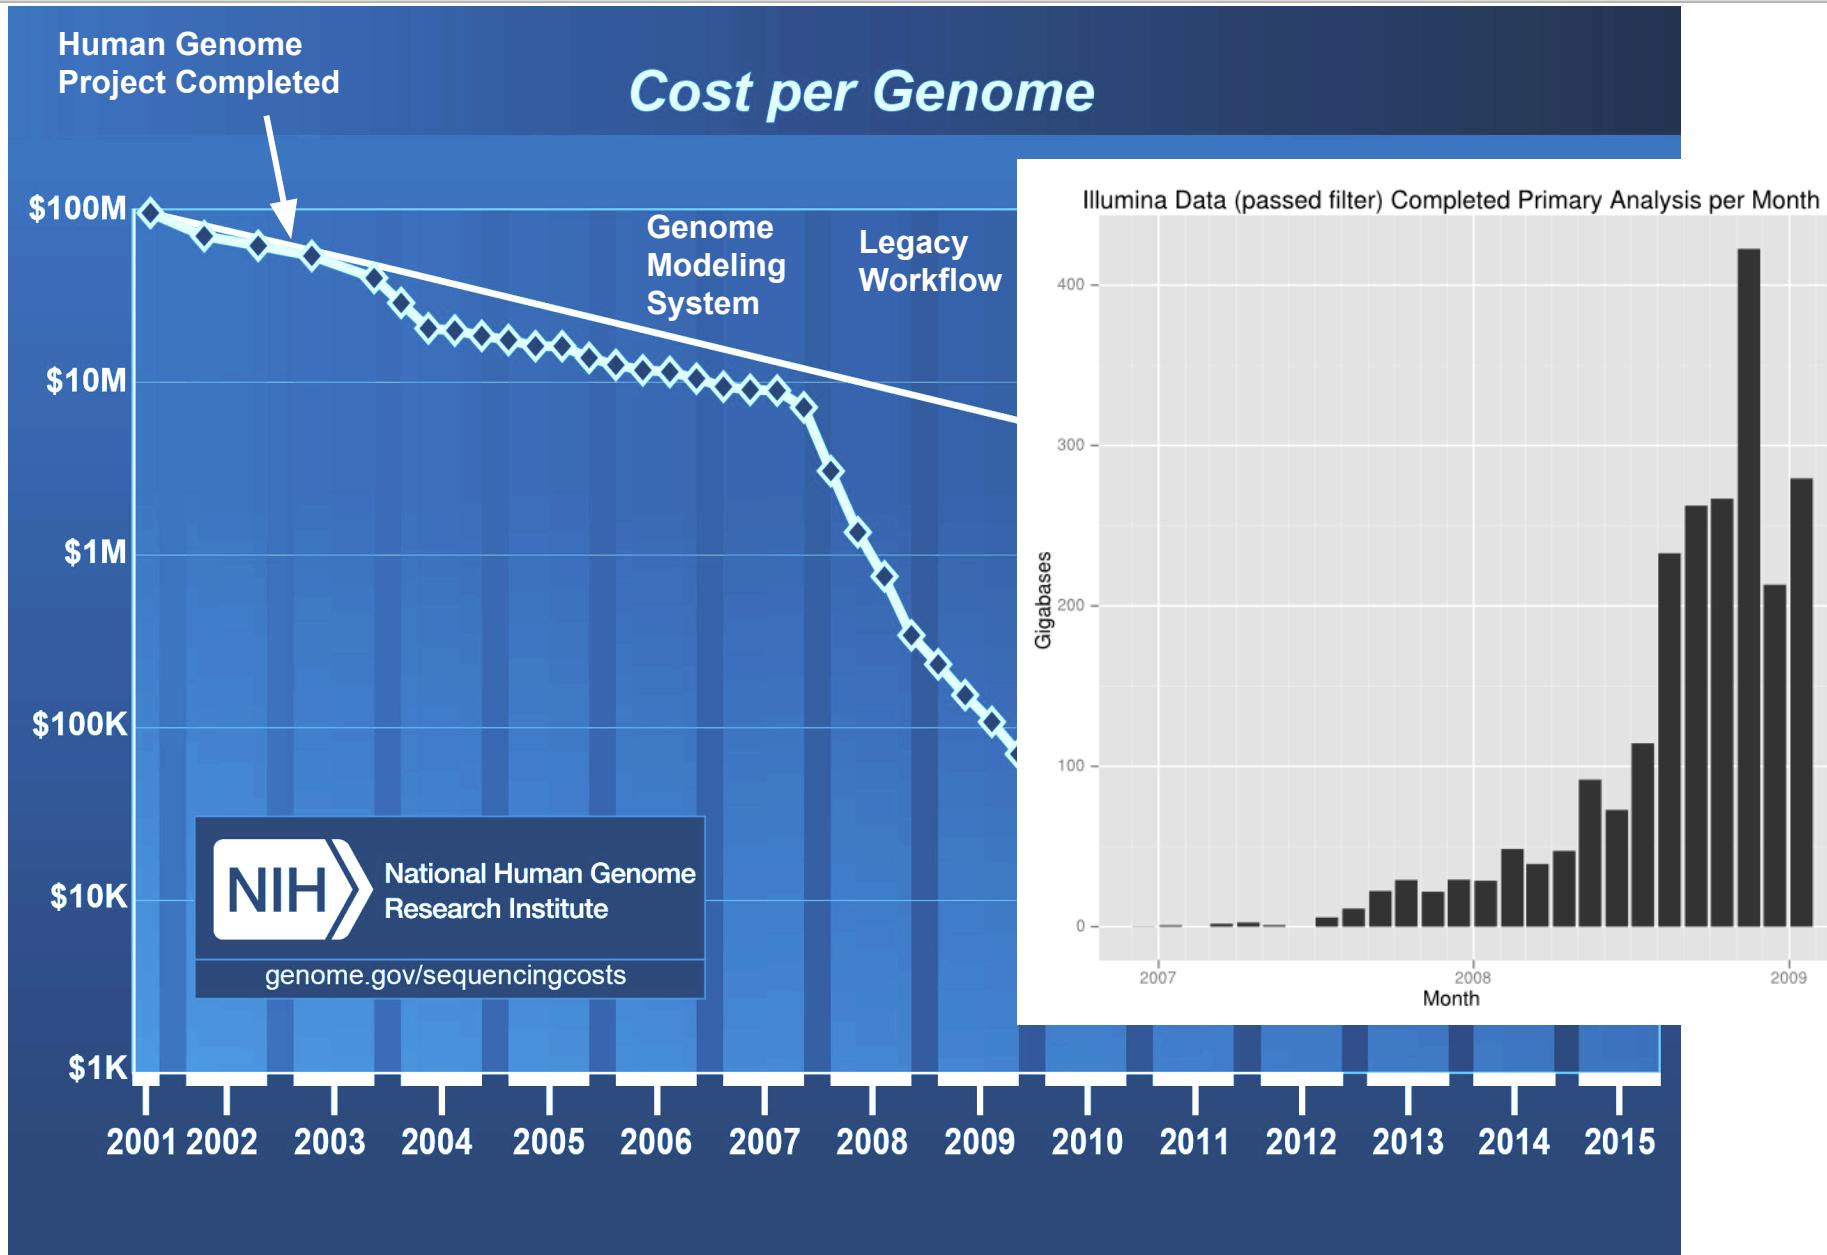


## Genome Modeling System

- Wrap Programs in Perl
- Expose Command Line Interface
- Create Automated Pipelines



# History and Background



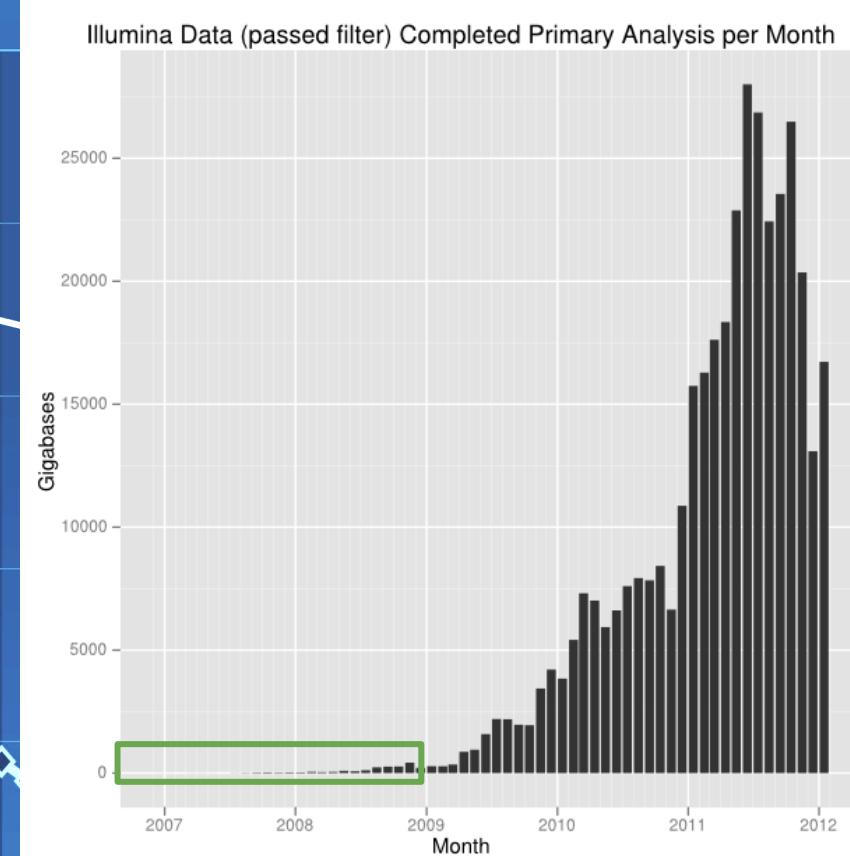
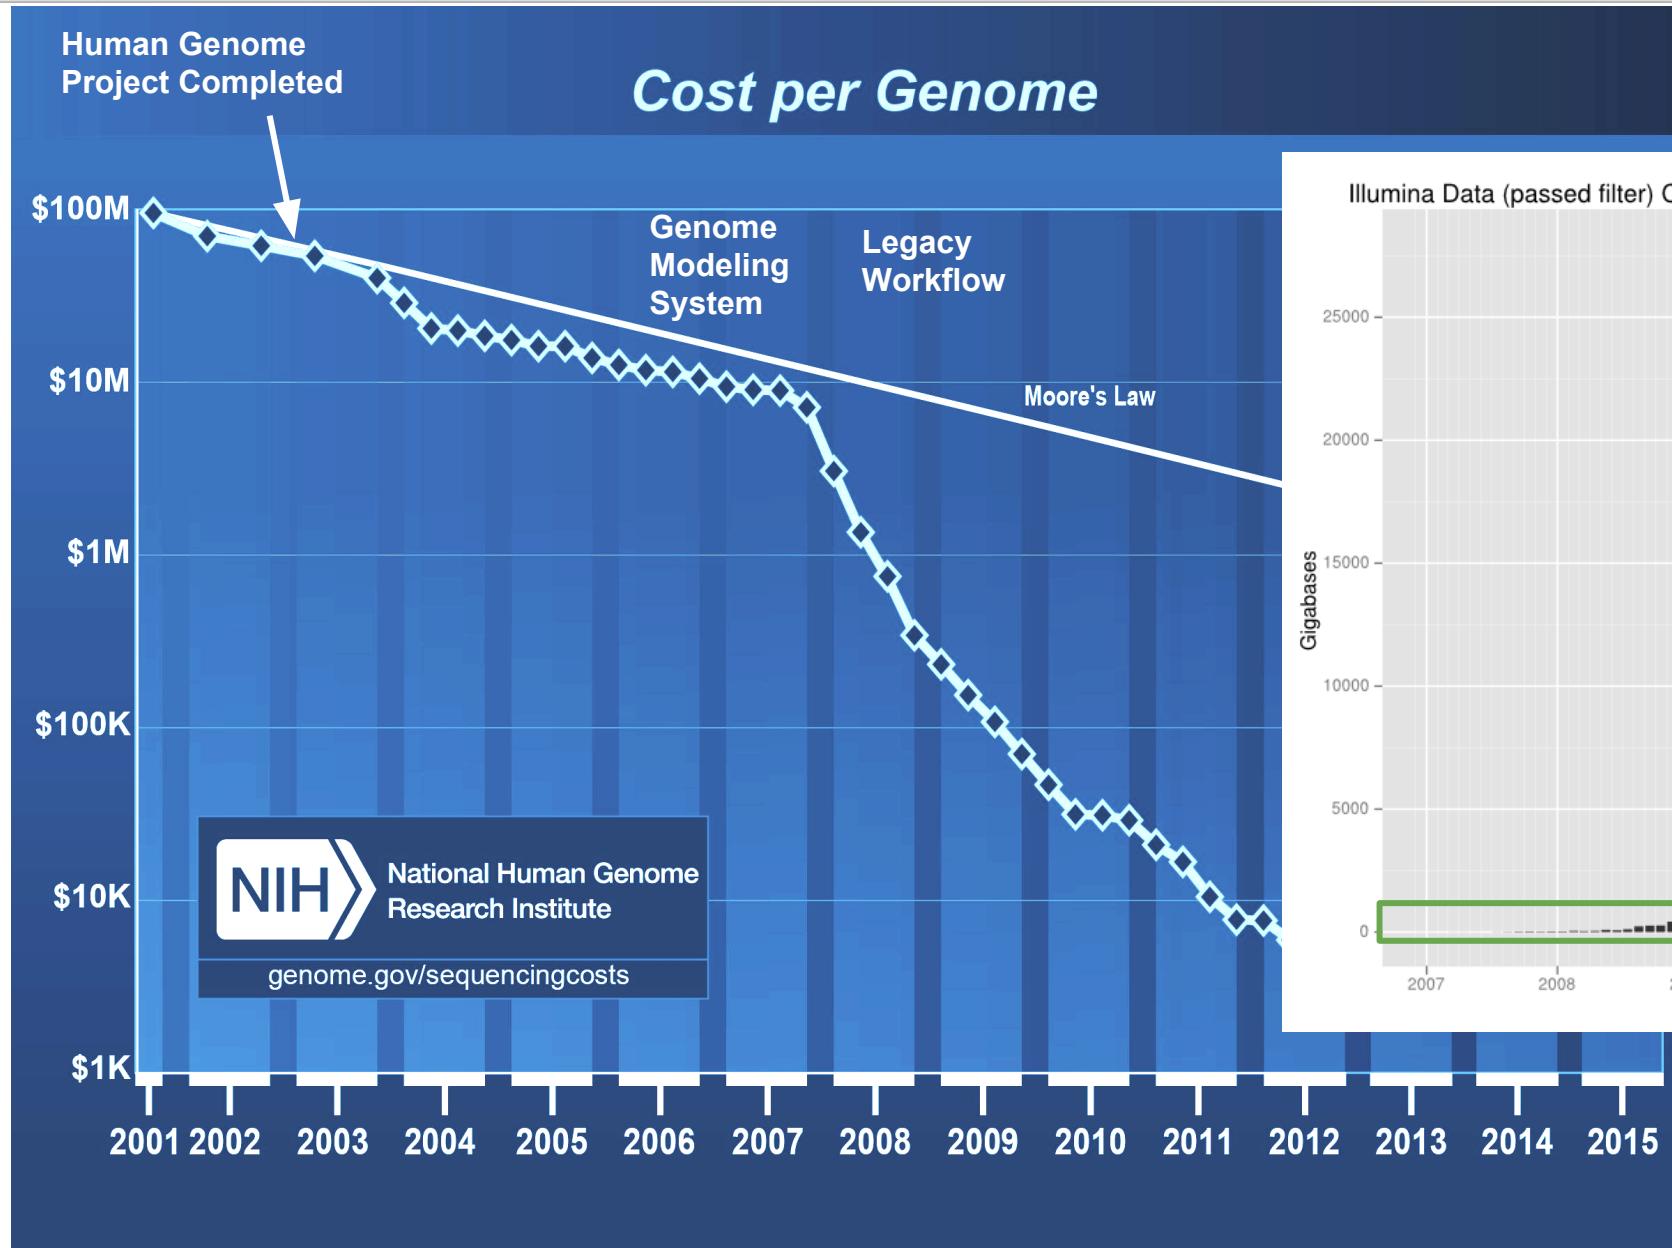
# History and Background

## Legacy Workflow

- Declarative (xml based workflow description)
- Integrates with datacenter management platform (IBM's LSF)
- Introduces concept of parallel-by
- Has centralized workflow tracking



# History and Background



# History and Background

## Legacy Workflow

### Features

- Declarative (xml based workflow description)
- Integrates with datacenter management platform (IBM's LSF)
- Introduces concept of parallel-by
- Has centralized workflow tracking



### Limitations

- Cannot scale up
- Highly coupled to Perl, UR and GMS
- Represents workflows as DAGs



# History and Background

|     | Dependency | Concurrency | Choice |
|-----|------------|-------------|--------|
| DAG | Yes        | Yes         | No     |
| FSM | Yes        | No          | Yes    |
| P/T | Yes        | Yes         | Yes    |

## Terms

DAG - Directed Acyclic Graph

FSM - Finite State Machine

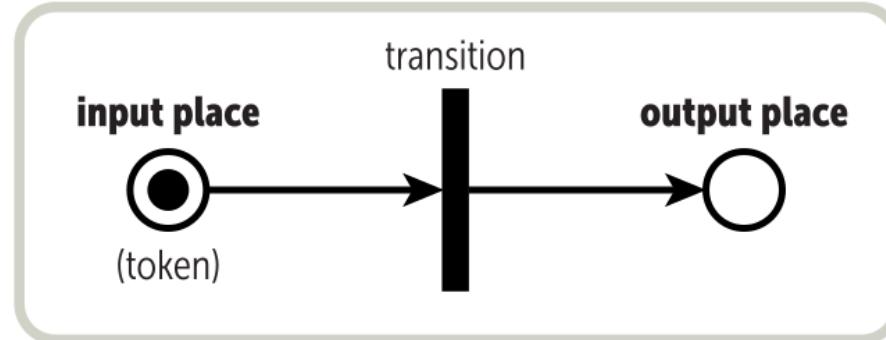
P/T - Place/Transition or Petri Net



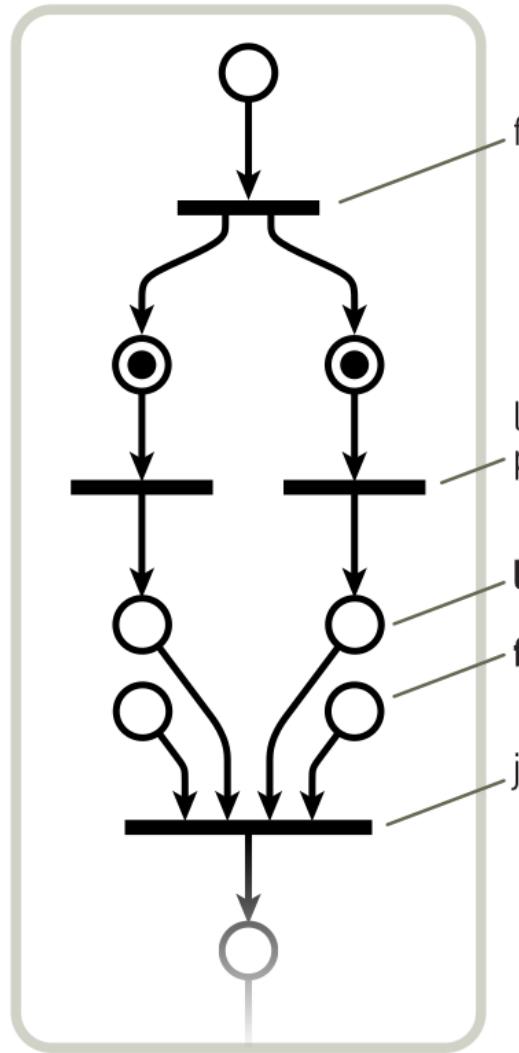
# History and Background

## Petri Net Basics

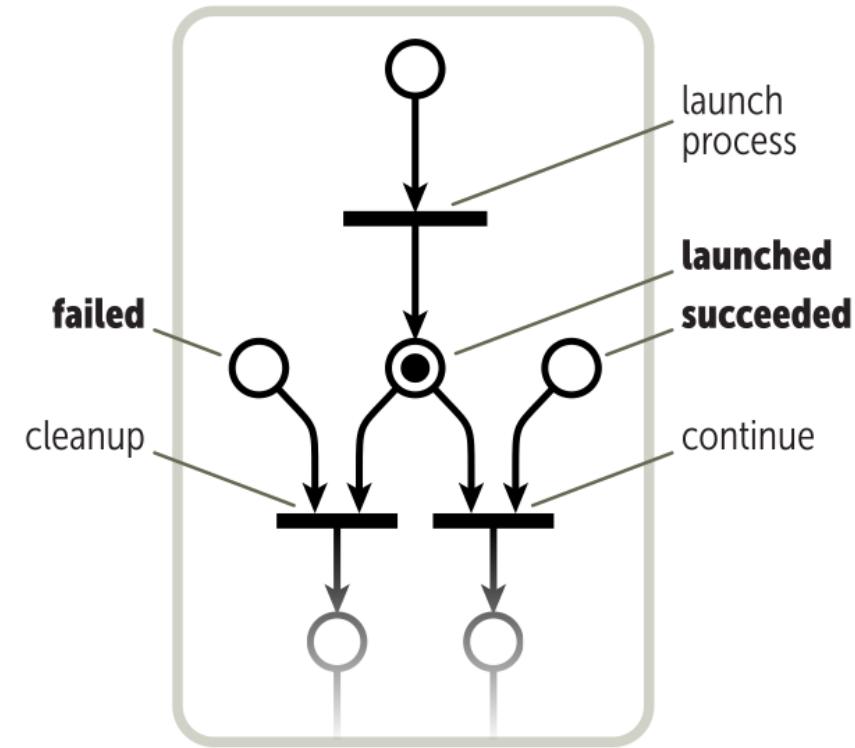
- Places represent **states**.
- Transitions perform **actions**.
- Transitions fire iff all input places have a token.
- After a transition fires, one token is placed in each output place.



# History and Background



Concurrency



Choice



# History and Background

## Flow

### Features

- Scales up nicely
  - Number of concurrent workflows
  - Number of tasks in a workflow
- Service Oriented
- High Reliability

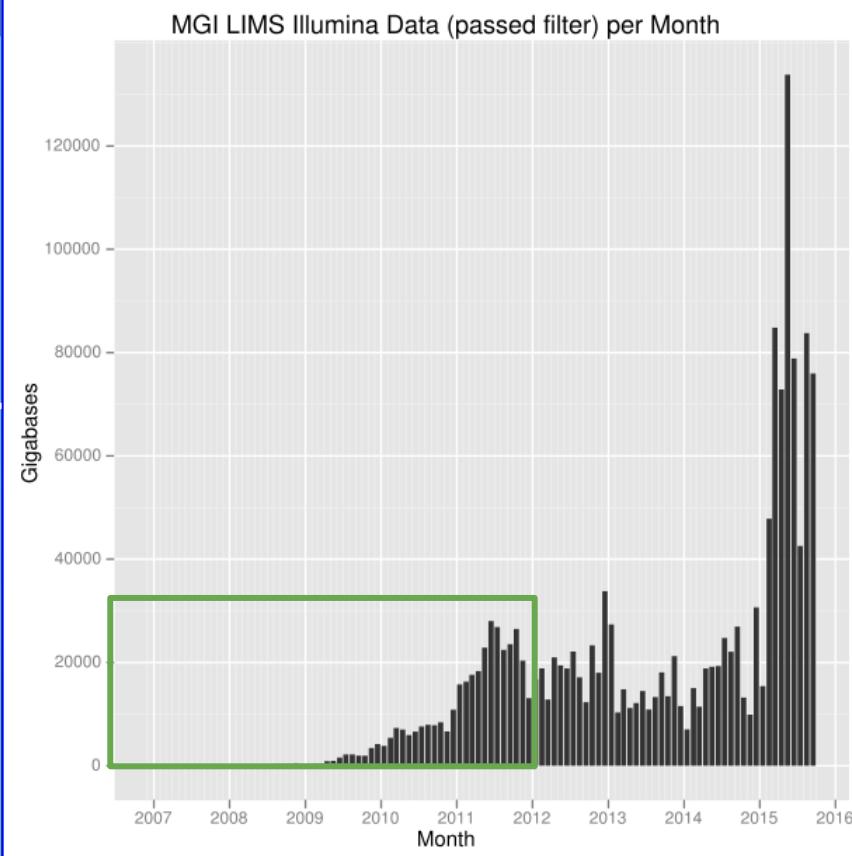
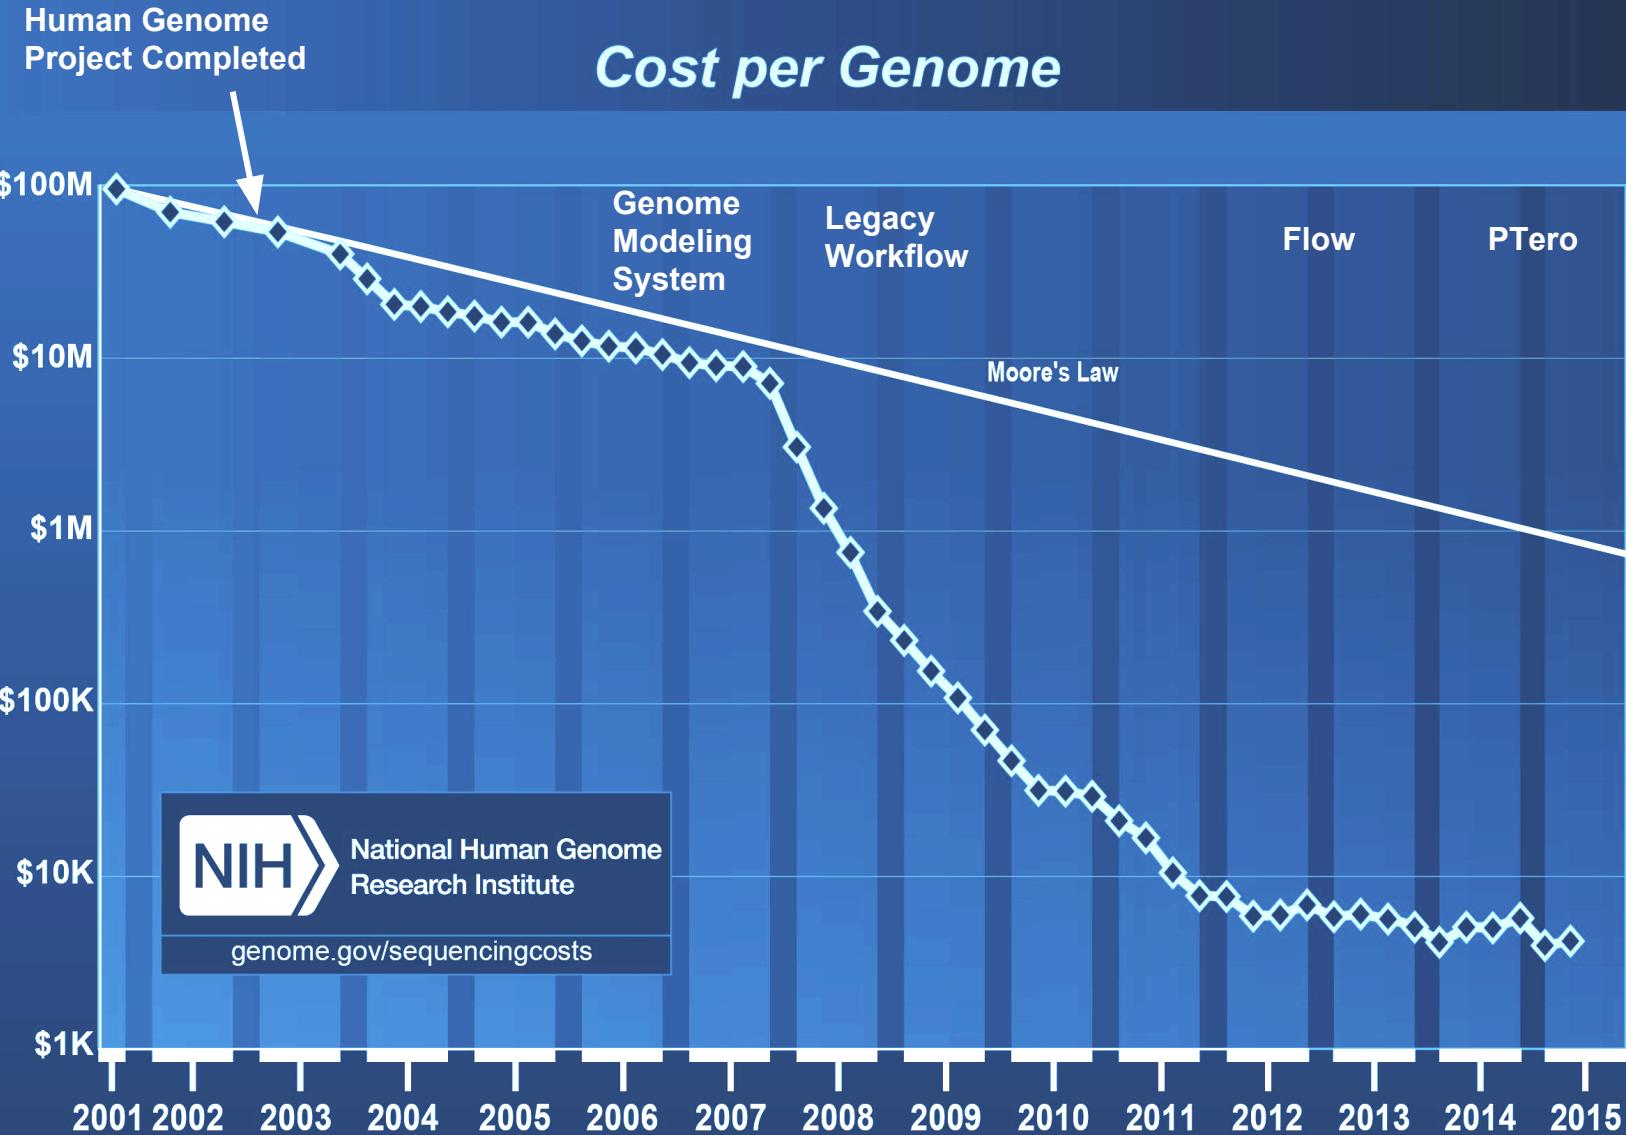


### Limitations

- Heterogeneous and complicated APIs
- Legacy workflow ties
  - Workflow status stored in old DB
  - Workflows represented with same xml

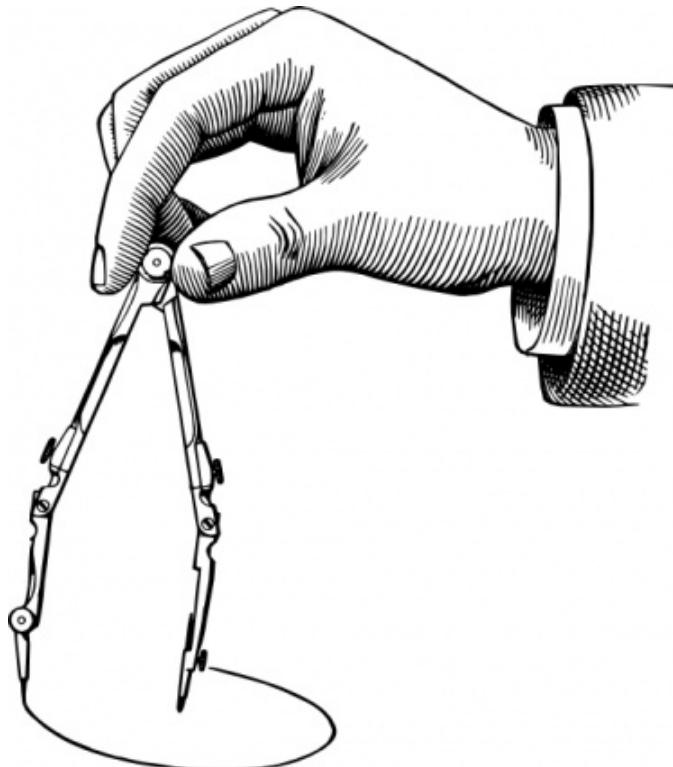


# History and Background



# History and Background

## PTero Design Goals



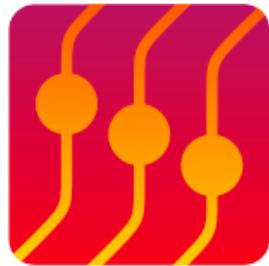
- Same ability to scale as Flow
- Service oriented with simple RESTful APIs
- JSON based workflow description
- Pluggable job execution services
- Outside adoption



# PTero Architecture



REST API



## Petri Service

- ★ Execute a Petri net



## Workflow Service

- ★ Execute a workflow
- ★ Instigate interactions with Petri service and Job services
- ★ Manage task inputs and outputs



## Job Services

### ShellCommand:

- ★ Execute a shell command given a user, working directory and environment

### LSF:

- ★ Submit a job as a given user specifying resources and other LSF options



# PTero Architecture

## Terminology:

**Task** - Something to complete

**Input/Output** - A JSON data structure

**Workflow** - A DAG of Tasks + inputs

**Method** - A way to finish a Task

come in two varieties:

Job - type of method that involves a job service

DAG - type of method that is a DAG of Tasks

**Execution** - A record of running a Method or Task

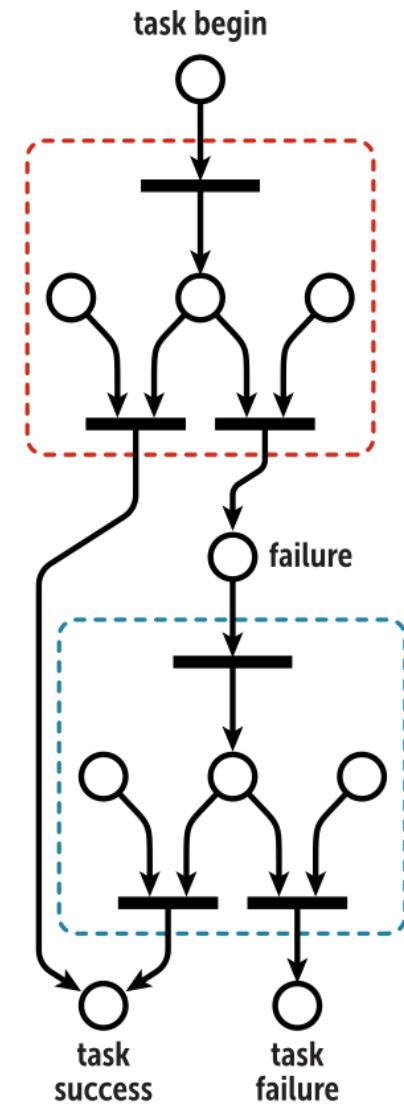
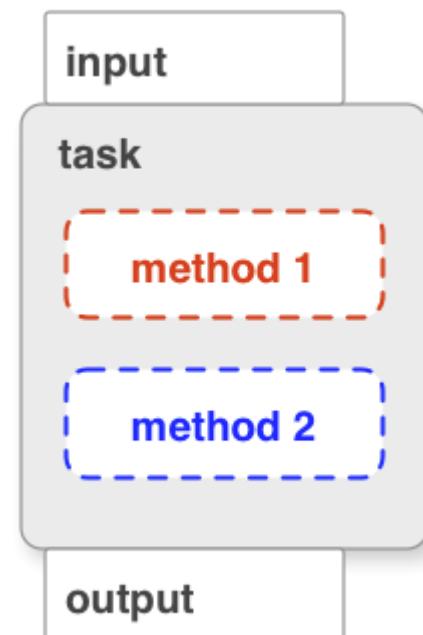


## Task Notation

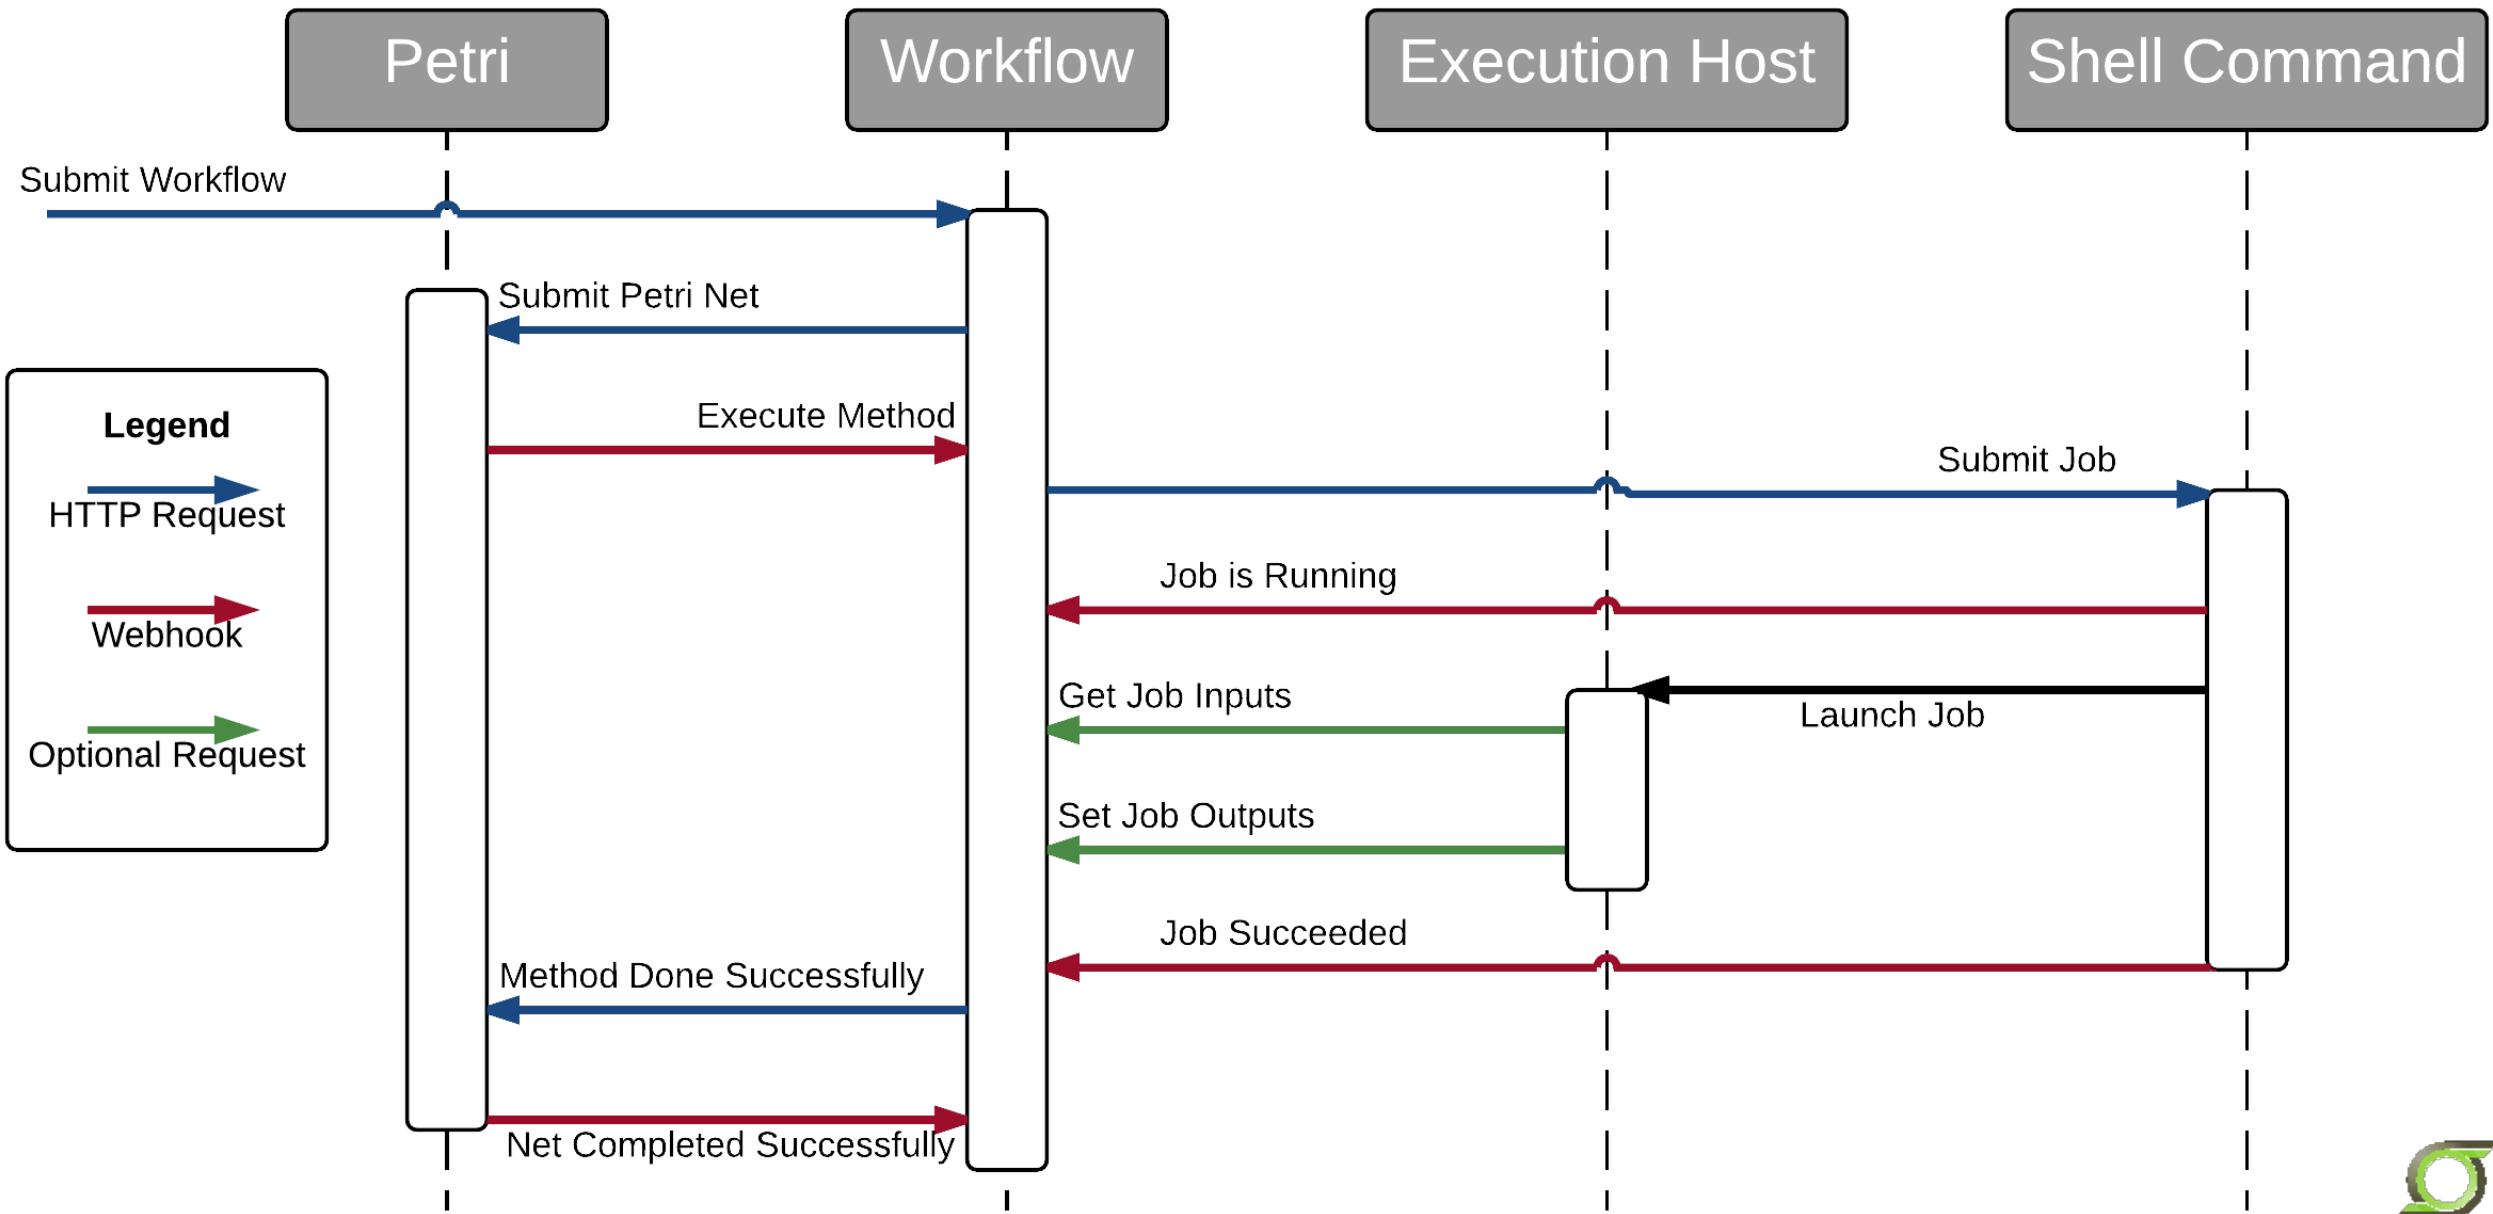


# PTero Architecture

## Translating a Task into a Petri Net

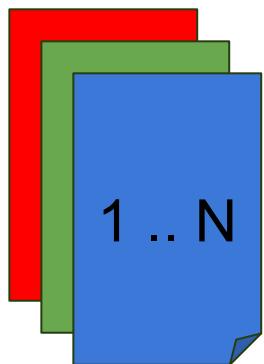


# PTero Architecture

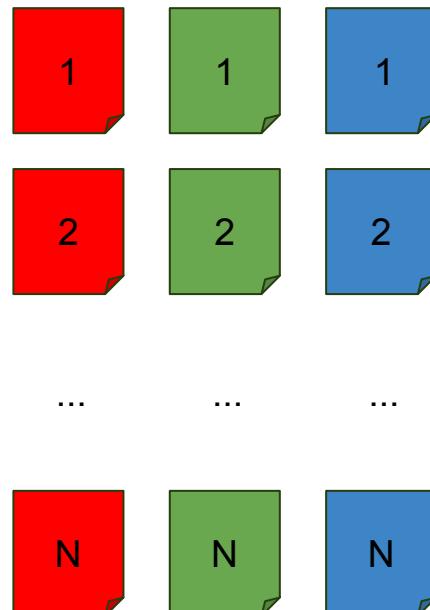


# PTero Architecture

One Sample Per File  
Many Chromosomes Per File



One Sample Per File  
One Chromosome Per File



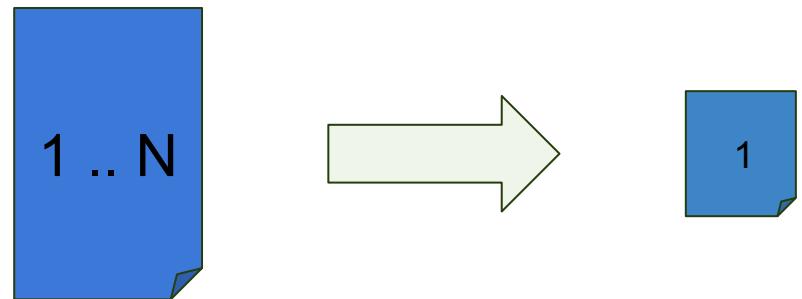
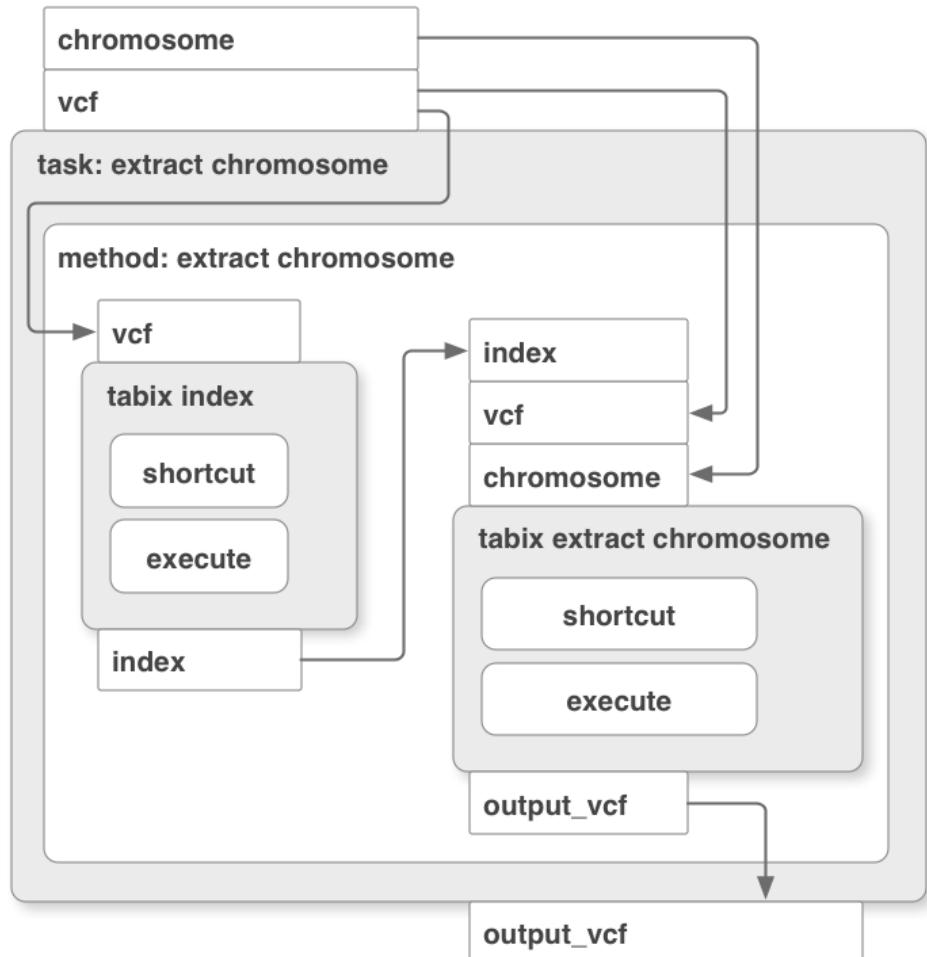
Many Samples Per File  
One Chromosome Per File



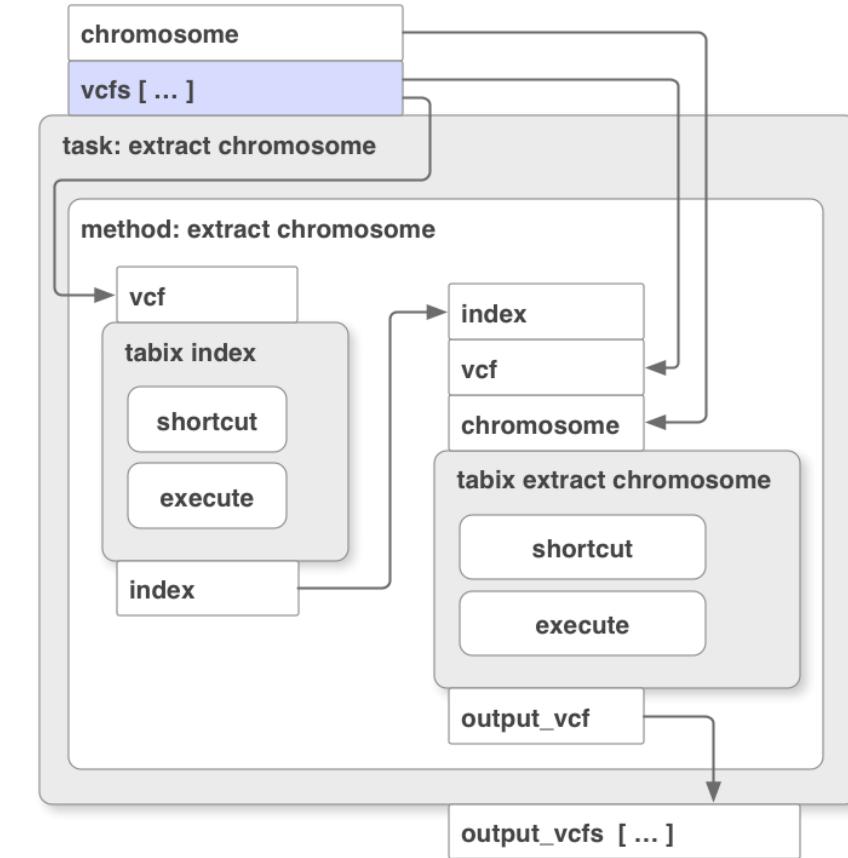
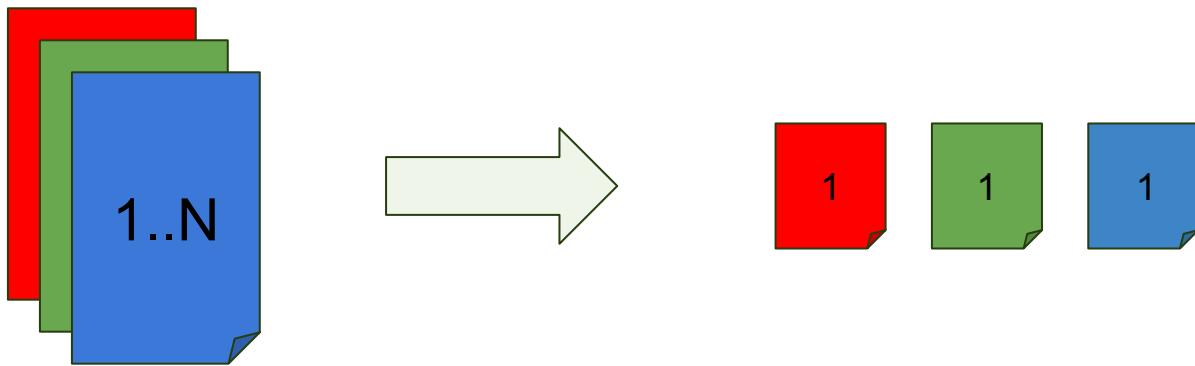
# PTero Architecture



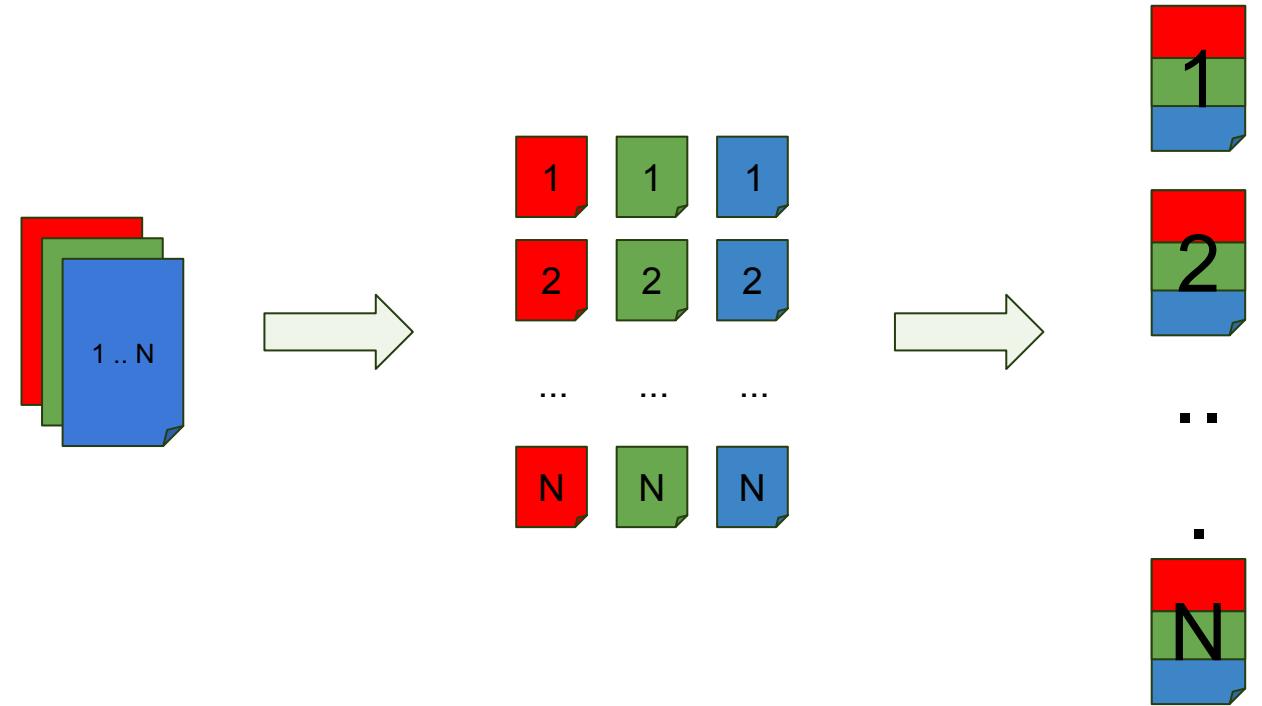
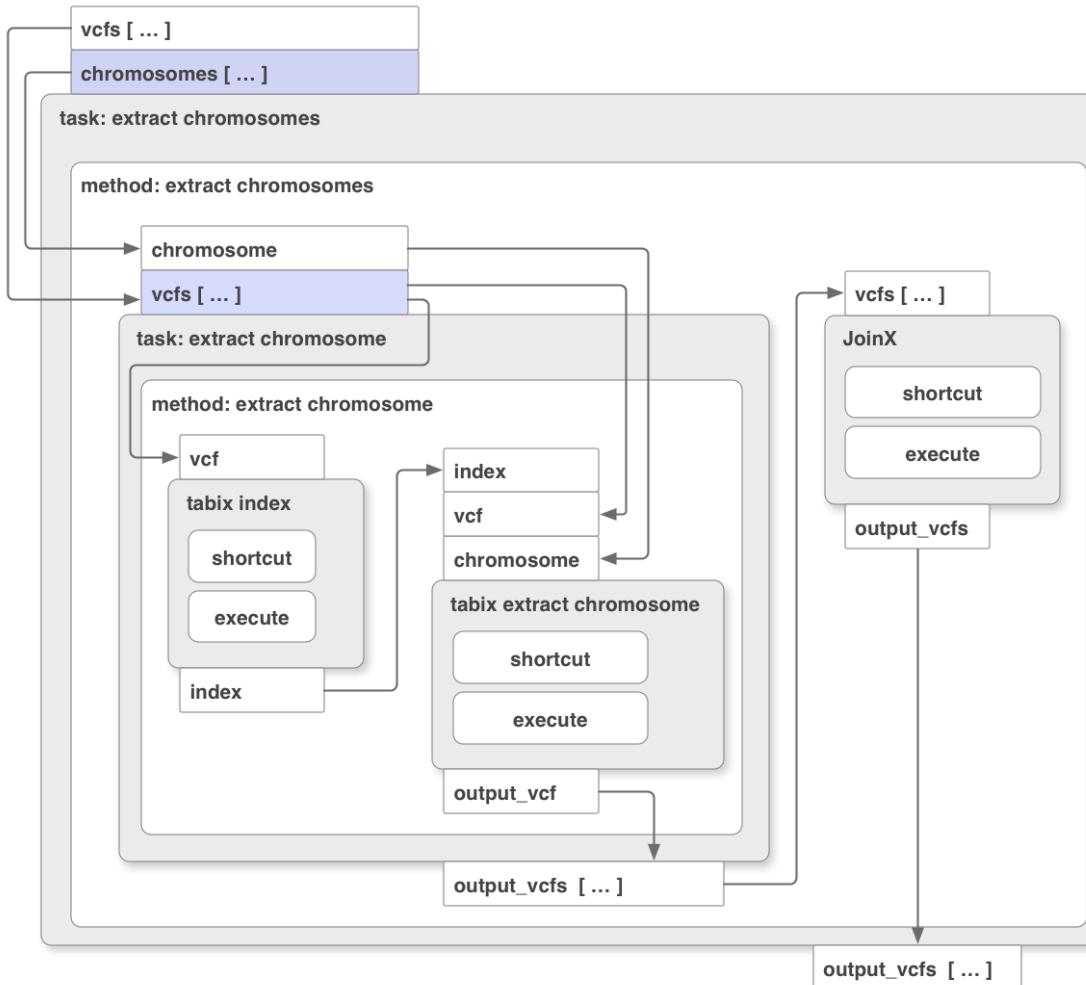
# PTero Architecture



# PTero Architecture



# PTero Architecture



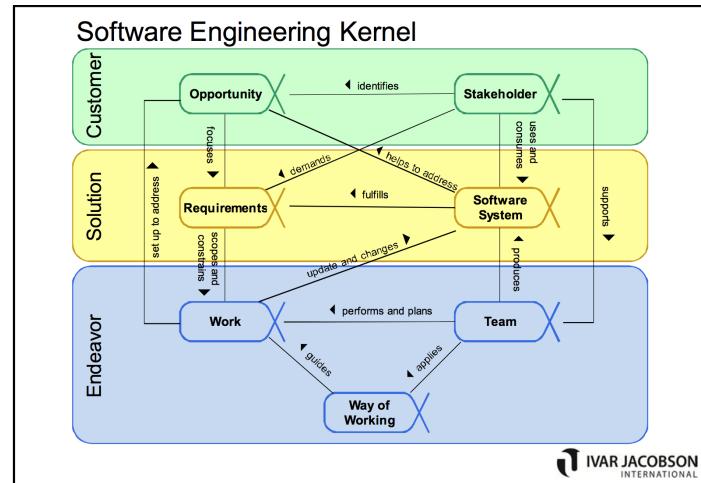
# PTero Architecture

```
1 $ ptero-view http://workflow-service/v1/workflows/5
2   TYPE      STATUS      STARTED      DURATION      NAME
3 Workflow succeeded 2015-09-23T23:30:55 00:00:18 FFFFJMALcRvKCKoGe7yPTzg
4 Task      succeeded 2015-09-23T23:30:57 00:00:16 merge chromosome vcfs [parallel-by: chromosome]
5 Task      succeeded 2015-09-23T23:30:57 00:00:16 . merge chromosome vcfs [0]
6 DAG      succeeded 2015-09-23T23:30:57 00:00:16 . . merge chromosome vcfs
7 Task      succeeded 2015-09-23T23:30:57 00:00:13 . . . extract chromosome from vcf [parallel-by: vcf]
8 Task      succeeded 2015-09-23T23:30:57 00:00:12 . . . . extract chromosome from vcf [0, 0]
9 DAG      succeeded 2015-09-23T23:30:58 00:00:11 . . . . . extract chromosome from vcf
10 Task     succeeded 2015-09-23T23:30:58 00:00:05 . . . . . . index_vcf task
11 Job      failed    2015-09-23T23:31:00 00:00:00 . . . . . . shortcut method
12 Job      succeeded 2015-09-23T23:31:01 00:00:02 . . . . . . index_vcf method
13 Task     succeeded 2015-09-23T23:31:03 00:00:06 . . . . . . extract_chromosome task
14 Job      failed    2015-09-23T23:31:04 00:00:02 . . . . . . shortcut method
15 Job      succeeded 2015-09-23T23:31:07 00:00:02 . . . . . . extract_chromosome method
16 Task     succeeded 2015-09-23T23:30:57 00:00:13 . . . extract chromosome from vcf [0, 1]
17 DAG      succeeded 2015-09-23T23:30:58 00:00:12 . . . . extract chromosome from vcf
18 Task     succeeded 2015-09-23T23:30:58 00:00:06 . . . . . . index_vcf task
19 Job      failed    2015-09-23T23:31:00 00:00:01 . . . . . . shortcut method
20 Job      succeeded 2015-09-23T23:31:03 00:00:01 . . . . . . index_vcf method
21 Task     succeeded 2015-09-23T23:31:04 00:00:05 . . . . . . extract_chromosome task
22 Job      failed    2015-09-23T23:31:06 00:00:01 . . . . . . shortcut method
23 Job      succeeded 2015-09-23T23:31:09 00:00:00 . . . . . . extract_chromosome method
24 Task     succeeded 2015-09-23T23:31:10 00:00:02 . . . join_vcfs task
25 Job      failed    2015-09-23T23:31:11 00:00:00 . . . . . . shortcut method
26 Job      succeeded 2015-09-23T23:31:12 00:00:00 . . . . . . join_vcfs method
27 Task     succeeded 2015-09-23T23:30:57 00:00:16 . merge chromosome vcfs [1]
28 DAG      succeeded 2015-09-23T23:30:57 00:00:16 . merge chromosome vcfs
29 Task     succeeded 2015-09-23T23:30:57 00:00:13 . . . extract chromosome from vcf [parallel-by: vcf]
30 Task     succeeded 2015-09-23T23:30:57 00:00:12 . . . . extract chromosome from vcf [1, 0]
31 DAG      succeeded 2015-09-23T23:30:58 00:00:11 . . . . extract chromosome from vcf
32 Task     succeeded 2015-09-23T23:30:58 00:00:05 . . . . . . index_vcf task
33 Job      succeeded 2015-09-23T23:31:00 00:00:00 . . . . . . shortcut method
34 Task     succeeded 2015-09-23T23:31:03 00:00:06 . . . . . . extract_chromosome task
35 Job      failed    2015-09-23T23:31:04 00:00:02 . . . . . . shortcut method
36 Job      succeeded 2015-09-23T23:31:07 00:00:02 . . . . . . extract_chromosome method
37 Task     succeeded 2015-09-23T23:30:57 00:00:13 . . . . extract chromosome from vcf [1, 1]
38 DAG      succeeded 2015-09-23T23:30:58 00:00:12 . . . . extract chromosome from vcf
39 Task     succeeded 2015-09-23T23:30:58 00:00:07 . . . . . . index_vcf task
40 Job      succeeded 2015-09-23T23:31:00 00:00:01 . . . . . . shortcut method
41 Task     succeeded 2015-09-23T23:31:05 00:00:04 . . . . . . extract_chromosome task
42 Job      failed    2015-09-23T23:31:06 00:00:01 . . . . . . shortcut method
43 Job      succeeded 2015-09-23T23:31:09 00:00:00 . . . . . . extract_chromosome method
44 Task     succeeded 2015-09-23T23:31:10 00:00:02 . . . join_vcfs task
45 Job      failed    2015-09-23T23:31:11 00:00:00 . . . . . . shortcut method
46 Job      succeeded 2015-09-23T23:31:12 00:00:00 . . . . . . join_vcfs method
```



# Development and Testing

## Development Methodology: SEMAT



## Software Engineering Method and Theory

<http://semat.org/>

<http://www.ivarjacobson.com/semat/>

## Software System Design: 12 Factor



### INTRODUCTION

In the modern era, software is commonly delivered as a service: called *web apps*, or *software-as-a-service*. The twelve-factor app is a methodology for building software-as-a-service apps that:

- Use declarative formats for setup automation, to minimize time and cost for new developers joining the project;
- Have a clean contract with the underlying operating system, offering maximum portability between execution environments;
- Are suitable for deployment on modern cloud platforms, obviating the need for servers and systems administration;
- Minimize divergence between development and production, enabling continuous deployment for maximum agility;
- And can scale up without significant changes to tooling, architecture, or development practices.

The twelve-factor methodology can be applied to apps written in any programming language, and which use any combination of backing services (database, queue, memory cache, etc).

### BACKGROUND

The contributors to this document have been directly involved in the development and deployment of hundreds of apps, and indirectly witnessed the development, operation, and scaling of hundreds of thousands of apps via our work on the Heroku platform.

<http://12factor.net/>

## Issue Tracking: GitHub and waffle.io

A screenshot of a GitHub issue board for the repository "genome/ptero". The board is divided into columns for Blocked (19 issues), Ready (39 issues), In Progress (2 issues), and Perl-Sdk (1 issue). Each issue card contains a title, a brief description, and labels indicating its type (e.g., enhancement, bug, workflow, deployment, apis, perl-sdk) and status (e.g., enhancement, bug, Phase 2 Deployment). The Perl-Sdk issue is labeled as a bug and has a note: "Relax parameter validation for job services". The Ready column includes an issue titled "Create Names/IDs for used services before submitting" and another titled "Make DB migrations separate command". The In Progress column includes an issue titled "Transition work from Adam to Josh". The Perl-Sdk column includes an issue titled "Implement additional validation checks".

<https://waffle.io/>

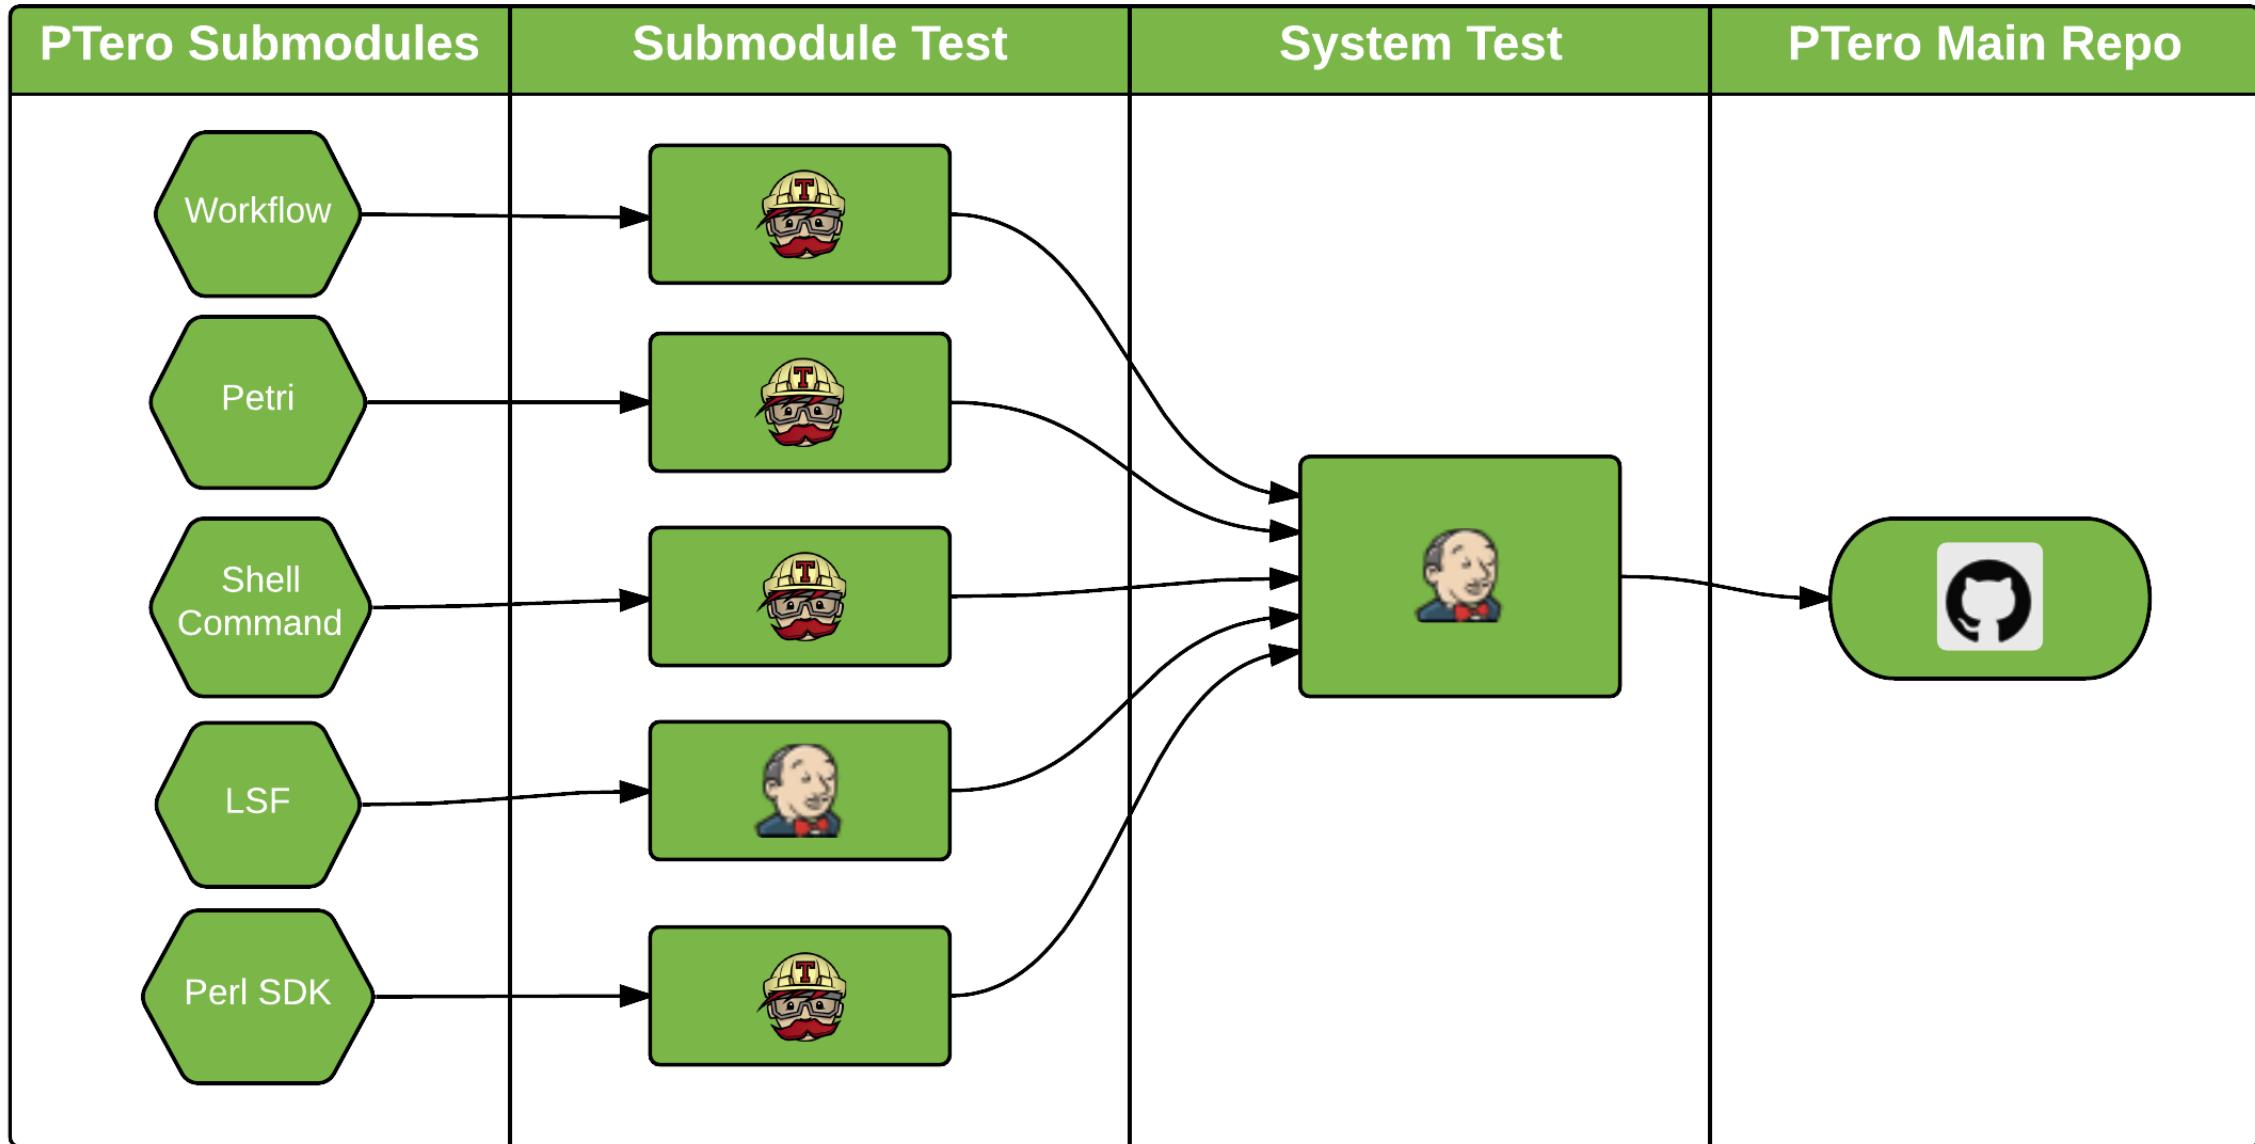


# Development

| Repo         | Language      | Sloccount    | Testing Sloccount | Test Coverage |
|--------------|---------------|--------------|-------------------|---------------|
| Workflow     | Python        | 2,965        | 1,957             | 93%           |
| Petri        | Python<br>lua | 1,765<br>222 | 1,873<br>-        | 95%<br>-      |
| ShellCommand | Python        | 291          | 647               | 83%           |
| LSF          | Python        | 486          | 452               | 76%           |
| PerlSDK      | Perl          | 1767         | 442               | 92%           |



# Development and Testing



# Deployment

## System tests as production monitoring

Jenkins

Project Ptero: Production Monitor

Project name: ptero\_production\_monitor\_perl\_sdk  
This runs tests against the production deployment of PTero.

Back to Dashboard Status Changes Workspace Build Now Delete Project Configure Rebuild Last Dependency Graph

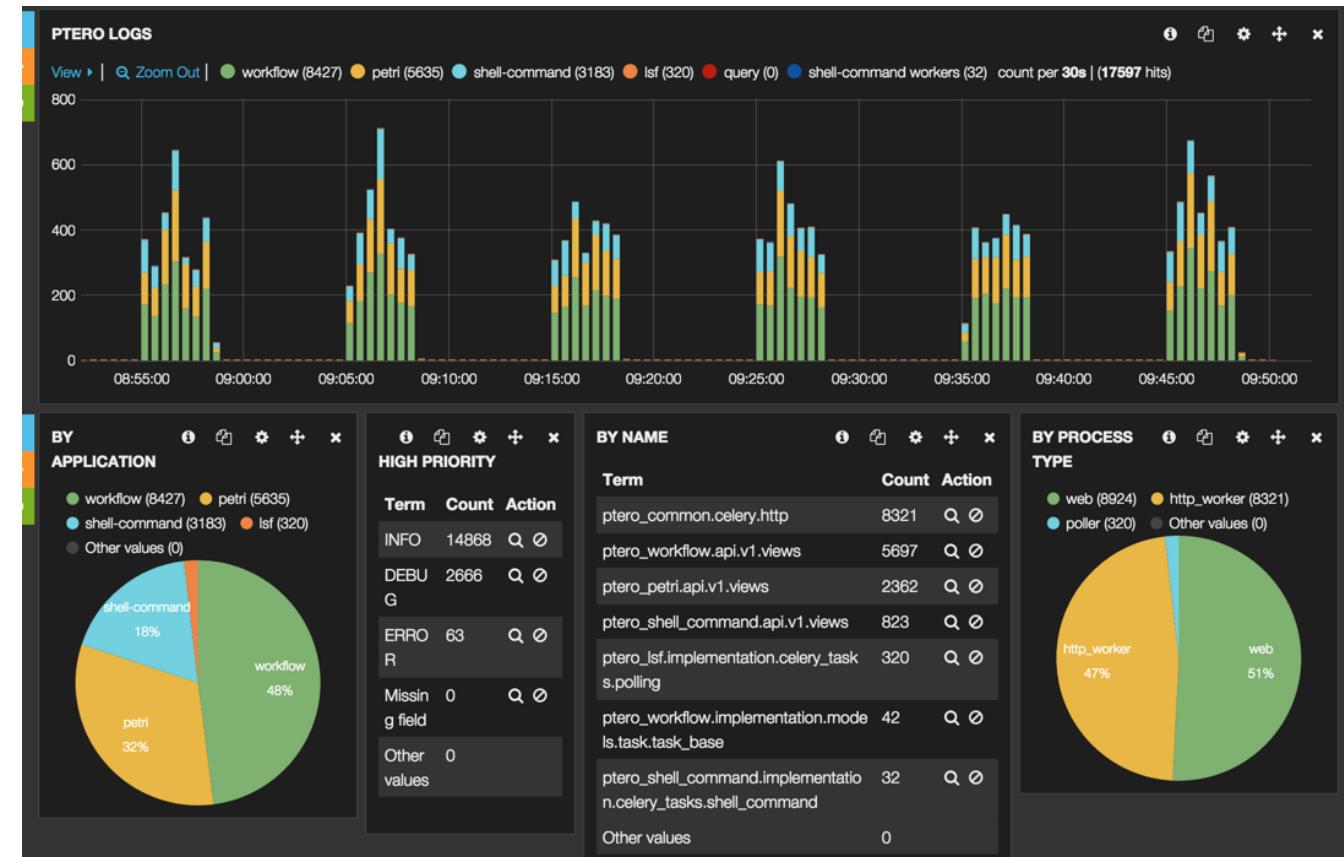
Build History trend

- Last build (#14679), 1 min 37 sec ago
- Last stable build (#14678), 11 min ago
- Last successful build (#14678), 11 min ago
- Last failed build (#14670), 1 hr 31 min ago
- Last unsuccessful build (#14670), 1 hr 31 min ago

Workspace Recent Changes

Permalinks

- Sep 23, 2015 8:45:01 PM #14679
- Sep 23, 2015 8:35:01 PM #14678 v0.1.10-lucid
- Sep 23, 2015 8:25:01 PM #14677 v0.1.10-lucid

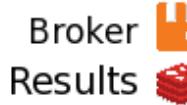
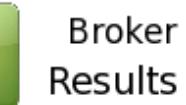
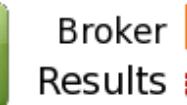
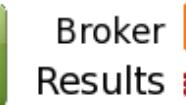
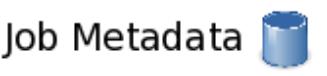


Logstash with Kibana Frontend



# Deployment

## PTero Processes

|         | Petri   | Workflow  | ShellCommand  | LSF   |
|---------|---|---|---|---|
| Backing |  Broker<br> Results<br><br> Petri Nets |  Broker<br> Results<br><br> Workflows,<br>Inputs/Outputs |  Broker<br> Results<br><br> Job Metadata |  Broker<br> Results<br><br> Job Metadata |
| Scaling | web<br>http_worker<br>worker<br><br>   | web<br>http_worker<br>worker  | web<br>http_worker<br>worker *  | web<br>http_worker<br>poller<br>scheduler<br>worker **<br>updater **  |



Celery



RabbitMQ



Redis



PostgreSQL

\* Deploy where you want shell-commands to be executed

\*\* Deploy where you have ability to submit and query for LSF jobs (must be LSF submit host)



# Development and Testing

```
git clone http://github.com/genome/ptero.git  
cd ptero  
git submodule update --init  
RUN_TESTS=1 vagrant up
```



# PTero Architecture

POST to <http://ptero-workflow/v1/workflows>

```
{  
  "inputs": {},  
  "links": [  
    {  
      "destination": "output connector",  
      "source": "Say Hello"  
    },  
    {  
      "destination": "Say Hello",  
      "source": "input connector"  
    }  
,  
  "tasks": {  
    "Say Hello": {  
      "methods": [  
        {  
          "name": "print hello",  
          "parameters": {  
            "commandLine": [  
              "bash",  
              "-c",  
              "echo 'hello'; echo 'world!' 1>&2"  
            ],  
            "environment": {},  
            "user": "vagrant",  
            "workingDirectory": "/home/vagrant/ptero/services/workflow/tests/scripts"  
          },  
          "service": "job",  
          "serviceUrl": "http://localhost:5000/v1"  
        }  
      ]  
    }  
  }  
}
```

- No inputs
- Control-flow links (no data to route)
- 1 Task named “Say Hello”
- 1 Method in that task named “print hello”



# Conclusion

---

## PTero is...

- designed to carry out large complex workflows
- service oriented with simple RESTful APIs
- focused in scope, heavily tested, extensible
- free and open source software
- under active development at

[github.com/genome/ptero](https://github.com/genome/ptero)



# Acknowledgements

## PTero Development Team



Travis Abbott



Mark Burnett



Charlotte Ferguson



Michael Kiwala  
[mkiwala@wustl.edu](mailto:mkiwala@wustl.edu)



Josh McMichael



David Morton  
[dmorton@wustl.edu](mailto:dmorton@wustl.edu)

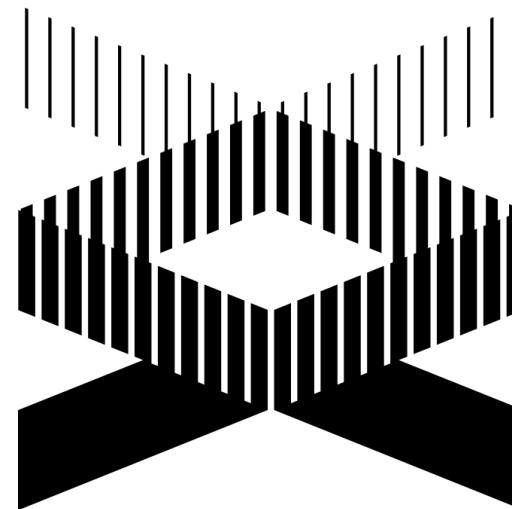


# People

We are the McDonnell Genome Institute at Washington University. Haley Abel, Benjamin Ainscough, Derek Albracht, Mark Ames, Lucinda Fulton, Elizabeth Appelbaum, Rachel Austin, Abhijit Badve, Matthew Bailey, Henry Bauer, Michael Becker, Edward Belter, Angela Bentlage, Jessica Bentlage, Thomas Biewald, Andrew Bohannon, Kirk Brege, Anthony Brummett, Theresa Caesar, Matthew Callaway, Song Cao, Darren Casimere, Wan Ching Chan, Gue Su Chang, Colby Chiang, Young-Jun Choi, Mari Jo Clark, Adam Coffman, Thomas Connell, Lisa Cook, Matthew Cordes, Laura Courtney, Amy D'Albora, John Dalbora, Indraniel Das, Anthony DeLuca, Tracie DeLuca, Sara Decarlo, Brandon Delgado, Ryan Demeter, Brian Derickson, Li Ding, Feiyu Du, James Eldred, Efrem Exum, Candace Farmer, Kelley Foyil, Rachel Frobel, Catrina Fronick, William Fronick, Robert Fugina, Bradley Fulton, Miranda Fulton, Robert Fulton, Kristi Futhey, Eric Geldmacher, Barbara Gillam, Chakravarthy Girda, Jennifer Godfrey, Susan Grasso, Julia Gray, Malachi Griffith, Obi Griffith, Jeanne Grigsby, David Gudermuth, Ira Hall, Terri Hall, Amy Hawkins, Todd Hepler, Jonathan Heusel, Jennifer Hodges, Kuan-lin Huang, Kimberly Hughes-Barnes, Stephanie Jackson, Reyka Jayasinghe, Rodney Jones, Krishna Kanchi, Chul Joo Kang, Kimberly Keen, Shamika Ketkar, Kyung Kim, Michael Kiwala, Daniel Koboldt, Sara Kohlberg, Colin Kremitzki, Milinn Kremitzki, Scott Kruchowski, Kilannin Krysiak, Jason Kunisaki, Nathan Kunkee, Sheila Lakanen, David Larson, Shawn Leonard, Shin Leong, Robert Lesurf, Andrew Levy, Shannon Lewis, Timothy Ley, Tiandao Li, Tina Lindsay, Charles Lu, Amy Ly, Sandra MacMillan, Vincent Magrini, Lenon Maguire, Christopher Maher, Jennifer Maher, Elaine Mardis, Christopher Markovic, John Martin, Robert Mashl, Judith McCart, Sean McGrath, Michael McLellan, Joshua McMichael, Samantha McNulty, Brian Meiningar, Karyn Meltz Steinberg, Kathie Mihindukulasuriya, Christopher Miller, Jordan Minx, Patrick Minx, Makedonka Mitreva, Deborah Moeller, Michael Montague, Thomas Mooney, David Morton, Pamela Nangle, William Nash, Joanne Nelsonkary, Ricky Ng, Mike Nhan, Thomas Nicholas, Jie Ning, David O'Brien, Gretchen O'Donnell, Michelle O'Laughlin, Kerri Ochoa, Bradley Ozenberger, Philip Ozersky, Stephanie Parrish, Josh Peck, Kymberlie Pepin, Allegra Petti, Craig Pohl, Robert Pokorney, Eric Ponce, Avinash Ramu, Allison Regier, Amy Reily, Susan Rock, Irina Ronko, Bruce Rosa, Kelsi Rotter, Ryan Rupp, Amy Sansone, Kyriena Schatzkamer, Debra Scheer, Heather Schmidt, Michael Schmidt, William Schroeder, Adam Scott, Sohini Sengupta, Cherilynn Shadding, Nicholas Sheehan, Susanna Siebert, Zachary Skidmore, Julie Smith, Steven Spargur, Nicholas Spies, Craig Steinmeyer, Scott Stevens, Bruce Striegel, Cynthia Strong, Dawn Sutter, Kenneth Swanson, Andrea Taylor, Thynn Thane, Brenda Theising, Brianne Thomeczek, Valerie Thompson, Chad Tomlinson, Lee Trani, Nikki Trapp, Ewanne Trevaskis, Richard Trippeer, Rusudan Turabelidze, Rahul Tyagi, Joelle Veizer, Tammi Vickery, Alex Wagner, Jason Waligorski, Jason Walker, Latricia Wallace, Grace Wang, Jiayin Wang, Wesley Warren, Brian Watts, James Weible, Matthew Weil, Michael Wendl, Brian White, Richard K. Wilson, Roxanne Wilson, Aye Wollam, Dee Wu, Song Wu, Matthew Wyczalkowski, Mingchao Xie, Kai Ye, Venkata Yellapantula, Xu Zhang.



# Funding



National Human  
Genome Research  
Institute

Grant: U54 HG003079 (Dr Wilson)



# Technologies

Postgresql

Redis

RabbitMQ

Python

Ubuntu

CoreOS

Celery

Sqlalchemy

Alembic

Flask

Perl5

tox

Travis CI

gunicorn

honcho

jsonschema

Moose

Deis

vmWare

OpenStack

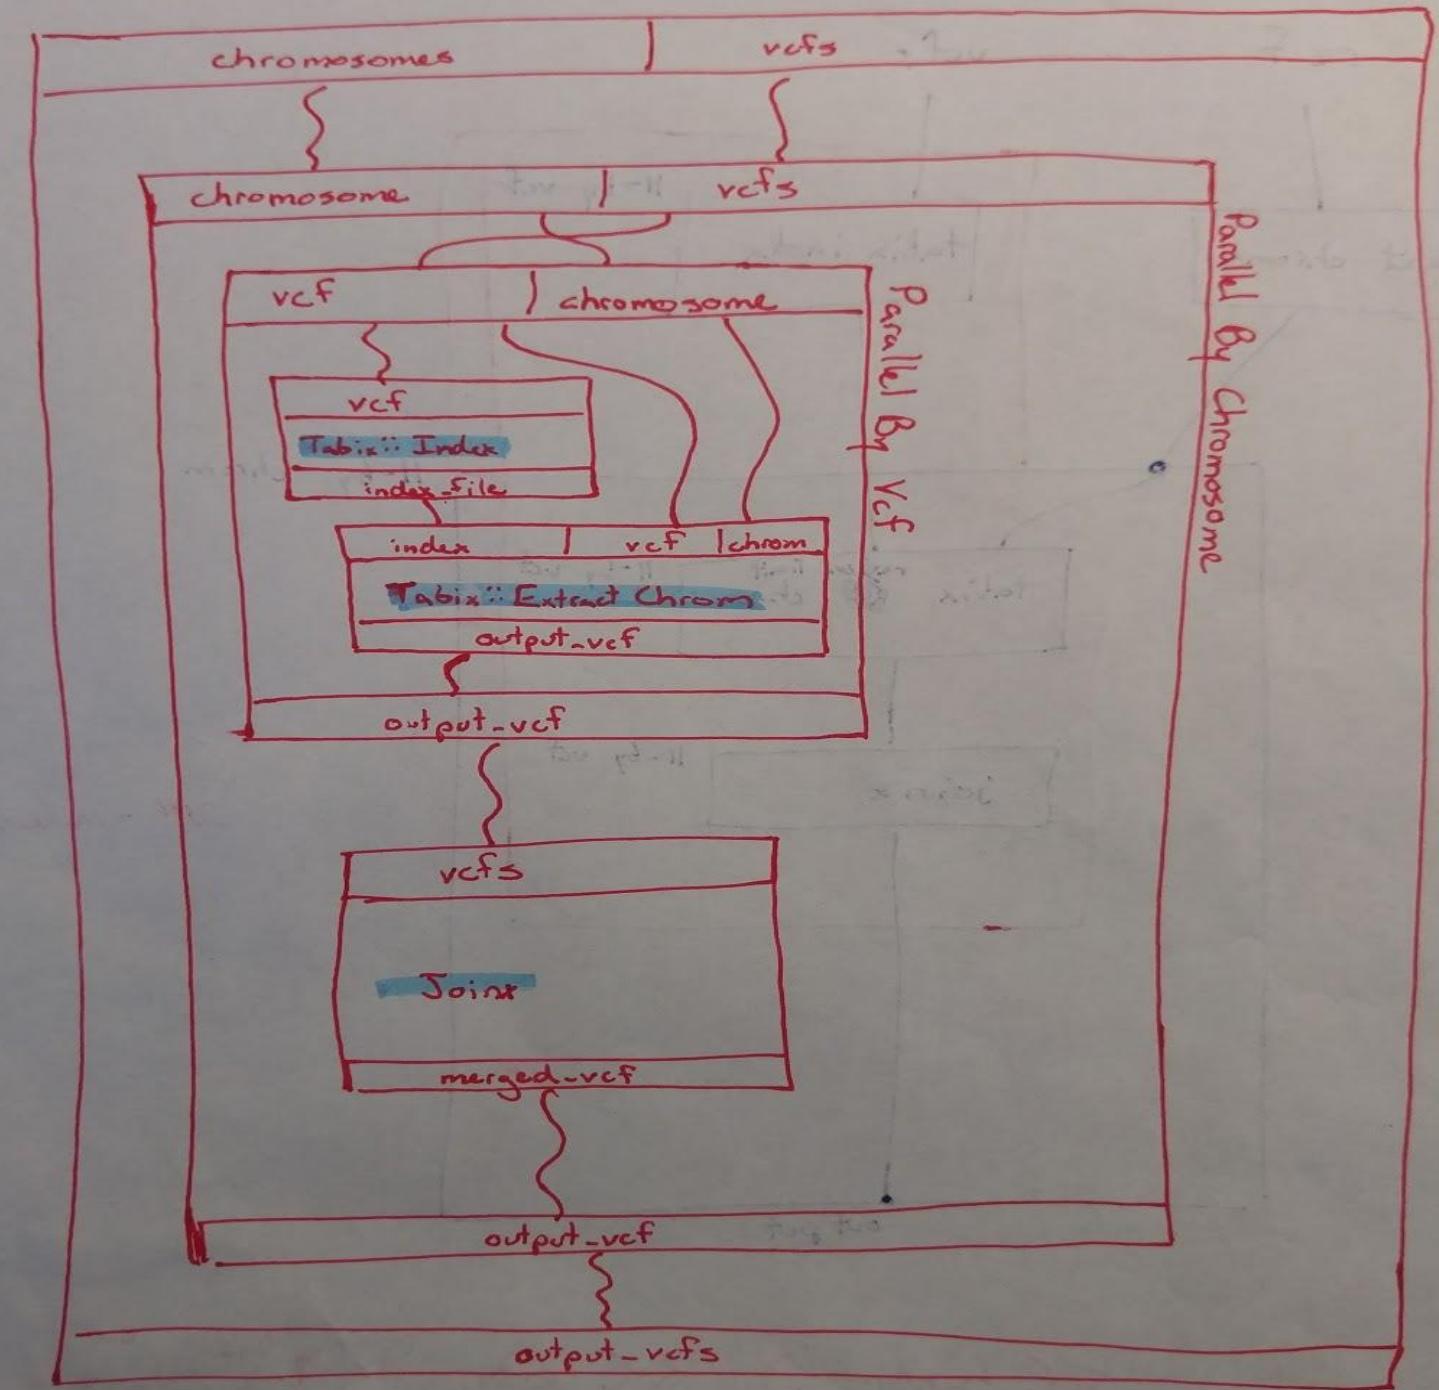
Heat Orchestration Template

VirtualBox

Docker







## History Context

### File Based Commands

with example command  
e.g. picard

These building block are not written by us.  
- wrapper  
- or use as-is

typically  
long running  
2h -  
2 weeks

### The Simplest WF: Scripted / Hard Coded

- Imperative
- On One Machine
- Manual Checkpointing
- failure/ completion tracking difficult (not centralized)

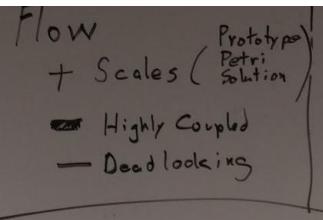
### Legacy WF (Pros)

DAG

- Describes work to be performed
- Integrates w/ LSF
- Centralized Tracking
- Shortcircuiting / Execute (Added Later)  
(FSM)

### Legacy WF (Limitations)

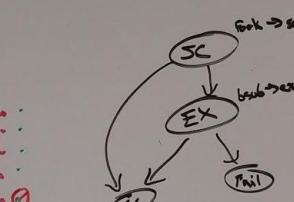
- Highly Coupled to Perl and to our GMS
- Scales poorly
  - + in width of a wf (phenotype correction; 10k's of tasks)
  - + # of concurrent WF's
  - + WF Deadlocking



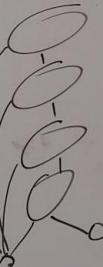
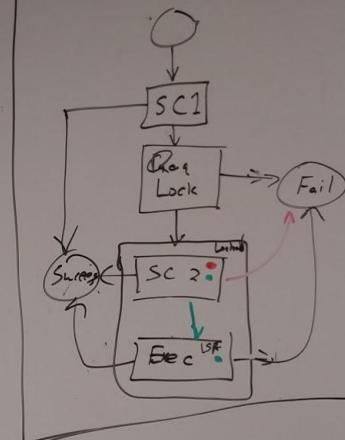
Shortcut

Execute

Set Result

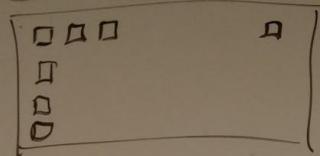


### Genome::Command



Development

Semant



12-factor

- heroku-style deployment
- deploy w/ git push
- Robust against ~~1 node~~ application failures
- Horizontal Scalability

Python

active community  
esp web apps  
team experience  
reuse code from Flow  
esp Petri code

GitHub PRs / Waffle.io Issue Tracking

1. to support open community
2. Waffle to aggregate issues and display kanban
3. Use of submodules
4. Travis CI integration/support  
Coveralls

Vagrant  
Develop and Test locally

Flake8

Test Coverage

How to get coverage reports  
while testing an SOA

100% Coverage - not the goal

Tests as monitors  
system tests; run against  
production deployment.

- Monitors:
- functionality
  - performance

Of not just the App, but the  
platforms and backing services.  
And app components are checked too.

Quick Comparison to Actual Prod Workflows

Ptero Itself:  
Architecture  
RESTful API  
Few Endpoints

Deployment

We use OS Heroku Clone  
Deis (was CoreOS)

We use a heat script to  
spin the whole thing up in OpenShift



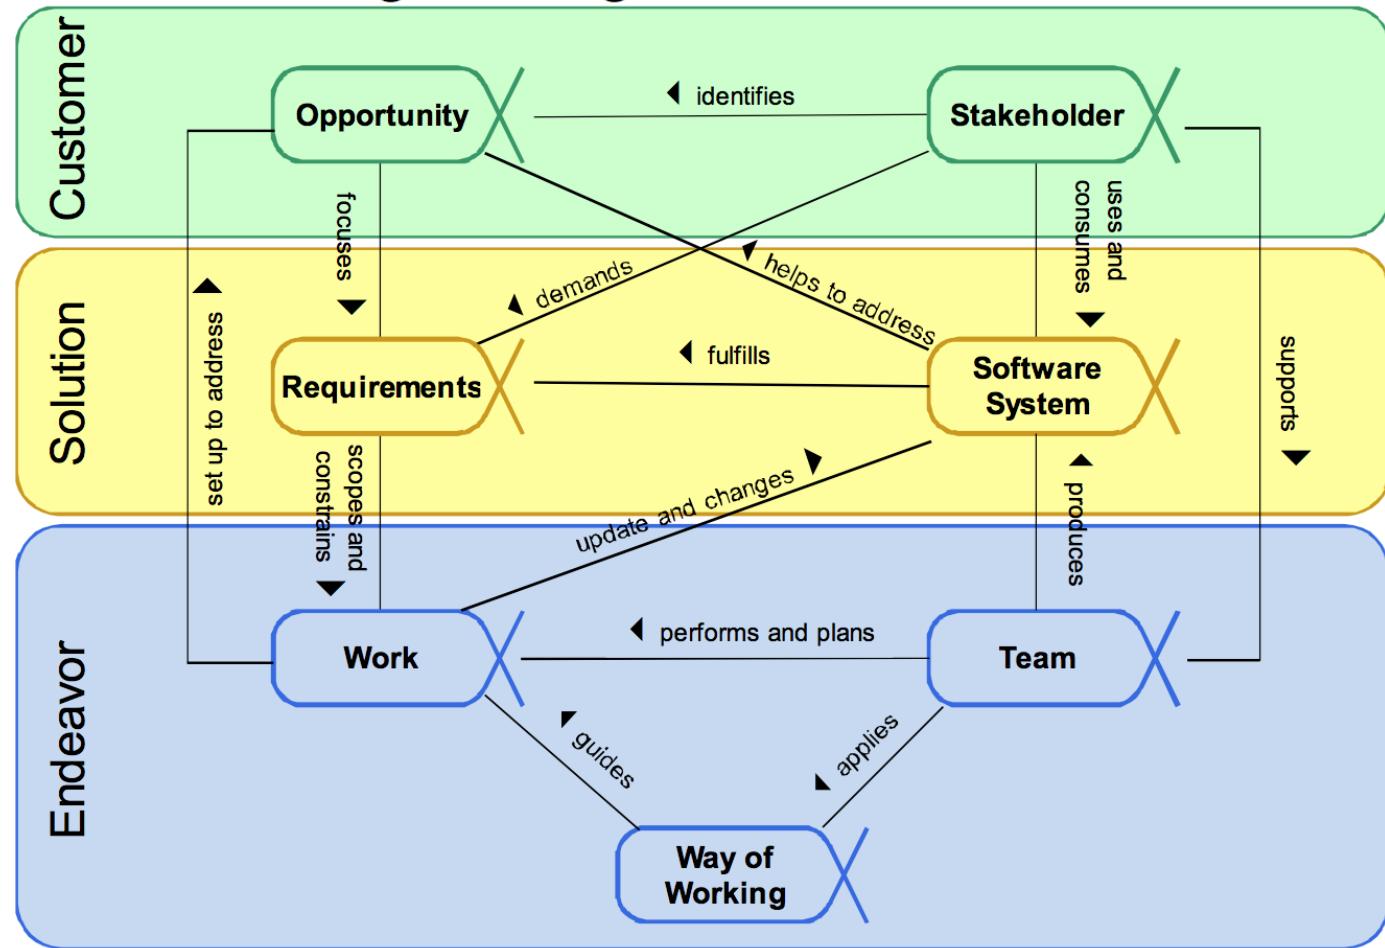
# Development Practices and Testing Choices

## SEMAT

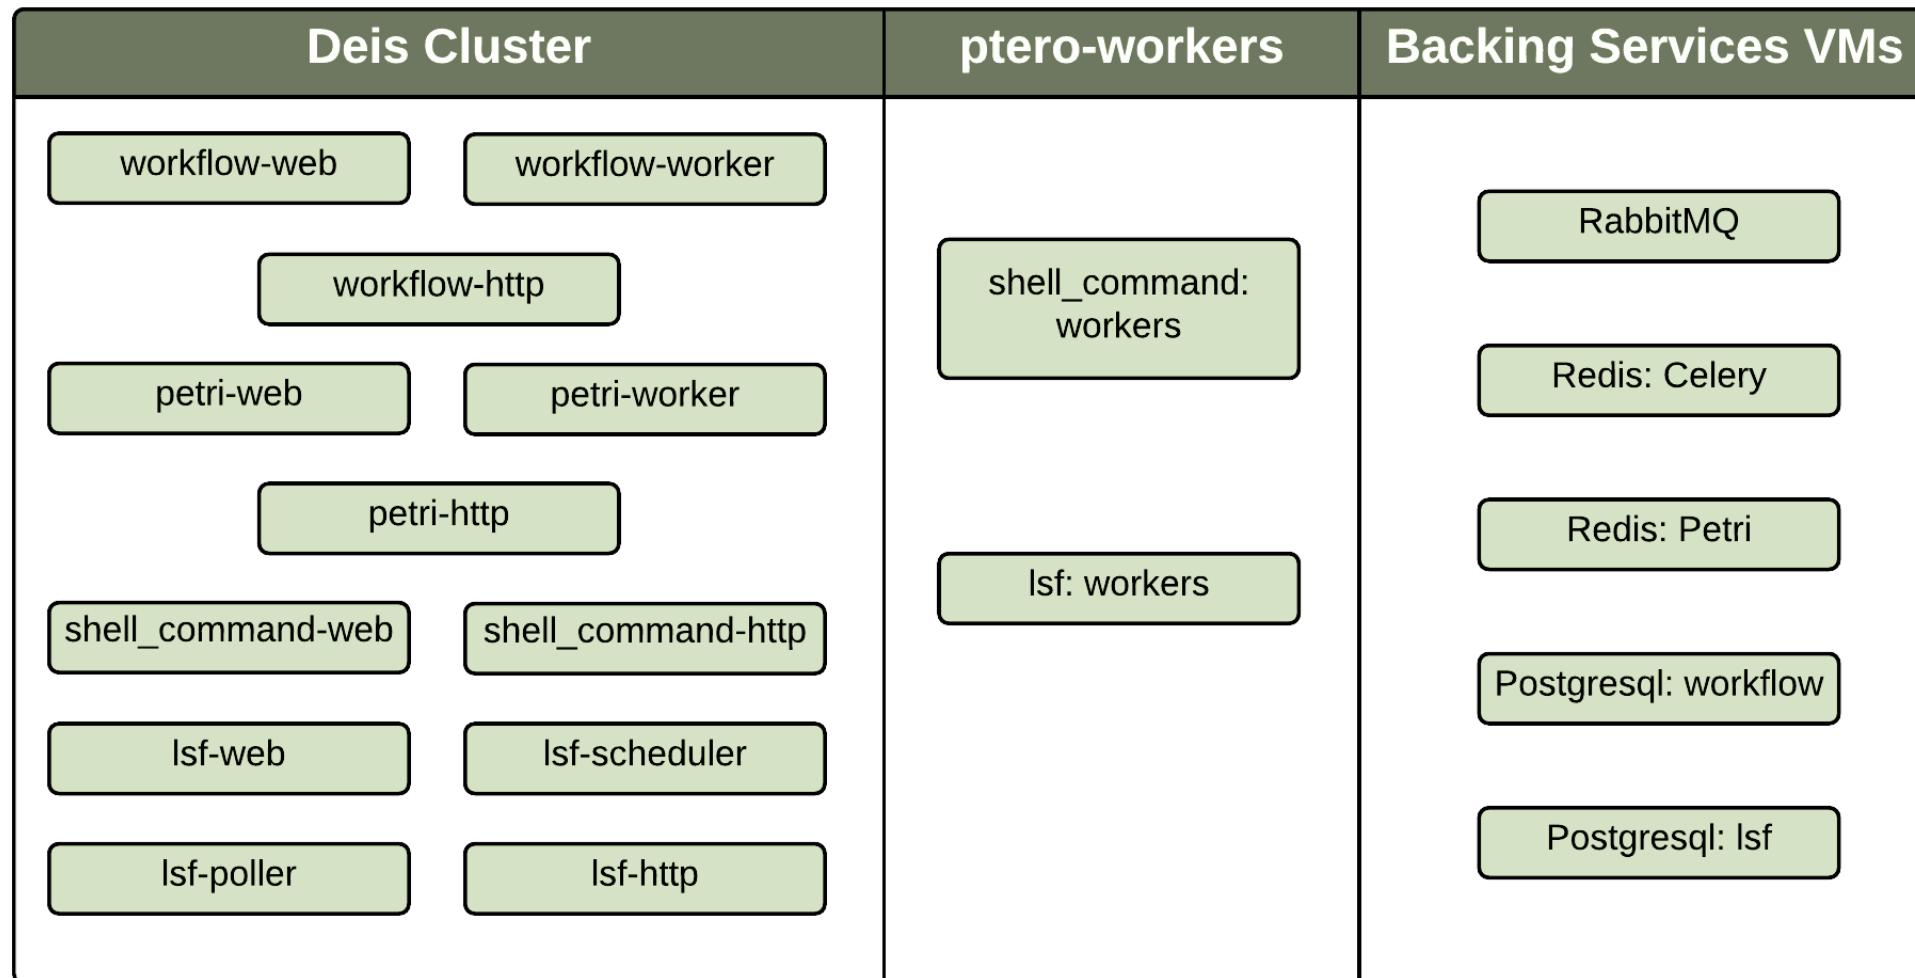
(Software Engineering Method and Theory)

- [semat.org](http://semat.org)
- [www.ivarjacobson.com/semat](http://www.ivarjacobson.com/semat)

## Software Engineering Kernel



# Deployment



# PTero Architecture

GET <http://ptero-workflow/v1/workflows/1>

-or-

GET <http://ptero-workflow/v1/workflows?name=3SYTDXt4Sp2eL4rtTdxjPg>

```
{  
  "name": "3SYTDXt4Sp2eL4rtTdxjPg",  
  "reports": {  
    "workflow-details": "http://localhost:7000/v1/reports/workflow-details?workflow_id=1",  
    "workflow-executions": "http://localhost:7000/v1/reports/workflow-executions?workflow_id=1",  
    "workflow-outputs": "http://localhost:7000/v1/reports/workflow-outputs?workflow_id=1",  
    "workflow-skeleton": "http://localhost:7000/v1/reports/workflow-skeleton?workflow_id=1",  
    "workflow-status": "http://localhost:7000/v1/reports/workflow-status?workflow_id=1",  
    "workflow-submission-data": "http://localhost:7000/v1/reports/workflow-submission-data?workflow_id=1"  
  },  
  "status": "succeeded"  
}
```

Workflow gets named automatically if one isn't specified at submission time.



# Development Practices and Testing Choices

