

Build Your own Distributed Framework on Mesos using Clojure

Presented By:

Edward Paget
Prasanna Gautam



Sept 24, 2015

Steps

1. Start Virtualbox VMs - Master and 2 Slaves
2. Generate a mesos-framework
3. Store State using Zookeeper
4. Recover from Failure
5. Leader Election
6. Web UI
7. Testing the Scheduler

What is Mesos



- Mesos is a cluster manager
- Handles scaling up resources
- Allows Dynamic Allocation
- Resource Isolation
- Works using resourceOffers to Frameworks

Mesos Essential Vocabulary

Term	Definition
Mesos Master	Coordinator, Offers Resources to Frameworks
Mesos Slaves	Where your tasks are run
Framework	Your Application on Mesos
Task	What Actually Runs
Scheduler	Handles Offers from mesos
Executor	Handles the launching of tasks
Zookeeper	Mesos uses for Coordination

Setup

Getting your Development Environment Ready

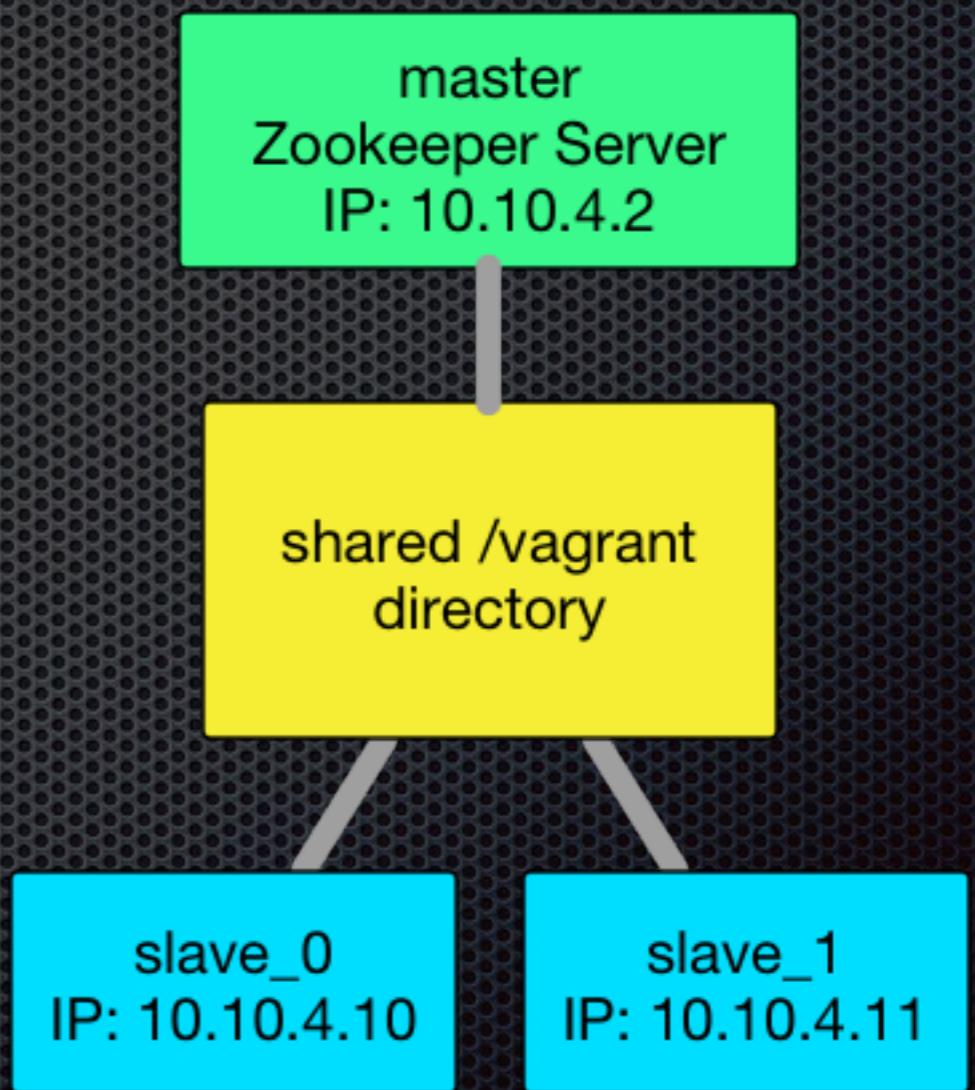
- Please insert the provided USB drives in your computers
- Requirements:
 - OS: Linux, OSX (Sorry we didn't test on Windows)
 - Latest Vagrant and Virtualbox
 - They're provided in the USB disk too

Setup

- Go to your USB Directory (/usb)
- * /Volumes/STL15MESOS on OSX
- * /media/STL15MESOS (most likely) on Linux
- Install Virtualbox from /usb/virtualbox
- Install Vagrant from /usb/vagrant/

Setup:Hello-Mesos

- cd /usb/code/hello-mesos
- vagrant up



Setup:Hello-Mesos

- run ./bin/diagnose-vagrant-setup.sh

```
→ hello-mesos git:(master) ✘ ./diagnose-vagrant-setup.sh
10.10.4.2 master
Success: Valid Master IP 10.10.4.2
imok
Success: Zookeeper is running on Master
2
Success: Valid number of slaves 2
10.10.4.2
Success: Master IP in the machine matches with 10.10.4.2
mesos-master start/running, process 6046
Success: mesos-master service is running
HTTP/1.1 200 OK
Success: mesos-master is listening on http://10.10.4.2:5050
Checking 10.10.4.10
HTTP/1.1 200 OK
Success: mesos-slave on 10.10.4.10 is reachable from Master
Checking 10.10.4.11
HTTP/1.1 200 OK
Success: mesos-slave on 10.10.4.11 is reachable from Master
Checking from slave_0
HTTP/1.1 200 OK
Success: mesos-master on http://10.10.4.2:5050 is reachable from slave_0
Checking from slave_1
HTTP/1.1 200 OK
Success: mesos-master on http://10.10.4.2:5050 is reachable from slave_1
```

Demo: (go)

- ❖ vagrant ssh master
- ❖ cd /vagrant
- ❖ lein repl
- ❖ (go)

```
user=> (go)
I0923 19:14:33.185351 2559 sched.cpp:157] Version: 0
2015-09-23 19:14:33,187:2534(0x7f117b676700):ZOO_INFO
2015-09-23 19:14:33,188:2534(0x7f117b676700):ZOO_INFO
2015-09-23 19:14:33,188:2534(0x7f117b676700):ZOO_INFO
2015-09-23 19:14:33,188:2534(0x7f117b676700):ZOO_INFO
2015-09-23 19:14:33,189:2534(0x7f117b676700):ZOO_INFO
2015-09-23 19:14:33,189:2534(0x7f117b676700):ZOO_INFO
2015-09-23 19:14:33,190:2534(0x7f117b676700):ZOO_INFO
2015-09-23 19:14:33,190:2534(0x7f117b676700):ZOO_INFO
2015-09-23 19:14:33,191:2534(0x7f117b676700):ZOO_INFO
2015-09-23 19:14:33,193:2534(0x7f117a674700):ZOO_INFO
2015-09-23 19:14:33,194:2534(0x7f117a674700):ZOO_INFO
I0923 19:14:33.195088 2590 group.cpp:313] Group proc
I0923 19:14:33.195560 2590 group.cpp:787] Syncing gr
I0923 19:14:33.195868 2590 group.cpp:385] Trying to
I0923 19:14:33.197350 2590 detector.cpp:138] Detecte
I0923 19:14:33.198221 2590 group.cpp:656] Trying to
W0923 19:14:33.199096 2590 detector.cpp:444] Leading
I0923 19:14:33.199683 2590 detector.cpp:481] A new l
I0923 19:14:33.201305 2591 sched.cpp:254] New master
I0923 19:14:33.202042 2591 sched.cpp:264] No credit
I0923 19:14:33.204704 2591 sched.cpp:448] Framework
:ready
```

Demo: (Web Interface)

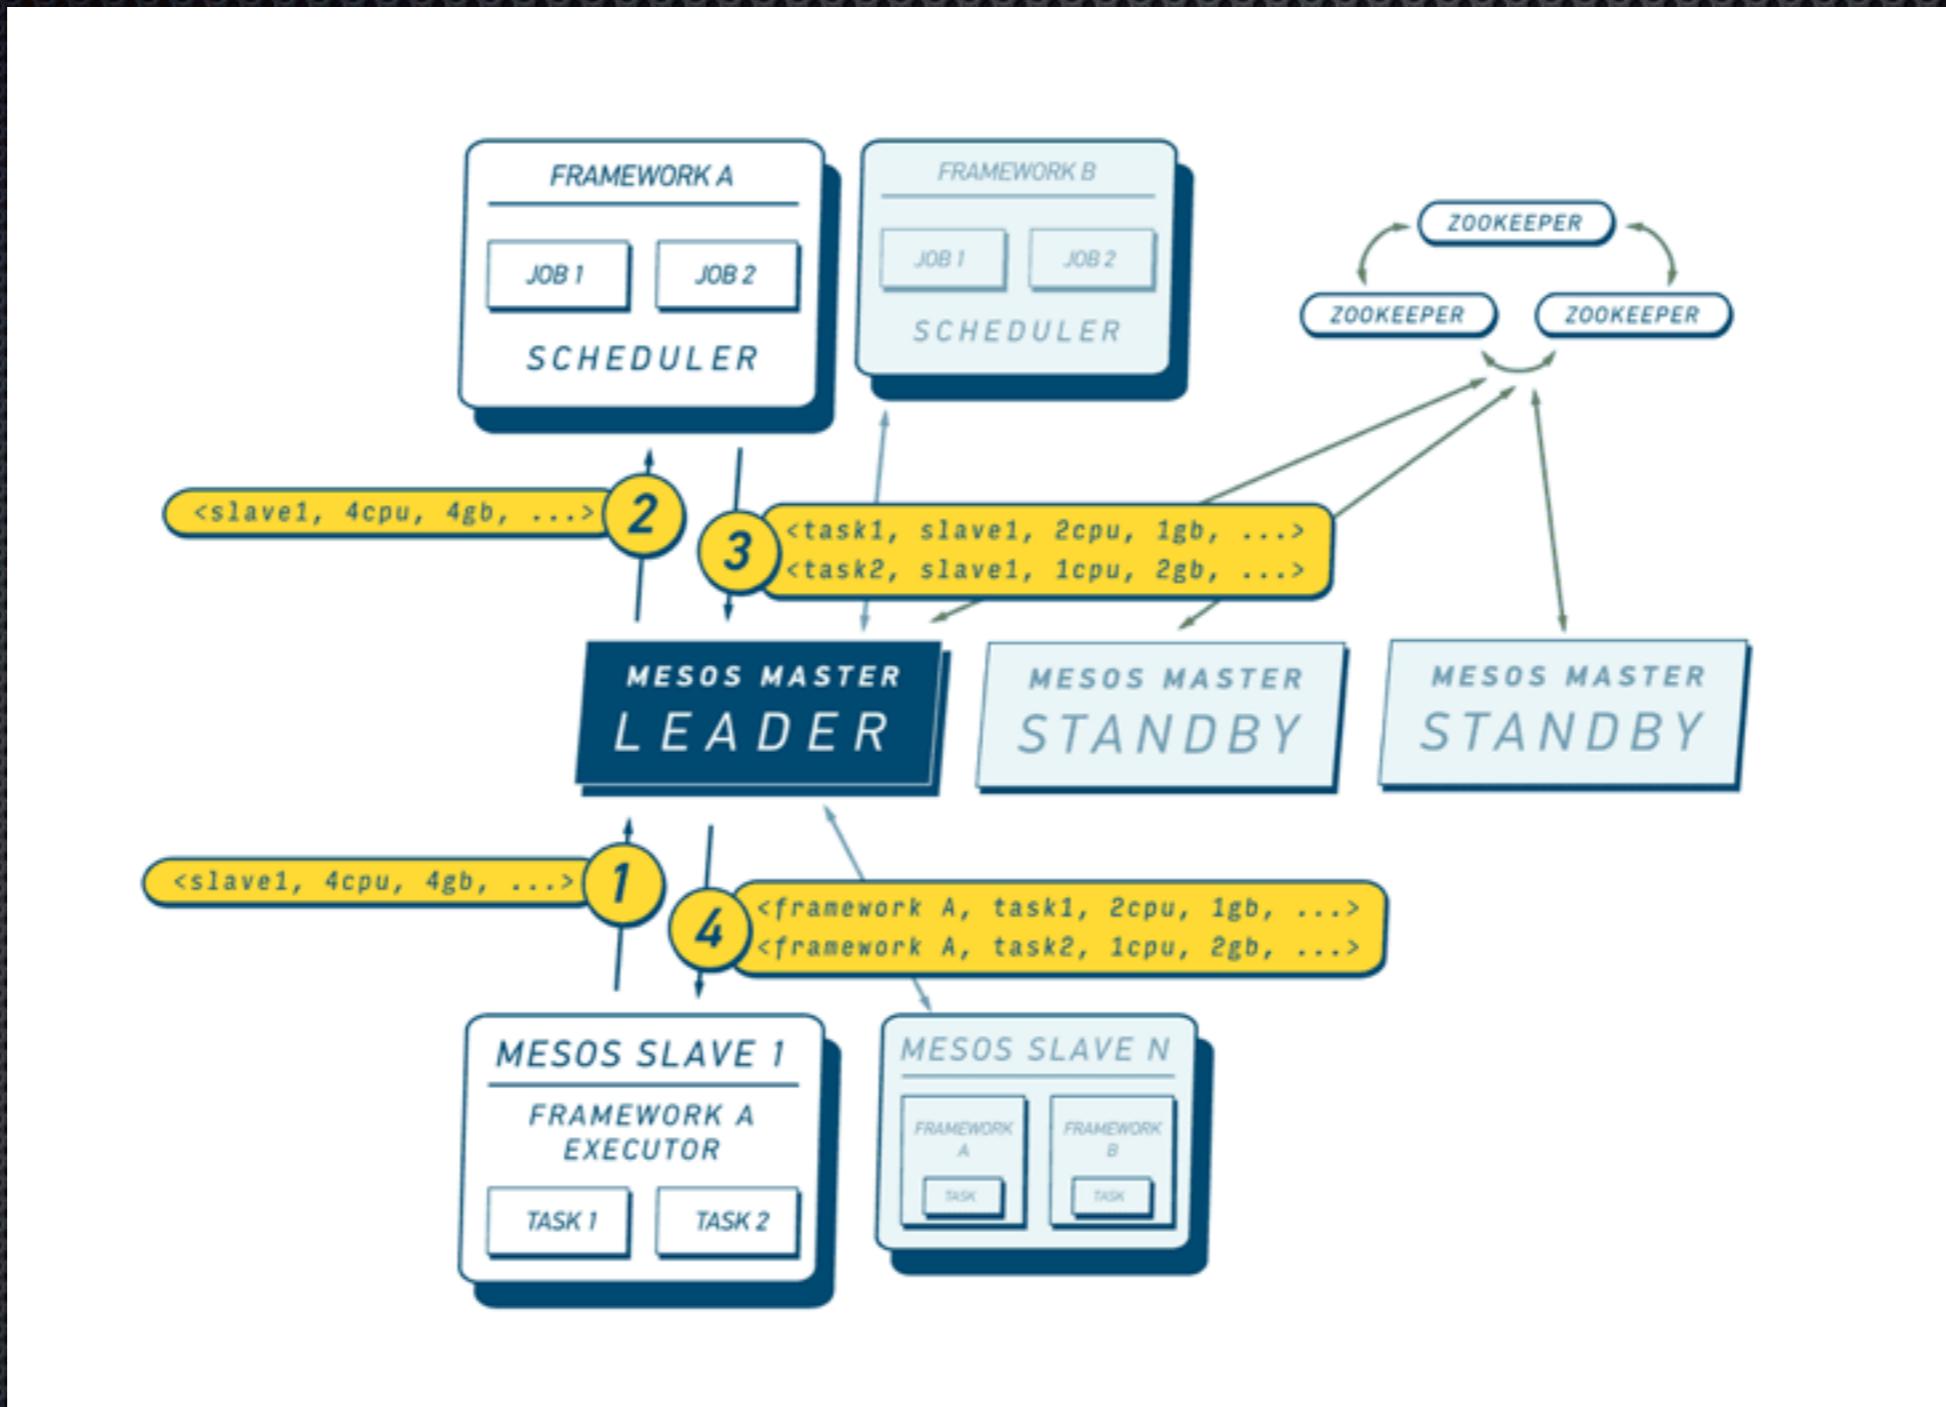
- <http://10.10.4.2:5050/#/>

Active Tasks					
ID	Name	State	Started ▼	Host	
f7bfe0fe-11a5-4670-a40f-858b2684d734	hello-mesos	RUNNING	2 minutes ago	10.10.4.10	Sandbox
748688f9-fcb0-4a8c-9e42-07eef2668278	hello-mesos	RUNNING	2 minutes ago	10.10.4.11	Sandbox

```
Registered executor on 10.10.4.10
Starting task f7bfe0fe-11a5-4670-a40f-858b2684d734
sh -c 'while true; do echo "Hey Mesos"; sleep 5; done'
Hey Mesos
Forked command at 2255
Hey Mesos
Hey Mesos
```

Active Frameworks										
ID ▼	Host	User	Name	Active Tasks	CPUs	Mem	Max Share	Registered	Re-Registered	
...5050-1817-0001	master	vagrant	hello-mesos	2	1	256 MB	50%	9 minutes ago	-	
...5050-1817-0000	master	root	marathon	0	0	0 B	0%	12 minutes ago	-	

How do Frameworks work?



Why Clojure?

- Clojure is a great example of “Making Easy Things Easy & Hard Things Possible”
- Distributed Systems are hard because of state management
 - You have to be very explicit
 - Clojure already forces you into this mindset
- Cleaner Code
- A REPL for your distributed system
- clj-mesos is really great!

Why Clojure? Java

```
Protos.TaskID taskId = Protos.TaskID.newBuilder()
    .setValue(Integer.toString(taskIDGenerator.incrementAndGet())).build();

Protos.ContainerInfo.DockerInfo.Builder dockerInfoBuilder = Protos.ContainerInfo.DockerInfo.newBuilder();
dockerInfoBuilder.setImage("busybox");

Protos.ContainerInfo.Builder containerInfoBuilder = Protos.ContainerInfo.newBuilder();
containerInfoBuilder.setType(Protos.ContainerInfo.Type.DOCKER);
containerInfoBuilder.setDocker(dockerInfoBuilder.build());

Protos.CommandInfo.Builder commandInfoBuilder = Protos.CommandInfo.newBuilder();
commandInfoBuilder.setShell(true);
commandInfoBuilder.setValue("while true; do echo \"Hey Mesos\"; sleep 5; done");

Protos.TaskInfo task = Protos.TaskInfo.newBuilder()
    .setName("task " + taskId.getValue())
    .setTaskId(taskId)
    .setSlaveId(offer.getSlaveId())
    .addResources(Protos.Resource.newBuilder()
        . setName("cpus")
        . setType(Protos.Value.Type.SCALAR)
        . setScalar(Protos.Value.Scalar.newBuilder().setValue(1)))
    .addResources(Protos.Resource.newBuilder()
        . setName("mem")
        . setType(Protos.Value.Type.SCALAR)
        . setScalar(Protos.Value.Scalar.newBuilder().setValue(128)))
    .setContainer(containerInfoBuilder)
    .setCommand(commandInfoBuilder)
    .build();
```

Why Clojure? clj-mesos

```
(defn docker-task-info
  [uuid {:keys [slave-id]}]
  [{:name "hello-mesos"
    :task-id uuid
    :slave-id slave-id
    :resources {:cpus min-cpu
                :mem min-mem}
    :command {:shell true
              :value "while true; do echo \"Hey Mesos\"; sleep 5; done"}
    :container {:type :docker
                :docker {:image "busybox"}}}])
```

Hello-World Repository

- git clone <https://github.com/clj-mesos-workshop/hello-mesos.git>

Step 1

Generating
a Mesos-Framework



Step 1: Commands

- Try
 - `lein new mesos-framework awesome-framework`
- Complete version
 - `git checkout step-one-generate-from-template`



Step 2
Storing State
using Zookeeper

Step 2: Commands

- git checkout step-two-add-zookeeper-for-state

Step 3

Recovering from Failure

FAILURE
IS ONLY THE
opportunity
to
BEGIN AGAIN,
ONLY THIS TIME MORE
Wisely.

~HENRY FORD

Step 3: Commands

- Complete version
- git checkout step-three-recover-from-failure



Step 4

Leader Election

Step 4: Commands

- git checkout step-four-leader-election

Step 5

Web UI

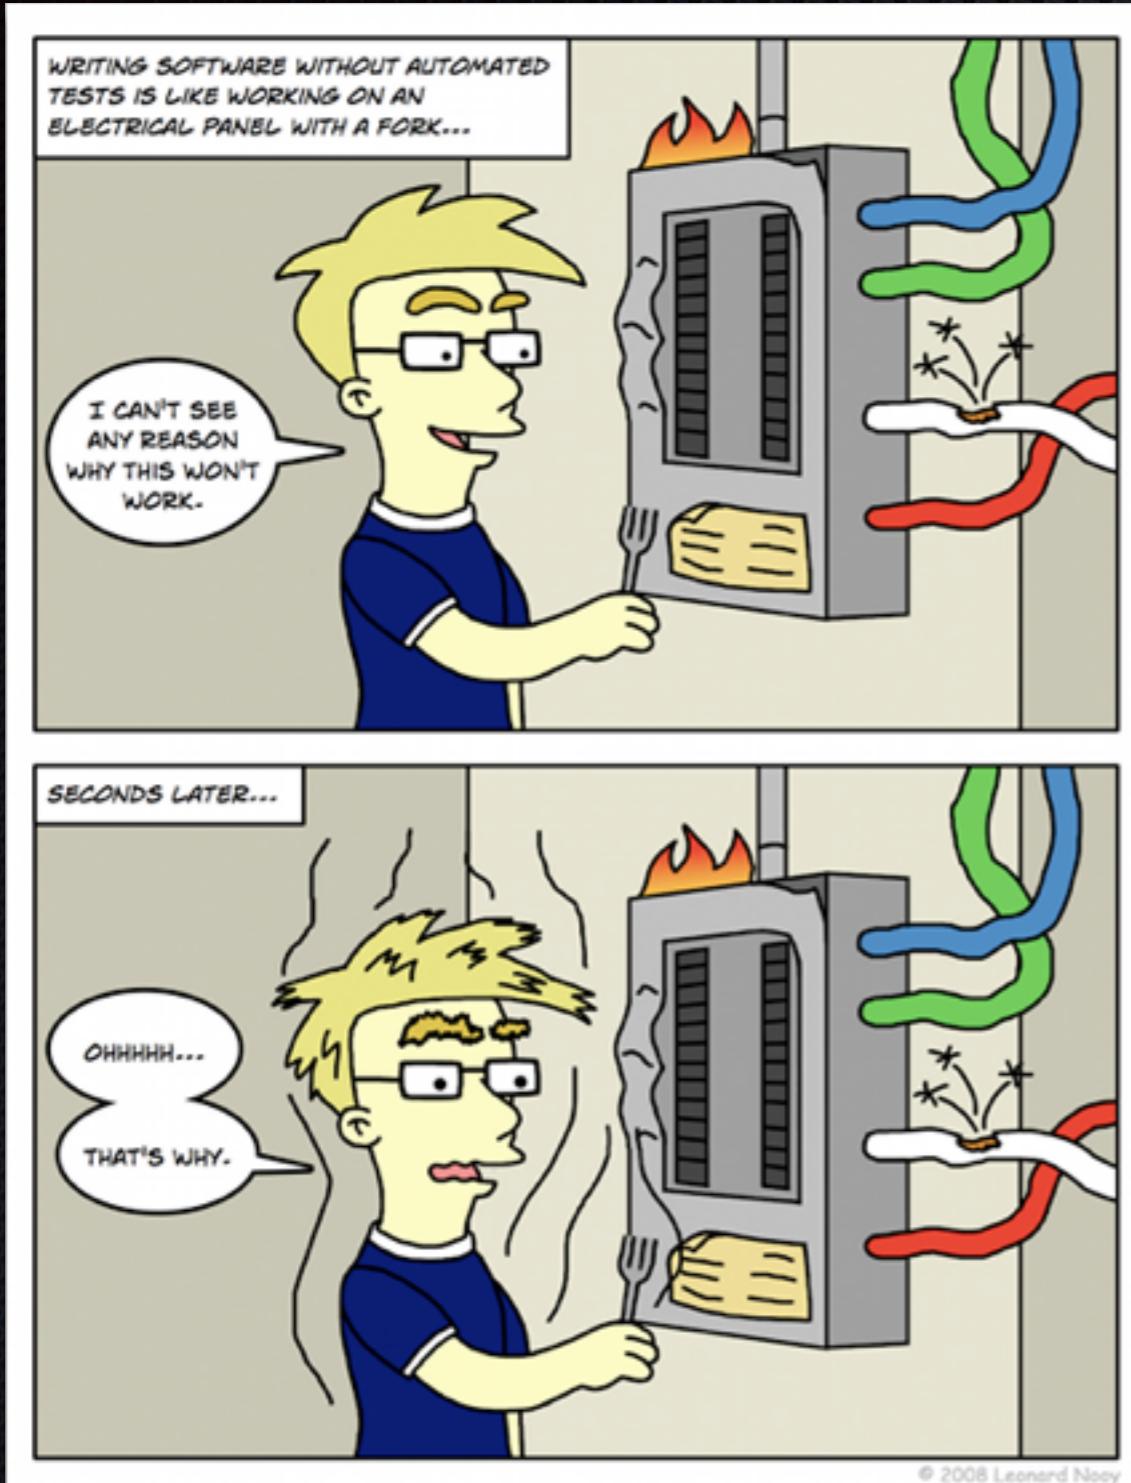
```
!DOCTYPE HTML PUBLIC  
<html>  
  <head>  
    <meta name="viewport" content="width=device-width, initial-scale=1.0, maximum-scale=1.0, user-scalable=0" />  
    <meta name="apple-mobile-web-app-capable" content="yes" />  
    <meta name="apple-mobile-web-app-status-bar-style" content="black" />  
    <link href="https://www.google.com" rel="stylesheet" type="text/css" />  
    <script src="https://www.google.com" type="text/javascript" />  
  </head>  
  <body>
```

Step 5: Commands

- git checkout step-five-web-ui

Step 6

Testing



Step 6: Commands

- `git checkout step-six-testing-the-scheduler`

**YOU THINK THIS IS THE
END?**



**THIS IS JUST THE
BEGINNING!**

memegenerator.net

Welcome to StrangeLoop '15